



solo.io

Multi Cluster Service Mesh Patterns, Operations, and Extensibility with WebAssembly

Idit Levine

Founder and CEO

[@udit_levine](https://twitter.com/udit_levine)

Christian Posta

Global Field CTO


[@christianposta](https://twitter.com/christianposta)

Challenges in Adopting Microservices

- Improve velocity of teams building and delivering code
- Decentralized implementations vs centralized operations
- Connect and include existing systems and investments
- Improve security posture
- Stay within regulations and compliance

A Solution is many, smaller clusters

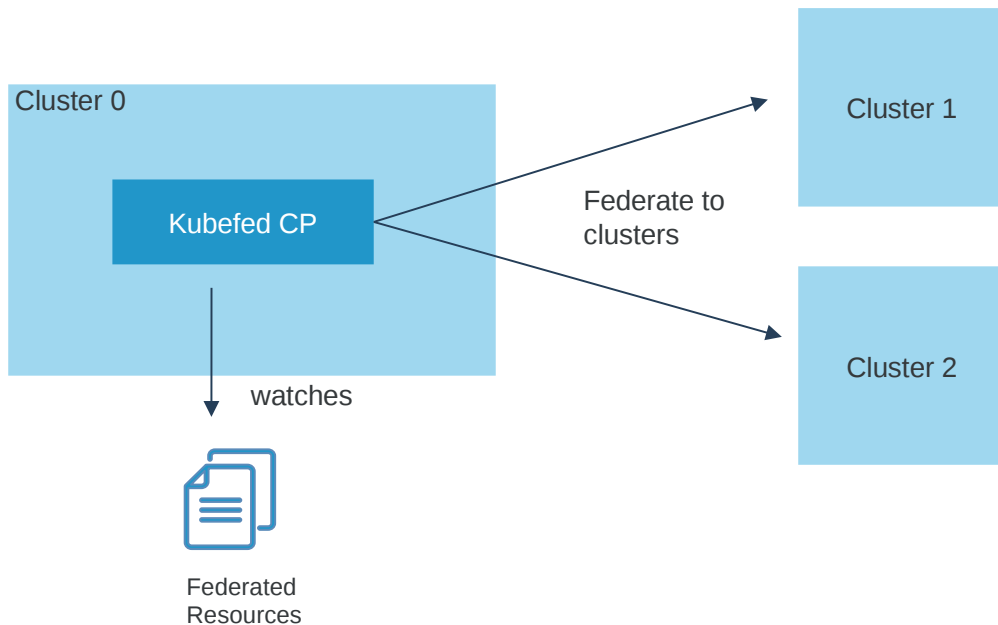
- High availability
- Compliance
- Isolation / Autonomy
- Scale
- Data locality, cost
- Public / DMZ / Private networks



Considerations for multiple clusters

- Exact replicas of each other, same fleet?
- Separate, non-uniform deployments?
- Single operational/administrative control
- Segmented by network? Segmented by team?
- Independent administration?
- Cluster federation? Independent clusters, which services are shared?

Example: Kubefed




```
apiVersion:
types.kubefed.io/v1beta1
kind: FederatedService
metadata:
  name: echo-server
spec:
  placement:
    clusterSelector:
      matchLabels: {}
  template:
    metadata:
      labels:
        app: echo-server
    spec:
      ports:
        - name: http
          port: 8080
      selector:
        app: echo-server
```


Services Need to Communicate with Each Other

Patterns within and outside of Kubernetes


Pattern: Flat network across pods



Pattern: Different network, expose all services



Pattern: Different network, controlled gateway



Forces to balance

- Security (AuthN/Z, encryption, identity)
- Service discovery
- Failover / traffic shifting / transparent routing
- Observability
- Separate networks
- Well-defined fault domains
- Building for scale



Can you build these patterns
just using Kubernetes?



AWS App Mesh

Service Mesh can help


Envoy Proxy is the magic behind service mesh



Envoy Implements the following:


- Zone aware, priority/locality load balancing
- Circuit breaking, outlier detection
- Timeouts, retries, retry budgets
- Traffic shadowing
- Request racing
- Rate limiting
- RBAC, TLS origination/termination
- Access logging, statistics collection

Envoy Proxy to do application networking heavy lifting



- Transparent client-side routing decisions
- TLS orig/termination
- Circuit breaking
- Stats collection

Envoy Proxy as backbone for multi-cluster communication federation



Other key Envoy Proxy features

- Request hedging
- Retry Budgets
- Load balancing priorities
- Locality weighted load balancing
- Zone aware routing
- Degraded endpoints (fallback)
- Aggregated clusters

Multi-Cluster Examples

Service Mesh examples using Envoy Proxy

Three Examples and Considerations

1

Shared Service Mesh
Control Plane with Flat
Network


2

Shared Service Mesh
Control Plane with
Separate Networks

3


Separate Service Mesh
Control Planes with
Separate Networks

1 - Istio shared control plane, flat network




- Simplest setup for multi-cluster
- No special Envoy routing (though may use zone-aware)
- Shared control plane increases the failure domain to multiple clusters
- Use simpler, flat networking if possible
- No special considerations for identity (identity domain is shared)
- Still need to federate telemetry collection

2 - Istio shared control plane, separate networks



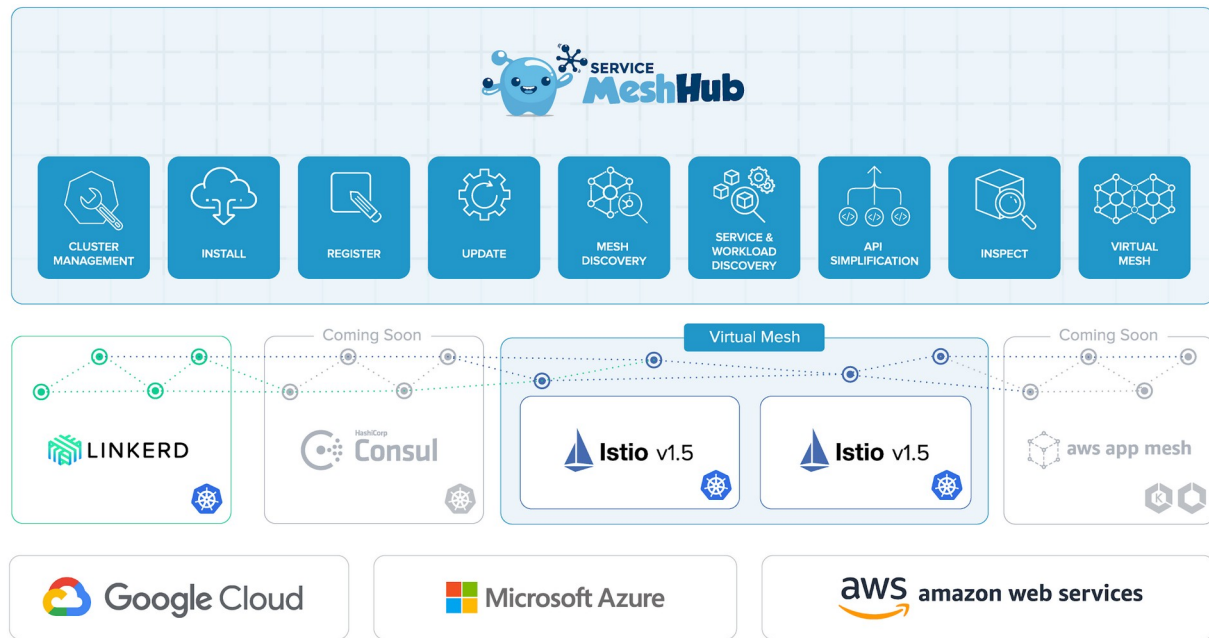
- Gateway allows communication between networks
- Envoy Locality Weighted LB (for the gateway endpoints). Istio calls this “split horizon EDS”.
- Shares same failure domain across all clusters
- Gateways facilitate communication AND control plane
- Slight increase in burden on operator to label networks and gateway endpoints correctly so Istio has that information

3 - Istio separate control plane, separate networks



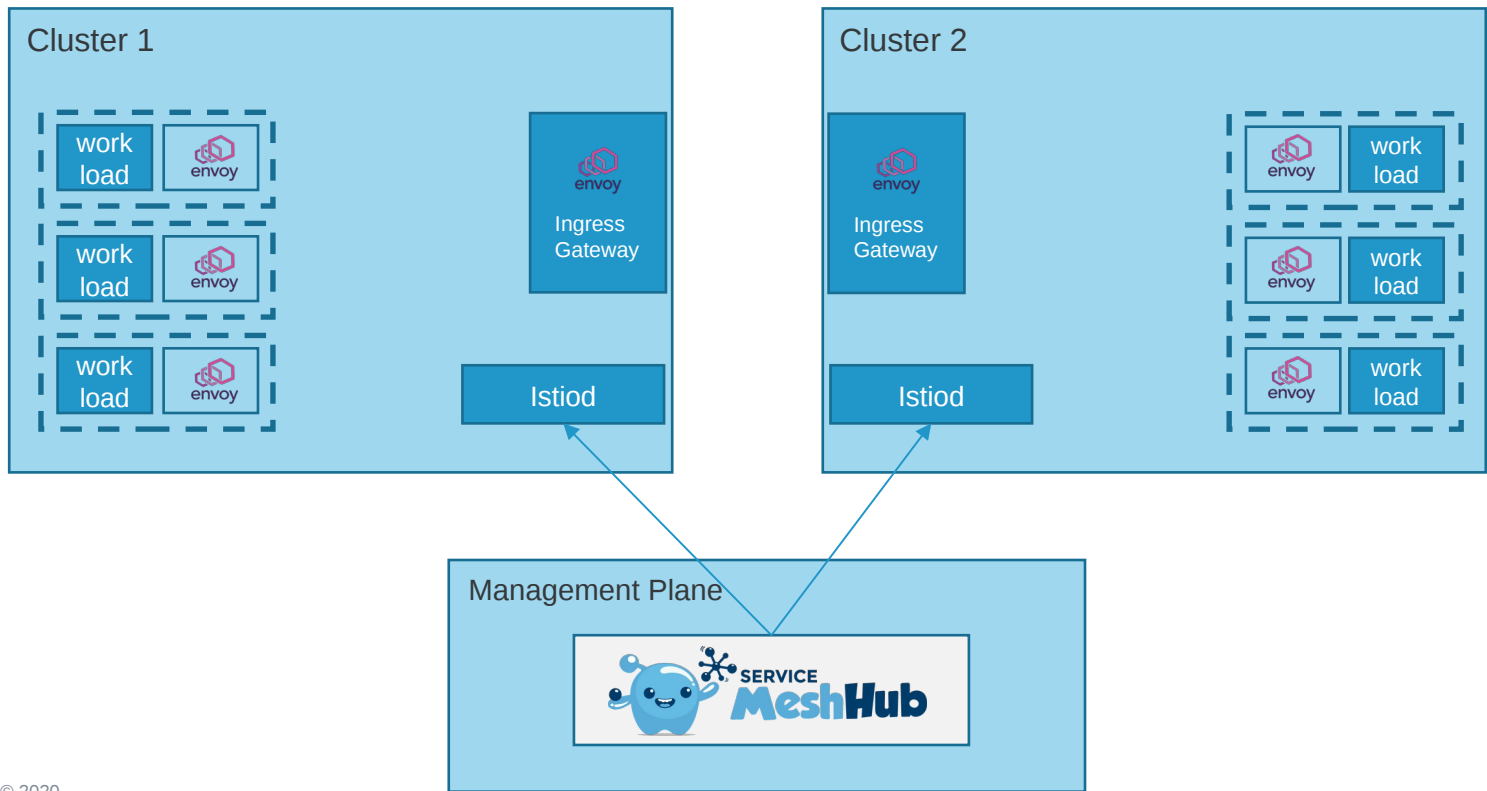
- Gateway allows communication between networks
- Uses Istio ServiceEntry to enable cross-network discovery
- Independent control planes
- Separate, independent failure domains
- Doesn't solve where trust domains MUST be separate (with federation at the boundaries)
- Increased operator burden to maintain service discovery, identity federation, and configuration across meshes

Open Source Service Mesh Hub



- Open Source
- Install & Configure
- Discovery
- Unified API
- Virtual Mesh
- Shared Identity
- Developer friendly tooling

Open Source Service Mesh Hub



Demo

Extending the Service Mesh Data Plane with WebAssembly

What is WebAssembly? It's neither web nor assembly

WebAssembly (abbreviated Wasm) is a binary instruction format for a stack-based virtual machine. Wasm is designed as a portable target for compilation of high-level languages like C/C++/Rust, enabling deployment on the web for client and server applications.

-> No Assembly -> Binary Instruction Format

-> No Web -> WASI

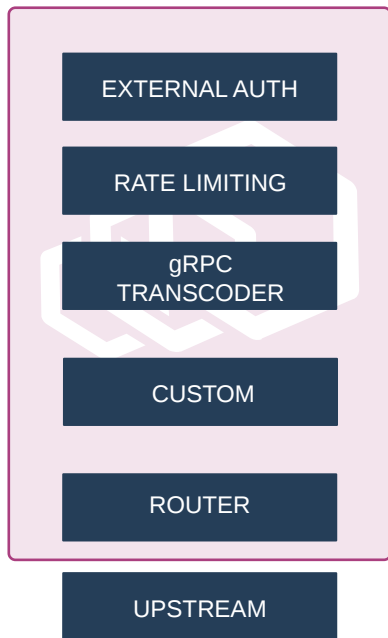
PORTABLE

SECURE

FAST


**ANY
LANGUAGE**

Challenges in customizing Envoy Proxy

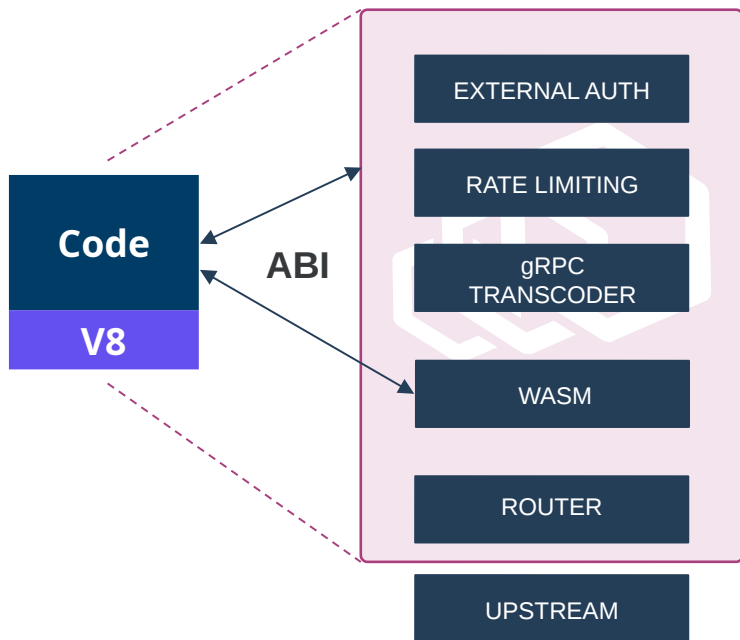


- Must maintain efficient performance for the data plane
- Writing custom filters is difficult and limited to C++
- Changes built into Envoy, maintain separate distro
- Hard dependencies, cascading failures
- Must stop and recompile

WebAssembly in Envoy Proxy



WebAssembly in Envoy Proxy



ABI: Application Binary Interface

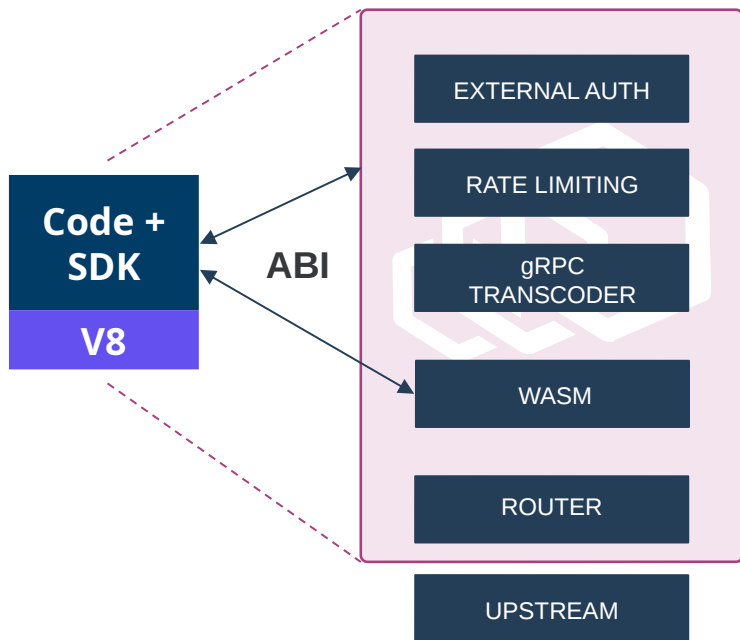
HTTP (L7) extensions

`proxy_on_http_request_headers`


- params:
 - `i32 (uint32_t) context_id`
 - `i32 (size_t) num_headers`
 - `i32 (bool) end_of_stream`
- returns:
 - `i32 (proxy_action_t) next_action`

Called when HTTP request headers are received from the client. Headers can be retrieved using `proxy_get_map` and/or `proxy_get_map_value`.

WebAssembly in Envoy Proxy



SDK Includes the following with more coming soon.



WebAssembly Hub



- Developer and operator workflow to build, share, and deploy
- Team and user management
- Operator uses declarative CRD
- Modules are OCI images
<https://github.com/solo-io/wasm-image-spec>

Demo



About Us - We're Hiring!

solo.io

WebAssembly Hub

webassemblyhub.io

Service Mesh Hub - Community meetings starting June 17th

github.com/solo-io/service-mesh-hub

Join the Community

slack.solo.io

