

# Join us for KubeCon + CloudNativeCon Virtual



Event dates: **August 17-20, 2020**

Schedule: [Now available!](#)

Cost: **\$75**

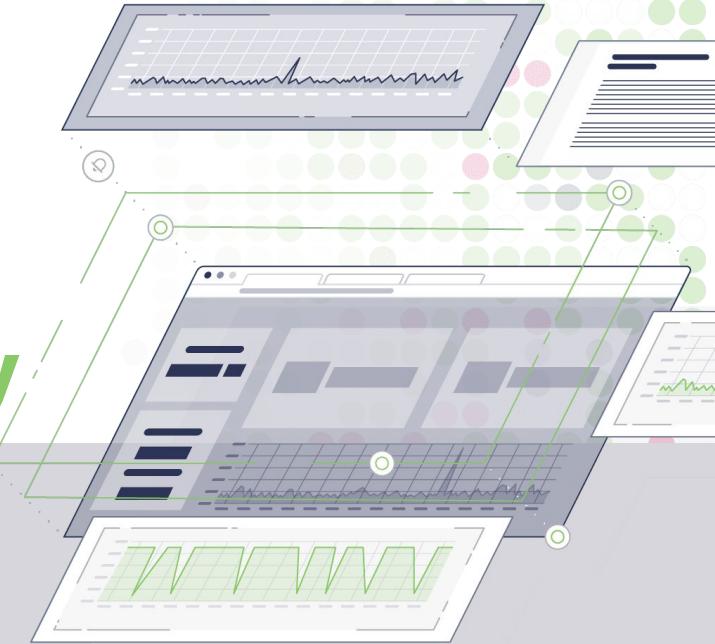
[\*\*Register now!\*\*](#)





# The top 7 most useful Kubernetes APIs for comprehensive cloud-native observability

Caleb Hailey, CEO & Co-founder  
Sensu, Inc. | [sensu.io](http://sensu.io)



# Overview

- Introduction
- Sensu ❤️ the CNCF
- The CNCF Cloud Native Trail Map
- Top 7 K8s APIs
- Getting to the “Why?” for observability



# Kubernetes APIs

## Introductions



# Sensu



# What is Sensu?

**Unified infrastructure and application monitoring at scale.** Gain deep visibility into servers, containers, services, applications, functions, and connected devices across any public or private cloud.

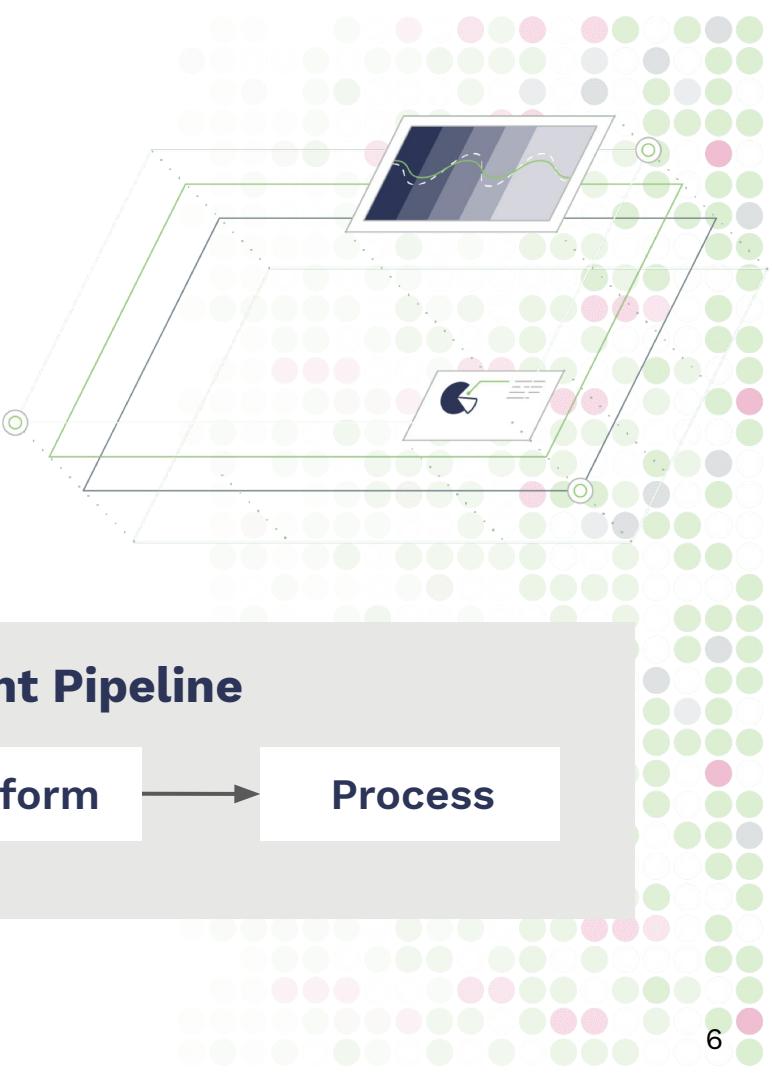
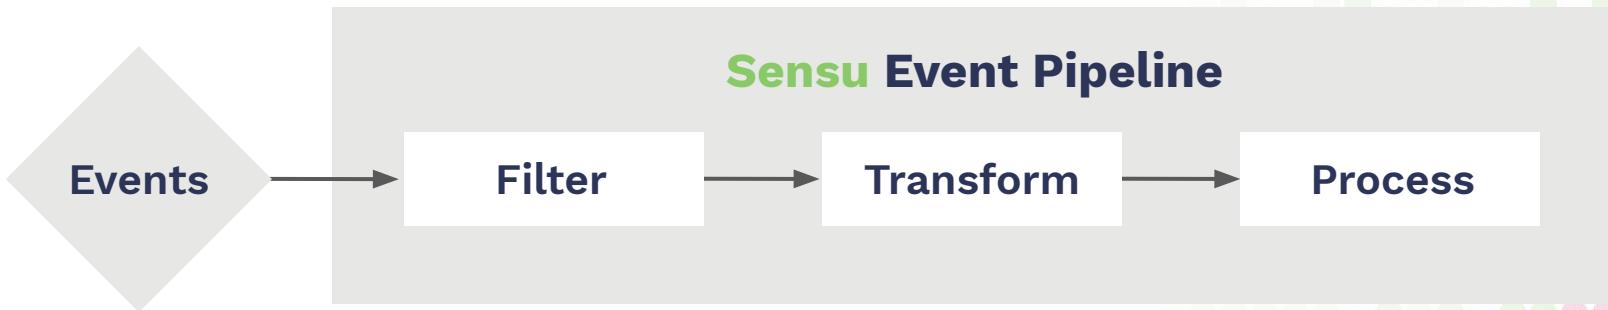
[sensu.io](https://sensu.io) | @sensu





# How does Sensu work?

Sensu is a flexible observability pipeline.

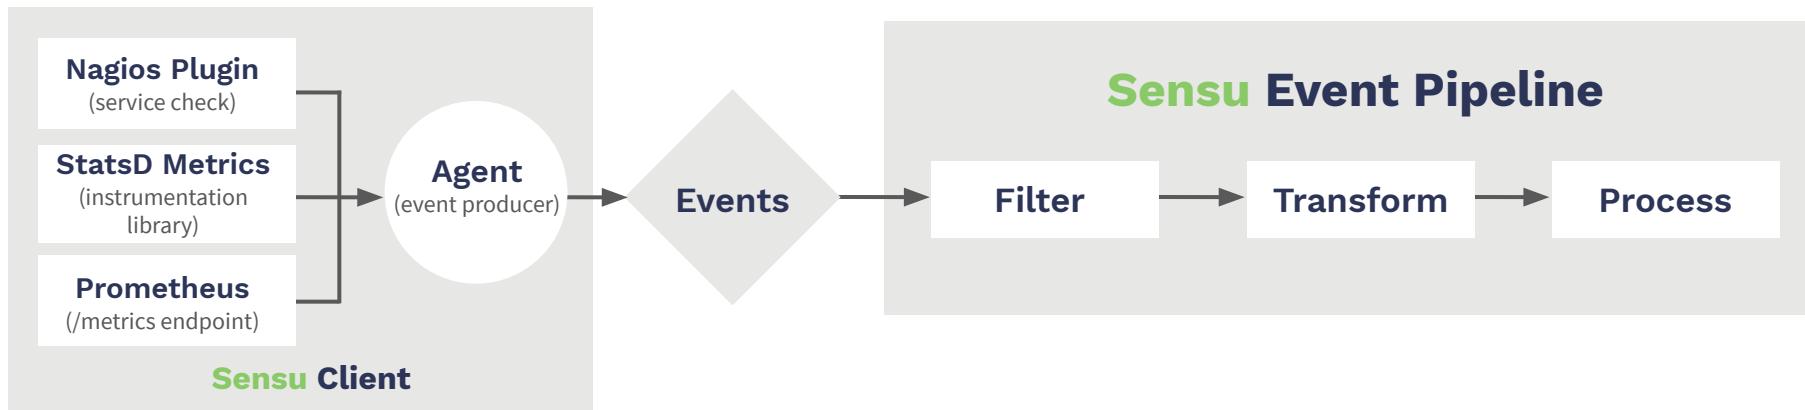




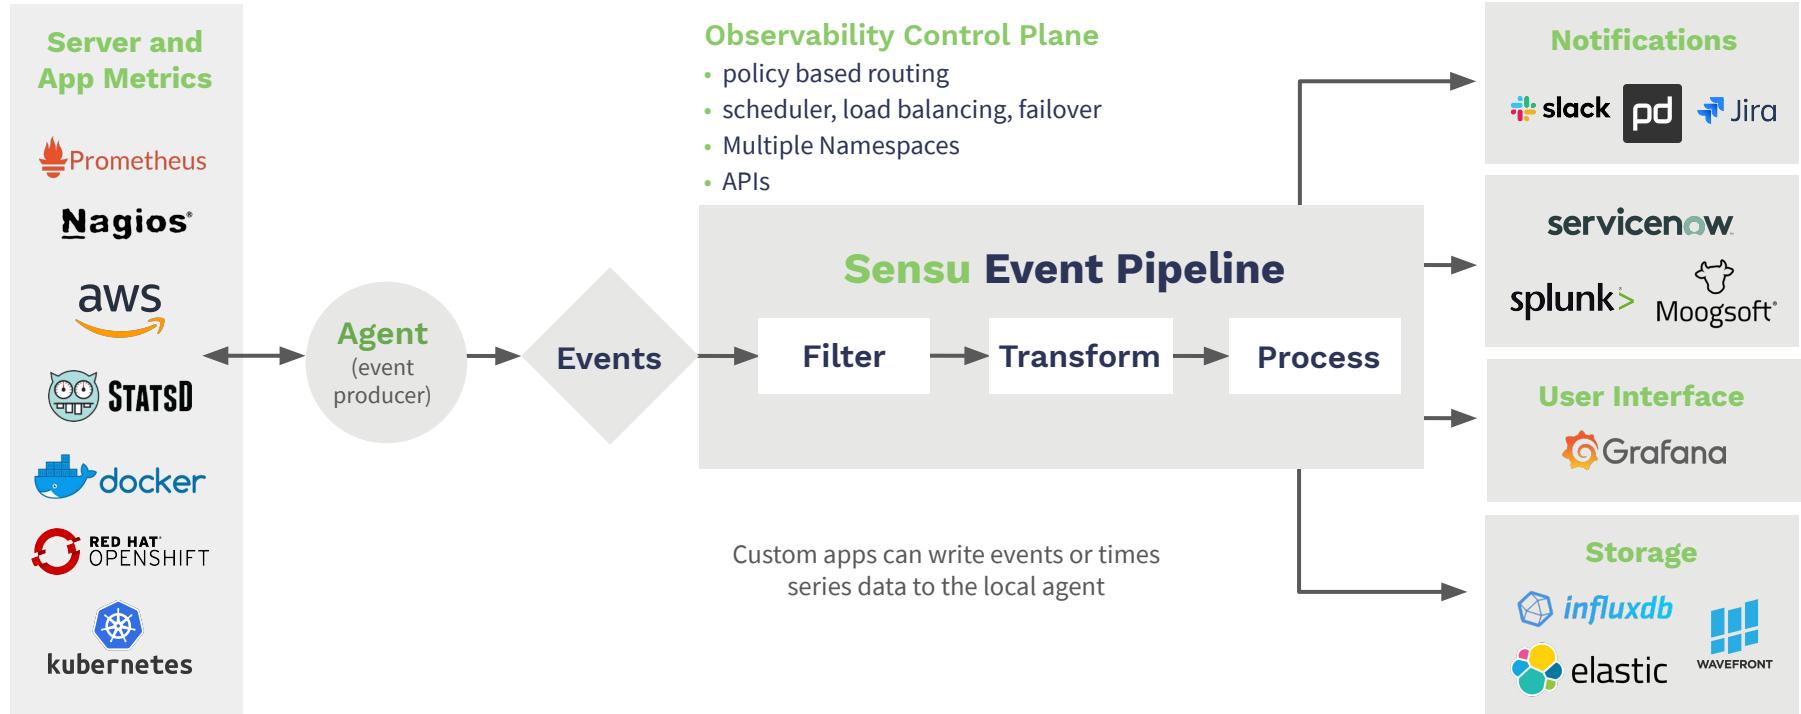
# Sensu Go

## The observability pipeline

The Sensu Go agent is a powerful cross-platform event producer.



# Observability Pipeline Architecture



# Sensu ❤️ the CNCF

- "Monitoring Kubernetes" (blog post)  
<https://www.cncf.io/blog/2019/01/09/monitoring-kubernetes-part-1-the-challenges-data-sources>
- "Monitoring Kubernetes Workloads: the Sidecar Pattern"  
<https://youtu.be/X14nPcOnUg0>
- "How Kubernetes works" (blog post)  
<https://www.cncf.io/blog/2019/08/19/how-kubernetes-works>
- KubeCon 2018-19 Sponsors



# Cloud Native Journey



"You can't perform any operation  
without proper observability."

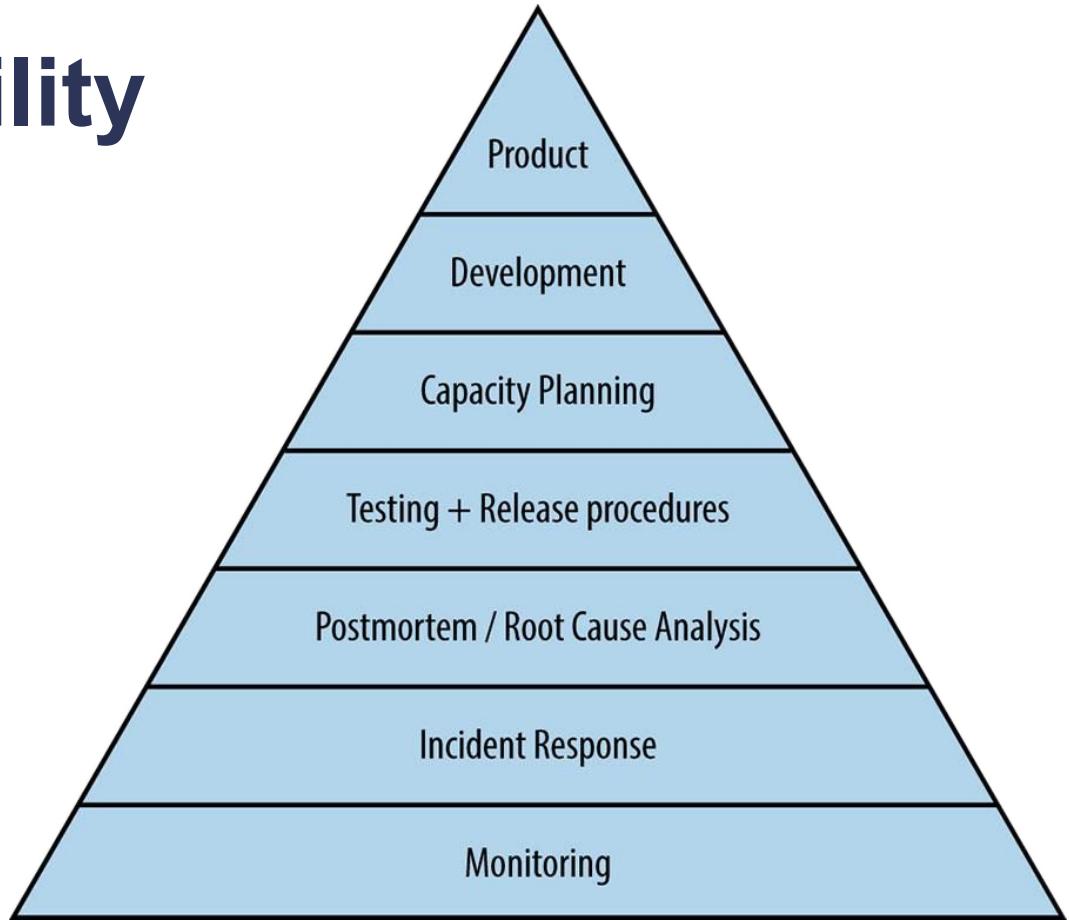
– Gene Kim

# Service Reliability Hierarchy

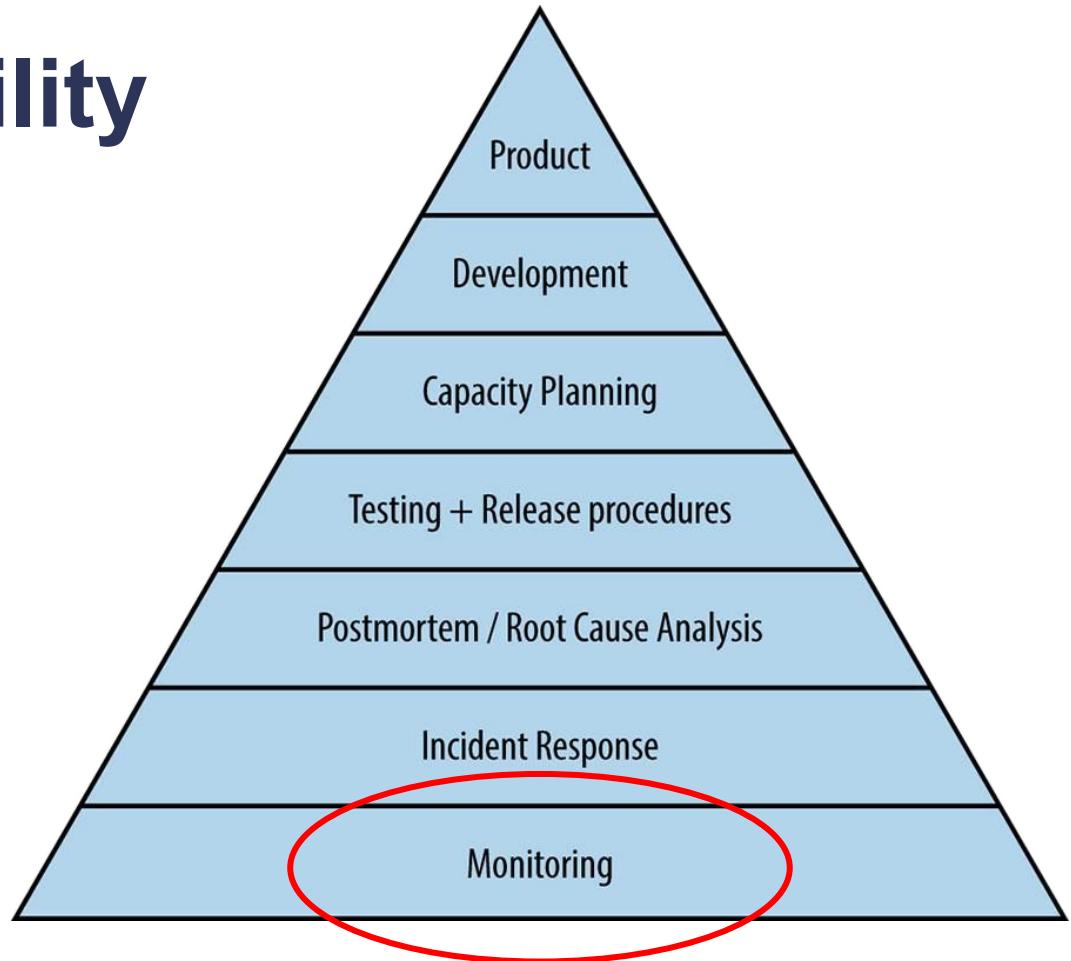
Part I. Introduction

Part II. Principles

Part III. Practices



# Service Reliability Hierarchy



# The CNCF Trail Map

## 1. CONTAINERIZATION

- Commonly done with Docker containers
- Any size application and dependencies (even PDP-11 code running on an emulator) can be containerized
- Over time, you should aspire towards splitting suitable applications and writing future functionality as microservices



## 3. ORCHESTRATION & APPLICATION DEFINITION

- Kubernetes is the market-leading orchestration solution
- You should select a Certified Kubernetes Distribution, Hosted Platform, or Installer: [cncf.io/ck](https://cncf.io/ck)
- Helm Charts help you define, install, and upgrade even the most complex Kubernetes application



## 5. SERVICE PROXY, DISCOVERY, & MESH

- CoreDNS is a fast and flexible tool that is useful for service discovery
- Envoy and Linkerd each enable service mesh architectures
- They offer health checking, routing, and load balancing



## 2. CI/CD

- Setup Continuous Integration/Continuous Delivery (CI/CD) so that changes to your source code automatically result in a new container being built, tested, and deployed to staging and eventually, perhaps, to production
- Setup automated rollouts, roll backs and testing
- Argo is a set of Kubernetes-native tools for deploying and running jobs, applications, workflows, and events using GitOps paradigms such as continuous and progressive delivery and MLOps



CNCF Incubating

## 4. OBSERVABILITY & ANALYSIS

- Pick solutions for monitoring, logging and tracing
- Consider CNCF projects Prometheus for monitoring, Fluentd for logging and Jaeger for Tracing
- For tracing, look for an OpenTracing-compatible implementation like Jaeger



# The CNCF Trail Map

## 1. CONTAINERIZATION

- Commonly done with Docker containers
- Any size application and dependencies (even PDP-11 code running on an emulator) can be containerized
- Over time, you should aspire towards splitting suitable applications and writing future functionality as microservices



## 3. ORCHESTRATION & APPLICATION DEFINITION

- Kubernetes is the market-leading orchestration solution
- You should select a Certified Kubernetes Distribution, Hosted Platform, or Installer: [cncf.io/ck](https://cncf.io/ck)
- Helm Charts help you define, install, and upgrade even the most complex Kubernetes application



kubernetes  
CNCF Graduated



HELM  
CNCF Graduated

## 5. SERVICE PROXY, DISCOVERY, & MESH

- CoreDNS is a fast and flexible tool that is useful for service discovery
- Envoy and Linkerd each enable service mesh architectures
- They offer health checking, routing, and load balancing



## 2. CI/CD

- Setup Continuous Integration/Continuous Delivery (CI/CD) so that changes to your source code automatically result in a new container being built, tested, and deployed to staging and eventually, perhaps, to production
- Setup automated rollouts, roll backs and testing
- Argo is a set of Kubernetes-native tools for deploying and running jobs, applications, workflows, and events using GitOps paradigms such as continuous and progressive delivery and MLOps



argo  
CNCF Incubating

## 4. OBSERVABILITY & ANALYSIS

- Pick solutions for monitoring, logging and tracing
- Consider CNCF projects Prometheus for monitoring, Fluentd for logging and Jaeger for Tracing
- For tracing, look for an OpenTracing-compatible implementation like Jaeger



Prometheus

CNCF Graduated



fluentd

CNCF Graduated



JAAGER

CNCF Graduated



OPENTRACING

CNCF Incubating



# CNCF's Cloud Native Trail Map



"This Cloud Native Trail Map is a recommended process for leveraging open source, cloud native technologies. At each step, you can choose a vendor-supported offering or do it yourself, and **everything after step #3 is optional based on your circumstances.**"

# Top 7 K8s APIs for O11y



## # 7: K8s Metrics APIs

- Most well-known APIs for monitoring
- Built-in Prometheus exporters
- kubelet-metrics (including cAdvisor)
- kube-state-metrics (cluster level)
- **Consumable by most observability tools**

# # 6: Service APIs

- Ingress, Endpoint, and Service resources
- Service **metadata**, **spec**, and **status**
- Service ports, internal and external IP addresses, load balancing, and label selector configuration details
- **kubectl describe service <service>**

# # 5: Container API



- **Pod API resource**

- Information about running containers
- Container **status** and **details** including image source, commands, networking, execution environment, and resource requirements

## # 4: Pod API

- Primary/core workload API resource
- `kubectl describe pod <pod name>`
- Pod **metadata**, **spec**, and **status**
- Controller references
- Read log API (misc operation)



Terminal — bash — 115x30

```
[homelab $ kubectl --namespace cncf-webinar logs nginx-example-5bd95679c8-bxqbh -c nginx --tail 10
10.24.1.1 -- [06/Jul/2020:18:21:34 +0000] "HEAD / HTTP/1.1" 200 0 "-" "curl/7.64.0" "-"
10.24.1.1 -- [06/Jul/2020:18:21:35 +0000] "HEAD / HTTP/1.1" 200 0 "-" "curl/7.64.0" "-"
10.138.15.223 -- [06/Jul/2020:18:21:38 +0000] "HEAD / HTTP/1.1" 200 0 "-" "curl/7.64.0" "-"
10.138.15.223 -- [06/Jul/2020:18:21:48 +0000] "HEAD / HTTP/1.1" 200 0 "-" "curl/7.64.0" "-"
10.24.1.1 -- [06/Jul/2020:18:23:33 +0000] "HEAD / HTTP/1.1" 200 0 "-" "curl/7.64.0" "-"
10.138.15.223 -- [06/Jul/2020:18:23:33 +0000] "HEAD / HTTP/1.1" 200 0 "-" "curl/7.64.0" "-"
10.138.15.223 -- [06/Jul/2020:18:23:34 +0000] "HEAD / HTTP/1.1" 200 0 "-" "curl/7.64.0" "-"
10.138.15.225 -- [06/Jul/2020:18:23:35 +0000] "HEAD / HTTP/1.1" 200 0 "-" "curl/7.64.0" "-"
10.24.1.1 -- [06/Jul/2020:18:23:38 +0000] "HEAD / HTTP/1.1" 200 0 "-" "curl/7.64.0" "-"
10.138.15.225 -- [06/Jul/2020:18:23:39 +0000] "HEAD / HTTP/1.1" 200 0 "-" "curl/7.64.0" "
[homelab $ curl -s "http://127.0.0.1:8888/api/v1/namespaces/cncf-webinar/pods/nginx-example-5bd95679c8-bxqbh/log?follow=false&tailLines=10&container=nginx"
10.24.1.1 -- [06/Jul/2020:18:21:34 +0000] "HEAD / HTTP/1.1" 200 0 "-" "curl/7.64.0" "-"
10.24.1.1 -- [06/Jul/2020:18:21:35 +0000] "HEAD / HTTP/1.1" 200 0 "-" "curl/7.64.0" "-"
10.138.15.223 -- [06/Jul/2020:18:21:38 +0000] "HEAD / HTTP/1.1" 200 0 "-" "curl/7.64.0" "-"
10.138.15.223 -- [06/Jul/2020:18:21:48 +0000] "HEAD / HTTP/1.1" 200 0 "-" "curl/7.64.0" "-"
10.24.1.1 -- [06/Jul/2020:18:23:33 +0000] "HEAD / HTTP/1.1" 200 0 "-" "curl/7.64.0" "-"
10.138.15.223 -- [06/Jul/2020:18:23:33 +0000] "HEAD / HTTP/1.1" 200 0 "-" "curl/7.64.0" "-"
10.138.15.223 -- [06/Jul/2020:18:23:34 +0000] "HEAD / HTTP/1.1" 200 0 "-" "curl/7.64.0" "-"
10.138.15.225 -- [06/Jul/2020:18:23:35 +0000] "HEAD / HTTP/1.1" 200 0 "-" "curl/7.64.0" "-"
10.24.1.1 -- [06/Jul/2020:18:23:38 +0000] "HEAD / HTTP/1.1" 200 0 "-" "curl/7.64.0" "-"
10.138.15.225 -- [06/Jul/2020:18:23:39 +0000] "HEAD / HTTP/1.1" 200 0 "-" "curl/7.64.0" "
[homelab $ ]
```

## # 3: Kubernetes Downward API



The Downward API enables Pods to expose information about themselves to Containers running in the Pod via files or environment variables.

# # 3: Kubernetes Downward API



## Motivation for the Downward API

It is sometimes useful for a Container to have information about itself, without being overly coupled to Kubernetes. The Downward API allows containers to consume information about themselves or the cluster without using the Kubernetes client or API server.

An example is an existing application that assumes a particular well-known environment variable holds a unique identifier. One possibility is to wrap the application, but that is tedious and error prone, and it violates the goal of low coupling. A better option would be to use the Pod's name as an identifier, and inject the Pod's name into the well-known environment variable.

<https://kubernetes.io/docs/~downward-api-volume-expose-pod-information/>

# # 3: Kubernetes Downward API

- Pod configuration directive
- Expose pod information via files or env vars
- **Alternative/complement to serviceAccount**
- Useful w/ monitoring & observability tools
- Enrich observations w/ rich context

# # 3: Kubernetes Downward API

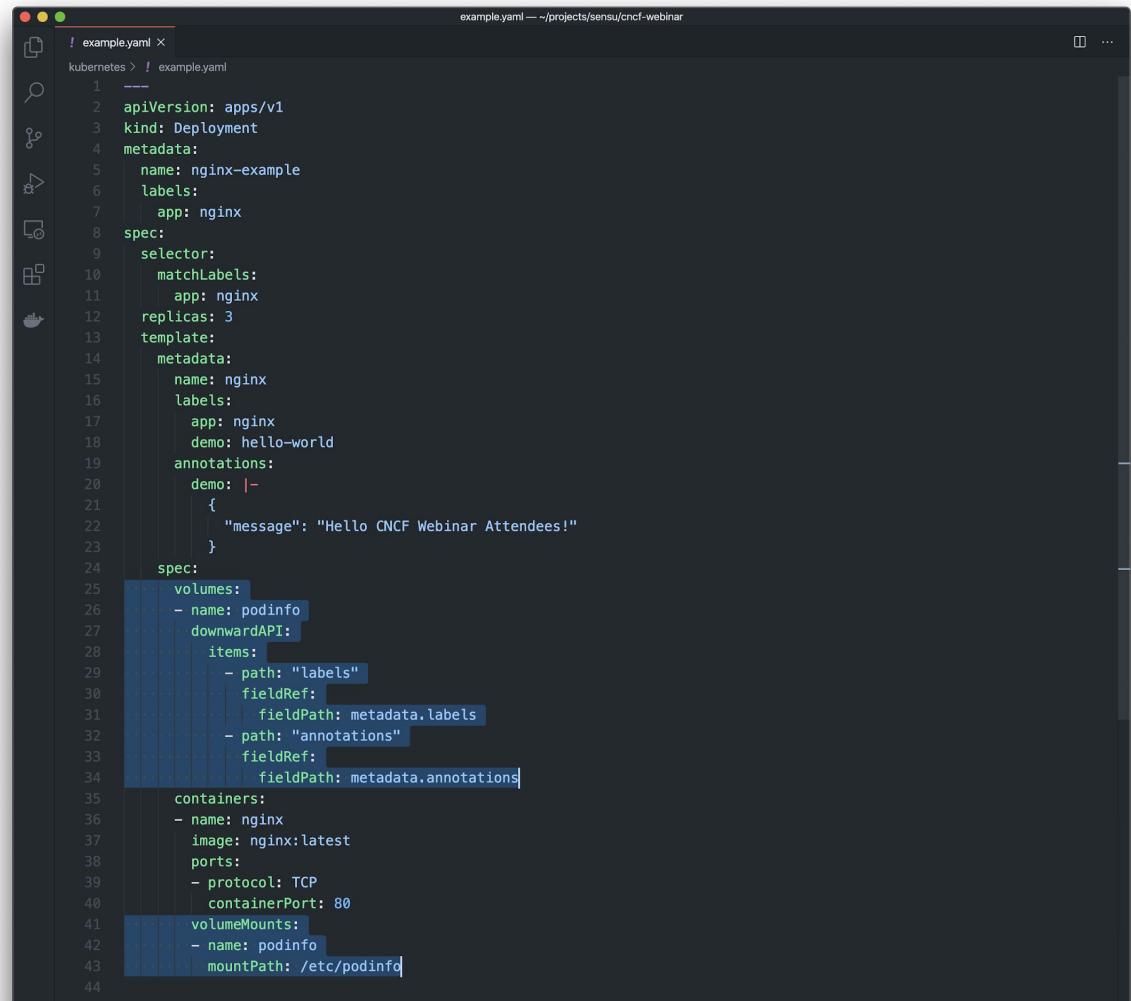
- Selectors: **fieldRef** & **resourceFieldRef**
- Available fields:
  - **metadata** (labels & annotations)
  - **spec** (nodeName, serviceAccountName)
  - **status** (status.podIP, status.hostIP)

# K8s Downward API

Configured as volume  
mount

-OR-

Configured as  
environment variable



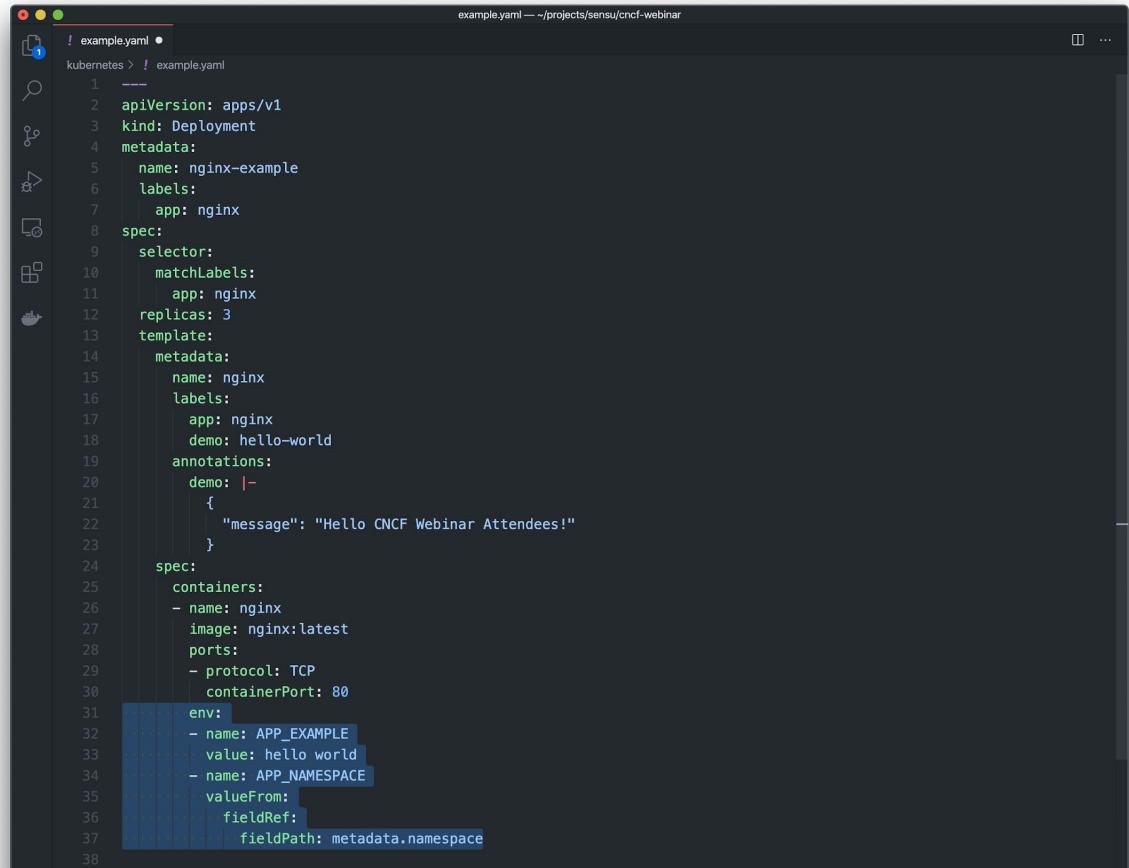
```
! example.yaml x
kubernetes > ! example.yaml
1
2  apiVersion: apps/v1
3  kind: Deployment
4  metadata:
5    name: nginx-example
6    labels:
7      app: nginx
8  spec:
9    selector:
10      matchLabels:
11        app: nginx
12    replicas: 3
13    template:
14      metadata:
15        name: nginx
16        labels:
17          app: nginx
18          demo: hello-world
19        annotations:
20          demo: |-
21            {
22              "message": "Hello CNCF Webinar Attendees!"
23            }
24    spec:
25      volumes:
26        - name: podinfo
27          downwardAPI:
28            items:
29              - path: "labels"
30                fieldRef:
31                  fieldPath: metadata.labels
32              - path: "annotations"
33                fieldRef:
34                  fieldPath: metadata.annotations
35      containers:
36        - name: nginx
37          image: nginx:latest
38          ports:
39            - protocol: TCP
40              containerPort: 80
41            volumeMounts:
42              - name: podinfo
43                mountPath: /etc/podinfo
44
```

# K8s Downward API

Configured as volume  
mount

—OR—

Configured as  
environment variable



```
! example.yaml •
kubernetes > ! example.yaml
1
2   apiVersion: apps/v1
3   kind: Deployment
4   metadata:
5     name: nginx-example
6     labels:
7       app: nginx
8   spec:
9     selector:
10       matchLabels:
11         app: nginx
12     replicas: 3
13     template:
14       metadata:
15         name: nginx
16         labels:
17           app: nginx
18           demo: hello-world
19         annotations:
20           demo: |-
21             {
22               "message": "Hello CNCF Webinar Attendees!"
23             }
24     spec:
25       containers:
26         - name: nginx
27           image: nginx:latest
28           ports:
29             - protocol: TCP
29               containerPort: 80
30       env:
31         - name: APP_EXAMPLE
32           value: hello world
33         - name: APP_NAMESPACE
34           valueFrom:
35             fieldRef:
36               fieldPath: metadata.namespace
37
38
```

Terminal — bash — 80x24

```
[root@nginx-example-5bd95679c8-bxqbh:/# export | grep APP_
declare -x APP_EXAMPLE="hello world"
declare -x APP_NAMESPACE="cnCF-webinar"
[root@nginx-example-5bd95679c8-bxqbh:/# ls /etc/podinfo
annotations  labels
[root@nginx-example-5bd95679c8-bxqbh:/# cat /etc/podinfo/labels && echo
app="nginx"
demo="hello-world"
pod-template-hash="5bd95679c8"
root@nginx-example-5bd95679c8-bxqbh:/# cat /etc/podinfo/]
```

Terminal — bash — 80x24

```
[sensu]$ kubectl exec -it nginx-6b8cb45687-crsgw -c sensu-agent /bin/ash
/] # echo $KUBE_KUBELET
gke-kubernetes-sensu-go-default-pool-33524e02-tq5r
/] # echo $SENSU_NAME
nginx-6b8cb45687-crsgw
/] # echo $SENSU_NAMESPACE
cnfcf-webinar
/] # echo $SENSU_LABELS | jq .
{
  "foo": "bar",
  "namespace": "cnfcf-webinar",
  "kubelet": "gke-kubernetes-sensu-go-default-pool-33524e02-tq5r"
}
/] #
```

nginx.yaml x

kubernetes > ! nginx.yaml

```
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
```

```
- name: sensu-agent
  image: sensu/sensu:5.21.0
  command: [
    "/opt/sensu/bin/sensu-agent", "start",
    "--config-file", "/etc/sensu/agent/agent.yaml",
    "--log-level", "debug"
  ]
  env:
    - name: KUBE_KUBELET
      valueFrom:
        fieldRef:
          fieldPath: spec.nodeName
    - name: SENSU_NAME
      valueFrom:
        fieldRef:
          fieldPath: metadata.name
    - name: SENSU_NAMESPACE
      valueFrom:
        fieldRef:
          fieldPath: metadata.namespace
    - name: SENSU_LABELS
      value: >-
        {
          "foo": "bar",
          "kube_namespace": "$(SENSU_NAMESPACE)",
          "kubelet": "$(KUBE_KUBELET)"
        }
    - name: SENSU_BACKEND_URL
      value: wss://sensu-backend-0.sensu.sensu-system.svc.cluster.local:443
    - name: SENSU_SUBSCRIPTIONS
      value: kubernetes linux nginx
    - name: SENSU_KEEPALIVE_INTERVAL
      value: "5"
    - name: SENSU_KEEPALIVE_WARNING_TIMEOUT
      value: "10"
```

```

Terminal — bash — 90x40
sensu $ kubectl get pods --namespace cncf-webinar
NAME           READY   STATUS    RESTARTS   AGE
nginx-6b8cb45687-crsgw  2/2    Running   0          51m
nginx-6b8cb45687-dsjbm  2/2    Running   0          51m
nginx-6b8cb45687-krwd8  2/2    Running   0          52m
sensu $ sensuctl entity list --namespace cncf-webinar --format json | jq '.[] | {
>   name: .metadata.name,
>   namespace: .metadata.namespace,
>   labels: .metadata.labels
> }'
{
  "name": "nginx-6b8cb45687-crsgw",
  "namespace": "cncf-webinar",
  "labels": {
    "foo": "bar",
    "kubelet": "gke-kubernetes-sensu-go-default-pool-33524e02-tq5r",
    "namespace": "cncf-webinar"
  }
}
{
  "name": "nginx-6b8cb45687-dsjbm",
  "namespace": "cncf-webinar",
  "labels": {
    "foo": "bar",
    "kubelet": "gke-kubernetes-sensu-go-default-pool-33524e02-mrk9",
    "namespace": "cncf-webinar"
  }
}
{
  "name": "nginx-6b8cb45687-krwd8",
  "namespace": "cncf-webinar",
  "labels": {
    "foo": "bar",
    "kubelet": "gke-kubernetes-sensu-go-default-pool-33524e02-64tz",
    "namespace": "cncf-webinar"
  }
}
sensu $

```

Sensu local-cluster > cncf-webinar > entities > nginx-6b8cb45687-crsgw

**Entity: nginx-6b8cb45687-crsgw**

Status	✓ success (0)	User
Last Seen	just now	Deregister
Subscriptions	kubernetes	Deregistration
	linux	Redacted Keys
	nginx	
	entity:nginx-6b8cb45687-crsgw	
Class	agent	OS
Hostname	nginx-6b8cb45687-crsgw	Platform
Cloud Provider	—	Architecture
VM System	—	ARM Version
VM Role	guest	libc Implementation
Adapter	eth0	
MAC	da:69:0e:07:a8:12	
IP Address	10.24.2.31/24	

**Labels**

- foo
- bar
- kubelet
- gke-kubernetes-sensu-go-d...

admin

## # 2: Kubernetes Events API

- First-class K8s resource
- Resource **state changes, errors, and other system messages**
- As seen in "kubectl describe" output
- Examples: image pull failures, volume mount failures, scheduling events, etc.

```
$ kubectl get events
```

```
Terminal — bash — 140x40
sensu $ kubectl -n default scale deployment nginx-example --replicas 2

Every 2.0s: kubectl -n default get pods                                         homelab: Wed Jul  8 21:35:29 2020
NAME                  READY   STATUS    RESTARTS   AGE
nginx-example-6d5bbc7d7-nl6rq   1/1     Running   0          61s
nginx-example-6d5bbc7d7-phz6t   1/1     Running   0          9m15s
nginx-example-6d5bbc7d7-ps49p   1/1     Running   0          61s
nginx-example-6d5bbc7d7-x59tx   1/1     Running   0          9m13s
nginx-example-6d5bbc7d7-zrs77   1/1     Running   0          61s

Every 2.0s: kubectl -n default get events                                         homelab: Wed Jul  8 21:35:29 2020
No resources found.
```

```
$ kubectl get events --field-selector involvedObject.kind!=Pod
```

```
Terminal — bash — 140x40
sensu $ kubectl -n default apply -f kubernetes/example-events-api.yaml

Every 2.0s: kubectl -n default get pods                                     homelab: Wed Jul  8 21:45:02 2020
No resources found.

Every 2.0s: kubectl -n default get events --field-selector involvedObject.kind!=Pod   homelab: Wed Jul  8 21:45:01 2020
No resources found.

[0] 0:kube-proxy- 1:webinar*
```

```
$ kubectl get events --field-selector type!=Normal
```



```
$ /api/v1/namespaces/cncf-webinar/events?fieldSelector%21%3DNormal
```

```
Terminal — bash — 140x40

sensu $ kubectl get events --field-selector type!=Normal
LAST SEEN    TYPE      REASON   OBJECT            MESSAGE
28s          Warning   Failed    pod/nginx-6c8997c749-88gj8   Failed to pull image "enginex:latest": rpc error: code = Unknown desc = Error response from daemon: pull access denied for enginex, repository does not exist or may require 'docker login'
28s          Warning   Failed    pod/nginx-6c8997c749-88gj8   Error: ErrImagePull
44s          Warning   Failed    pod/nginx-6c8997c749-88gj8   Error: ImagePullBackOff
sensu $ █

sensu $ curl -XGET -s "http://127.0.0.1:8888/api/v1/namespaces/cncf-webinar/events?fieldSelector=type%21%3DNormal" | jq '.items[] | {
>     type: .type,
>     message: .message,
>     reason: .reason,
>     kind: .involvedObject.kind,
>     name: .involvedObject.name
> }'
{
  "type": "Warning",
  "message": "Failed to pull image \"enginex:latest\": rpc error: code = Unknown desc = Error response from daemon: pull access denied for enginex, repository does not exist or may require 'docker login'",
  "reason": "Failed",
  "kind": "Pod",
  "name": "nginx-6c8997c749-88gj8"
}
{
  "type": "Warning",
  "message": "Error: ErrImagePull",
  "reason": "Failed",
  "kind": "Pod",
  "name": "nginx-6c8997c749-88gj8"
}
{
  "type": "Warning",
  "message": "Error: ImagePullBackOff",
  "reason": "Failed",
  "kind": "Pod",
  "name": "nginx-6c8997c749-88gj8"
}
sensu $ █

[0] 0:kube-proxy_ 1:webinar*           "homelab" 09:27 09-Jul-20
```

# # 1: Kubernetes API Watchers

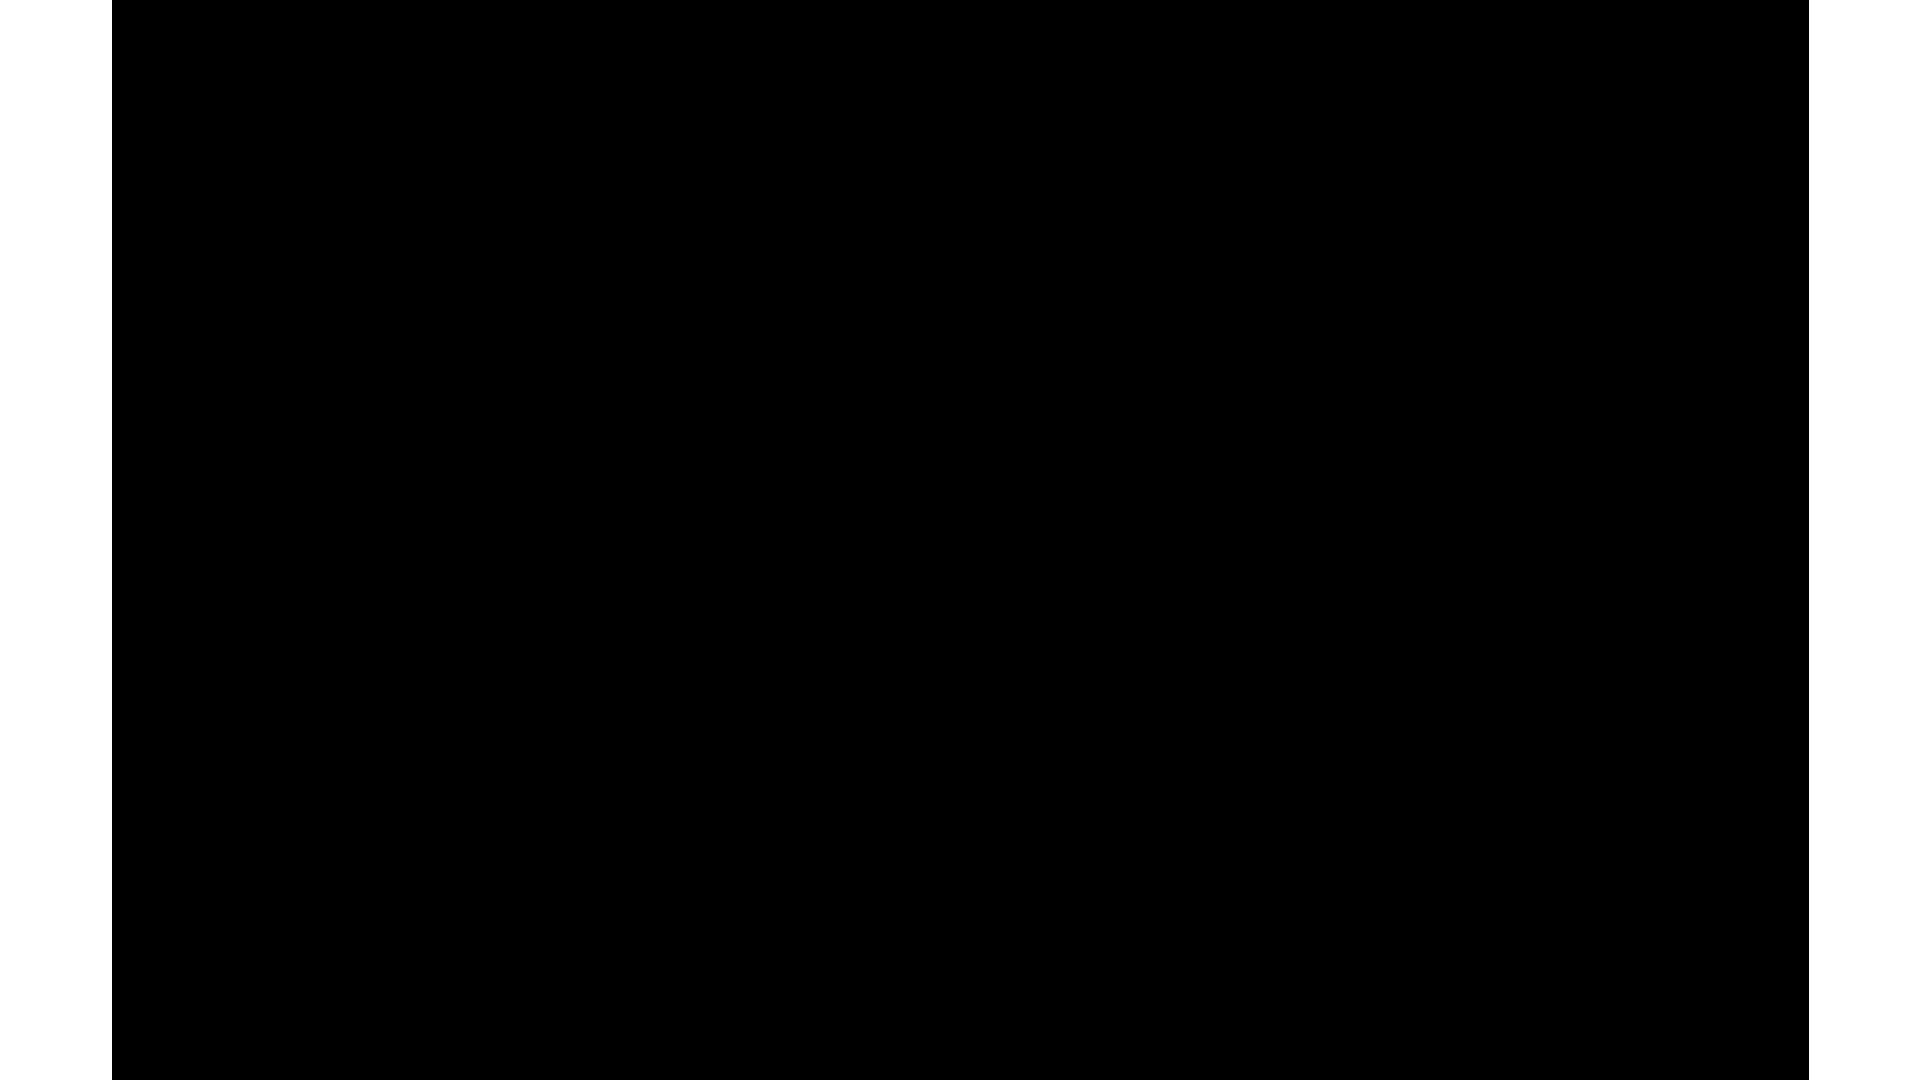
- Change notifications for **all K8s resources**
- Supports resource **instances & collections**

GET /api/v1/namespaces/{namespace}/pods

GET /api/v1/watch/namespaces/{namespace}/pods?resourceVersion=123456789

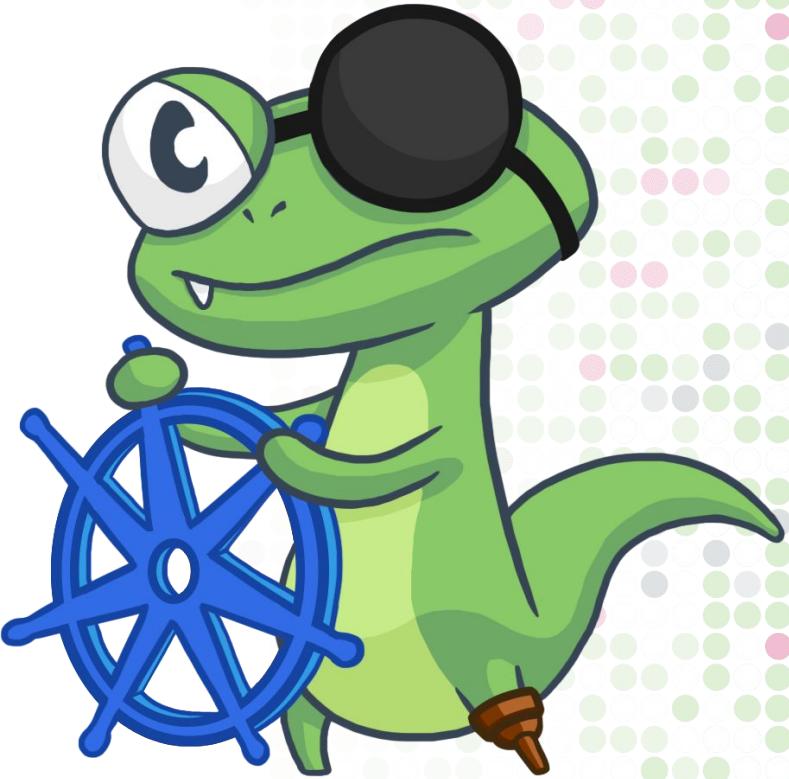
GET /api/v1/watch/namespaces/{namespace}/pods/{name}?resourceVersion=123

- Returns **ADDED**, **MODIFIED**, and **DELETED** change notifications



# Top 7 Kubernetes APIs:

7. Kubernetes Metrics APIs
6. Service API
5. Container API
4. Pod API
3. Kubernetes Downward API
2. Kubernetes Events API
1. Kubernetes API Watchers



# Closing remarks:

## Getting to the why

- Holistic view of K8s health
- Richer observations
- Additional context needed to complement metrics
- **Observability should tell the whole story**



# Going further: join our webinar!



- Filling gaps in K8s observability
- Getting richer context with an o11y pipeline
- New K8s integrations for Sensu
- Live webinar: **August 20, 2020 @ 10:30 AM PT**
- RSVP: [bit.ly/k8s-o11y](https://bit.ly/k8s-o11y)

# Thank you!



Learn more:  
[sensu.io](https://sensu.io)

Join our community:  
[discourse.sensu.io](https://discourse.sensu.io)

Join our webinar next month!  
[bit.ly/k8s-o11y](https://bit.ly/k8s-o11y)

