




AWS Controllers for Kubernetes

The AWS universe of services, now Kubeified.

Jay Pipes, Principal Open Source Engineer
@jaypipes



A familiar story...

Alice is a web developer and a huge Kubernetes fan.

She's developed a node.js application and is building her application into a Docker image.

The application uses a SQLite database for simple storage.

Alice goes to deploy on Kubernetes...

```
kubectl apply -f deployment.yml
```

```
kubectl apply -f service.yml
```

and probably...

```
kubectl apply -f ingress-nginx.yml
```

Everything is great. Until...

Ten users try using the site at once.

SQLite falls over.

Alice needs to set up a real database.

So, Alice sets up a real database...

```
kubectl apply -f postgres-secret.yml
```

```
kubectl apply -f postgres-volume-claim.yml
```

```
kubectl apply -f postgres-deployment.yml
```

```
kubectl apply -f postgres-service.yml
```

Hmm...

Now Alice is in the RDBMS administration game.

Definitely not what Alice had in mind.

So, what to do?

AWS to the rescue!

Alice finds out about Amazon Relational Database Service (RDS).

w00t! No more Alice the DBA!

But, there's a problem....

Where is Alice's cozy Kubernetes experience?

Create database


Choose a database creation method [Info](#)


☒ **Standard Create**
You set all of the configuration options, including ones for availability, security, backups, and maintenance.


☐ **Easy Create**
Use recommended best-practice configurations. Some configuration options can be changed after the database is created.


Engine options


Engine type [Info](#)


☐ Amazon Aurora


☐ MySQL



☐ MariaDB


☐ PostgreSQL


☐ Oracle


☒ Microsoft SQL Server


© 2020, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Alice doesn't have to use the AWS console...

What about the aws CLI tool?

What about CloudFormation?

What about Terraform?

But those aren't Kubernetes.

And Alice loves her Kubernetes.

But not enough to be a DBA.

Alice wants RDS, but doesn't want to leave her Kubernetes experience.

What can Alice do?

kubectl apply -f db.yml

```
apiVersion: rds.services.k8s.aws/v1alpha1
kind: DBInstance
metadata:
  name: mydb
spec:
  dbInstanceClass: db.m1.large
  dbInstanceIdentifier: mydb
  engine: PostgreSQL
...
```

AWS Controllers for Kubernetes

Solving Alice's problems. And yours too, hopefully.

A Kubernetes experience for AWS services

- Resources for AWS managed services are just another Kubernetes manifest
- Kubernetes stores the desired resource state
- ACK service controller handles the lifecycle of AWS managed service resources
- No CloudFormation behind the scenes

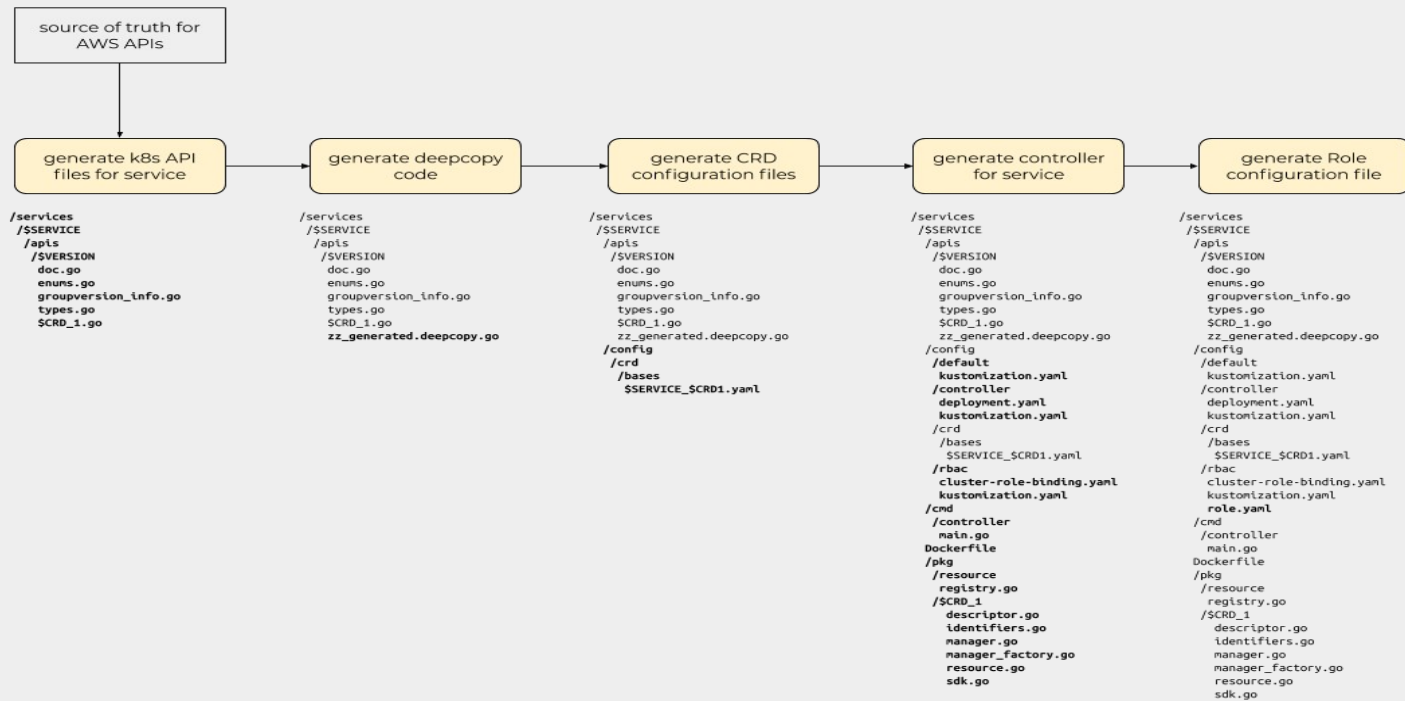
Design of ACK

- Each AWS managed service has a separate ACK service controller
- Install using Helm, static manifests or helper script
- Everything, including controller implementation, is code-generated
- Consult with AWS service teams to ensure API calling and behaviour is correct
- Not EKS specific, runs on any Kubernetes

Code generation

- Multiple phases
- Reads API models from aws-sdk-go, figures out which resources are CRDs (top-level)
- Creates Kubernetes API type definitions
- Generates resource manager for each CRD
- Generates linkage between AWS SDK and ACK runtime
- Generates config and build files

Code generation



Authorization and access control

- Kubernetes RBAC for custom resources (CRs)
- IAM Roles for AWS service and account permissions
- Each ACK service controller in own Deployment
- ACK service controller supplied with environment variables for AWS access credentials
- Use IAM Roles for Service Accounts to automate
- <https://aws.github.io/aws-controllers-k8s/user-docs/authorization/>

Cross-account resource management (soon)

- Avoid installing ACK controllers in many clusters
- Kubernetes cluster admin associates an AWS account ID to a Namespace
- All ACK CRs must be in a Namespace
- Application developer creates AWS managed resources by creating CR in a Namespace
- ACK controller looks up Role ARN
- Controller calls `STS::AssumeRole` to pivot client

What about secret stuff?

```
apiVersion: rds.services.k8s.aws/v1alpha1
kind: DBInstance
metadata:
  name: mydb
spec:
  dbInstanceClass: db.m1.large
  dbInstanceIdentifier: mydb
  engine: PostgreSQL
  masterUserPassword: UhmPlainText!?
...
```

Ah, that's better. And more Kubernetes-like.

```
apiVersion: rds.services.k8s.aws/v1alpha1
kind: DBInstance
metadata:
  name: mydb
spec:
  dbInstanceClass: db.m1.large
  dbInstanceIdentifier: mydb
  engine: PostgreSQL
  masterUserPassword:
    name: dbsecrets
    key: masterUserPassword
...
```

Other things coming soon

- Standardized AWS tag representation for all ACK resources
- Control tags that all CRs (in a Namespace) should have
- Common rate limiting and throttling support
- “Adopting” pre-existing resources

Developer preview

- S3 Bucket
- SNS Topic
- SQS Queue
- ECR Repository
- DynamoDB [Global]Table
- API Gateway V2

Soon

- RDS DBInstance, DBCluster
- Elasticache CacheCluster
- CloudFront Distribution
- EC2 VPC Subnet, SecurityGroup, InternetGateway
- EKS
- Release roadmap:
 - <https://github.com/aws/aws-controllers-k8s/projects/1>

Interested in contributing?

<https://github.com/aws/aws-controllers-k8s>

<https://github.com/aws/containers-roadmap/issues/456>