



Security of the mesh and in-mesh security

Alcide x CNCF, March 2019
Gadi Naor, CTO



about# gadi.naor

--from **tel_aviv**

--enjoy **skateboarding**

--engineering @**check_point**

--engineering @**altor_networks**

--engineering @**juniper_networks**

--cto **alcide**

--since 07.16

run

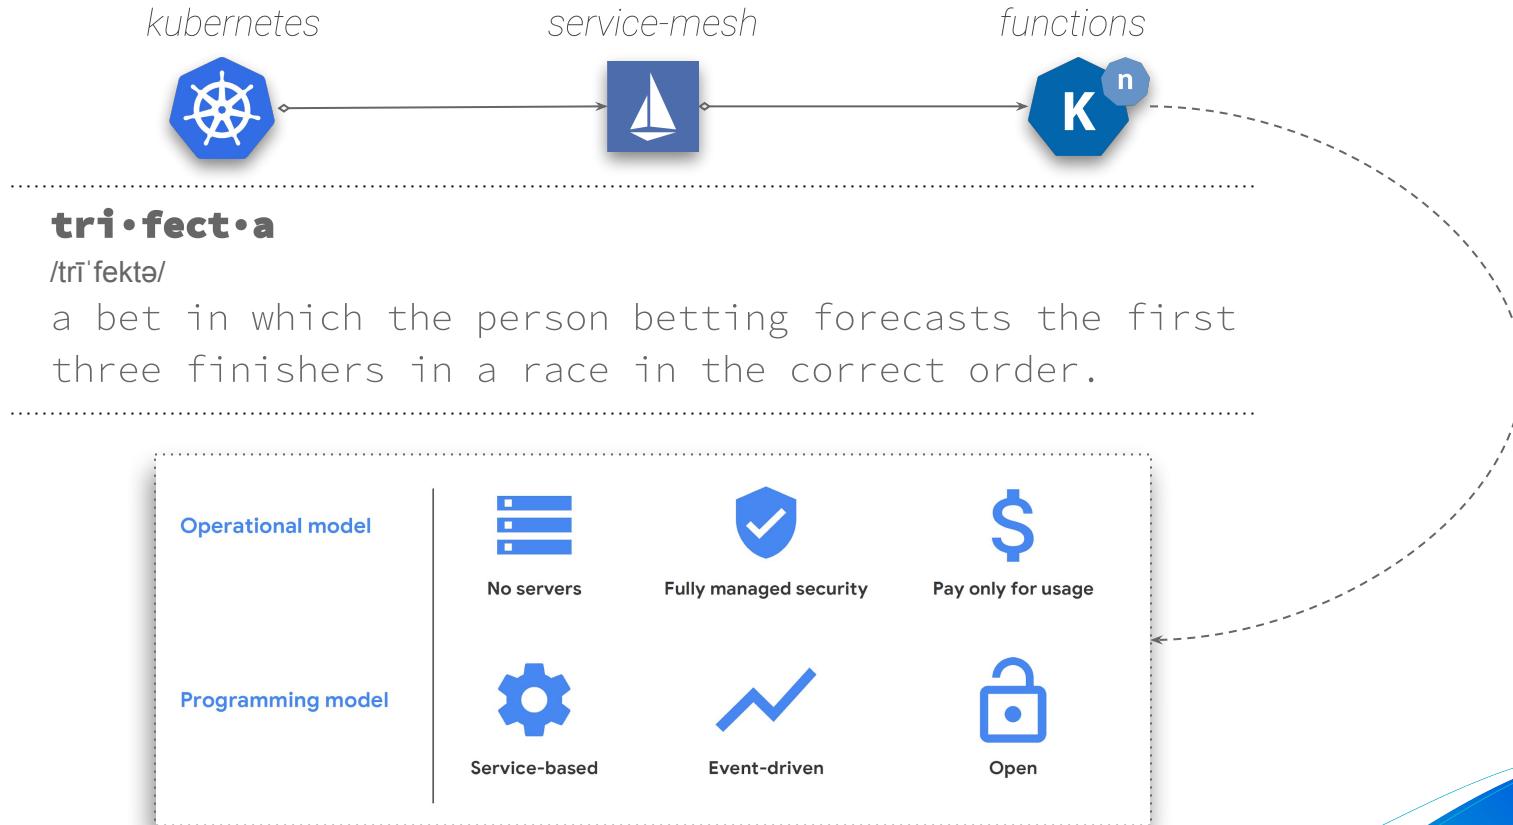


In This Webinar

In-Mesh Security & Off-Mesh Security

- Cloud Native Application - The New Stack
- Security & Operations of Service Mesh
- Beyond Zero Trust - Security Challenges In Cloud Native Applications
- Defense & Segmentation Strategies
- Understanding Data Leak Vectors

Cloud Native Applications - The New Stack



The Cloud Native App Stack



\$ Primitives for deploying and managing container workloads at scale



\$ How workloads interact with each other over the network in a secure, reliable, observable way



\$ “PaaS” abstractions for specifying, running, and modifying applications

Istio - The Promised Land

Connect, manage and secure microservices

Connect

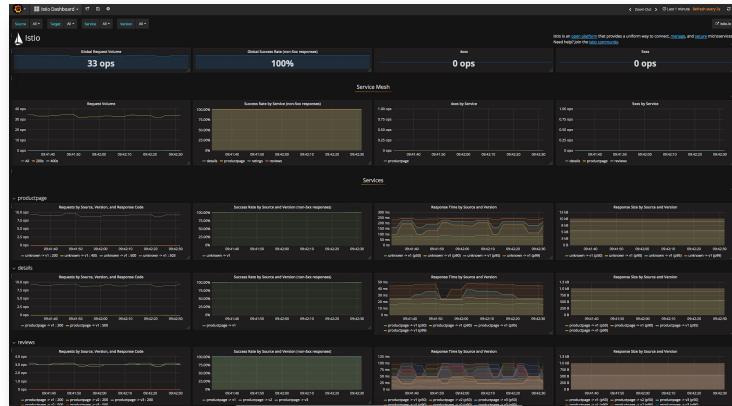
- Layer 7 path-based routing
- Traffic shaping
- Load balancing - A/B testing, canarying

Manage

- Telemetry
- Fleet-wide Visibility - Zipkin, Prometheus & Grafana

Secure

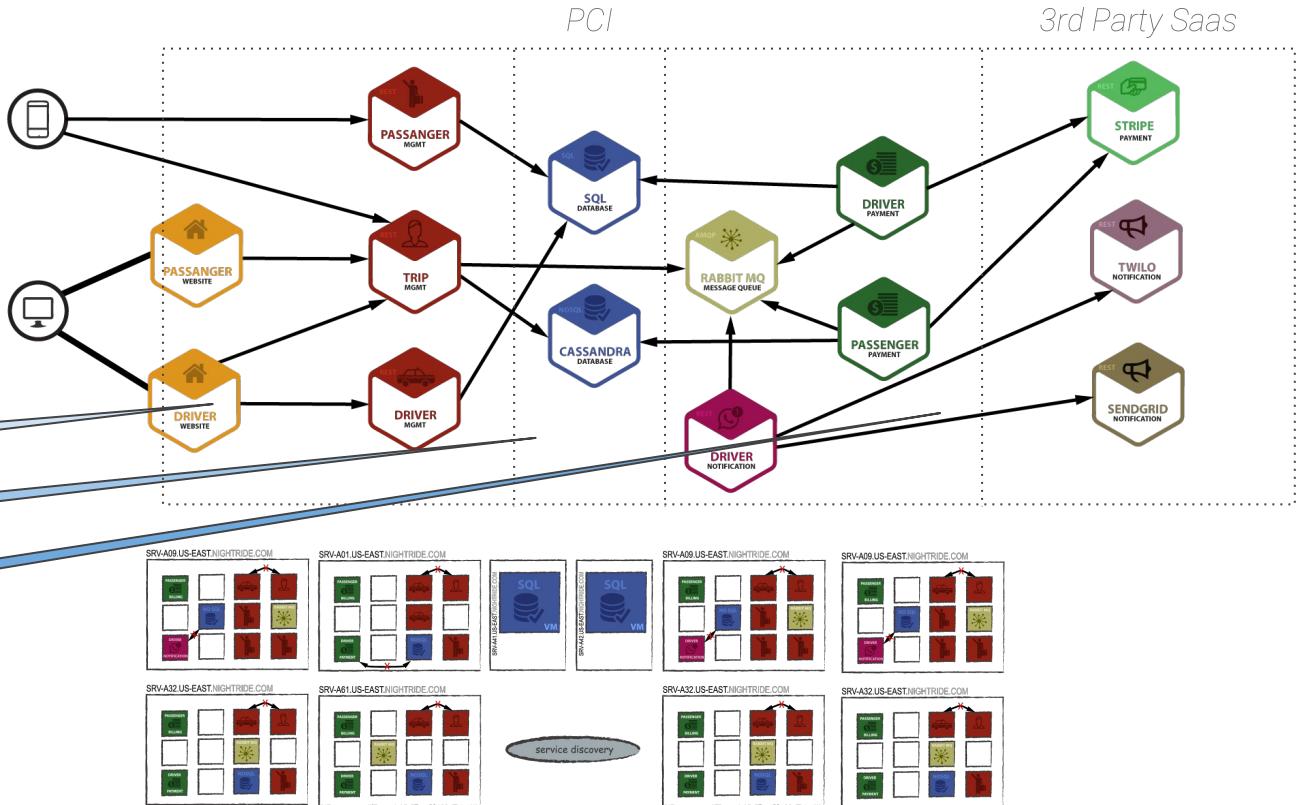
- Identity based service access control
- Service authorization - API level access control
- Service-Service encryption with TLS (mTLS)



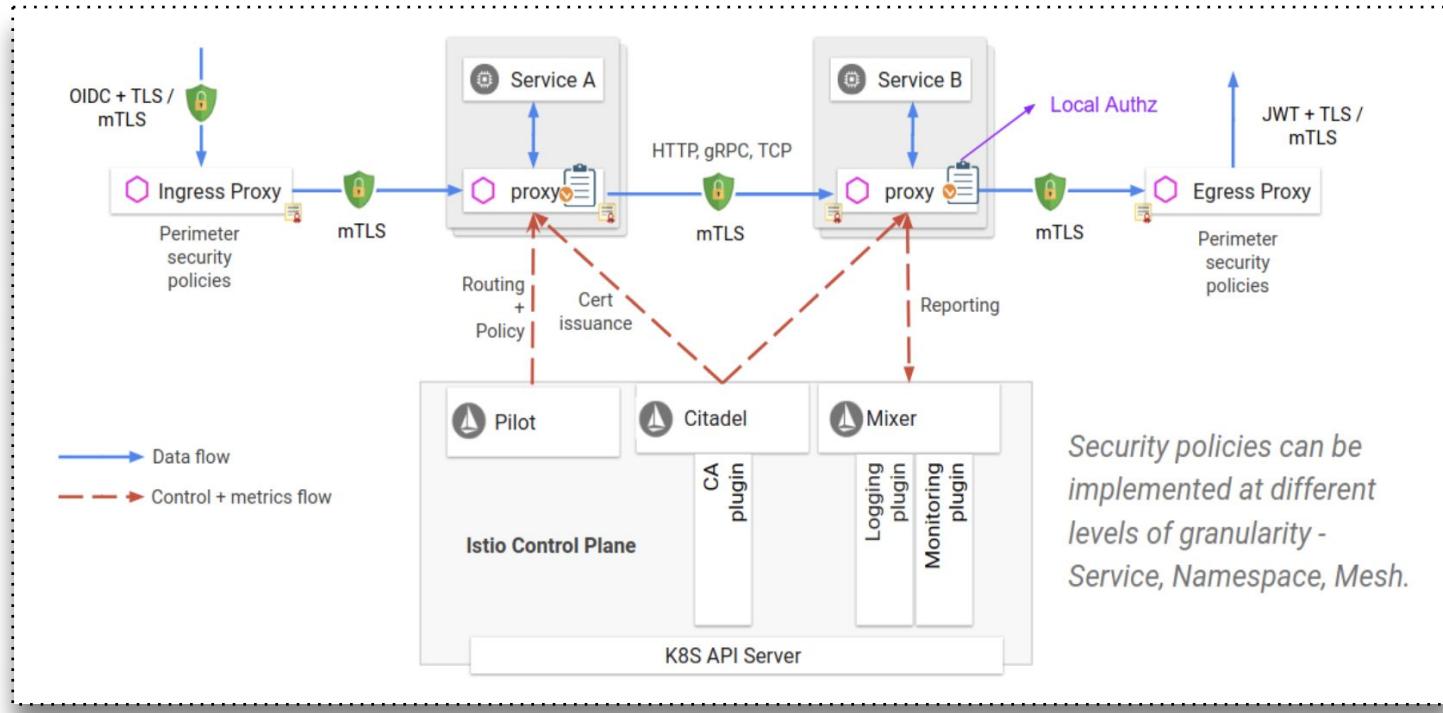
The Service Mesh

Building Blocks

Ride The {{ Microservices }} Lightning

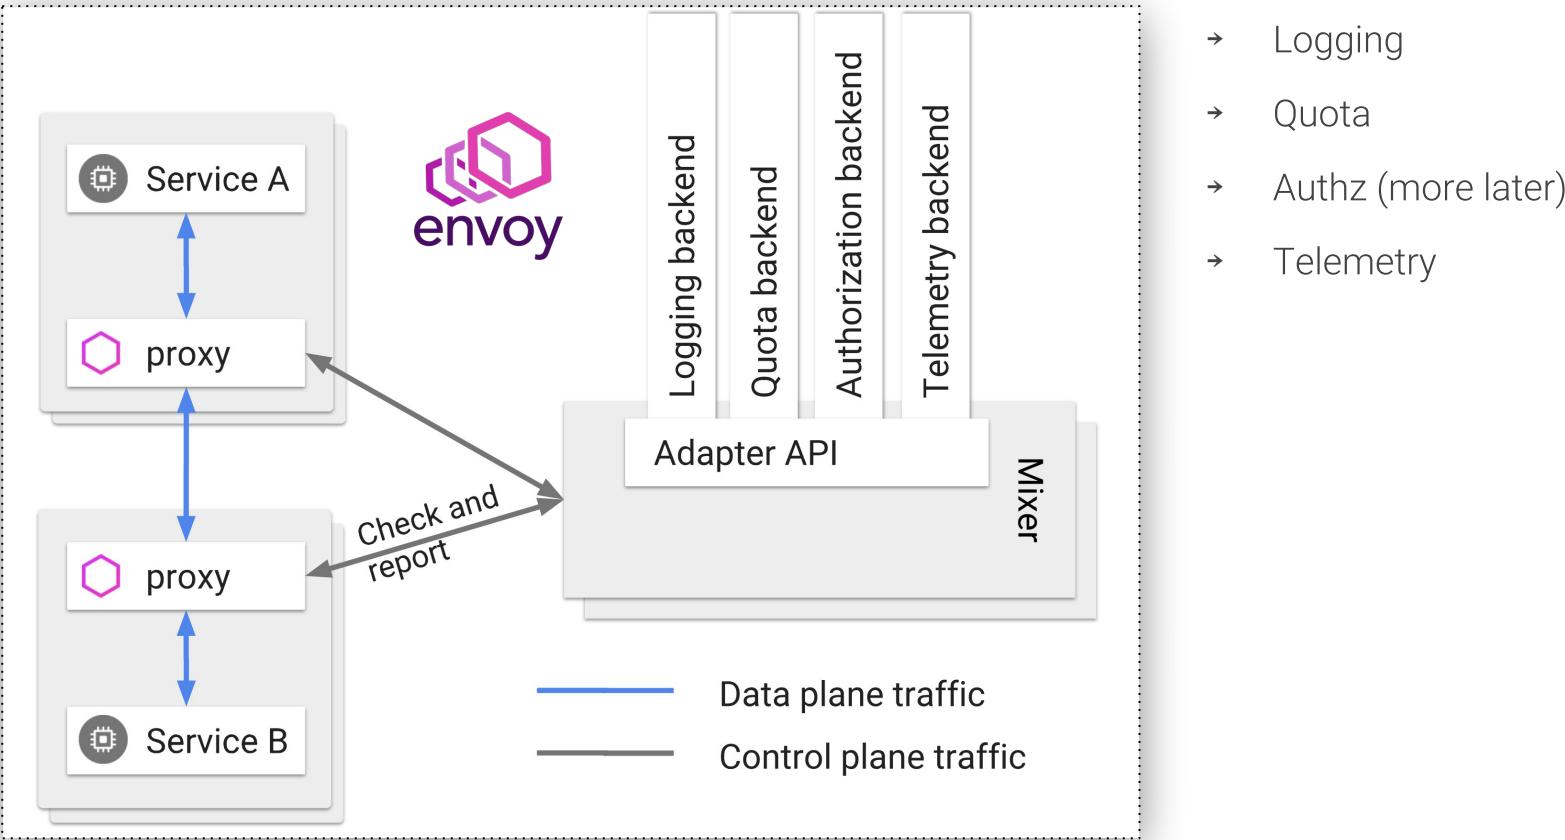


ATLS - “Istio”

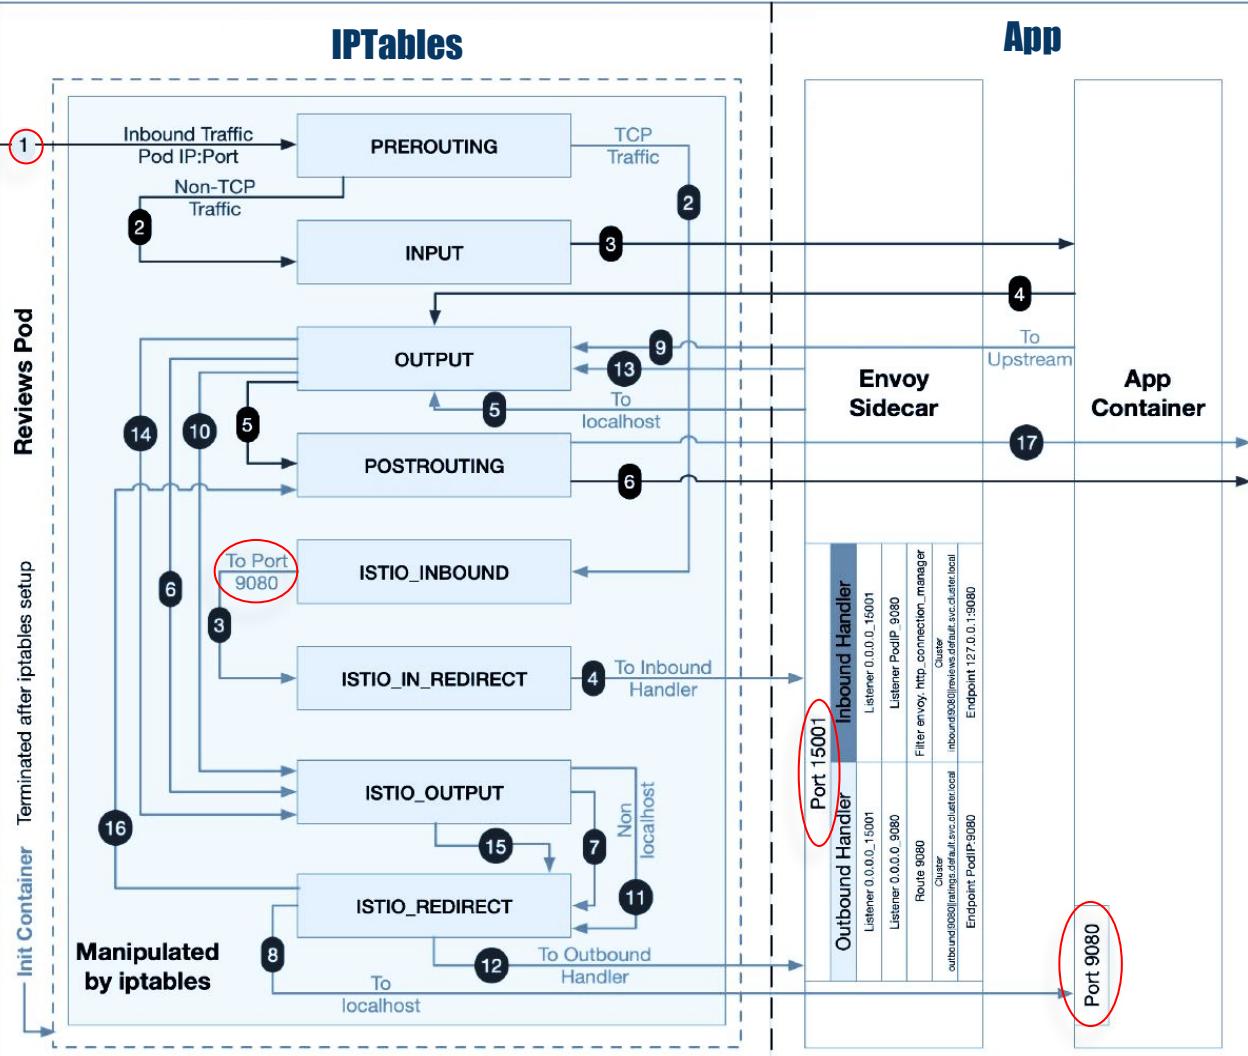


<https://cloud.google.com/security/encryption-in-transit/application-layer-transport-security/>

Controlling Envoy Dataplane



IPTables



Service Mesh Dataplane



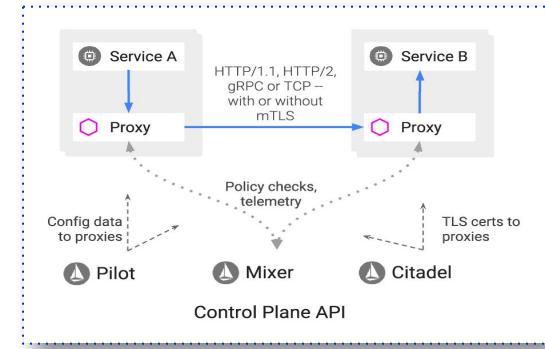
- Auto Inject @ Admission Controller
 - xDS API
 - Observability Pantheon
 - Rich Protocol Support
 - H2, gRPC, MongoDB, DynamoDB,..
 - Side Car Tax
 - Network Latency
 - CPU
 - Memory

Mesh Escape 101

```
$ kubectl -n default exec -it ratings-v1-77f657f55d-zq248 -c istio-proxy -- bash -c "id -u"  
1337
```

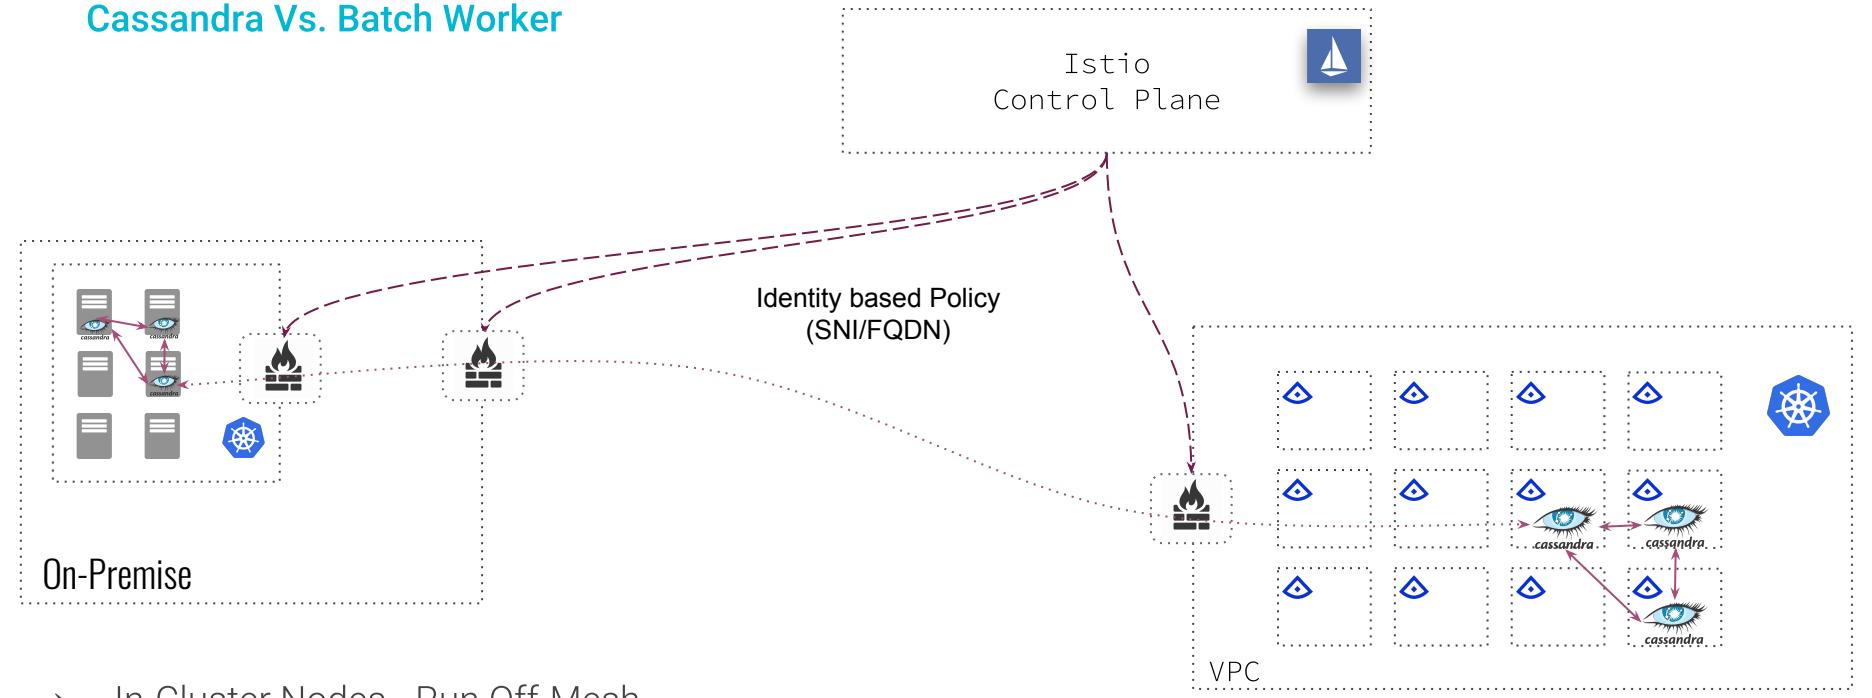
```
$ useradd envoyuser -u 1337 and su - envoyuser  
bye bye mesh
```

- IPv6 - east-west communications
- Tunneling (more later)
- 'Do not inject' annotation
- Does not cover infra services, control plane, ...
- Use Pod Security Policy and admission controllers for better control



Performance // Not All Services Are Born Equal

Cassandra Vs. Batch Worker



- In-Cluster Nodes - Run Off-Mesh
- Cassandra - DC Aware datastore
- Spirit: use the right tool for the right task

In-Mesh & Off-Mesh

Service Mesh Security Building Blocks

Application Segmentation aka Zero Trust

- Bring security into the heart of both network and application mesh
- Enable protection at the workload level, as well as service level
- Wrapping granular security policies around an individual workload (or group of workloads).
- Identify and prevent threats moving laterally through the network
- Add east-west protection to the traditional perimeter security model.
- Contain and quarantine threats, within the micro-segment - prevent propagation.

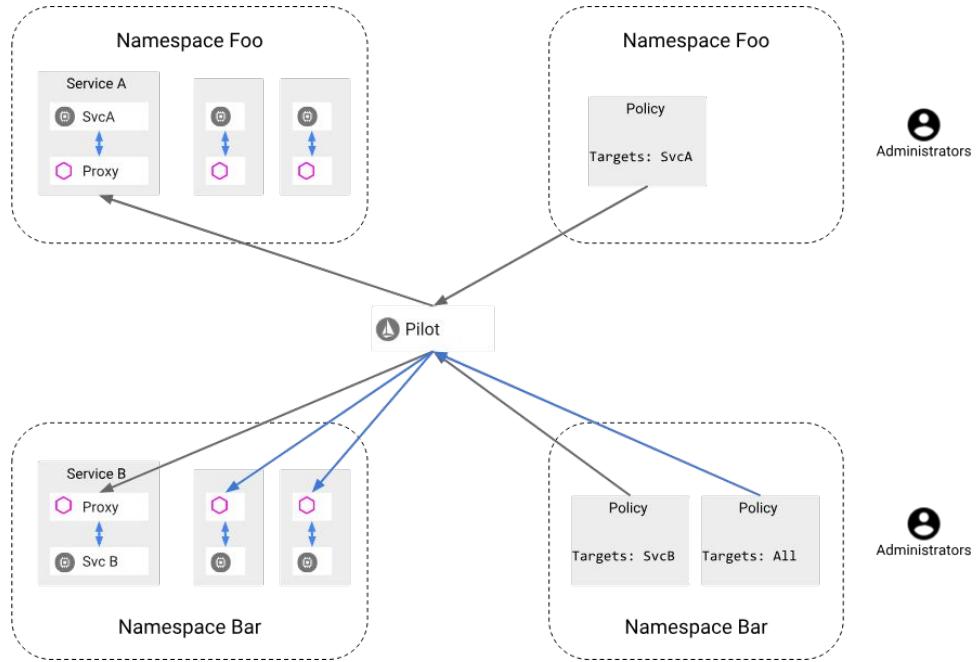
Good Start ... Not Enough

- DevSecOps & CI/CD
- Application & Infra convergence → Have an Isolation Game Plan?!
- Threat Protection - AppSec, Scans, Exploits (cpu,...), Container Escape

Istio Service Authentication

You need to:

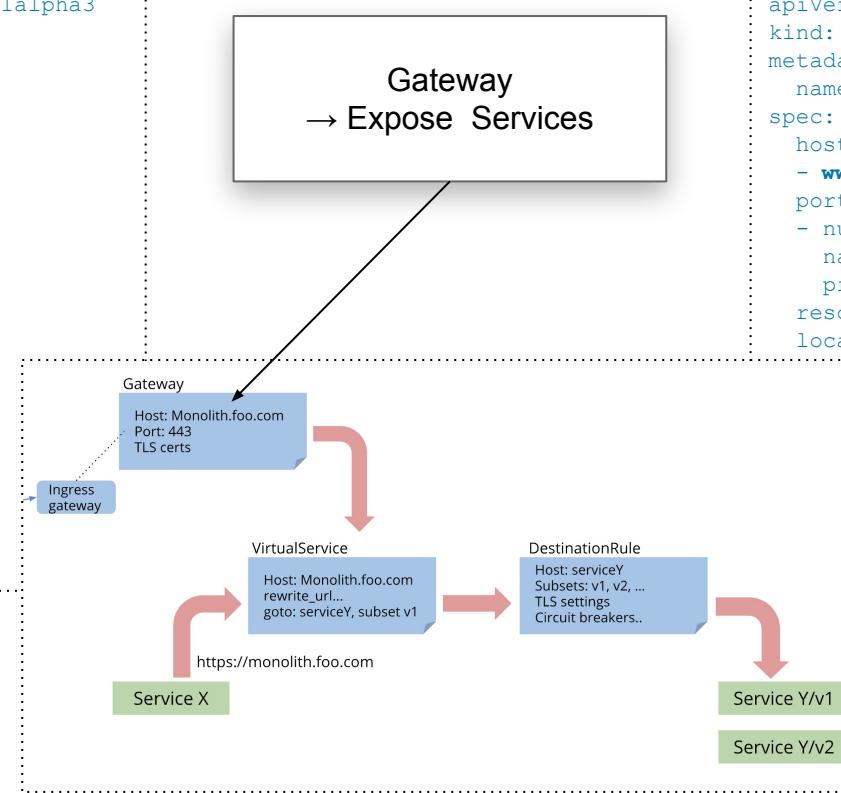
- Istio Mixer
 - Policy decision APIs
 - Service-to-Service Authentication
 - mTLS - identity provisioning
 - key life-cycle management
- End-user Authentication
 - Origin Authentication
 - Request level JWT (Auth0, Google Auth,...)



Istio Ingress & Egress Policy

Example: Grant egress access <https://www.google.com>

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: google
spec:
  hosts:
    - www.google.com
  tls:
    - match:
        - port: 443
        sni_hosts:
          - www.google.com
    route:
      - destination:
          host: www.google.com
          port:
            number: 443
          weight: 100
```



```
apiVersion: networking.istio.io/v1alpha3
kind: ServiceEntry
metadata:
  name: google
spec:
  hosts:
    - www.google.com
  ports:
    - number: 443
      name: https
      protocol: HTTPS
  resolution: DNS
  location: MESH_EXTERNAL
```

Istio Service Authorization (RBAC)

Example: Grant authenticated users with read access to version v1,v2 of products service

```
apiVersion: "rbac.istio.io/v1alpha1"
kind: RbacConfig
metadata:
  name: default
spec:
  mode: 'ON_WITH_INCLUSION'
  inclusion:
    namespaces: ["default"]
```

```
apiVersion: "rbac.istio.io/v1alpha1"
kind: ServiceRole
metadata:
  name: products-viewer-version
  namespace: default
spec:
  rules:
    - services: ["products.default.svc.cluster.local"]
      methods: ["GET", "HEAD"]
      constraints:
        - key: request.headers[version]
          values: ["v1", "v2"]
```

read access

condition

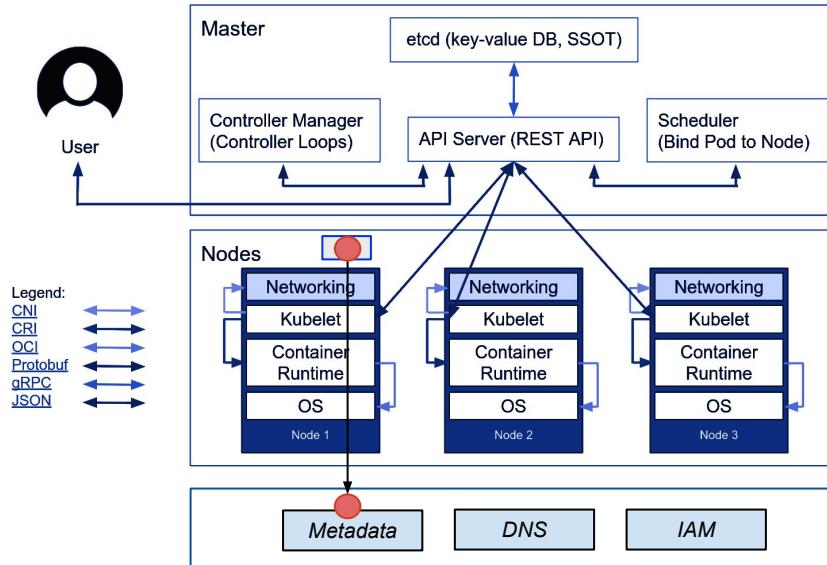
```
apiVersion: "rbac.istio.io/v1alpha1"
kind: ServiceRoleBinding
metadata:
  name: binding-products-all-authenticated-users
  namespace: default
spec:
  subjects:
    - properties:
        source.principal: "*"
  roleRef:
    kind: ServiceRole
    name: "products-viewer-version"
```

authenticated
users

Exploit Chain Example

Cloud Metadata Service

1. Server Side Request Forgery (SSRF) App vulnerability in screenshotting functionality of **Shopify Exchange**
2. Gain access from pod to Google Cloud Metadata
3. Obtain kube-env from cloud metadata
4. Obtain kubelet certificates from kube-env
5. Run arbitrary command with kubectl using the certs from kube-env



Off Mesh //Kubernetes Network Policies 101

With security, no news ... is good news

History

pre-work Q4'15 → alpha v1.2 (Mar'16) → stable v1.7 (June'17)

Anatomy & Spirit

- Pods are not-isolated by default
- Require CNI policy support (calico, canal, weavenet, cilium)
- Multiple policies attached to Pods
- Policies “associated” to pods based on labels
- Policy rules for ingress, egress or both
 - Ingress : “Who can connect to this Pod?”
 - Egress : “Who this Pod can connect to?”

Off Mesh // Kubernetes Network Policies 101

Rules Of Engagement

- Rules are allowing traffic (whitelist)
- Rules are additive
- Container Ports only (not service ports)
- Empty groups [] implement deny
- API access level control → Istio/Api Gateway
- Policy management - hard over time
- Separate from Istio network policies

YAML Traps

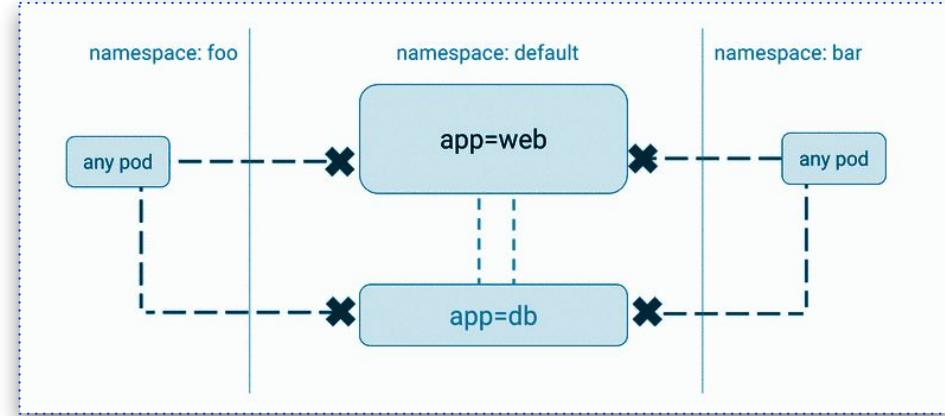
```
...  
ingress:  
- from:  
  - namespaceSelector:  
    matchLabels:  
      user: alice  
podSelector:  
    matchLabels:  
      role: client  
...
```

```
...  
ingress:  
- from:  
  - namespaceSelector:  
    matchLabels:  
      user: alice  
  - podSelector:  
    matchLabels:  
      role: client  
...
```

Off Mesh // Microservice Per Namespace

Namespace Isolation

```
kind: NetworkPolicy
apiVersion: networking.k8s.io/v1
metadata:
  namespace: my-namespace
  name: deny-from-other-namespaces
spec:
  podSelector:
    matchLabels:
  ingress:
    - from:
      - podSelector: {}
```



- Whitelist ingress traffic from namespace pods
 - Deny the rest
- Traffic dropped at the destination
 - Noisy/Hostile Neighbour

Off Mesh // Segment & Conquer

Node Taints & Pod Tolerations

- Manipulate Kubernetes Scheduler
- Force Workload placement on certain nodes

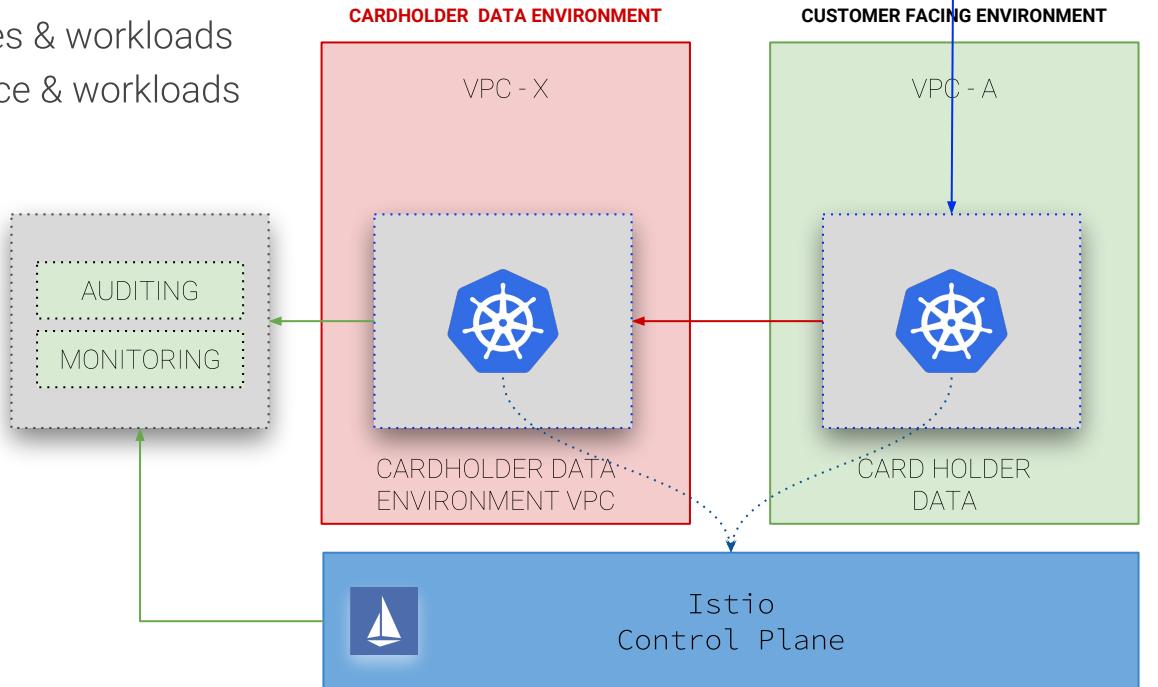
```
$ kubectl taint nodes es-node elasticsearch=false:NoExecute
```

```
...
tolerations:
  - key: elasticsearch
    operator: Equal
    value: false
    effect: NoExecute
```

Commerce Application On Istio

Use Case: PCI DSS Compliance

- Cluster X: PCI Cardholder services & workloads
 - Cluster A: Customer facing service & workloads
 - In-Cluster special segmentation
 - Cross Clusters & VPCs policies
 - Istio Network Policies
-
- DevSecOps overhead
 - Compute Resource overhead

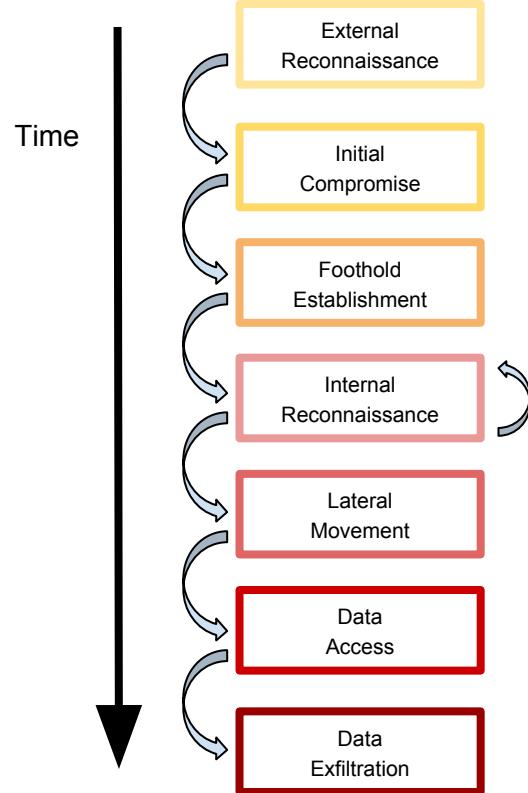


Attack Model

Istio & Service Discovery

Model of Attack on Istio & Kubernetes

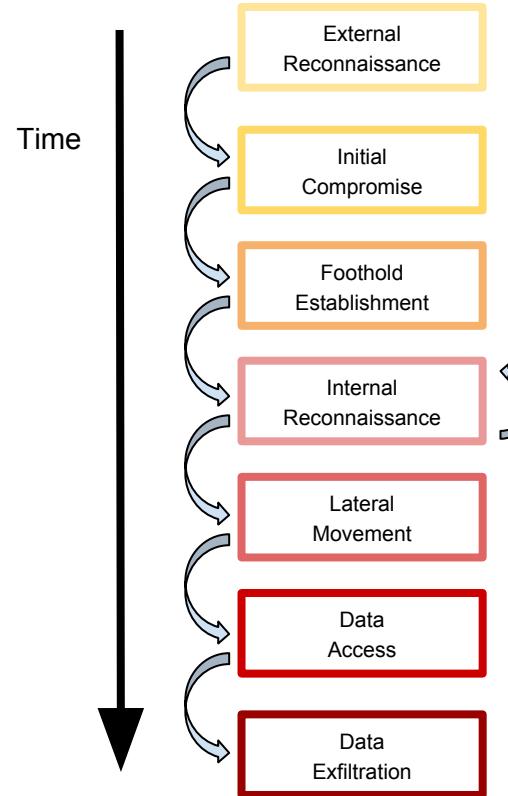
1. Collect information to be used in preparing and executing the attack
2. Initial access to a deployed component
 - o compromise externally-exposed vulnerable workload; deploy rogue workload
3. Prepare the compromised infiltration point as long term proxy for the attack
 - o connect to Command and Control (C&C); persist
4. Identify targets, enablers or obstacles to expanding the scope of the infiltration
 - o internal scanning; infrastructure API (ab)use
5. Lateral movement: infiltrate additional resources from the established foothold or from other previously compromised resources
 - o pivoting, escalation; bypass security policies, firewalls...
6. Often, the target of the attack is to access sensitive information stored in data repositories
 - o DBs, message-queues, files, APIs; within cluster or accessible from cluster
7. Often, data extraction is the end goal of an attack.
 - o This phase is centered around perimeter-facing network connections.



Model of Attack on Istio & Kubernetes

1. Collect information about the target and executing the attack
2. Initial access to a destination system
 - compromise externally-exposed vulnerable workload; deploy rogue workload
3. Prepare the compromised infiltration point as long term proxy for the attack
 - connect to Command and Control (C&C); persist
4. Identify targets, enablers or obstacles to expanding the scope of the infiltration
 - internal scanning; infrastructure API (ab)use
5. Lateral movement: infiltrate additional resources from the established foothold or from other previously compromised resources
 - pivoting, escalation; bypass security policies, firewalls...
6. Often, the target of the attack is to access sensitive information stored in data repositories
 - DBs, message-queues, files, APIs; within cluster or accessible from cluster
7. Often, data extraction is the end goal of an attack.
 - This phase is centered around perimeter-facing network connections.

Using DNS Tunneling



DNS Tunneling

- Tunneling: embedding messages of one network protocol within messages of a different network protocol.
- In Security Context: bypass security boundaries
 - Encapsulating malicious traffic within permitted, ubiquitous traffic
 - e.g. bypass perimeter firewall, Kubernetes/Istio network security policies
- Attacker sets up a tunnel from Internet to a compromised Pod (Workload)
 - Pass commands to it
 - Exfiltrate data harvested in the cluster through its service discover infra

DNS Tunneling // Attack and Protection

- DNS tunneling may be used by attacker to connect compromised WL in a Kubernetes cluster to Internet
 - covert channel for C&C and data exfiltration
 - bypassing perimeter firewall, cloud/k8s/istio network security policies
- Challenges in detection and mitigation
 - high-resolution monitoring: deep packet inspection per WL
 - scalable and fast analysis
 - detection using WL context
 - automatically adapt to changes of WLs and WLs activity in cluster
 - fine-grained mitigation

Zero Trust API Access Control

Definition

- Each API is protected (God Bless Envoy & Istio)
- East-West API Calls
- Authentication API Keys, Bearer Tokens, TLS Client Certificates, or HTTP Basic Authorization
- Authorization - Attribute Based Access Control

Challenges

- GraphQL, Non HTTP protocols (Kafka, C*, ...)
- Egress Controls - Workload Segmentation
- Off service mesh
- API Sprawl
- Detect Anomaly & Drifts

Istio Security

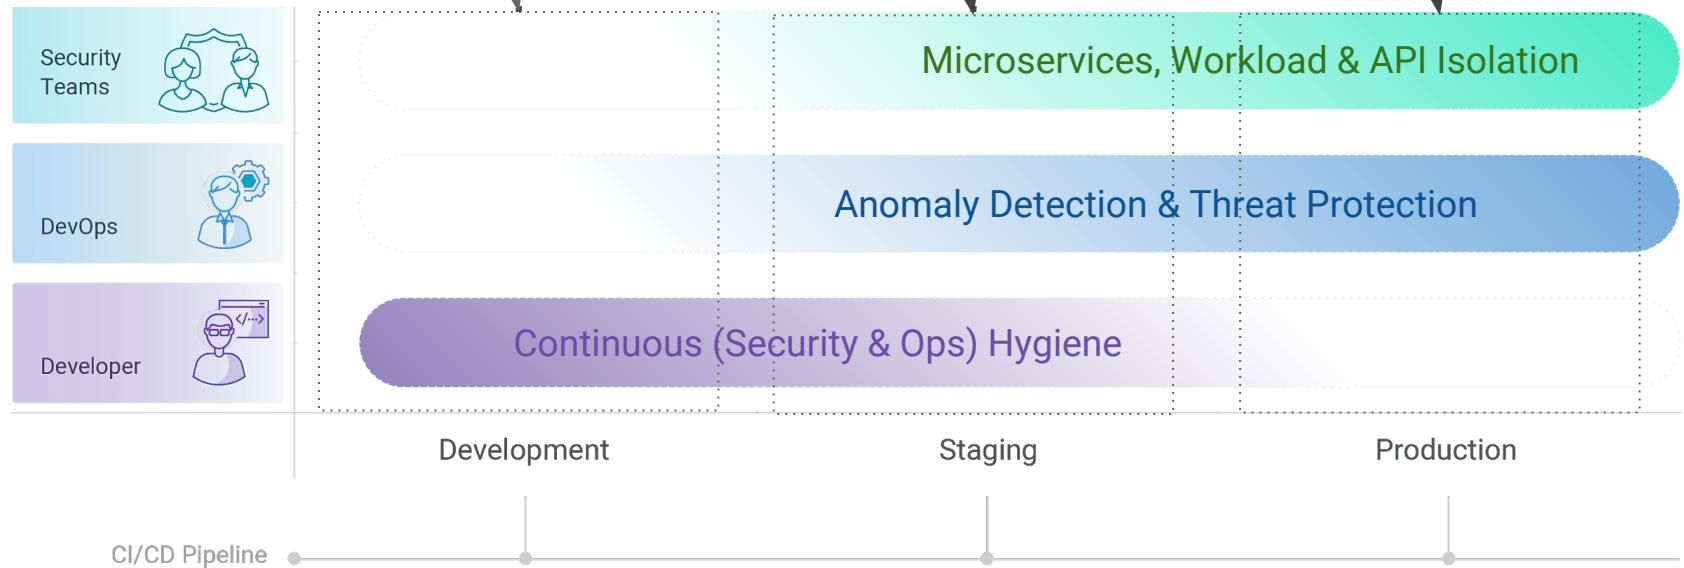
The Good, Bad & Ugly

- Pods are not-isolated by default
- Basic east-west workload segmentation between cluster workloads
- Decouple microservices and policy rules - with healthy flexibility
- Cluster nodes network policies covered by the underlying cloud security groups
- Policy management - extremely challenging over time.
- You can't embed policy into workloads/microservices
- No explicit way for drop rules
- Dropped traffic is not logged anywhere
- Full network security stack (cryptomining, scans, tunneling, spoofing, application attacks,...)

Complexity Is The Worst Enemy For Security

SMeshOps - Mitigate Challenges Collectively

Software Supply
Chain Hygiene





Q & A

gadi@alcide.io

get.alcide.io/request-demo



Thank You

gadi@alcide.io

get.alcide.io/request-demo