

Cluster API

A Kubernetes subproject to simplify
cluster lifecycle management

Katie Gamanji - CNCF Technical
Oversight Committee & Cloud
Platform Engineer, American Express

Naadir Jeewa - Maintainer, Cluster API
Provider AWS & Senior Member of
Technical Staff, VMware

11 June 2020



Katie Gamanji

Cloud Platform Engineer

@American Express

Technical Oversight Committee

@CNCF



@k_gamanji



katie-gamanji

Naadir Jeewa

Senior Member of Technical Staff
VMware

Maintainer, Kubernetes Cluster API Provider AWS



[randomvariable](mailto:jeewan@vmware.com) jeewan@vmware.com

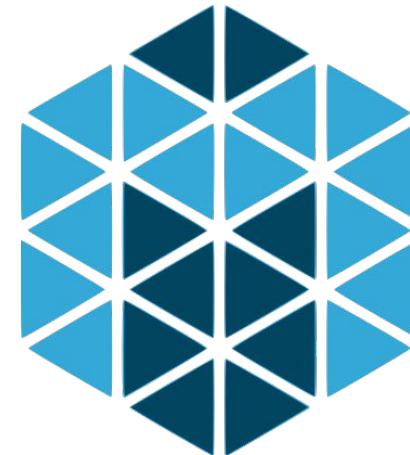
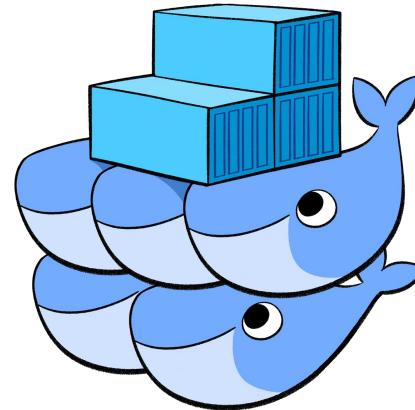


Agenda

Kubernetes Lifecycle
Management with
Cluster API

- Cluster API overview
- How Cluster API delivers cloud independence
- How controllers enable declarative infrastructure
- Demo
- Roadmap
- Cluster API and you

Container Orchestrators



kubernetes

AMERICAN EXPRESS



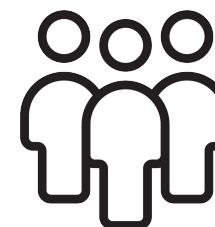
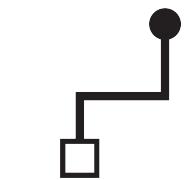


58%

—
Production
Systems

2K

—
Contributors



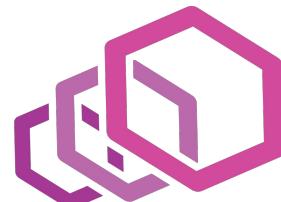
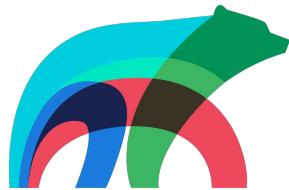
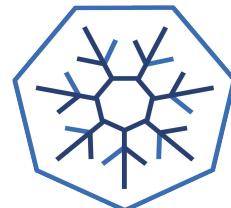
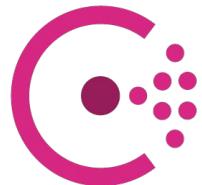
23K

—
KubeCon
Attendees

AMERICAN EXPRESS

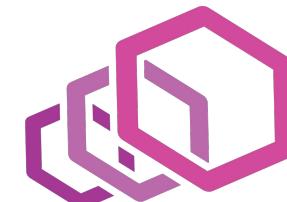
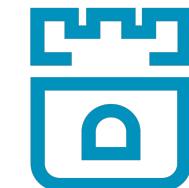
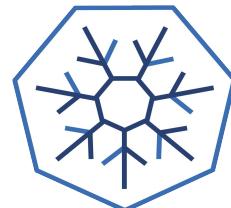
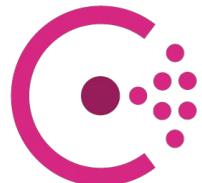


AMERICAN EXPRESS



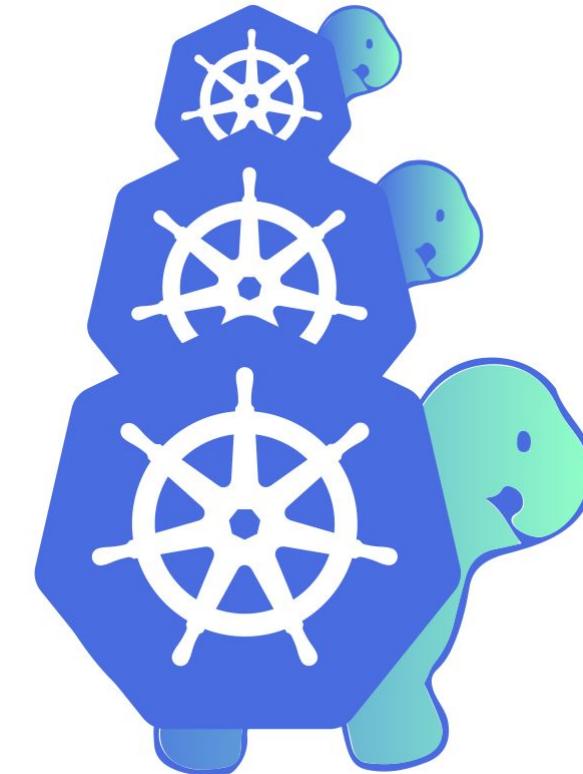


CLOUD NATIVE Landscape

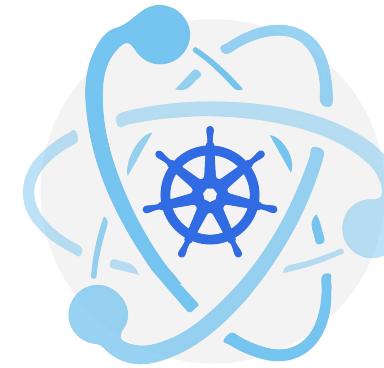


ClusterAPI

Cluster Provisioning



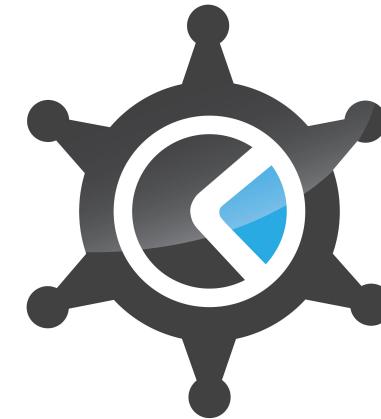
Bootstrap Tools



kubeadm



KUBESPRAY



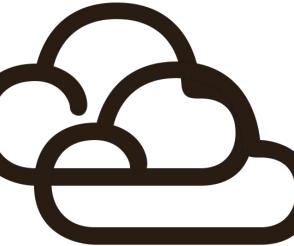
kops



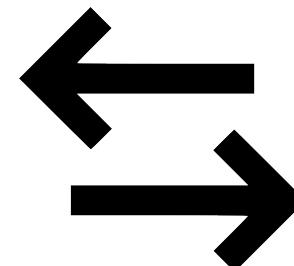
TECTONIC

Challenges

Bootstrap Providers



Cross-cloud



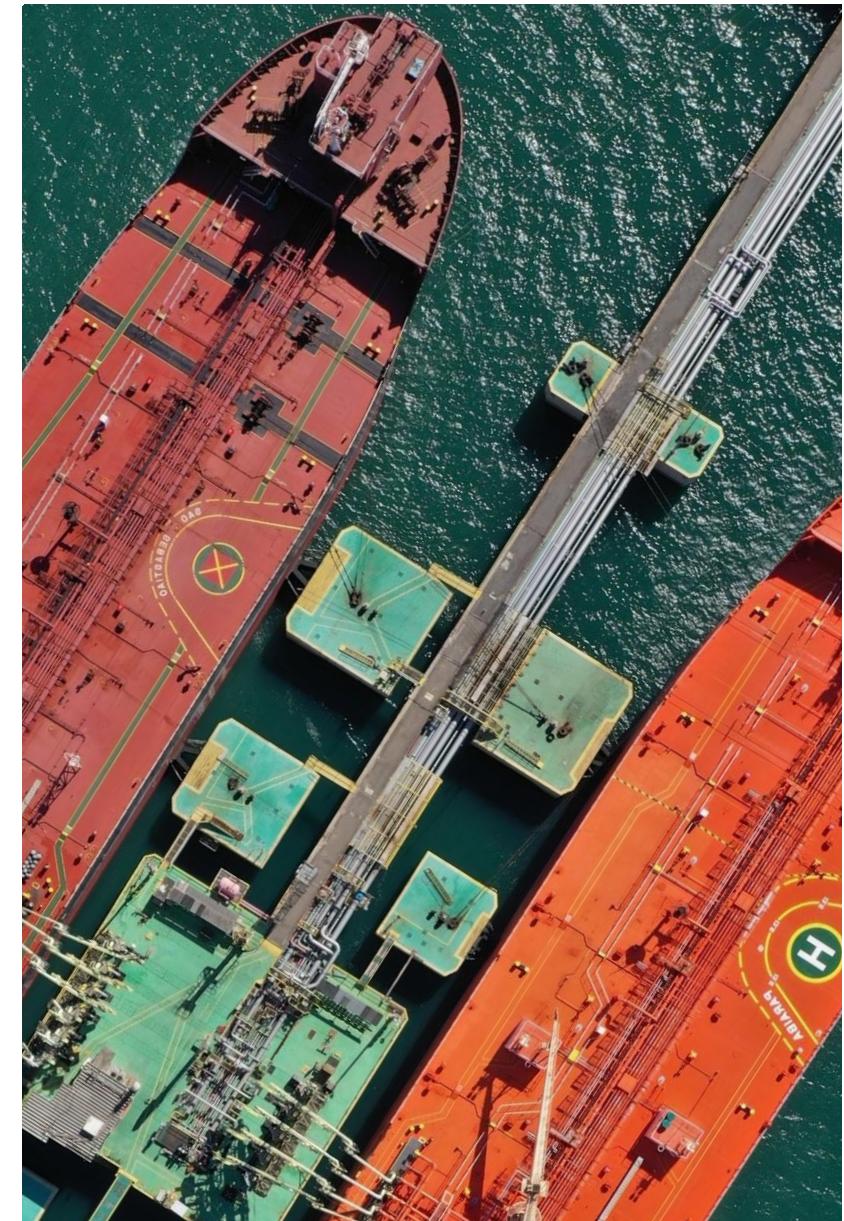
Tool migration



China

Cluster API

“ Provision of declarative APIs for cluster creation, configuration, and management. ”



ClusterAPI

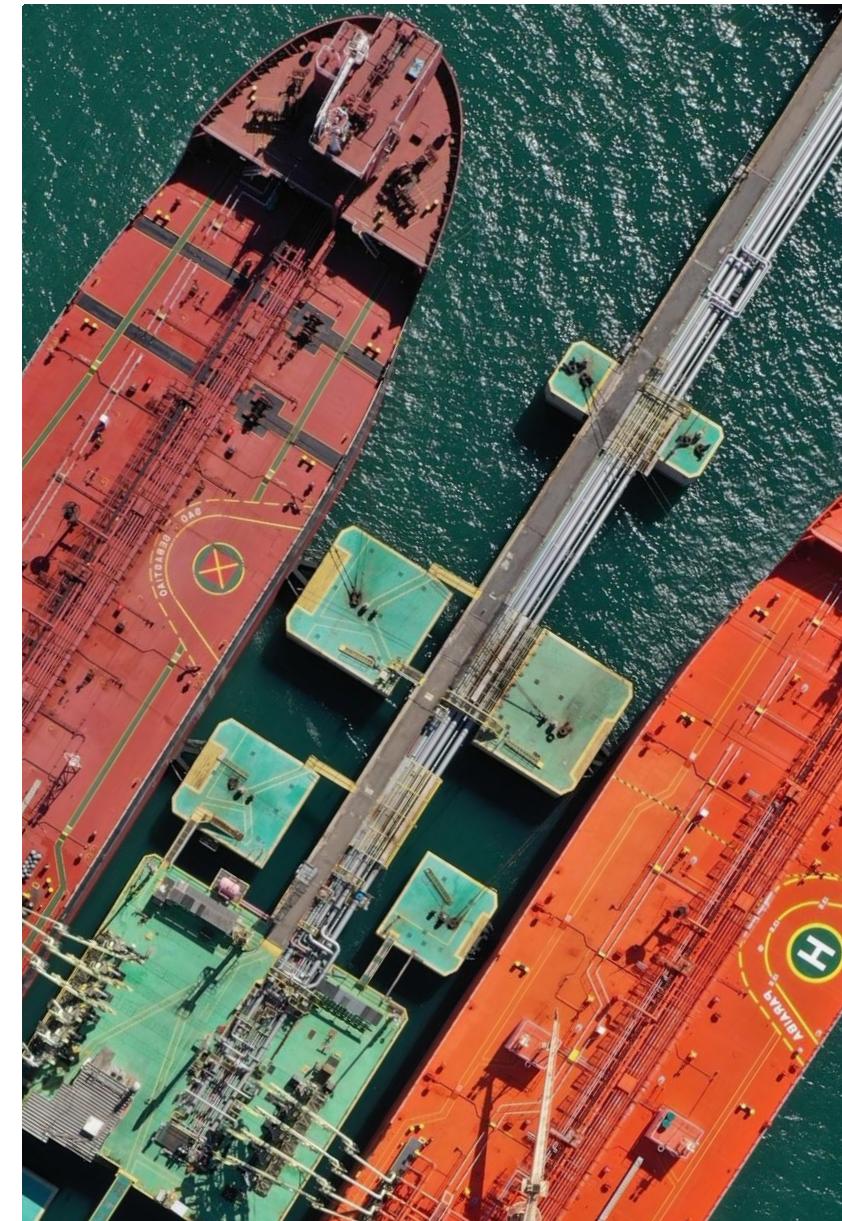
SIG-cluster-lifecycle:

- Initial release: April 2019
- API: v1alpha3

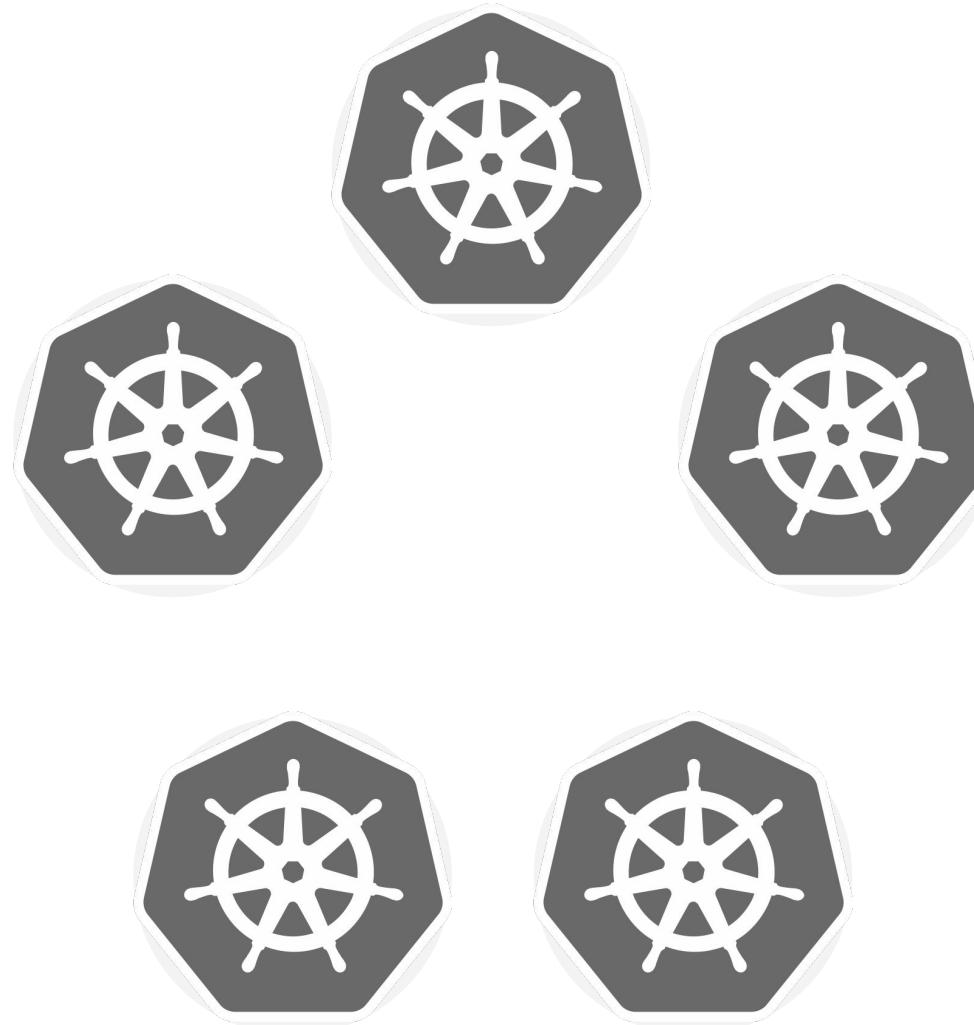


A grid of logos for various cloud providers and infrastructure companies, including AWS, Azure, packet, DigitalOcean, Google Cloud, TALOS, openstack, IBM Cloud, Tencent, BAIDU AI CLOUD, Alibaba Cloud, VMware, and EXOSCALE.

cluster-api.sigs.k8s.io/reference/providers.html

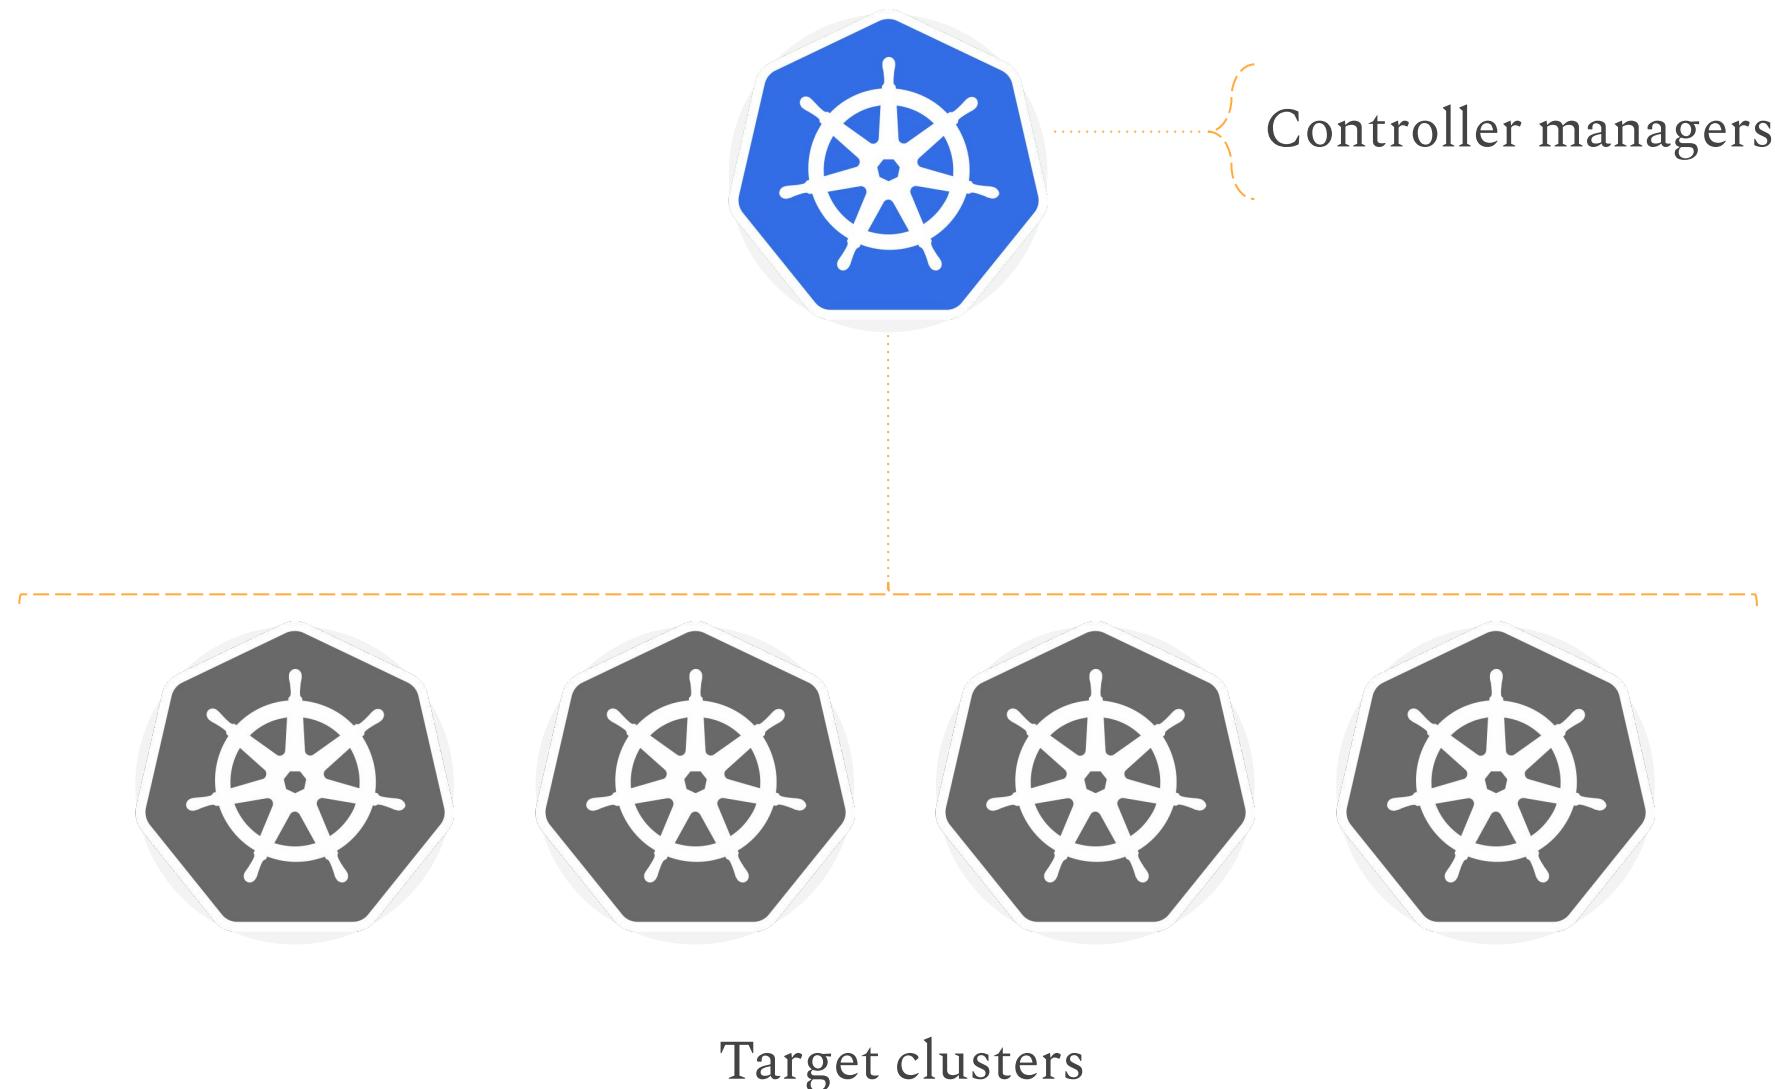


1 / 4

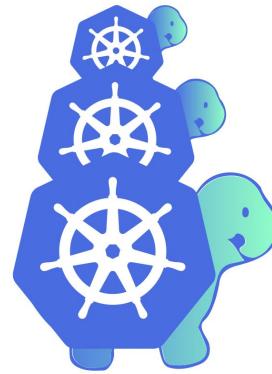


Management cluster

2 / 4

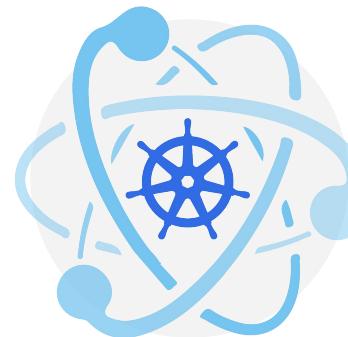


3 / 4: Controller Managers



ClusterAPI
CRDs

+



Bootstrap
Providers

+

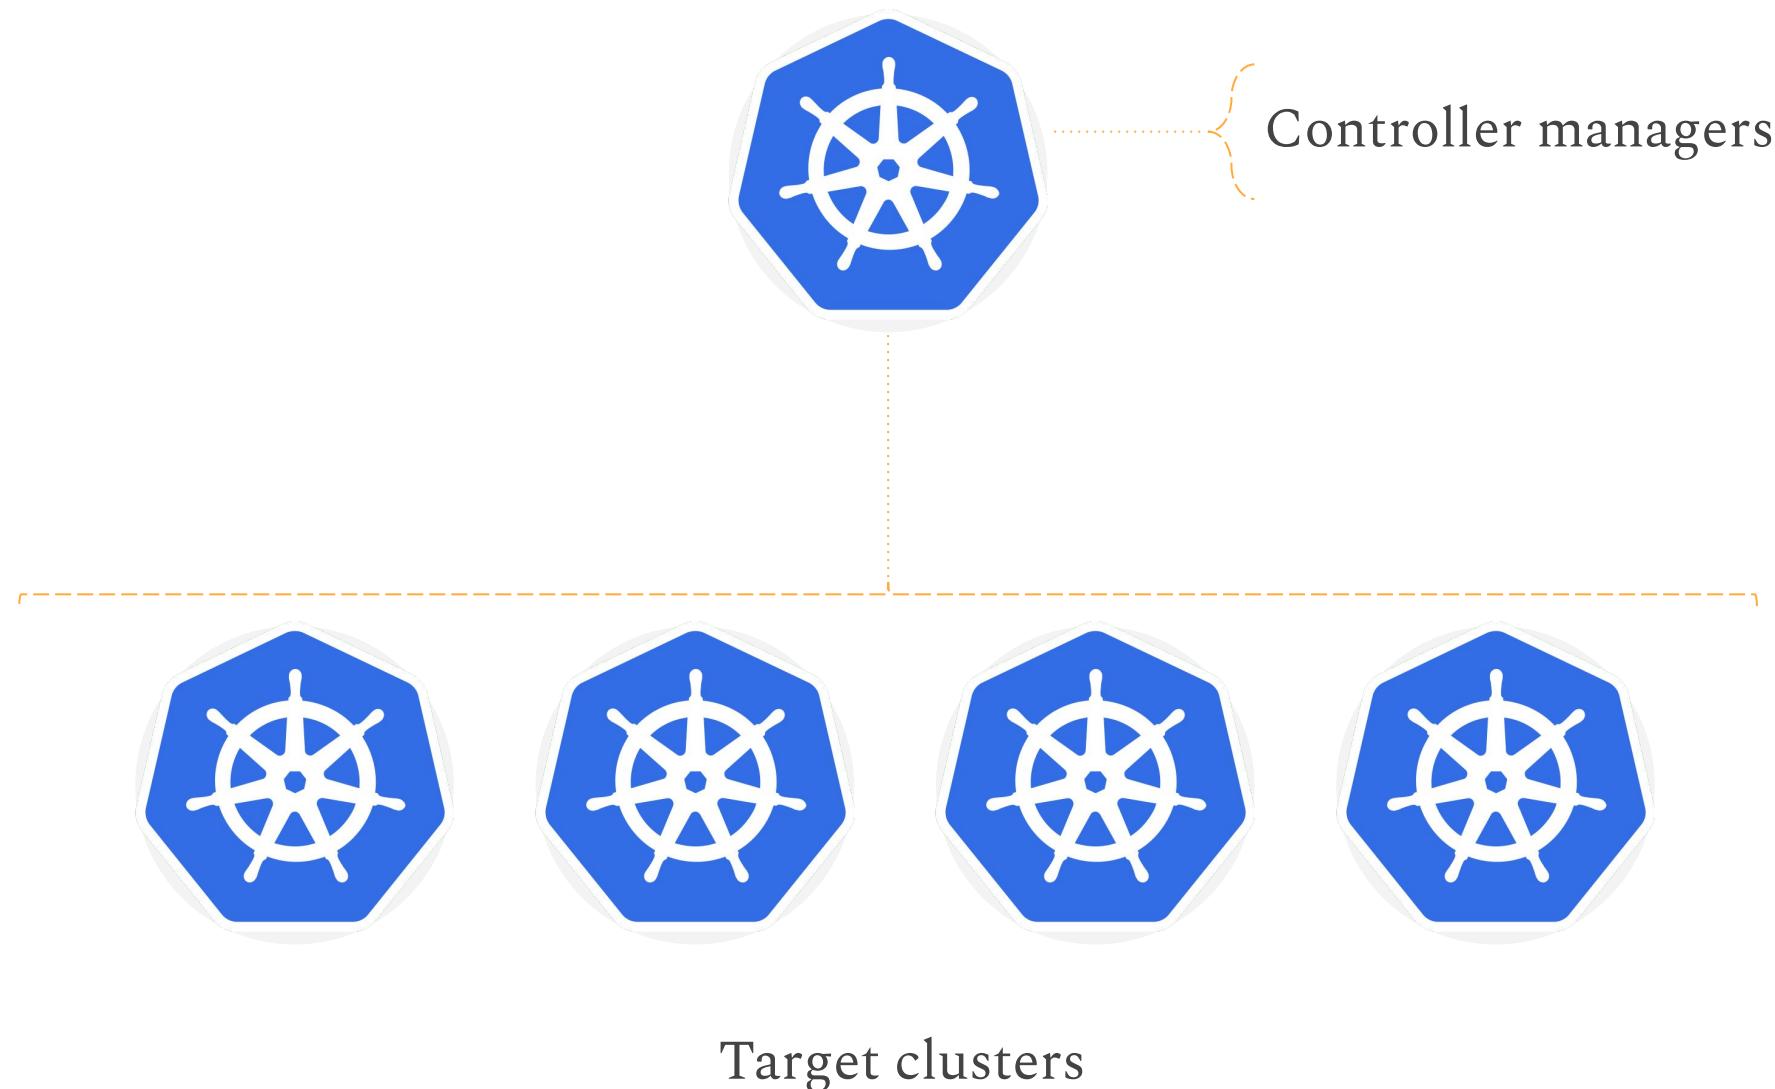


Infrastructure
Providers



Management cluster

4 / 4

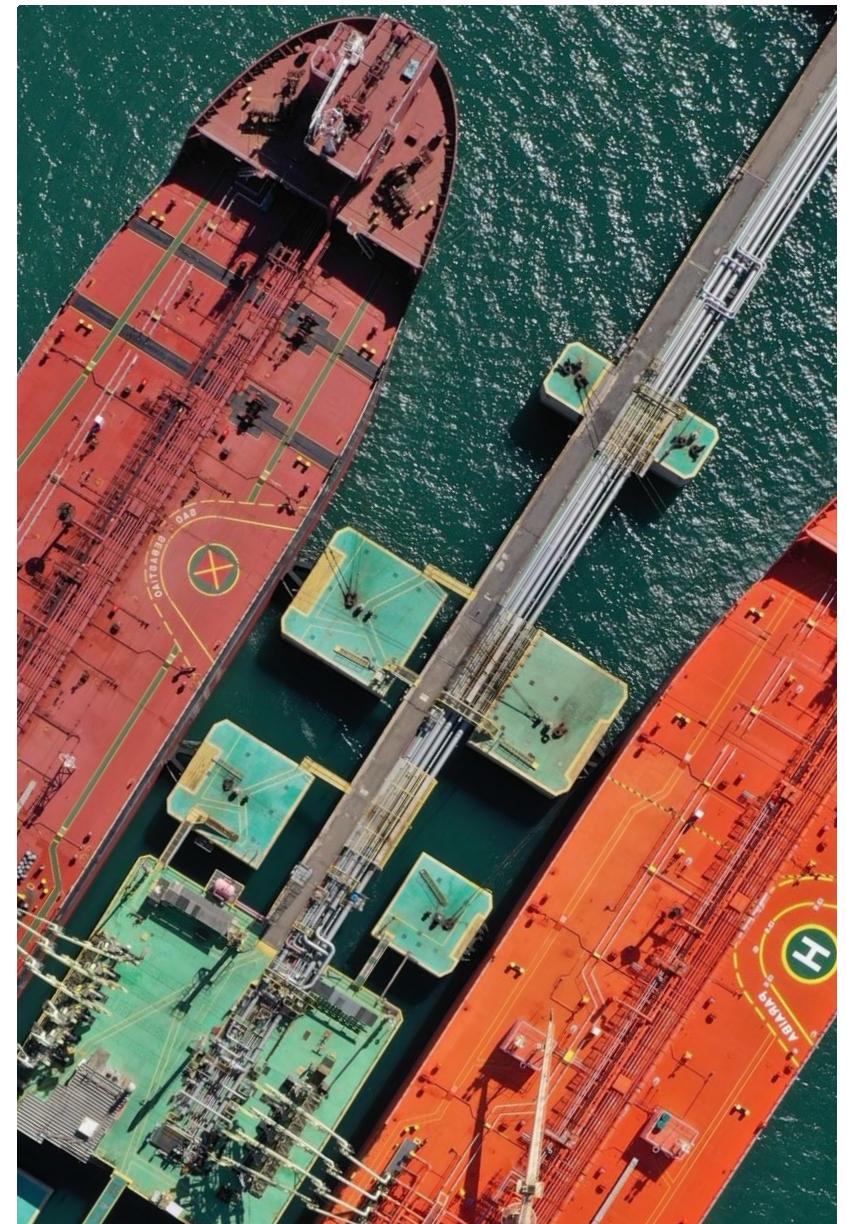




```
├── bastion-ssh-config
│   ├── tasks
│   └── templates
├── bootstrap-os
│   ├── defaults
│   ├── files
│   └── tasks
├── container-engine
│   ├── containerd
│   ├── cri-o
│   ├── docker
│   ├── docker-storage
│   └── meta
├── etcd
│   ├── defaults
│   ├── handlers
│   ├── meta
│   ├── tasks
│   └── templates
└── kubernetes
    ├── client
    ├── kubeadm
    ├── master
    ├── node
    ├── node-label
    ├── preinstall
    ├── tokens
    ├── kubernetes-apps
    │   ├── ansible
    │   ├── cloud_controller
    │   ├── cluster_roles
    └── container_engine_accelerator
        ├── csi_driver
        ├── external_provisioner
        ├── helm
        ├── ingress_controller
        └── meta
    └── metrics_server
    └── network_plugin
    └── persistent_volumes
    └── policy_controller
    └── registry
    └── rotate_tokens
└── kubespray-defaults
    ├── defaults
    ├── meta
    └── tasks
└── network_plugin
    ├── calico
    ├── canal
    ├── cilium
    ├── cloud
    ├── cni
    ├── flannel
    └── kube-router
    └── [...]
```

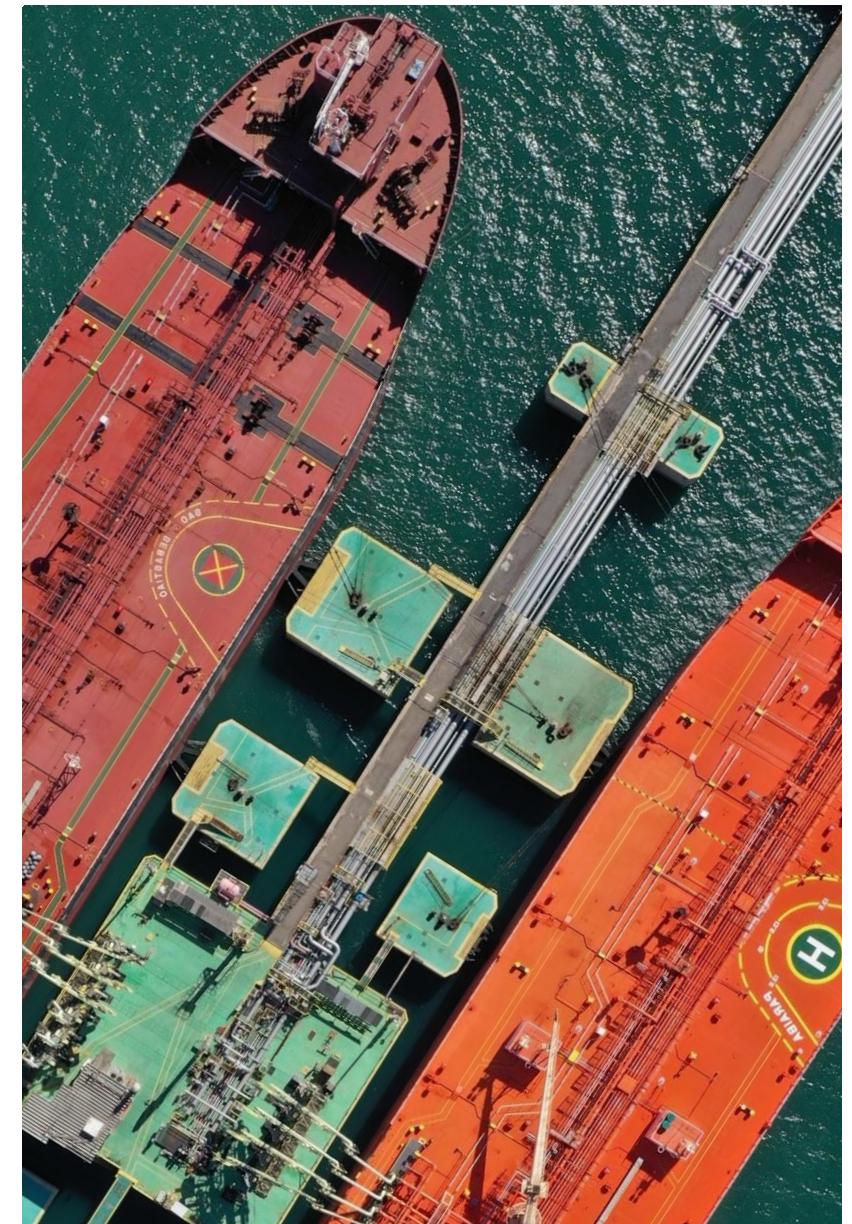


```
apiVersion: ${{ v1alpha3 }}
kind: Cluster
metadata:
  name: test-cluster
spec:
  clusterNetwork:
    pods:
      cidrBlocks: ["192.168.0.0/16"]
  infrastructureRef:
    apiVersion: ${{ v1alpha3 }}
    kind: AWSCluster
    name: test-cluster
  controlPlaneRef:
    kind: KubeADMControlPlane
  [...]
```





```
apiVersion: ${{ v1alpha3 }}
kind: Cluster
metadata:
  name: test-cluster
spec:
  clusterNetwork:
    pods:
      cidrBlocks: ["192.168.0.0/16"]
  infrastructureRef:
    apiVersion: ${{ v1alpha3 }}
    kind: AWSCluster
    name: test-cluster
  controlPlaneRef:
    kind: KubeADMControlPlane
  [...]
```





```
apiVersion: ${{ v1alpha3 }}
```

```
kind: Cluster
```

```
metadata:
```

```
  name: test-cluster
```

```
spec:
```

```
  clusterNetwork:
```

```
    pods:
```

```
      cidrBlocks: ["192.168.0.0/16"]
```

```
infrastructureRef:
```

```
  apiVersion: ${{ v1alpha3 }}
```

```
  kind: AWSCluster
```

```
  name: test-cluster
```

```
controlPlaneRef:
```

```
  kind: KubeADMControlPlane
```

```
[...]
```

```
apiVersion: ${{ v1alpha3 }}
```

```
kind: AWSCluster
```

```
metadata:
```

```
  name: test-cluster
```

```
spec:
```

```
  region: eu-central-1
```

```
  sshKeyName: default
```



```
apiVersion: ${{ v1alpha3 }}
```

```
kind: Cluster
```

```
metadata:
```

```
  name: test-cluster
```

```
spec:
```

```
  clusterNetwork:
```

```
    pods:
```

```
      cidrBlocks: ["192.168.0.0/16"]
```

```
infrastructureRef:
```

```
  apiVersion: ${{ v1alpha3 }}
```

```
  kind: GCPCluster
```

```
  name: test-cluster
```

```
controlPlaneRef:
```

```
  kind: KubeADMControlPlane
```

```
[...]
```

```
apiVersion: ${{ v1alpha3 }}
```

```
kind: GCPCluster
```

```
metadata:
```

```
  name: test-cluster
```

```
spec:
```

```
  region: europe-west3
```

```
  project: CAPI
```

```
  network:
```

```
    name: default-capi
```

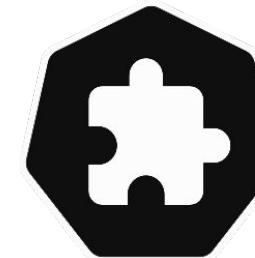


```
apiVersion: ${{ v1alpha3 }}
kind: Cluster
metadata:
  name: test-cluster
spec:
  clusterNetwork:
    pods:
      cidrBlocks: ["192.168.0.0/16"]
  infrastructureRef:
    apiVersion: ${{ v1alpha3 }}
    kind: AzureCluster
    name: test-cluster
  controlPlaneRef:
    kind: KubeADMControlPlane
  [...]
```

```
apiVersion: ${{ v1alpha3 }}
kind: AzureCluster
metadata:
  name: test-cluster
spec:
  location: westeurope
  networkSpec:
    vnet:
      name: default-capi
  resourceGroup: capi
```

ClusterAPI

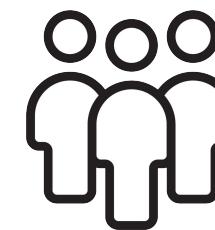
Takeaways



Building blocks



Cloud agnostic



Community

What is a **Kubernetes** Cluster?

Infrastructure

- Networking
- Firewall Rules
- Servers
- Load balancer

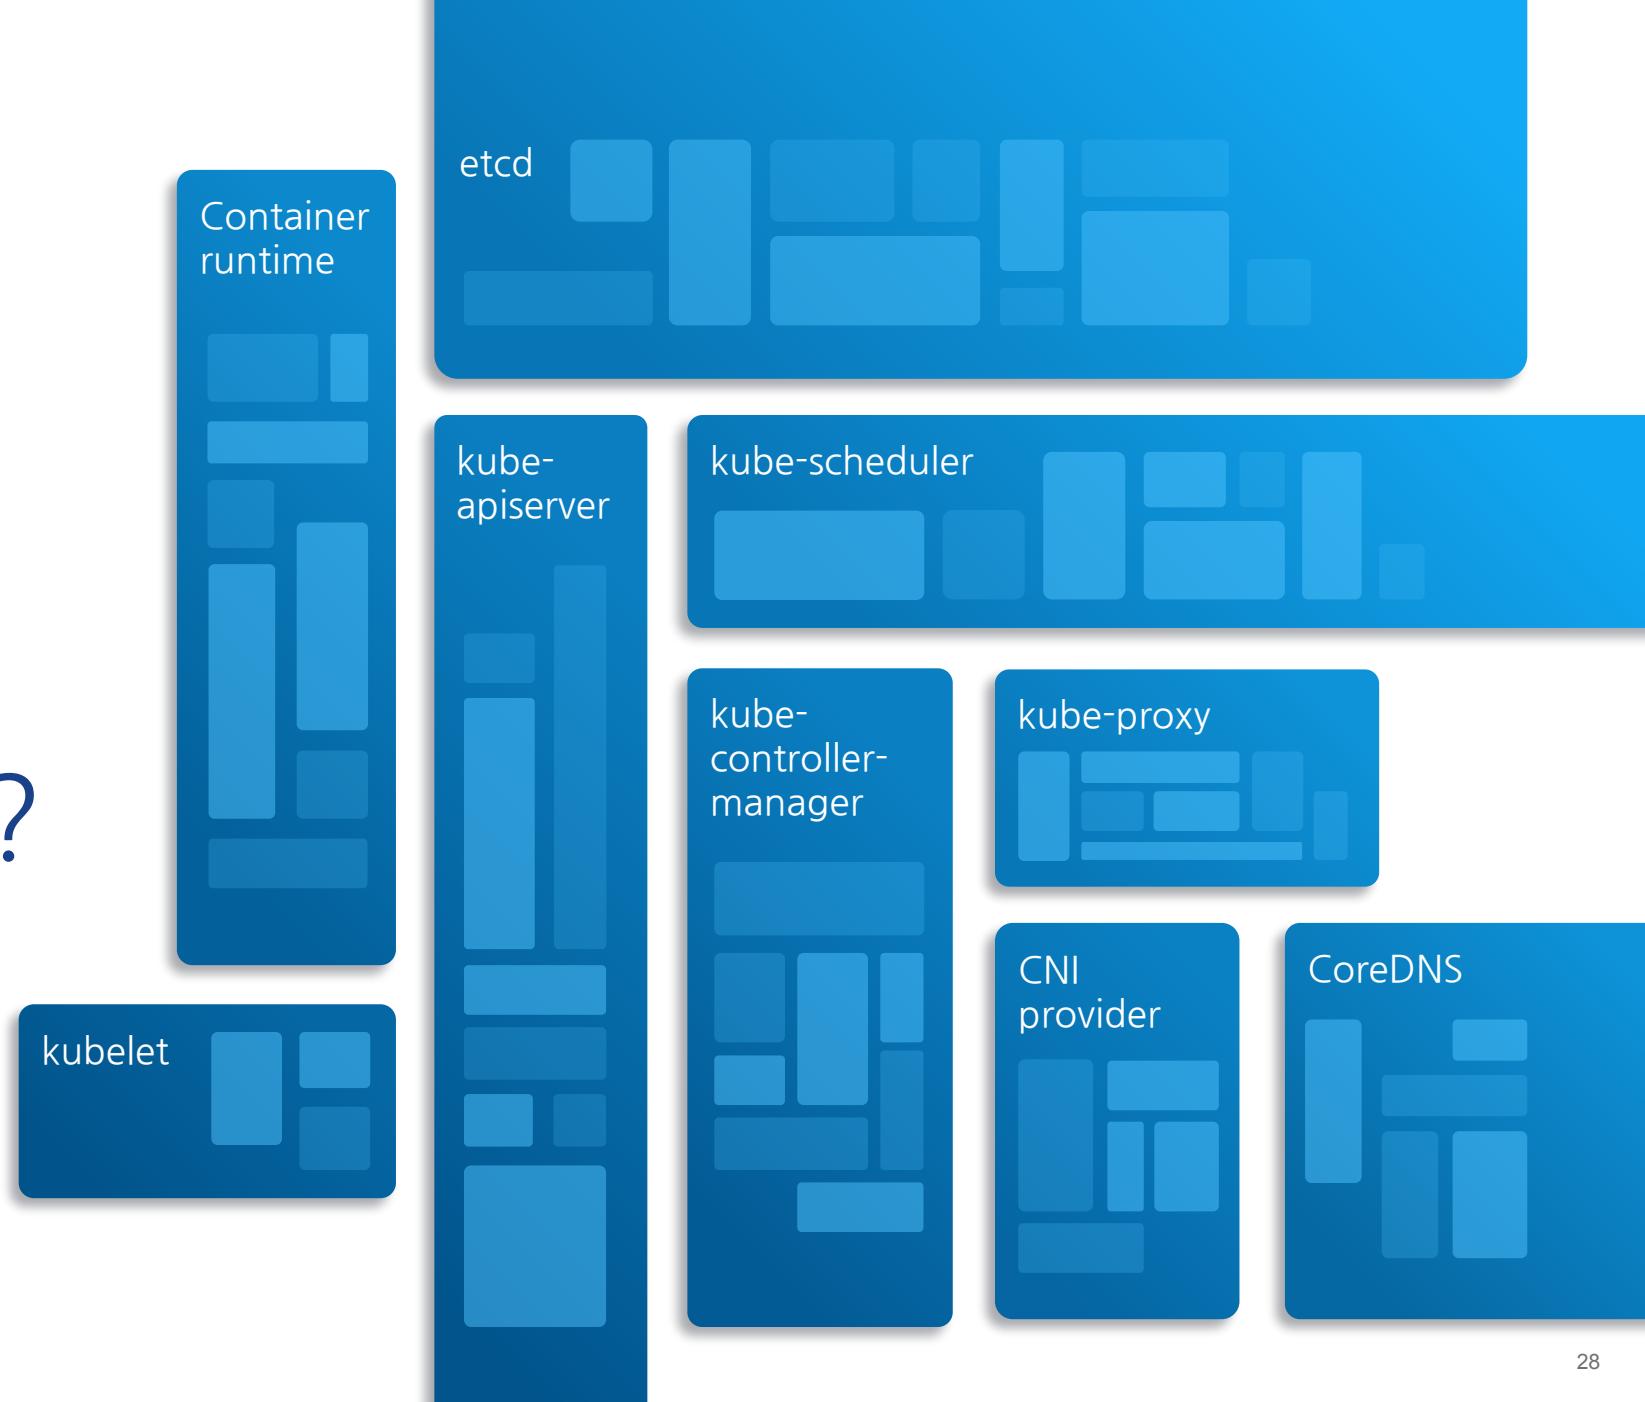
Kubernetes components

Add-ons

- Disaster recovery (Velero)
- Ingress (Contour, Gimbal)
- Registry (Harbor)
- Conformance (Sonobuoy)

User Apps

Kubernetes Components: Bootstrapping?





A PROJECT OF
SIG CLUSTER LIFECYCLE

Kubeadm

Consistent Kubernetes
cluster bootstrapping

github.com/kubernetes/kubeadm

Kubernetes Cluster Provisioning?

Infrastructure

- Networking
- Firewall Rules
- Servers
- Load balancer

Kubernetes components

Add-ons

- Disaster recovery (Velero)
- Ingress (Contour)
- Registry (Harbor)
- Conformance (Sonobuoy)

User Apps

Kubernetes in Layers

Add-ons

Ecosystem

Cloud Integrations

Runtime: Kubernetes

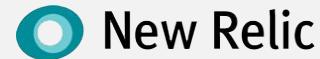
Bootstrapping: kubeadm

Provisioning: Cluster API

API Driven Infrastructure

Cluster API

Community Participation



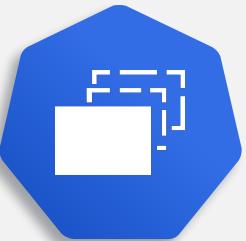
bit.ly/capidevstats

Cluster API Abstractions

Kubernetes



Pod



ReplicaSet



Deployment

Cluster API



Machine



MachineSet



MachineDeployment



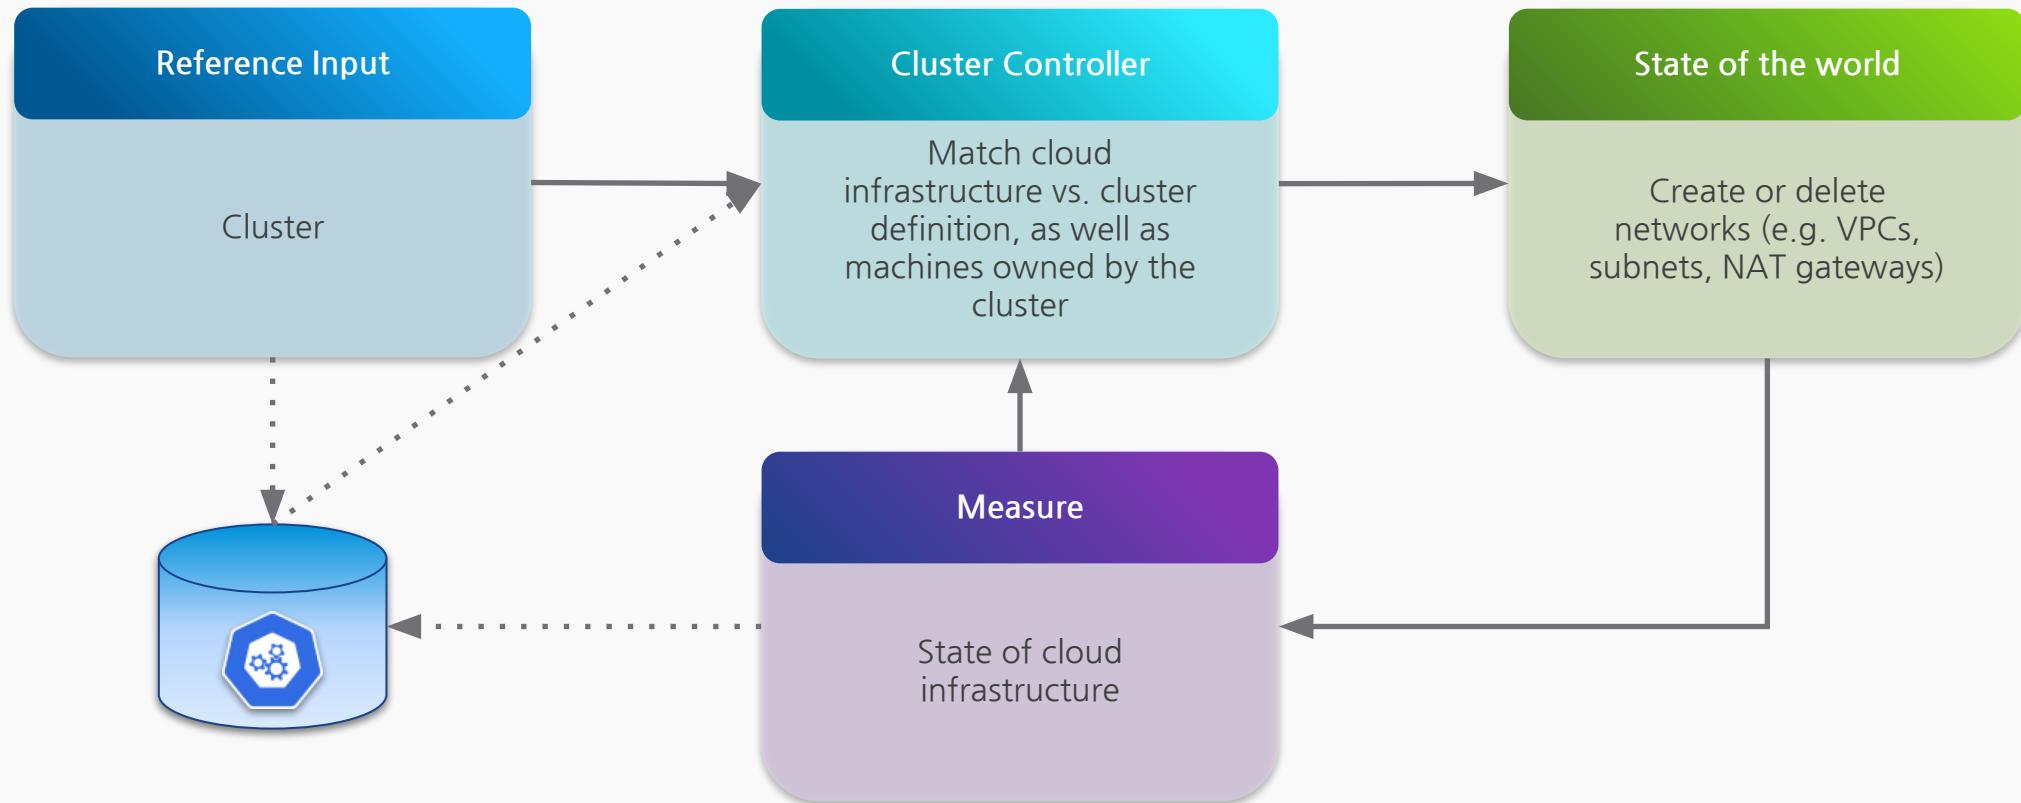
Cluster



ControlPlane

Cluster Controller

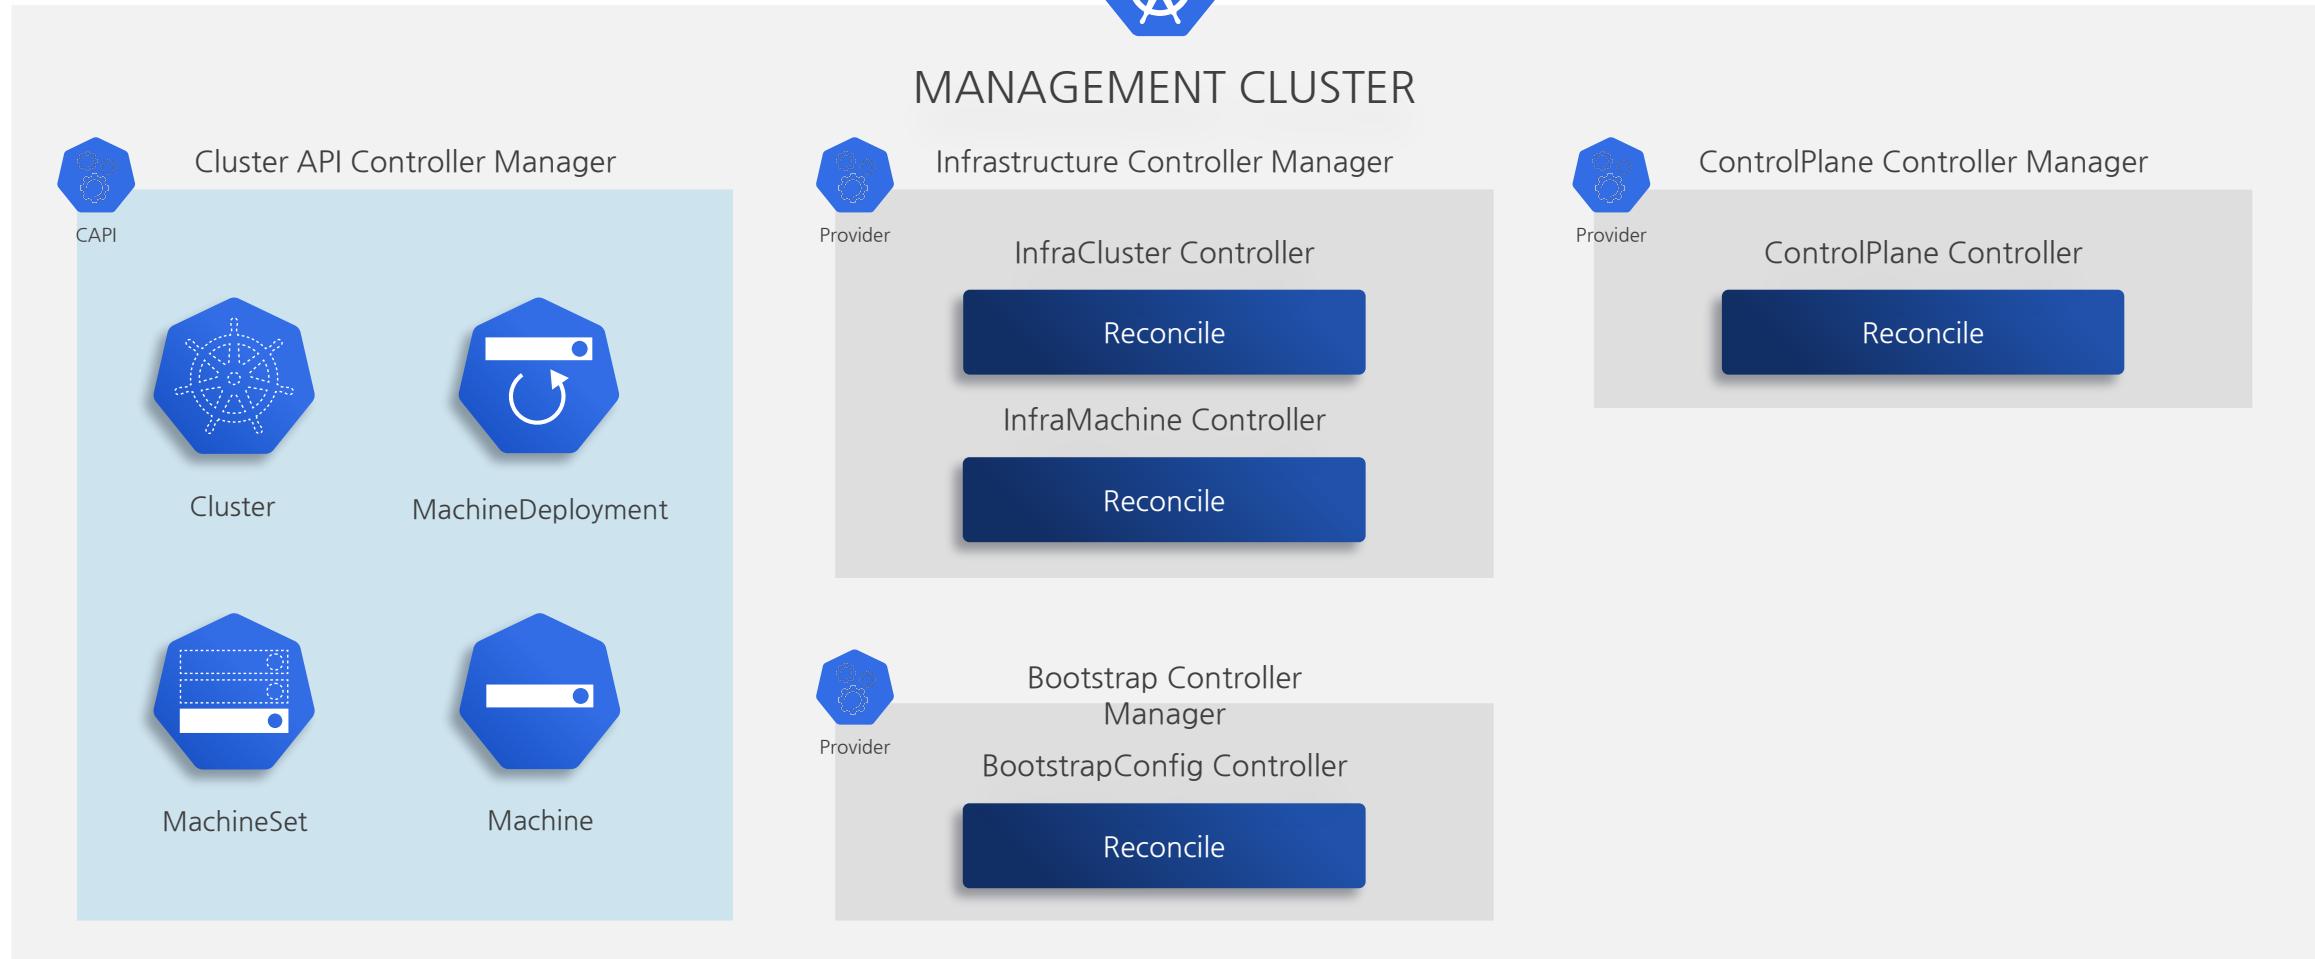
How Cluster API uses continuous reconciliation



Cluster API

Current Architecture

Kubebuilder
Cluster API
Provider



Cluster Creation

Management Cluster

Install CRDs, controllers, providers into management cluster

Create objects

Cluster

ControlPlane

Add-ons (CNI)

MachineDeployment(s) for worker nodes

Cluster API creates infrastructure and bootstraps target cluster



Cluster Creation Workflow

Start with “bootstrap” cluster (kind)

Install CRDs and controllers
into bootstrap cluster

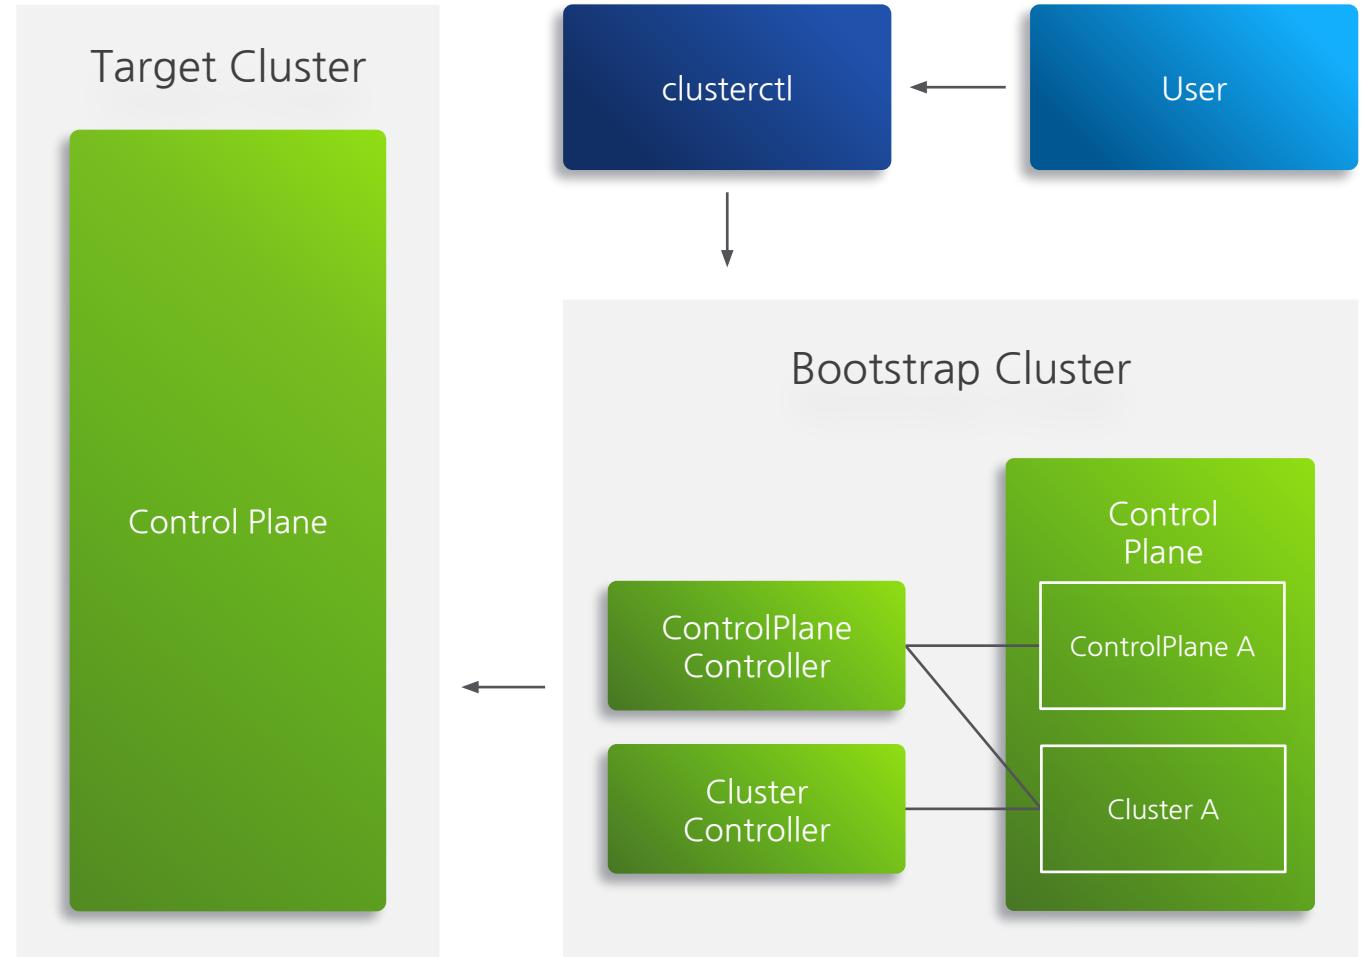
Create objects
(cluster, control panel, etc.)

Creates infrastructure and
bootstraps target cluster

Copy Cluster API and objects
into target cluster

Delete kind bootstrap cluster

Use target cluster for all
future management



clusterctl

Command-line interface for
unified lifecycle management

clusterctl init

Installs the cluster API components in target cluster to make it
into a management cluster

clusterctl upgrade

Upgrades cluster API and provider components installed in the
management cluster

clusterctl delete

Deletes provider components from the management cluster

clusterctl config cluster

Generates workload cluster manifest based on provided cluster
template

clusterctl move

Moves Cluster API objects between management clusters

Cluster Templates

Predefined workload
cluster configurations

Enables infrastructure providers and Cluster API operators to provide recommended deployment configurations for creating workload clusters.

Example templates could include:

- Highly available control plane for production clusters
- Development setup with single node control plane
- GPU-enabled nodes for machine learning workloads
- High memory nodes for memory intensive workloads

Demo

Cluster API Roadmap

Long Term Priorities

Extending control plane lifecycle management to offer more flexibility around upgrade strategies, making configuration changes

Improving observability and visibility into Cluster API for end-users

Working towards Cluster API beta to stabilize the API

Project Governance

Building community muscle for the long term

Release cadence

Reduce scope to iterate faster

At least every 6 months with public planning

v0.3.0 released in March 2020

Communication and collaboration through KEPs

Many alphas and betas until APIs stabilize
(similar to kubeadm)

Underlying controllers and providers can be robust, though!

What's New in v0.3

Control Plane

Control Plane managed as a single entity instead of individual machines

Upgrade of the Control Plane nodes

User Experience

Lifecycle Management of Cluster API components via clusterctl

Developer and end-user documentation improvements

Testing

Testing framework to enable infrastructure and bootstrap providers to validate Cluster API behaviors

What's New in v0.3

Machine Pools

Infrastructure-provider specific implementations for managed group of nodes

Machine Health Checks

Automated monitoring and removal of unhealthy nodes
Does not currently support KubeADMControlPlane; work to add support is under way

Failure Domains

Distribution of Control Plane nodes across infrastructure-specific failure domains

Cluster API and You

How to Get Involved

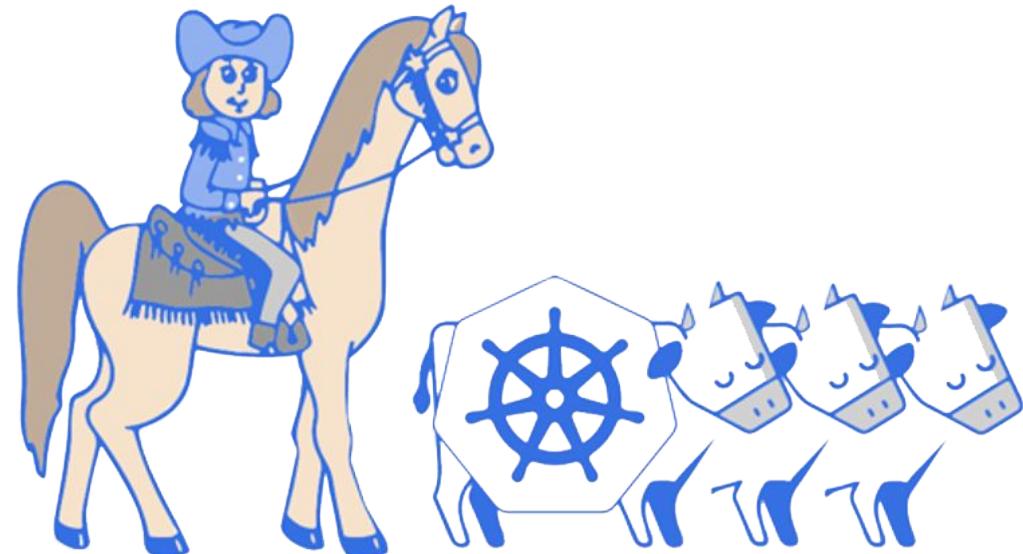
Join Us!

[Cluster API discuss forum](#)

[sig-cluster-lifecycle](#) Google Group
to gain access to documents and
calendars

Cluster API working group sessions
[Weekly on Wednesdays](#)
[10:00 PT on Zoom](#)

Chat with us on [Slack](#): #cluster-api



sigs.k8s.io/cluster-api

Thank You

Appendix

Cluster API CRDs

Cluster API: CRDs



Cluster



ControlPlane



Machine



MachineSet



MachineDeployment

Spec

pod CIDR

service CIDR

service DNS suffix

Infrastructure provider-specific cluster reference

Status

API server endpoints

Error details

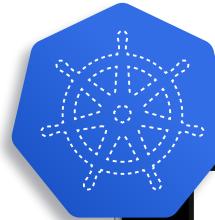
Cluster Example

Cluster

Cluster-wide configuration

Generic networking concepts like pod and service ranges or DNS domain

Providers can modify and override behavior where needed



```
apiVersion: cluster.x-k8s.io/v1alpha3
kind: Cluster
metadata:
  name: cluster-api-demo
spec:
  clusterNetwork:
    services:
      cidrBlocks: ["10.96.0.0/12"]
    pods:
      cidrBlocks: ["192.168.0.0/16"]
      serviceDomain: "cluster.local"
  infrastructureRef:
    kind: AWSCluster
    apiVersion:
infrastructure.cluster.x-k8s.io/v1alpha3
  name: cluster-api-demo
  namespace: default
```

InfraCluster Example

InfraCluster (AWSCluster)

Provider-specific cluster configuration



```
apiVersion: infrastructure.cluster.x-k8s.io/v1alpha3
kind: AWSCluster
metadata:
  name: cluster-api-demo
spec:
  region: us-east-2
  sshKeyName: default
```

Cluster API: CRDs



Cluster



ControlPlane



Machine



MachineSet



MachineDeployment

Spec

Desired replica count

Desired Kubernetes version

Infrastructure provider-specific template reference

Kubeadm Configuration for initializing & joining nodes to the control plane

Timestamp after which the control plane should be upgraded

Status

Number of replicas in the control plane

Number of replicas updated to match desired version

Number of replicas ready for the control plane

Number of unavailable replicas

Indication whether control plane has been initialized

Indication whether control plane is ready

ControlPlane Example

KubeadmControlPlane

Declarative control plane lifecycle management with Kubeadm

Replicas has the desired number of control plane machines

InfrastructureTemplate provides pluggable provider-specific machine definitions for control plane machines

KubeadmConfig provides means for configuring initialization, cluster and join configuration for control plane machines



```
apiVersion: controlplane.cluster.x-k8s.io/v1alpha3
kind: KubeadmControlPlane
metadata:
  name: capi-demo-control-plane
spec:
  replicas: 3
  infrastructureTemplate:
    kind: AWSMachineTemplate
    apiVersion: infrastructure.cluster.x-k8s.io/v1alpha3
    name: capi-demo-control-plane
  kubeadmConfigSpec:
    initConfiguration:
      ...
    clusterConfiguration:
      ...
    joinConfiguration:
      ...
  version: 1.17.4
```

ControlPlane Example

AWSMachineTemplate

Defines a template for AWS-specific machine creation



```
apiVersion: infrastructure.cluster.x-k8s.io/v1alpha3
kind: AWSMachineTemplate
metadata:
  name: capi-demo-control-plane
spec:
  template:
    spec:
      instanceType: m5.large
      iamInstanceProfile:
      "control-plane.cluster-api-provider-aws.sigs.k8s.io"
      sshKeyName: capi-demo-keys
```

Cluster API: CRDs



Cluster



ControlPlane



Machine



MachineSet



MachineDeployment

Spec

Kubernetes version

Infrastructure provider-specific reference

Provider ID

Bootstrap provider-specific configuration reference

Status

Node reference*

Last updated time

Kubernetes version

Error details

IP addresses

Phase

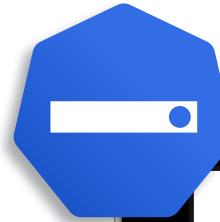
Machine Example

Machine

Configuration for a specific machine

Spec has the desired kubelet version

Providers can modify and override behavior where needed

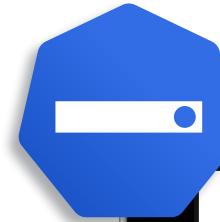


```
apiVersion: cluster.x-k8s.io/v1alpha3
kind: Machine
metadata:
  name: capi-demo-machine
  labels:
    cluster.k8s.io/cluster-name: cluster-api-demo
spec:
  bootstrap:
    configRef:
      kind: KubeADMConfig
      apiVersion: bootstrap.cluster.x-k8s.io/v1alpha3
      namespace: default
      name: capi-demo-machine
  infrastructureRef:
    kind: AWSMachine
    apiVersion: infrastructure.cluster.x-k8s.io/v1alpha3
    namespace: default
    name: capi-demo-machine
  version: "v1.17.4"
```

InfraMachine Example

InfraMachine (AWSMachine)

Infrastructure provider-specific machine configuration

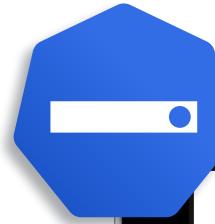


```
apiVersion: infrastructure.cluster.x-k8s.io/v1alpha3
kind: AWSMachine
metadata:
  name: capi-demo-machine
  labels:
    cluster.k8s.io/cluster-name: cluster-api-demo
spec:
  instanceType: m5.large
  iamInstanceProfile:
  "controllers.cluster-api-provider-aws.sigs.k8s.io"
  sshKeyName: default
```

Bootstrap Configuration Example

BootstrapConfig (KubeadmConfig)

Bootstrap configuration for
a specific machine



```
apiVersion: bootstrap.cluster.x-k8s.io/v1alpha3
kind: KubeadmConfig
metadata:
  name: capi-demo-machine
spec:
  initConfiguration:
  nodeRegistration:
    name: '{{ ds.meta_data.hostname }}'
    kubeletExtraArgs:
      cloud-provider: aws
```

Cluster API: CRDs



Cluster



ControlPlane



Machine



MachineSet



MachineDeployment

Spec

Replica count

Minimum ready seconds

Deletion policy

Label selector

Machine template

Status

Replicas (observed, fully labeled, ready, available)

Observed generation

Error details

Cluster API: CRDs



Cluster



ControlPlane



Machine



MachineSet



MachineDeployment

Spec

Replica count
Label selector
Machine template
Replacement strategy
Minimum ready seconds

Revision history limit
Paused
Progress deadline seconds

Infrastructure provider-specific InfraMachineTemplate
Bootstrap provider-specific BootstrapConfigTemplate

Status

Observed generation
Replicas (total, updated, ready, available, unavailable)

MachineDeployment Example

MachineDeployment

Declarative updates for Machines via MachineSets

Update strategy allows control of the rate at which a change is applied



```
apiVersion: cluster.x-k8s.io/v1alpha3
kind: MachineDeployment
metadata:
  name: nodepool-0
  labels: {cluster.k8s.io/cluster-name: cluster-api-demo}
spec:
  replicas: 3
  selector:
    matchLabels:
      cluster.x-k8s.io/cluster-name: cluster-api-demo
      nodepool: nodepool-0
  template:
    metadata:
      labels:
        cluster.x-k8s.io/cluster-name: cluster-api-demo
        nodepool: nodepool-0
    spec:
      version: v1.17.4
      bootstrap:
        configRef:
          name: nodepool-0
          apiVersion: bootstrap.cluster.x-k8s.io/v1alpha3
          kind: KubeadmConfigTemplate
      infrastructureRef:
        name: nodepool-0
        apiVersion: infrastructure.cluster.x-k8s.io/v1alpha3
        kind: AWSMachineTemplate
```

InfraMachineTemplate Example

InfraMachineTemplate (AWSMachineTemplate)

Template used for creating
Provider-specific InfraMachine instances
for the MachineDeployment



```
apiVersion: infrastructure.cluster.x-k8s.io/v1alpha3
kind: AWSMachineTemplate
metadata:
  name: nodepool-0
spec:
  template:
    spec:
      instanceType: m5.large
      iamInstanceProfile:
      "nodes.cluster-api-provider-aws.sigs.k8s.io"
      sshKeyName: default
```

BootstrapConfigTemplate Example

BootstrapConfigTemplate (KubeadmConfigTemplate)

Template used for creating BootstrapConfig objects for configuring the InfraMachine instances for the MachineDeployment



```
apiVersion: bootstrap.cluster.x-k8s.io/v1alpha3
kind: KubeadmConfigTemplate
metadata:
  name: nodepool-0
spec:
  template:
    spec:
      joinConfiguration:
        nodeRegistration:
          name: '{{ ds.meta_data.hostname }}'
        kubeletExtraArgs:
          cloud-provider: aws
```

Cluster API Deployment Topology

5 separate deployments

Deployment 1

Common Cluster
API controllers

Cluster
Machine

MachineSet

MachineDeployment

Deployment 2

Cluster API
webhooks

Deployment 3

Control Plane Provider
-specific controllers
KubeadmControlPlane

KubeadmControlPlane

Deployment 4

Infrastructure Provider
-specific controllers

InfraCluster
InfraMachine

Deployment 5

Bootstrap Provider
-specific controllers

BootstrapConfig