



KubeDirector: Open Source Project for Stateful Applications on Kubernetes



Today's Speakers



Tom Phelan

Co-Founder and Chief Architect
BlueData

 @tapbluedata



Joel Baxter

Distinguished Engineer
BlueData

 @joel_k_baxter

Agenda

- Kubernetes (K8s) and Stateful / Stateless Applications
- Complex Stateful Applications on Kubernetes
- BlueData, BlueK8s, and KubeDirector
- KubeDirector Deep Dive
- KubeDirector Demonstration
- Key Takeaways



What is Kubernetes (K8s?)



- Open source “**platform**” for container orchestration
- Platform **building blocks** vs. turnkey platform
 - <https://kubernetes.io/docs/concepts/overview/what-is-kubernetes/#what-kubernetes-is-not>
- Top use case is **stateless / microservices** deployments
- Evolving for **stateful** applications

Stateless Applications on K8s

- **Stateless**
 - Each application service instance is configured identically
 - All information stored remotely
 - “Remotely” refers to some persistent storage that has a life span different from that of the container
 - Frequently referred to as “*cattle*”



Stateful Applications on K8s?

- **Stateful**

- Each application service instance is configured differently
- Critical information stored locally
- “Locally” means that the application running in the container accesses the information via file system reads/writes rather than some remote access protocol
- Frequently referred to as “*pets*”



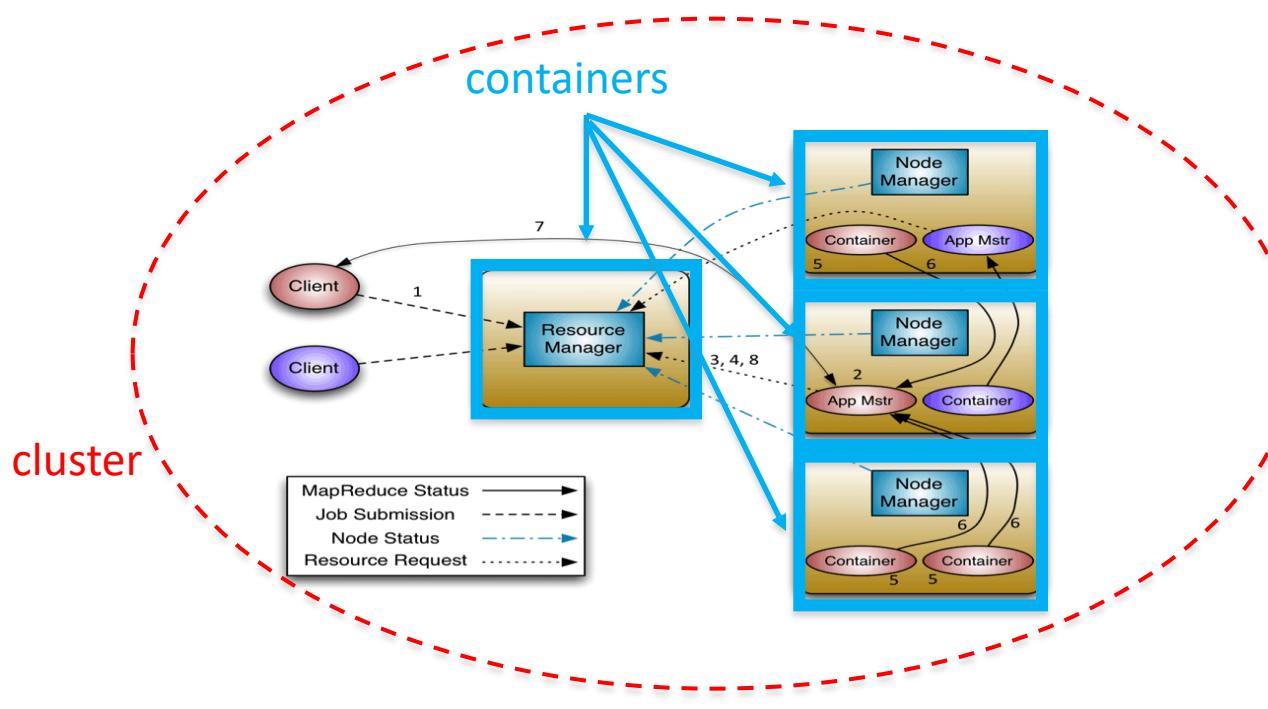
Complex Stateful Applications

- Big Data / AI / Machine Learning / Deep Learning
- What do all these applications have in common?
 - Require large amounts of data
 - Use distributed processing, multiple tools / services
 - When on-prem, typically deployed on bare-metal
 - Do *not* have a cloud native architecture
 - No microservices
 - Application instance-specific state



Example: Hadoop in Containers

Running Hadoop clusters in containers:



Kubernetes – Components

- Objects
- Pods
- Statefulsets
- PersistentVolumes
- Operators
- Custom Resource Definitions



kubernetes

Kubernetes – Operators



- Operator
 - A way of packaging, deploying, and managing a given application
- Operator Framework
 - A set of developer and runtime tools to help accelerate the writing of a Operator
- Operator SDK
 - An SDK that further hides the complexities of the Kubernetes API

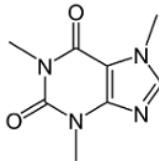
Source: <https://coreos.com/operators>



Kubernetes – Operators



- Application-specific means a new operator needs to be written for each application



theano



What to Do?

- There needs to be an easier way to deploy and manage clusters running complex stateful applications



BlueK8s and KubeDirector

- BlueK8s is an Apache open source initiative focused on bringing enterprise support for complex stateful applications to Kubernetes
- A series of open source projects will be rolled out under the BlueK8s umbrella
 - The first major project is “KubeDirector”:
<https://github.com/bluek8s/kubedirector>

Source: www.bluedata.com/blog/2018/07/operation-stateful-bluek8s-and-kubernetes-director



Operation: Stateful BlueK8s and KubeDirector



kubernetes

Motivation

- Why create KubeDirector? Why use it?
 - E.g. why not app-specific operators, Helm, Kubeflow...
- Reframed: which architecture enables features we want (current or future)?
- Find sweet spot for users between two extremes:
 - Direct use of K8s APIs & “generic” deployment
 - Hardcoded application-specific solutions
- Abstractions + features guided by domain focus

Domain Focus

- Interested in best supporting apps that:
 - Are scale-out
 - May have “non cloud-native” service architecture
 - Have stateful cluster members
 - Need to access data lakes
 - Have user roles w/ distinct workflows and privileges
 - Integrate w/ enterprise services for authentication, certificate and license management, etc.

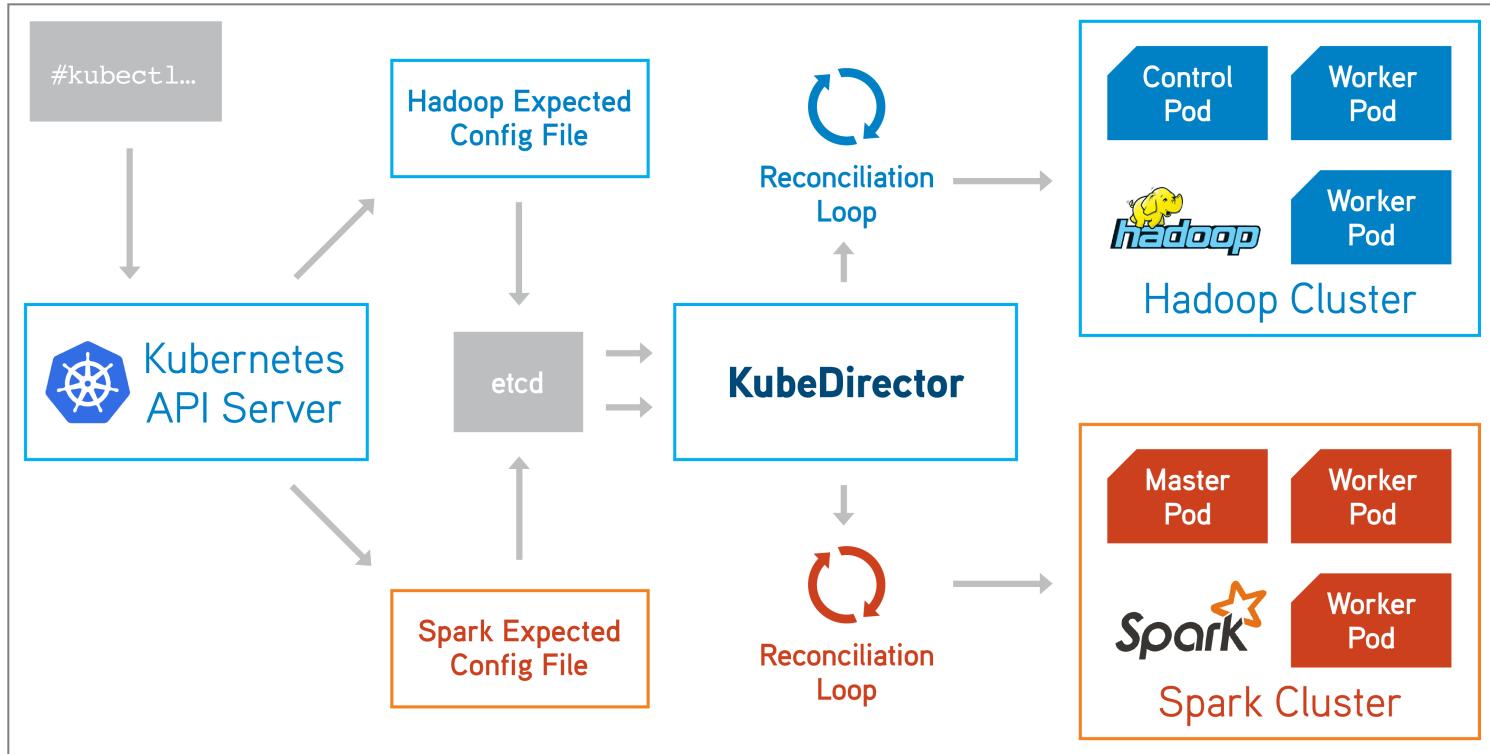
KubeDirector Overview

- KubeDirector is a K8s “custom controller”
- Watches for custom resources (CRs) to appear/change
- Creates/modifies standard K8s resources (StatefulSets etc.) in response, to implement specs from CRs
- Differs from normal Kubernetes Operator pattern:
 - No app-specific logic in KubeDirector code
 - App deployment is data-driven from external app definitions
 - Supports interactions among different apps + other objects



Deploy KubeDirector to K8s

kubectl create -f kubedirector/deployment.yaml



Separation of Concerns

- Application experts (on-site or elsewhere)
 - Responsible for making app images/metadata/configscripts
 - No need to write Go code or understand Operator concepts
- Administrators (on-site)
 - Select which apps are available to end users
 - Change app versions independently of KubeDirector upgrade
- End users
 - Pick from menu of applications and config choices

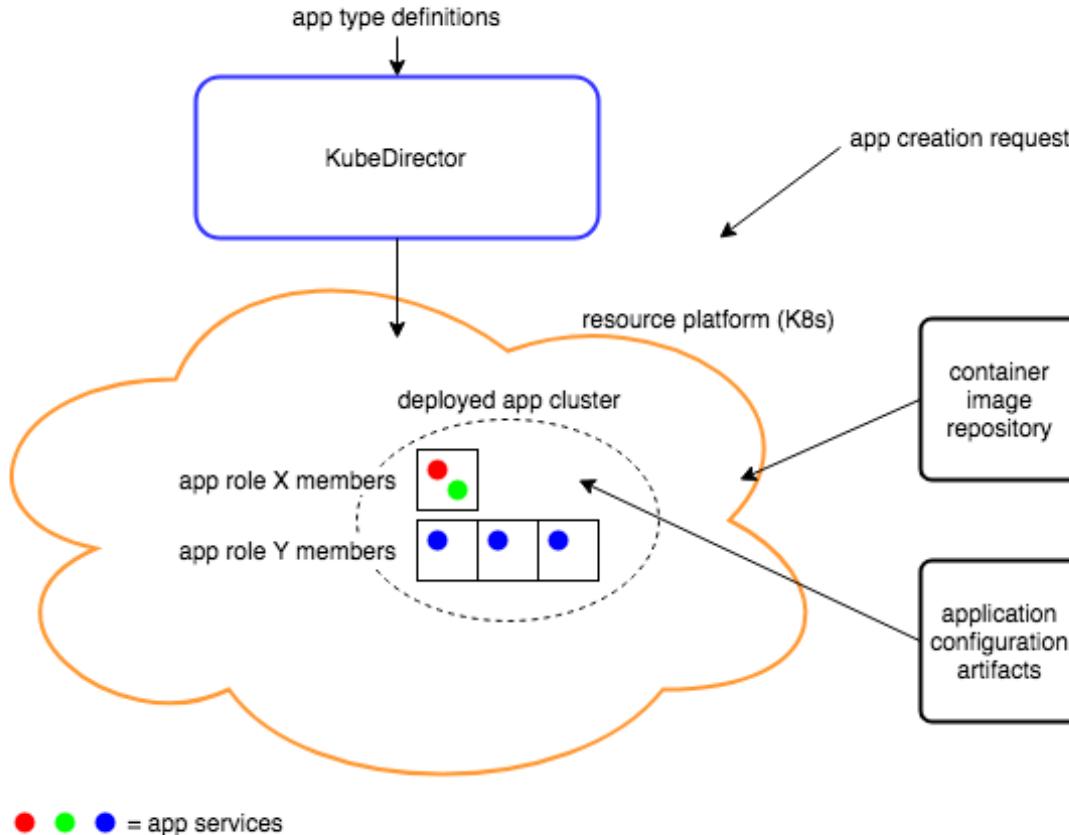
Alternatives Comparison 1/2

- Support distinctions between IT, app expert, project manager, and data scientist
 - Unlike Helm 2/3 & Kubeflow
- Integrate with K8s user authentication and ACLs
 - Unlike Helm 2 (Tiller)
- Support post-deployment autoremediation, autoscale, and other lifecycle events w/ app-specific logic
 - Unlike Helm 3 & Kubeflow

Alternatives Comparison 2/2

- Also a couple of behaviors not found in app-specific operators, and not a picnic in other solutions:
 - Support end-user import of new application types
 - Apply common features across multiple application types from different developers

KubeDirector Concepts



Custom Resource Definitions

- Primary CRD: KubeDirectorCluster
 - Models any kind of application instance launchable by KubeDirector
- Other CRDs for related objects, e.g.
 - App definitions (KubeDirectorApp)
 - DataTaps and other shared storage
 - Config sets for AD/LDAP integration for containers
 - Machine Learning models
- This talk will concentrate on KubeDirectorCluster/App

- KubeDirector Administration
- Application Preparation
- Application Instance Deployment

Deployment

- Create custom resource definitions (CRDs) in your K8s cluster
- Deploy KubeDirector
 - Normally runs in a Pod on same K8s cluster
 - Authenticates to K8s API w/ privileged service account
- Configure KubeDirector global settings
 - E.g. supply app definitions, set types of service & storage



- KubeDirector Administration
- Application Preparation
- Application Instance Deployment

App Definition Metadata

- App identifier/description/version
- Service endpoints
- Available “roles”, and container image per role
- Available deploy-time choices, and their effects on services per role
- Info for optional runtime setup package
- And more!

App Definition Example 1/3

```
apiVersion: kubedirector.bluedata.io/v1alpha1
kind: KubeDirectorApp
metadata:
  name: spark221e2

spec:
  label:
    name: Spark 2.2.1 on centos7x with Jupyter
  default_image_repo_tag: docker.io/bluedata/sparkbase:2.0
  default_config_package:
    package_url: https://s3.amazonaws.com/mybucket/spark221e2/appconfig.tgz
```



App Definition Example 2/3

```
roles:  
- id: controller  
  cardinality: 1  
- id: worker  
  cardinality: 0+  
  
services:  
- id: spark  
  label:  
    name: Spark master  
  endpoint:  
    port: 7077  
  
          - id: spark_master_ui  
            label:  
              name: Spark master (web UI)  
            endpoint:  
              port: 8080  
              is_dashboard: true  
              url_scheme: http  
- id: spark_worker_ui  
  label:  
    name: Spark worker (web UI)  
  endpoint:  
    port: 8081  
    is_dashboard: true  
    url_scheme: http
```



App Definition Example 3/3

```
config:  
  
  selected_roles:  
    - controller  
    - worker  
  
  role_services:  
    - role_id: controller  
      service_ids:  
        - spark  
        - spark_master_ui  
    - role_id: worker  
      service_ids:  
        - spark_worker_ui
```

Application Setup Package

- Optional tgz injected into each container, contains:
 - Entrypoint script
 - Standard script functions for reading deployment info
 - Any artifacts (config file templates etc.) required for setup
- Entrypoint script will be invoked at lifecycle events:
 - This container has just been created
 - Some other member(s) added to or removed from the cluster

Setup Script Actions

- Perform any setup that requires runtime info
 - E.g. FQDNs of other member(s) of the cluster
- Enable and start appropriate services
 - Can query the role of current node
 - Services-to-start depend on role and deploy-time choices
- Can use features of KubeDirector Agent in future

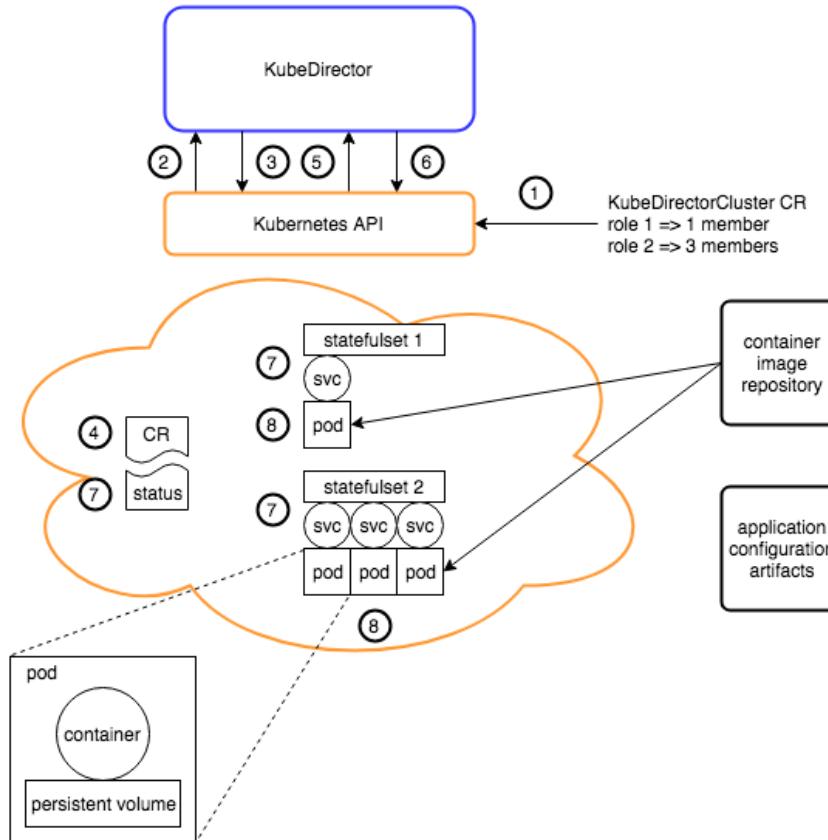
- KubeDirector Administration
- Application Preparation
- **Application Instance Deployment**

Custom Resource Creation

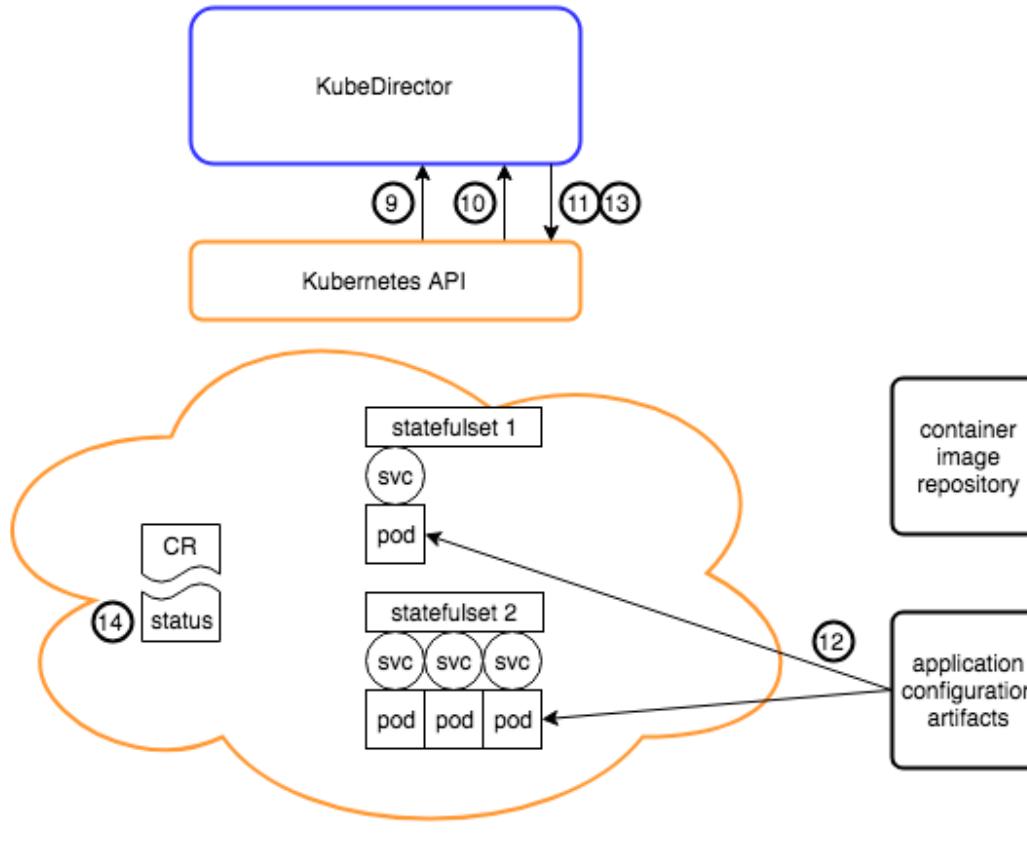
```
apiVersion: "kubedirector.bluedata.io/v1alpha1"
kind: "KubeDirectorCluster"
metadata:
  name: "spark-instance"
spec:
  app: spark221e2
  roles:
    - id: controller
      resources:
        limits:
          memory: "4Gi"
    - id: worker
      members: 2
```



Cluster Creation Sequence 1/2



Cluster Creation Sequence 2/2



Other Operations

- Shrink & expand of role member count is handled similarly
- All resources are automatically cleaned up if CR is deleted (because CR is their “owner”)
- End user can read the CR to see current status, service objects, event history, etc.

Key Takeaways

- Running complex stateful applications on Kubernetes is challenging today
- The goal of BlueK8s and KubeDirector is to make it easier to run such applications on Kubernetes
- Learn more about KubeDirector:
 - <https://github.com/bluek8s/kubedirector/wiki>



Join the KubeDirector Community!



Tom Phelan



@tapbluedata

Joel Baxter



@joel_k_baxter

<https://github.com/bluek8s>

