



Network Service Mesh Webinar





Agenda

- Housekeeping
- NSM Vision
- State of the NSM
- NSM Future
- Deep Dive:
 - How the Magic Works
 - Interdomain
 - HW NICs





Housekeeping



<https://networkservicemesh.io>



NSMCon

Nov 18, 2019 | San Diego, California
Colocated with Kubecon+CloudNativeCon 2019



These slides



Network Service Mesh

NSM Vision





NSM Vision

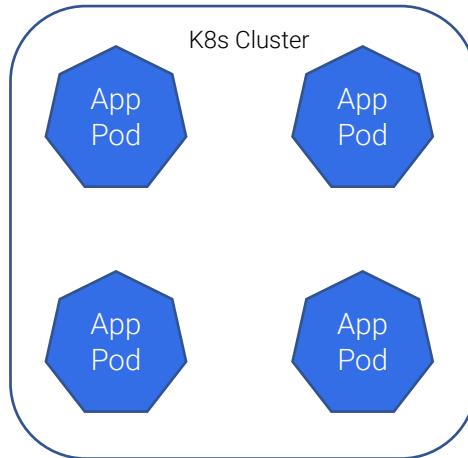
- The Problems
- The Non-Solutions
- The NSM Solutions





Runtime Domain

K8s is a 'Runtime Domain'...



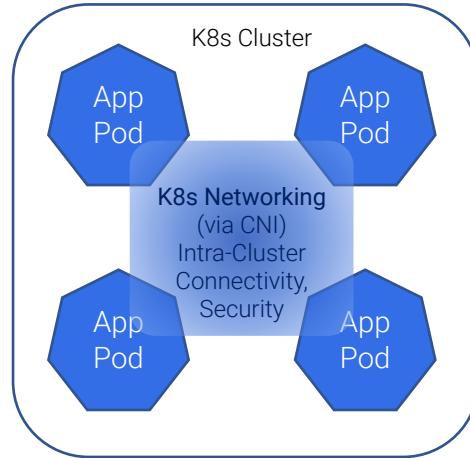


Connectivity Domain

K8s is a ‘Runtime Domain’...

With a ‘Connectivity Domain’...

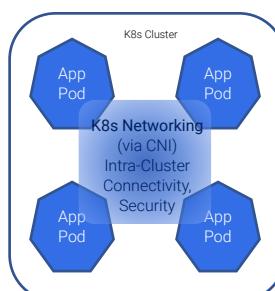
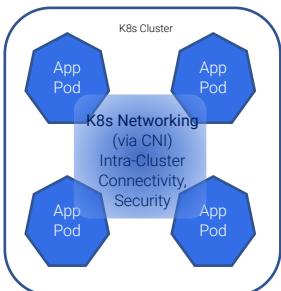
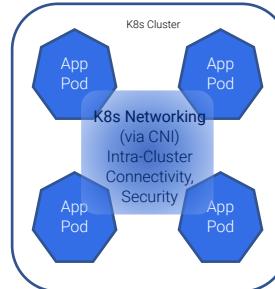
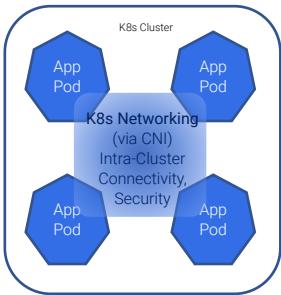
- Pure L3
- + **Service Discovery/Routing:** K8s Services
- + **Isolation:** K8s Network Policies (Isolation)
- (Optionally) + L7 Service Mesh(Istio etc)
- Intra Cluster





The Problems

What about East/West traffic between workloads (Pods) in different clusters?



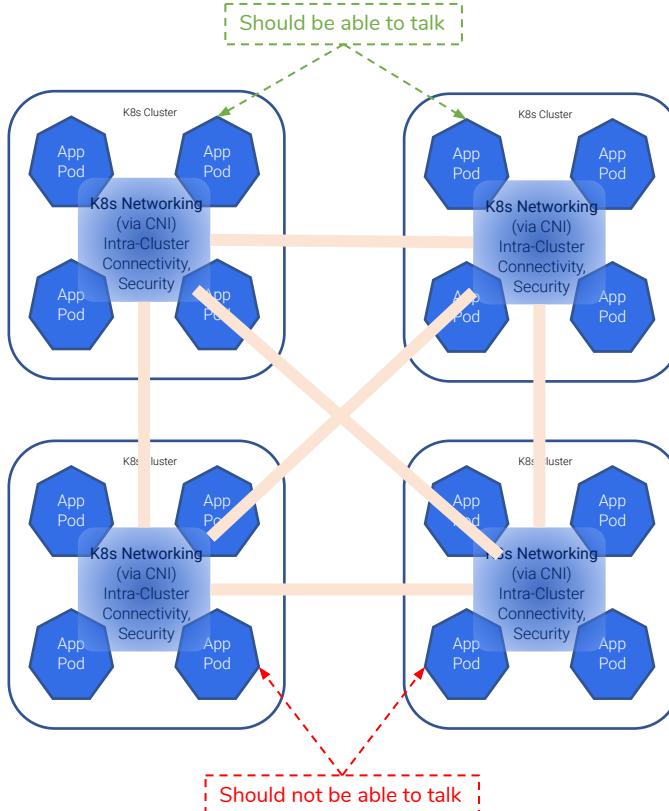


The Non-Solutions

All non-NSM attempts involve cluster-to-cluster networking:

Problems:

- **Inter-cluster Workload Isolation**



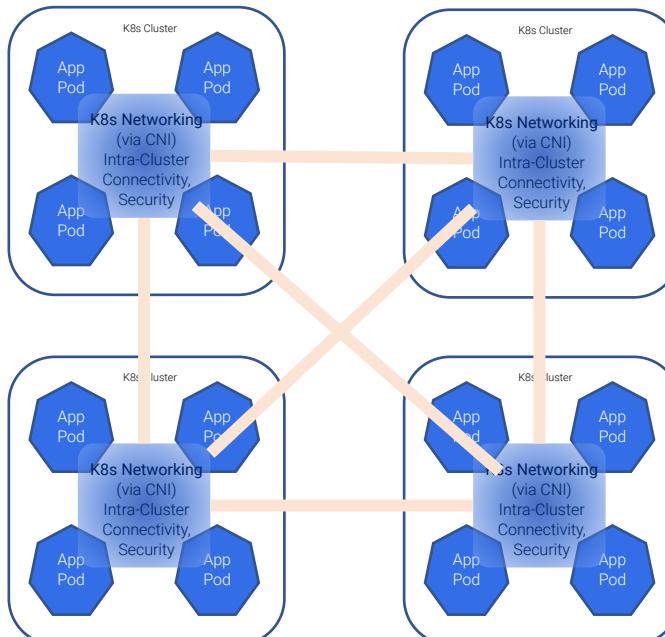


The Non-Solutions

All non-NSM attempts involve cluster to cluster networking:

Problems:

- Inter-cluster Workload Isolation
- **Full Mesh between clusters explodes combinatorics – number of links scales like number of clusters choose 2 (ie: factorially)**



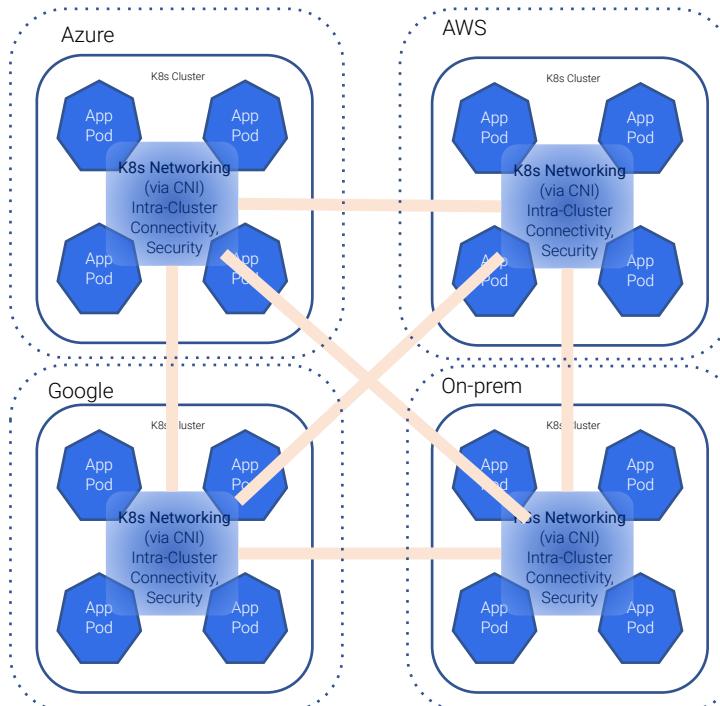


The Non-Solutions

All non-NSM attempts involve cluster to cluster networking:

Problems:

- Inter-cluster Workload Isolation
- Full Mesh between clusters explodes combinatorics – number of links scales like number of clusters choose 2 (ie: factorially)
- **Complex often manual cluster to cluster link setup between different public/private cloud providers (possibly involving complex firewall rules depending on how its done).**



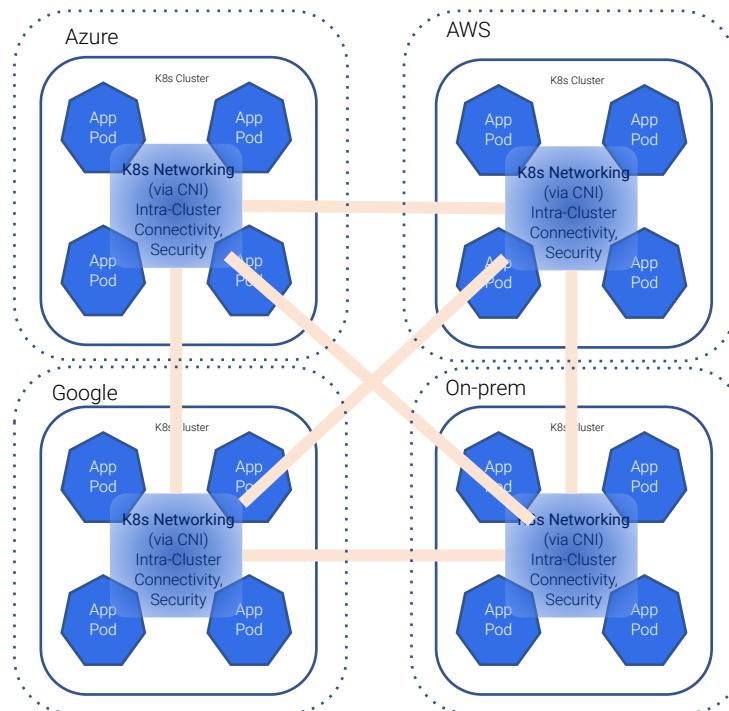


The Non-Solutions

All non-NSM attempts involve cluster to cluster networking:

Problems:

- Inter-cluster Workload Isolation
- Full Mesh between clusters explodes combinatorics – number of links scales like number of clusters choose 2 (ie: factorially)
- Complex often manual cluster to cluster link setup between different public/private cloud providers (possibly involving complex firewall rules depending on how its done).
- **Inter-cluster Service Discovery/Routing**

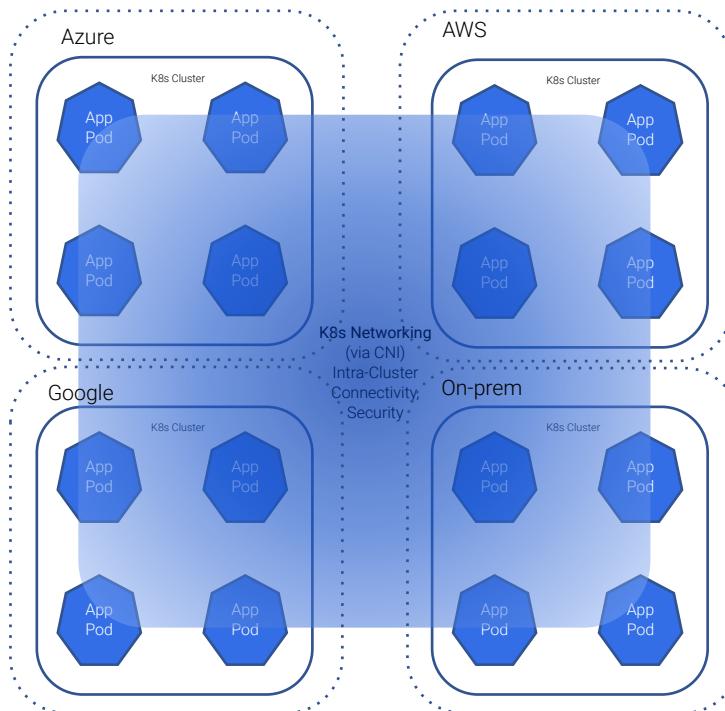


The Federation non-Solution



Attempt to 'Federate' multiple clusters:

- Hides rather than fixes inter-cluster link combinatorics/complexity/manualness
- Doesn't scale:
 - Services/Network Policies have enough trouble scaling in a single cluster, with low latency updates

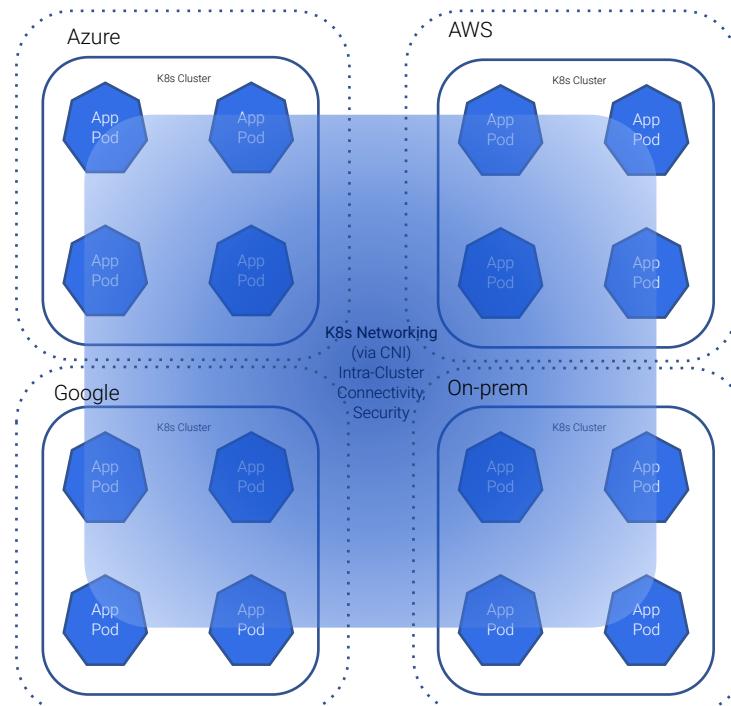
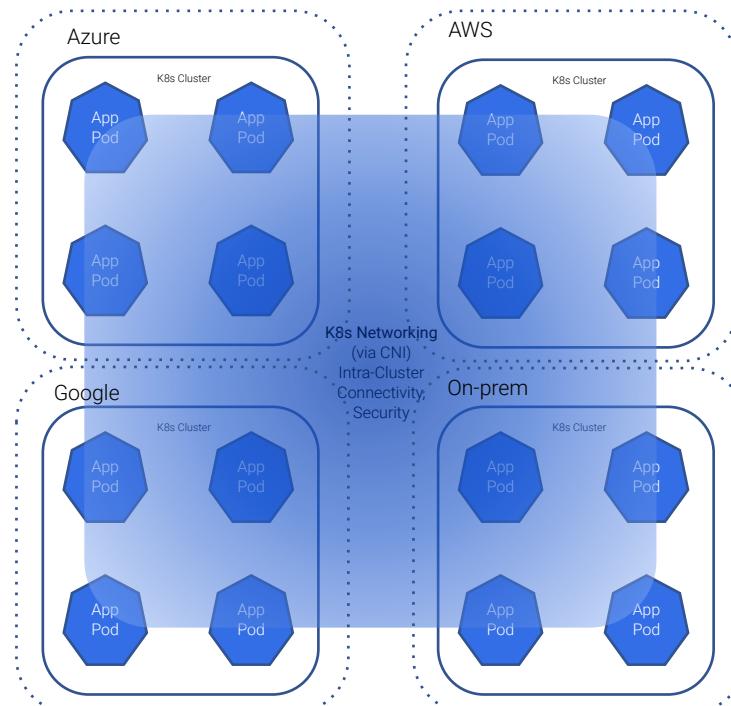
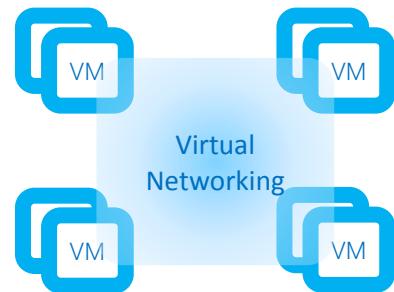
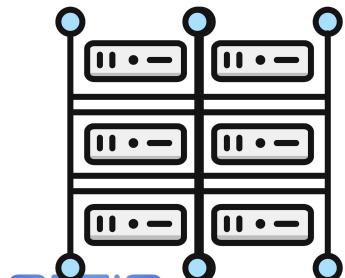




The Federation non-Solution

Attempt to 'Federate' multiple clusters:

- Also incompatible with non-K8s runtime domains
 - Semantics of VM domain different than K8s
 - Semantics of on-prem server networking different than K8s





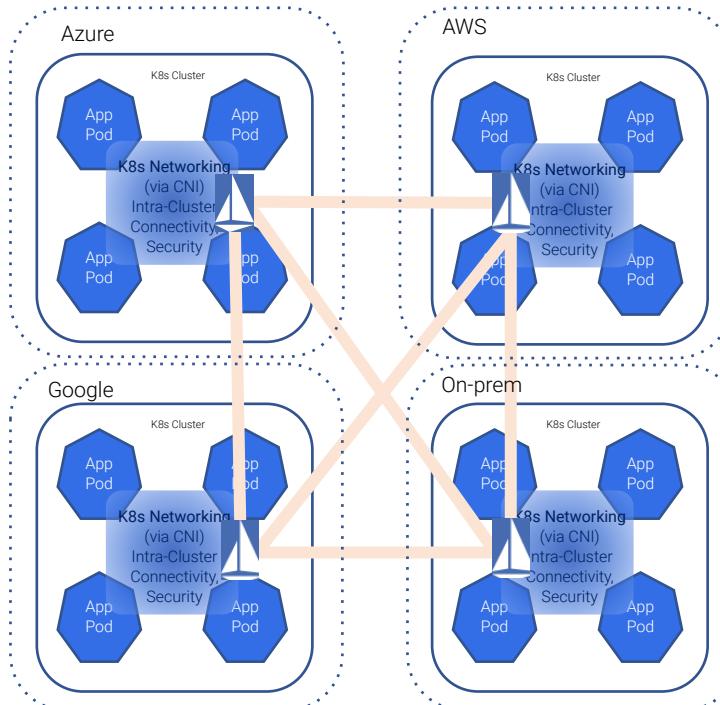
The Service Mesh Intercluster Gateway non-Solution

Solving with Service Mesh(Istio etc)

Inter-cluster gateways

Problems:

- Only works for L7, not L3
- Same full mesh combinatorics problems
- Same complex often manual cluster to cluster link setup between different public/private cloud providers (possibly involving complex firewall rules depending on how its done).

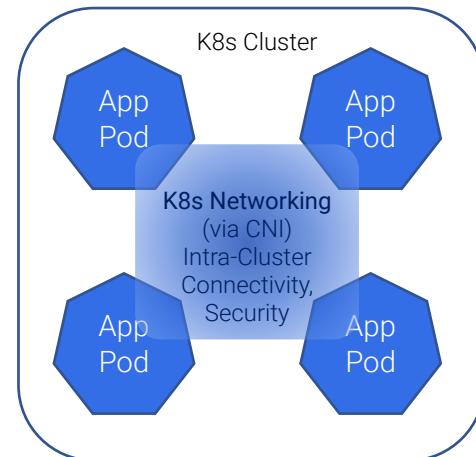




The NSM Realization

'Connectivity Domain' Independence:

- Welding your 'connectivity domain' to your 'runtime domain' (cluster) is mistake



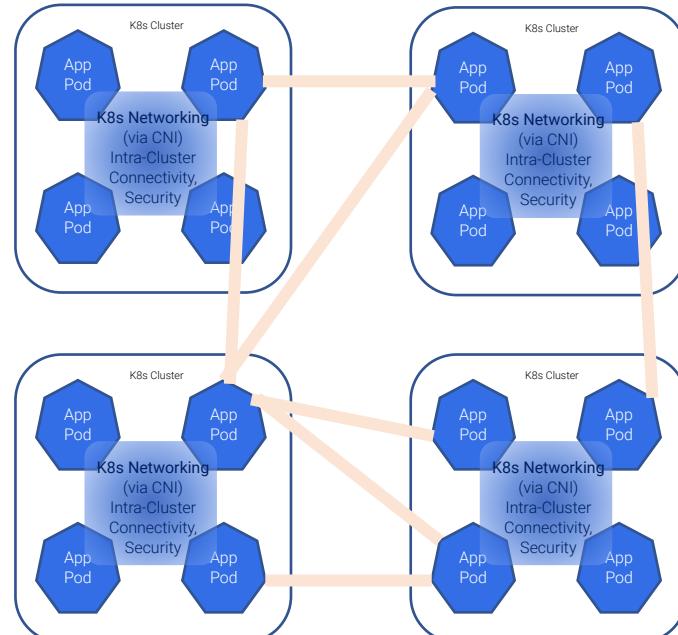


The NSM Realizations

'Connectivity Domain' Independence:

- Welding your 'connectivity domain' to your 'runtime domain' (cluster) is mistake

**What you really care about is
workload to workload connectivity:
independent of runtime domain.**

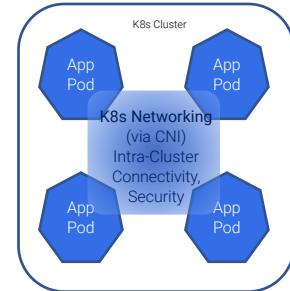
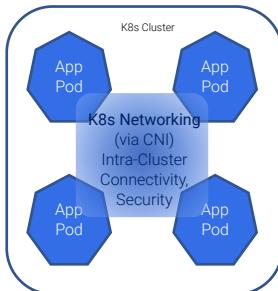
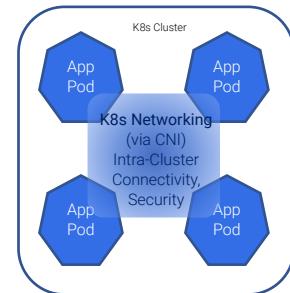
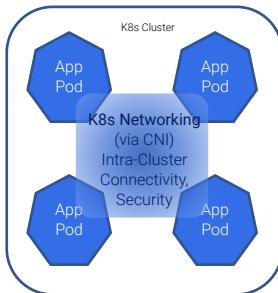




The NSM Solution

Leave Intra-Cluster Networking Alone:

- Orthogonal to CNI
- Harmless to existing K8s Networking





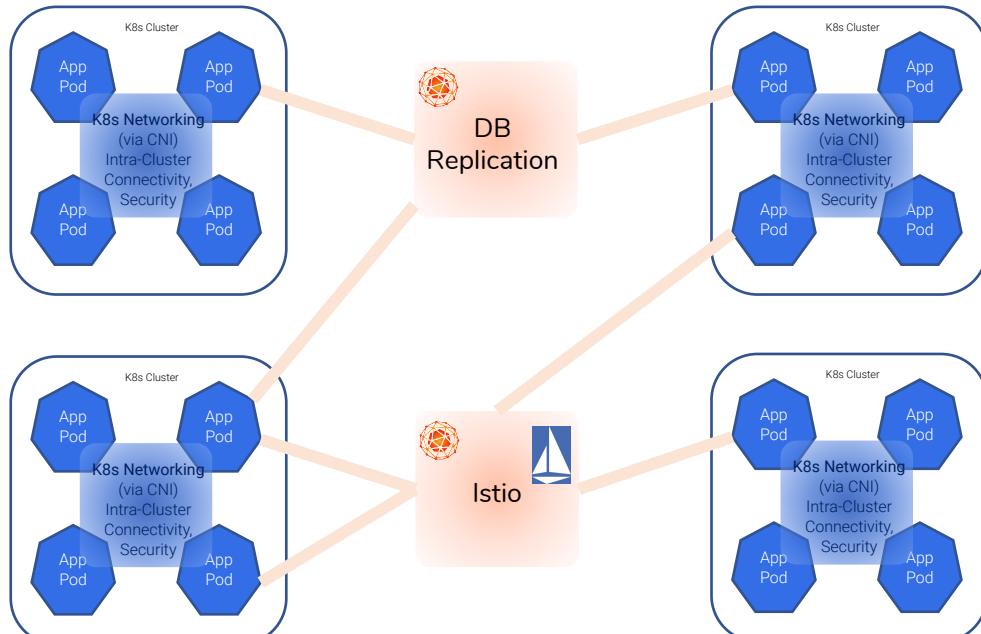
The NSM Solution

Leave Intra-Cluster Networking Alone:

- Orthogonal to CNI
- Harmless to existing K8s Networking

Allow workloads to connect to new
'connectivity domains':

- With the
Connectivity/Security/Observability
features needed in that connectivity
domain





The NSM Solution

Leave IntraCluster Networking Alone:

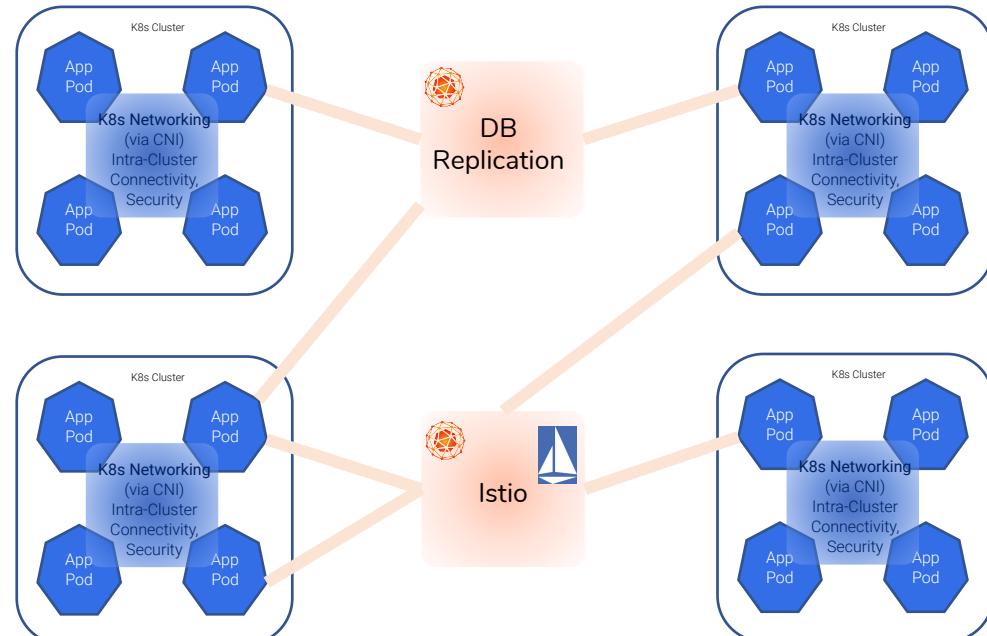
- Orthogonal to CNI
- Harmless to existing K8s Networking

Allow workloads to connect to new
'connectivity domains':

- With the
Connectivity/Security/Observability
features needed in that connectivity
domain

Examples:

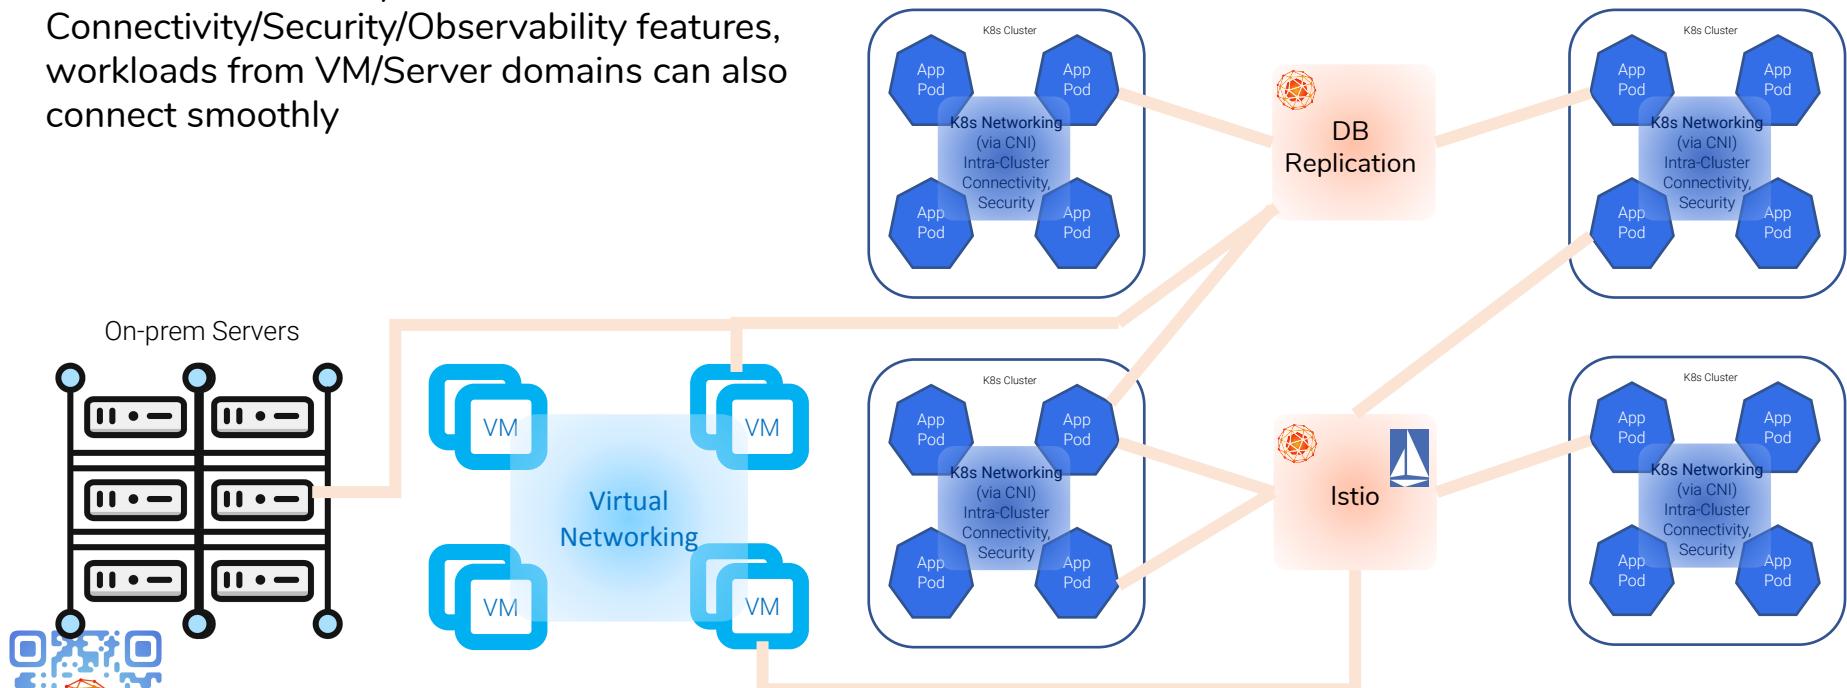
- **DB Replication Connectivity Domain**
 - Pure vL3 domain between DB replicas: where-ever they may be
- **Istio Connectivity Domain**
 - Single Istio instance serving workloads wherever they may be over vL3 domain.





The NSM Solution

Because connectivity domain has its own Connectivity/Security/Observability features, workloads from VM/Server domains can also connect smoothly





Network Service Mesh

State of the NSM





State of the NSM - Community

- CNCF Project: NSM is now a CNCF project



**CLOUD NATIVE
COMPUTING FOUNDATION**





State of the NSM - Early Use

- **CNF Testbed:** NSM is used in the CNF Testbed project for Cloud Native NFV for Telco.



CNF Testbed





State of the NSM

- CNCF Project
- **Multi-cloud CI:** NSM runs CI on AKS/EKS/GKE/Vanilla K8s (100+ tests each)

```
Elapsed total: 1h22m36.004136867s
Tests time: 1h17m50.120297751s
Tasks Completed: 449
Remaining: 11.037871128s (1).

Running:
    Example-helm-icmp on cluster azure-2 elapsed: 1m59.043478331s

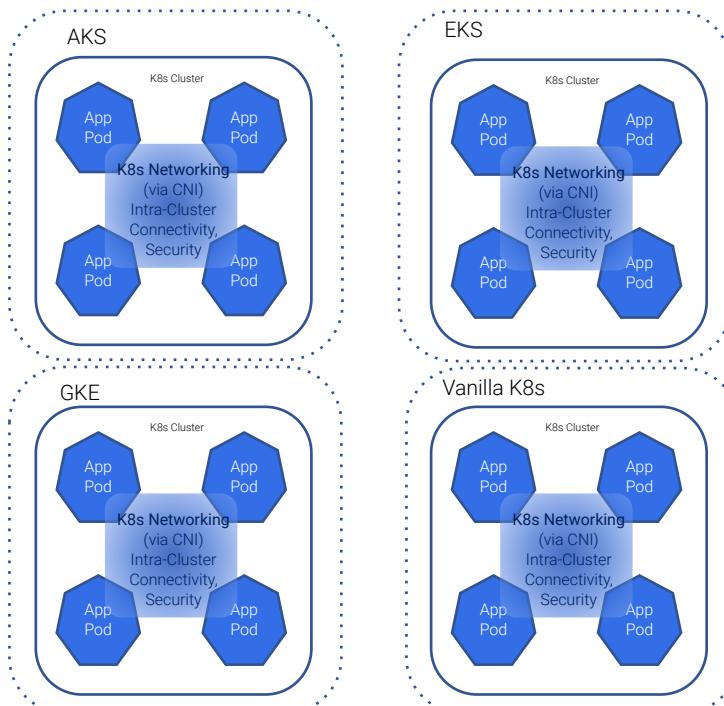
Clusters:
Cluster: packet Tasks left: 0
packet-1 shutdown uptime: 1h16m41.78334579s
packet-2 shutdown uptime: 1h16m33.555938821s

Cluster: gke Tasks left: 0
gke-1 shutdown uptime: 1h17m50.120404233s
gke-2 shutdown uptime: 1h17m46.403253797s

Cluster: aws Tasks left: 0
aws-1 shutdown uptime: 1h8m48.426188875s
aws-2 shutdown uptime: 1h7m59.154977848s

Cluster: azure Tasks left: 1
azure-1 ready uptime: 1h10m2.826986099s
azure-2 running test uptime: 1h12m15.346359264s

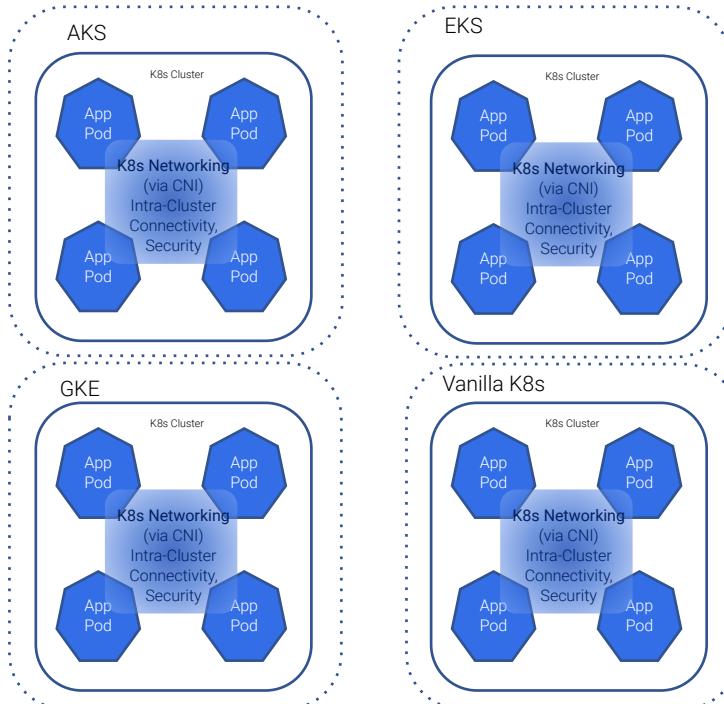
Status Passed: 449
Status Failed: 0
Status Timeout: 0
Status Skipped: 0
```





State of the NSM

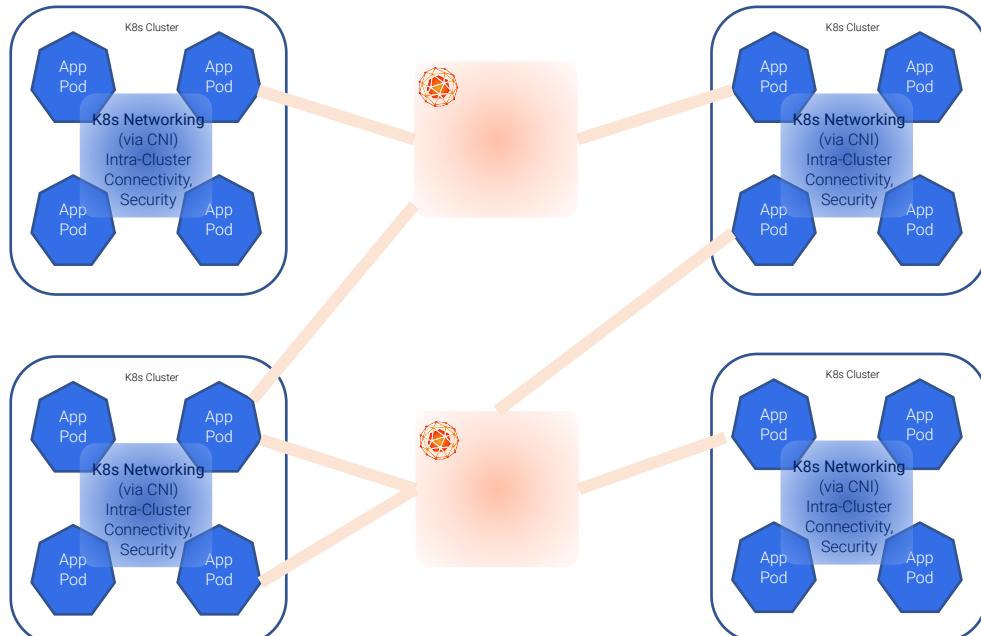
- CNCF Project
- Multi-cloud CI
- **Resiliencyv1 (AutoHealing)**: Can auto heal connections between Pods and Network Services if various system elements restart or NSE providing Network Service dies without disturbing client Pod.





State of the NSM

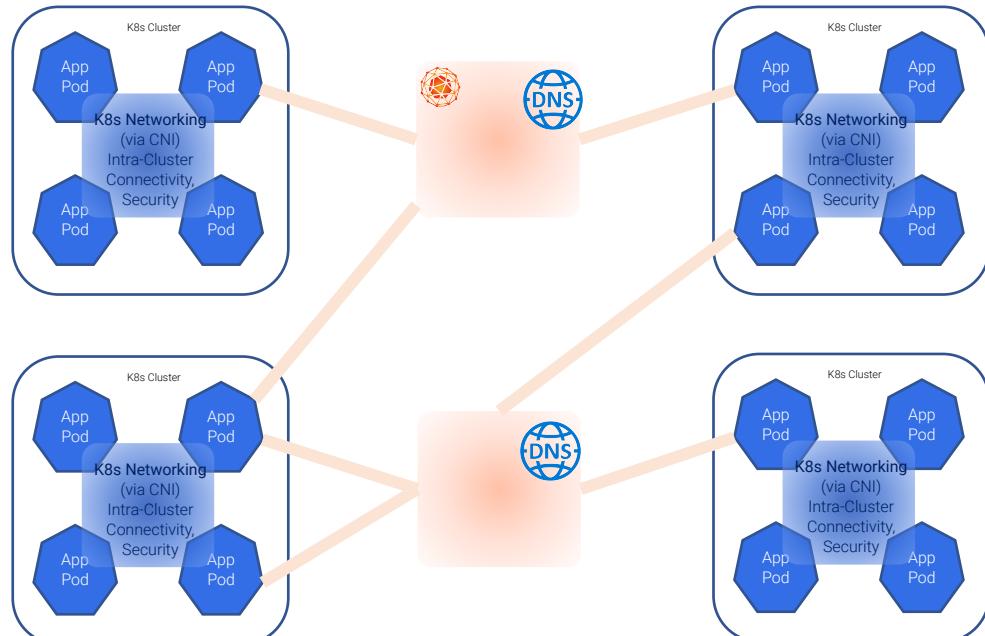
- CNCF Project
- Multi-cloud CI
- Resiliencyv1 (AutoHealing)
- **Inter-domain:** Initial Inter-domain support merged





State of the NSM

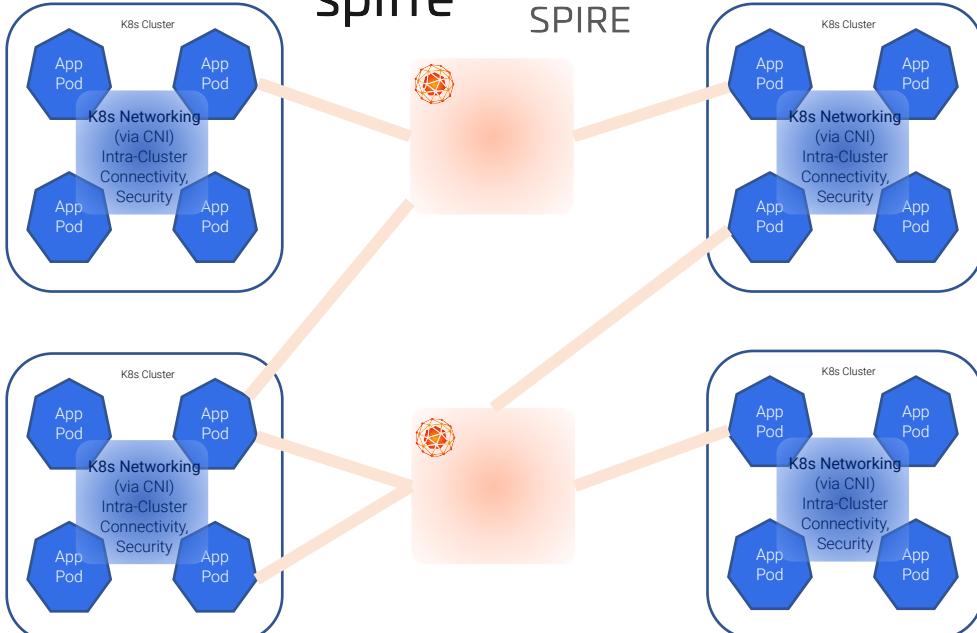
- CNCF Project
- Multi-cloud CI
- Resiliencyv1 (AutoHealing)
- Inter-domain
- **DNS:** Each Network Service (Connectivity Domain) can provide DNS to workload additivitly (ie: without breaking K8s DNS).





State of the NSM

- CNCF Project
- Multi-cloud CI
- Resiliencyv1 (AutoHealing)
- Inter-domain
- DNS
- **Security:** Spiffe/Spire based security – initial work done.





Network Service Mesh

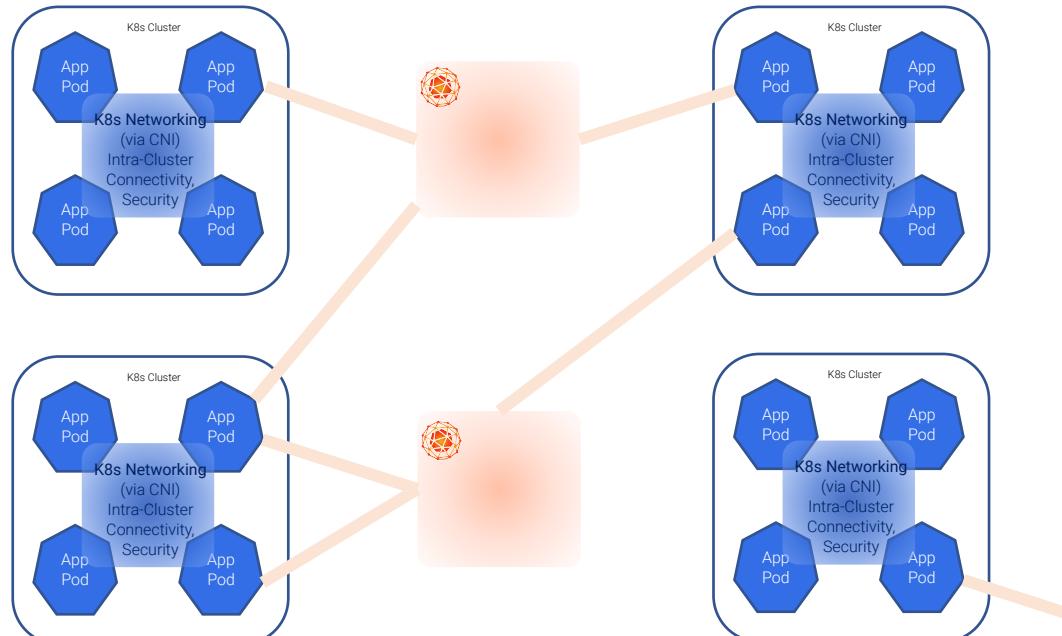
NSM Future





Resiliencyv2

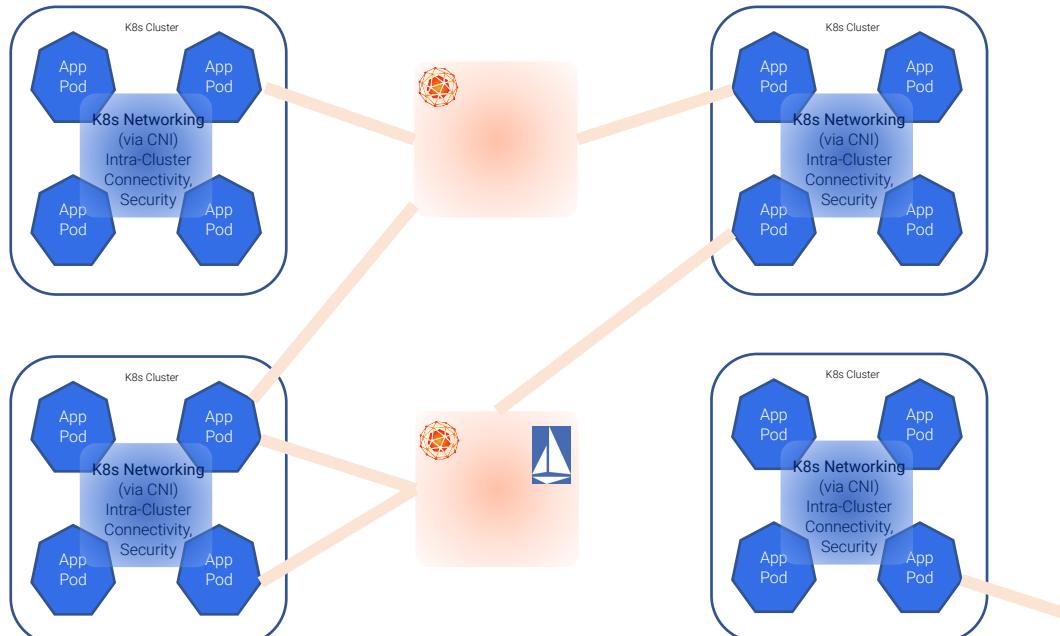
Autoheal connections from Client
Pods to Network Service
Endpoints (NSEs) they are
connected to even if *all*
non-client elements of the
system restart simultaneously
and NSE dies without impacting
Client.





Istio on NSM

Run an Istio domain over an
NSM Network Service



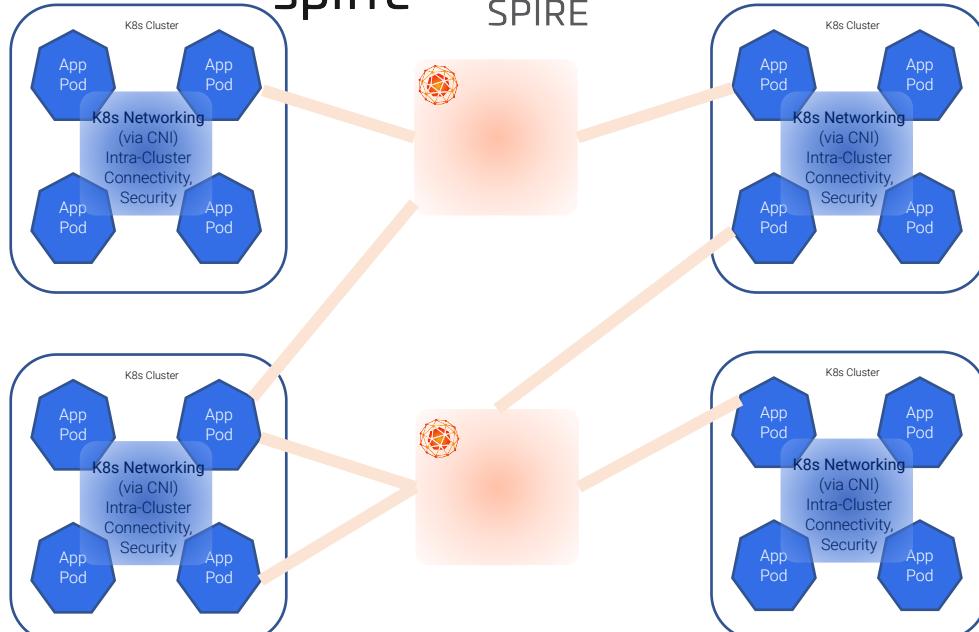


Open Policy Agent: Authz

Using OPA to allow the Network Service Mesh to enforce admissions policy based on Spiffe/Spire identities



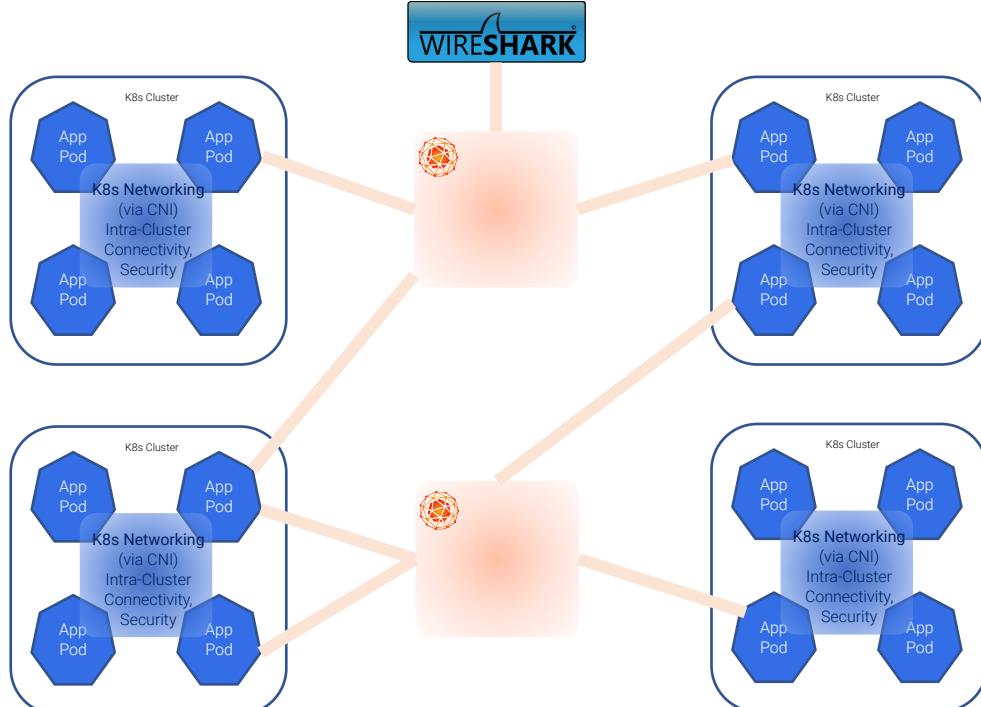
Open Policy Agent





Packet Capture Observability

Make it simple for Network Services (Connectivity Domains) to allow developers to securely get packet capture observability at per workload granularity.



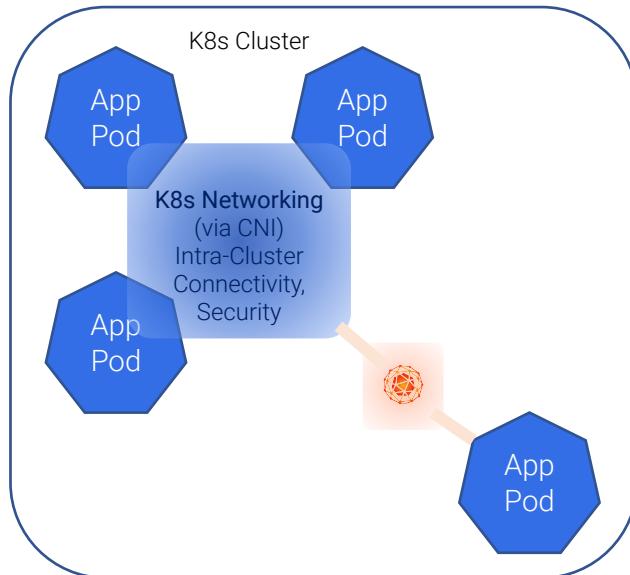


CNI Intercept

Allow (safe) insertion of Network Service between Pod and its CNI interface.

Would allow adding features to IntraCluster Networking with any CNI

Could be used for inserting Envoy Sidecar for Istio via NSM.



NSMCon @Kubecon NA



First [NSMCon](#) Nov 18 @KubeCon NA in San Diego



NSMCon





Network Service Mesh

How the Magic Works

Network Service Registry Domain

Network Service Registry

Registry of:

- NetworkServices
- NetworkServiceEndpoints
- NetworkServiceManagers
 - (more later on this)

Network Service Registry Domain

Network Service Registry

Network
Service
Manager
(NSMgr)

...

Network
Service
Manager
(NSMgr)

Network Service Registry Domain

Network Service Registry

Network Service Manager Domain

Network Service Client (NSC)

⋮

Network Service Client (NSC)

Network Service Manager (NSMgr)

Network Service Endpoint (NSE)

⋮

Network Service Endpoint (NSE)

NSM Forwarder

Network Service Manager Domain

Network Service Client (NSC)

⋮

Network Service Client (NSC)

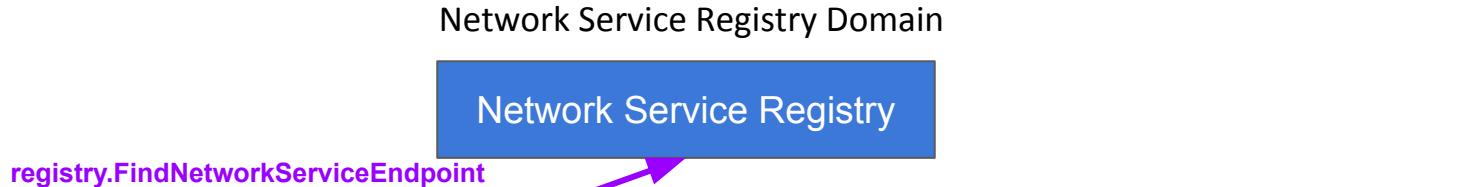
Network Service Manager (NSMgr)

Network Service Endpoint (NSE)

⋮

Network Service Endpoint (NSE)

NSM Forwarder



Network Service Manager Domain

Network Service Client (NSC)

Network Service Manager (NSMgr)

Network Service Client (NSC)

Network Service Endpoint (NSE)

Network Service Endpoint (NSE)

NSM Forwarder

Network Service Manager Domain

Network Service Client (NSC)

Network Service Manager (NSMgr)

Network Service Client (NSC)

Network Service Endpoint (NSE)

Network Service Endpoint (NSE)

NSM Forwarder

Network Service Registry Domain

Network Service Registry

Network Service Manager Domain

Network Service Client (NSC)

Network Service Manager (NSMgr)

Network Service Client (NSC)

Network Service Endpoint (NSE)

Network Service Endpoint (NSE)

remote.NetworkService.Request

NSM Forwarder

Network Service Manager Domain

Network Service Client (NSC)

Network Service Manager (NSMgr)

Network Service Endpoint (NSE)

Network Service Endpoint (NSE)

⋮

Network Service Endpoint (NSE)

Network Service Endpoint (NSE)

⋮

NSM Forwarder

Network Service Registry Domain

Network Service Registry

Network Service Manager Domain

Network Service Client (NSC)

Network Service Manager (NSMgr)

Network Service Client (NSC)

Network Service Endpoint (NSE)

Network Service Endpoint (NSE)

NSM Forwarder

vWire (L2/L3 connection (tunnel))

Network Service Manager Domain

Network Service Client (NSC)

Network Service Manager (NSMgr)

Network Service Client (NSC)

Network Service Endpoint (NSE)

Network Service Endpoint (NSE)

NSM Forwarder



Network Service Mesh

Interdomain



Examples of ‘Domains’

Public Cloud1

K8s Cluster1

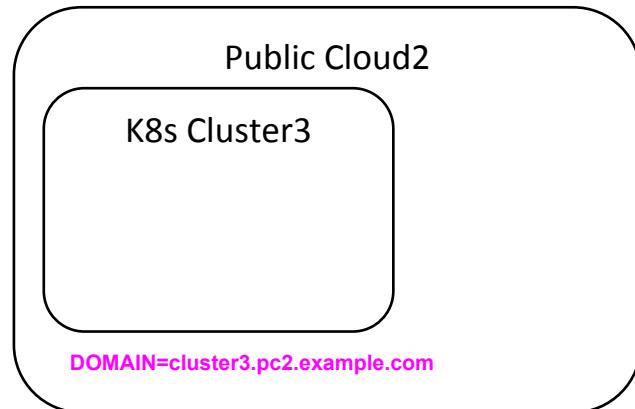
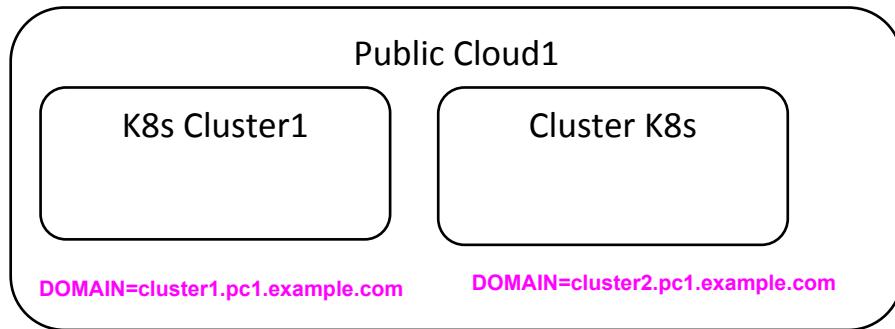
Cluster K8s

DOMAIN=cluster1.pc1.example.com

DOMAIN=cluster2.pc1.example.com

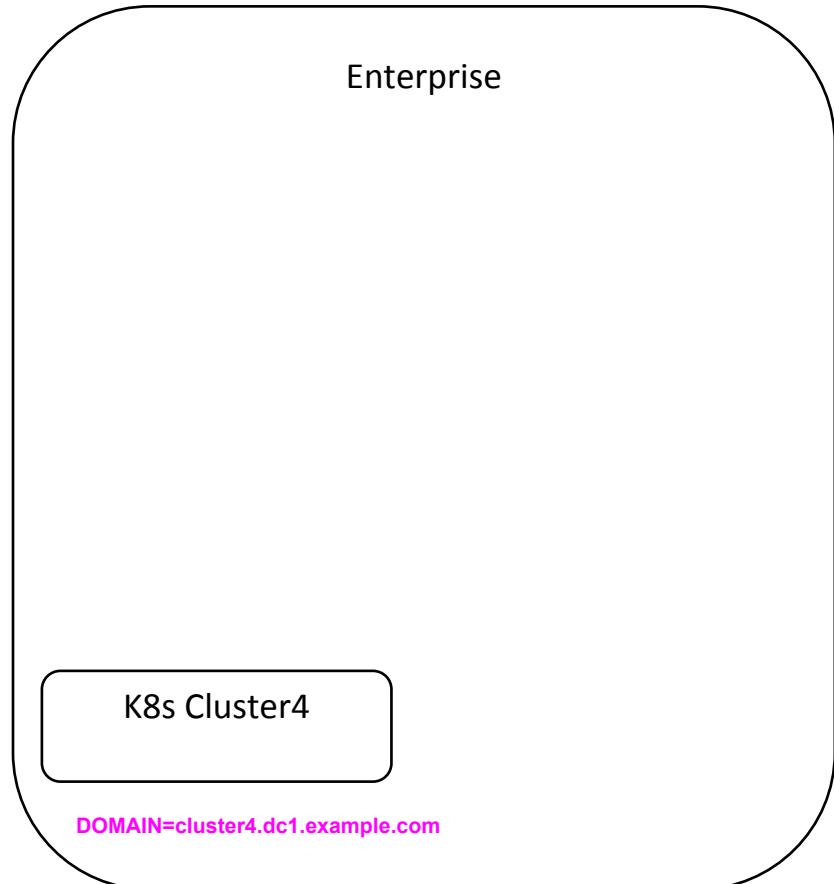
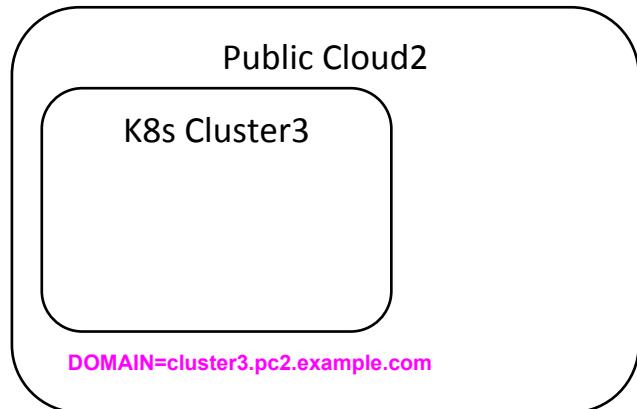
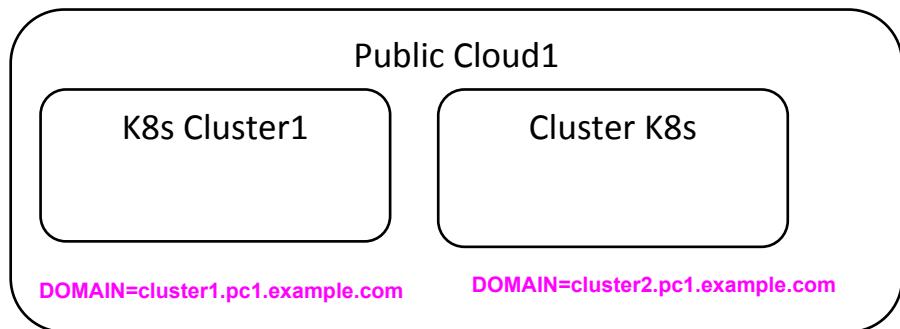


Examples of ‘Domains’



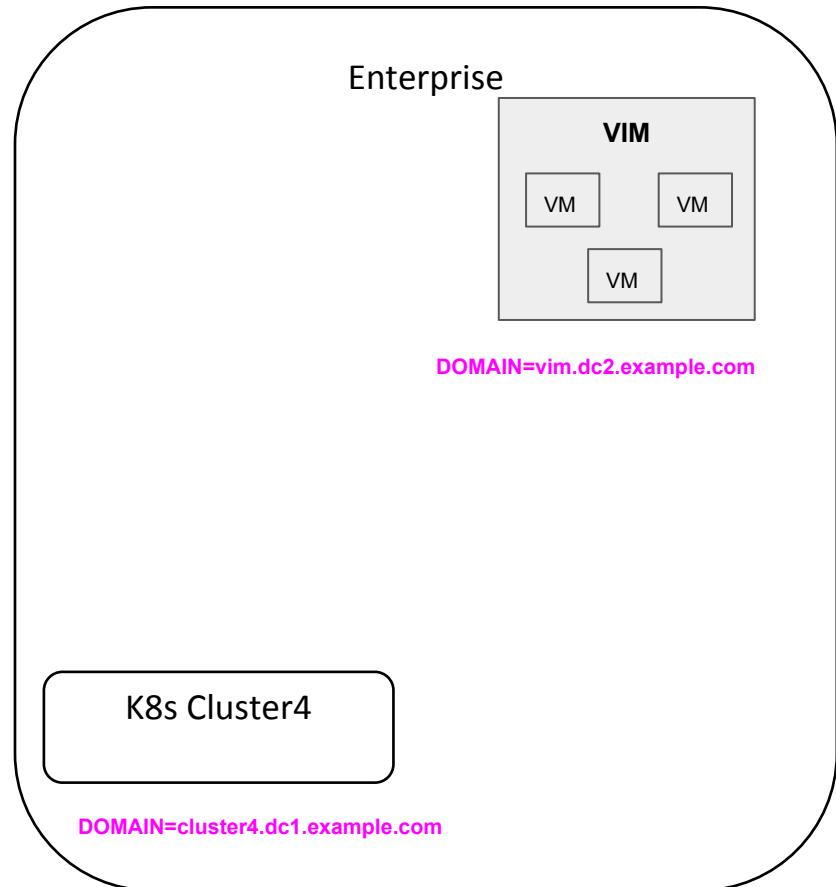
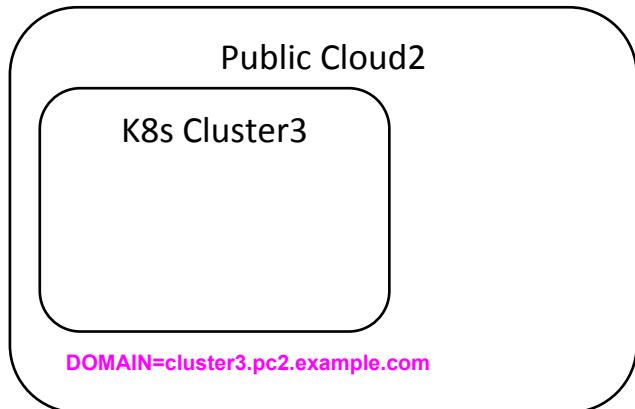
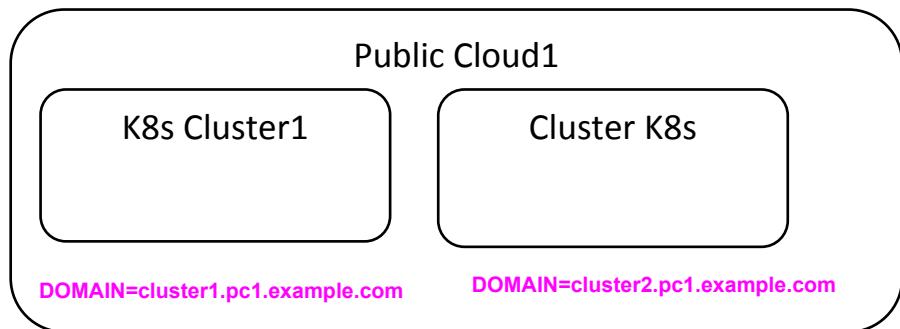


Examples of 'Domains'



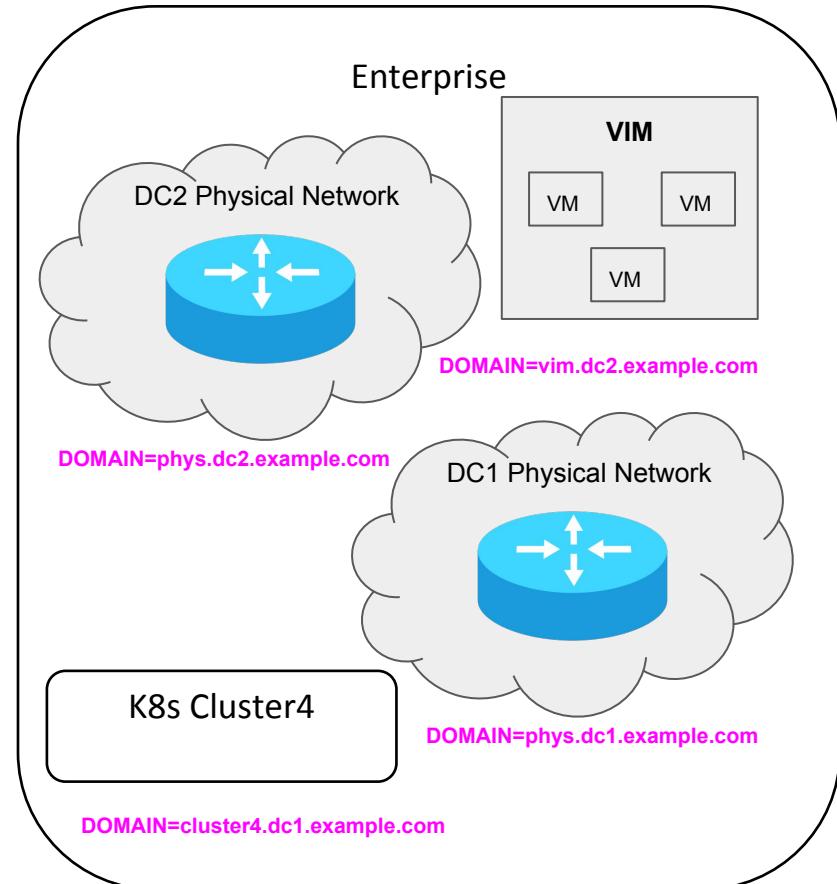
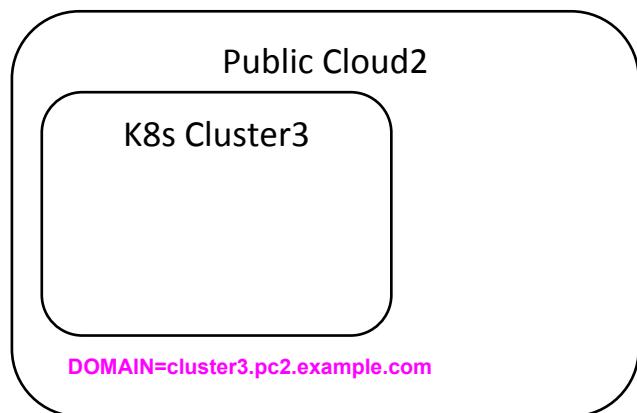
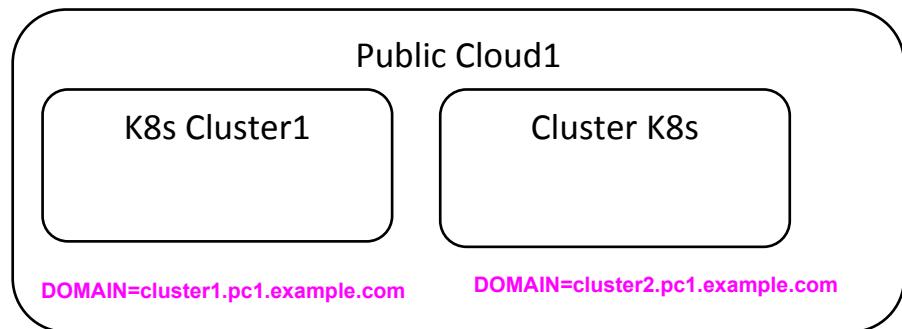


Examples of 'Domains'





Examples of ‘Domains’



Kubernetes Cluster

Kubernetes API Server
(Network Service Registry
via CRDs)

Node(Network Service Manager Domain)

Network Service
Client (NSE)
(Pod)

⋮
Network Service
Endpoint (NSE)
(Pod)

Network
Service
Manager
(NSMgr)
(Daemonset)

NSM Forwarder
(kernel/vswitch)

Network Service Registry Domain (example.com)

Network Service Registry

Network Service Manager Domain

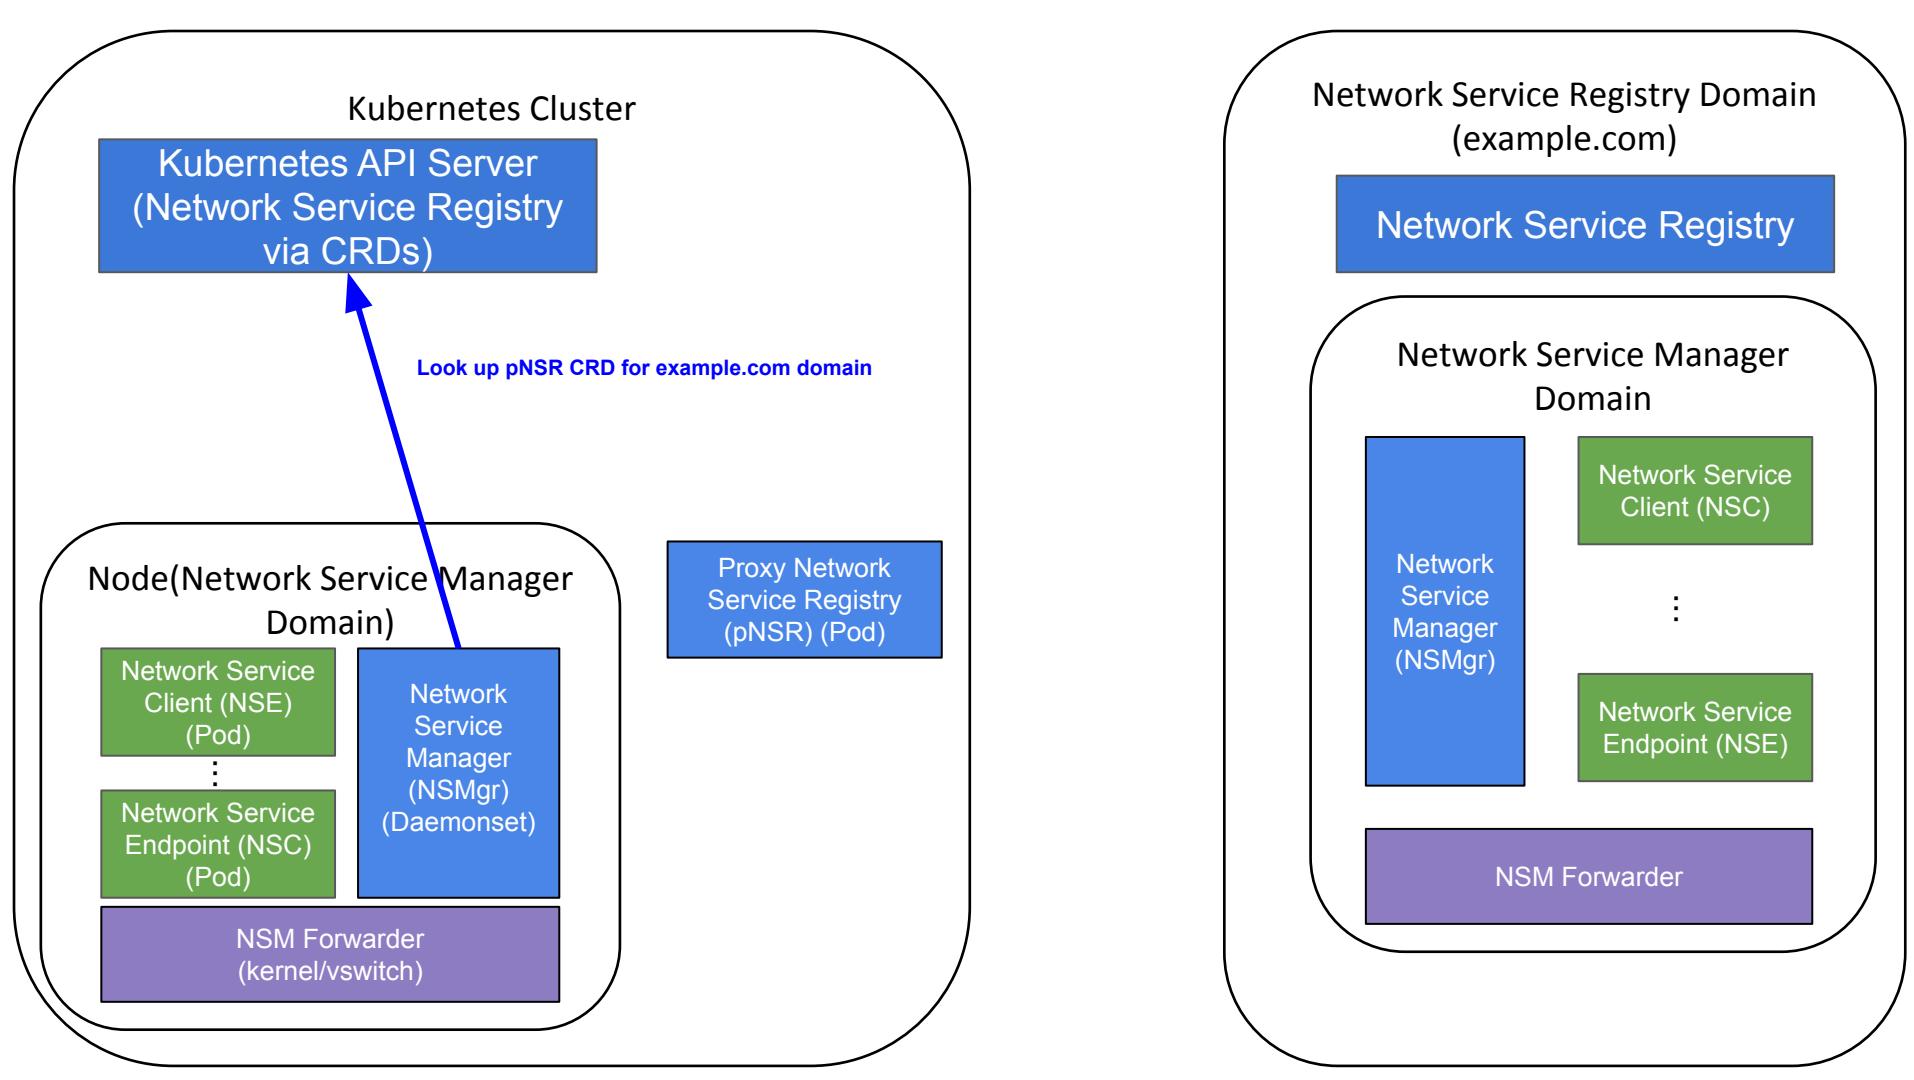
Network
Service
Manager
(NSMgr)

Network Service
Client (NSC)

⋮

Network Service
Endpoint (NSE)

NSM Forwarder



Kubernetes Cluster

Kubernetes API Server
(Network Service Registry
via CRDs)

Node(Network Service Manager Domain)

Network Service Client (NSE) (Pod)

⋮

Network Service Endpoint (NSE) (Pod)

Network Service Manager (NSMgr) (Daemonset)

NSM Forwarder (kernel/vswitch)

Proxy Network Service Registry (pNSR) (Pod)

`registry.FindNetworkServiceEndpoint`

Network Service Registry Domain (example.com)

Network Service Registry

Network Service Manager Domain

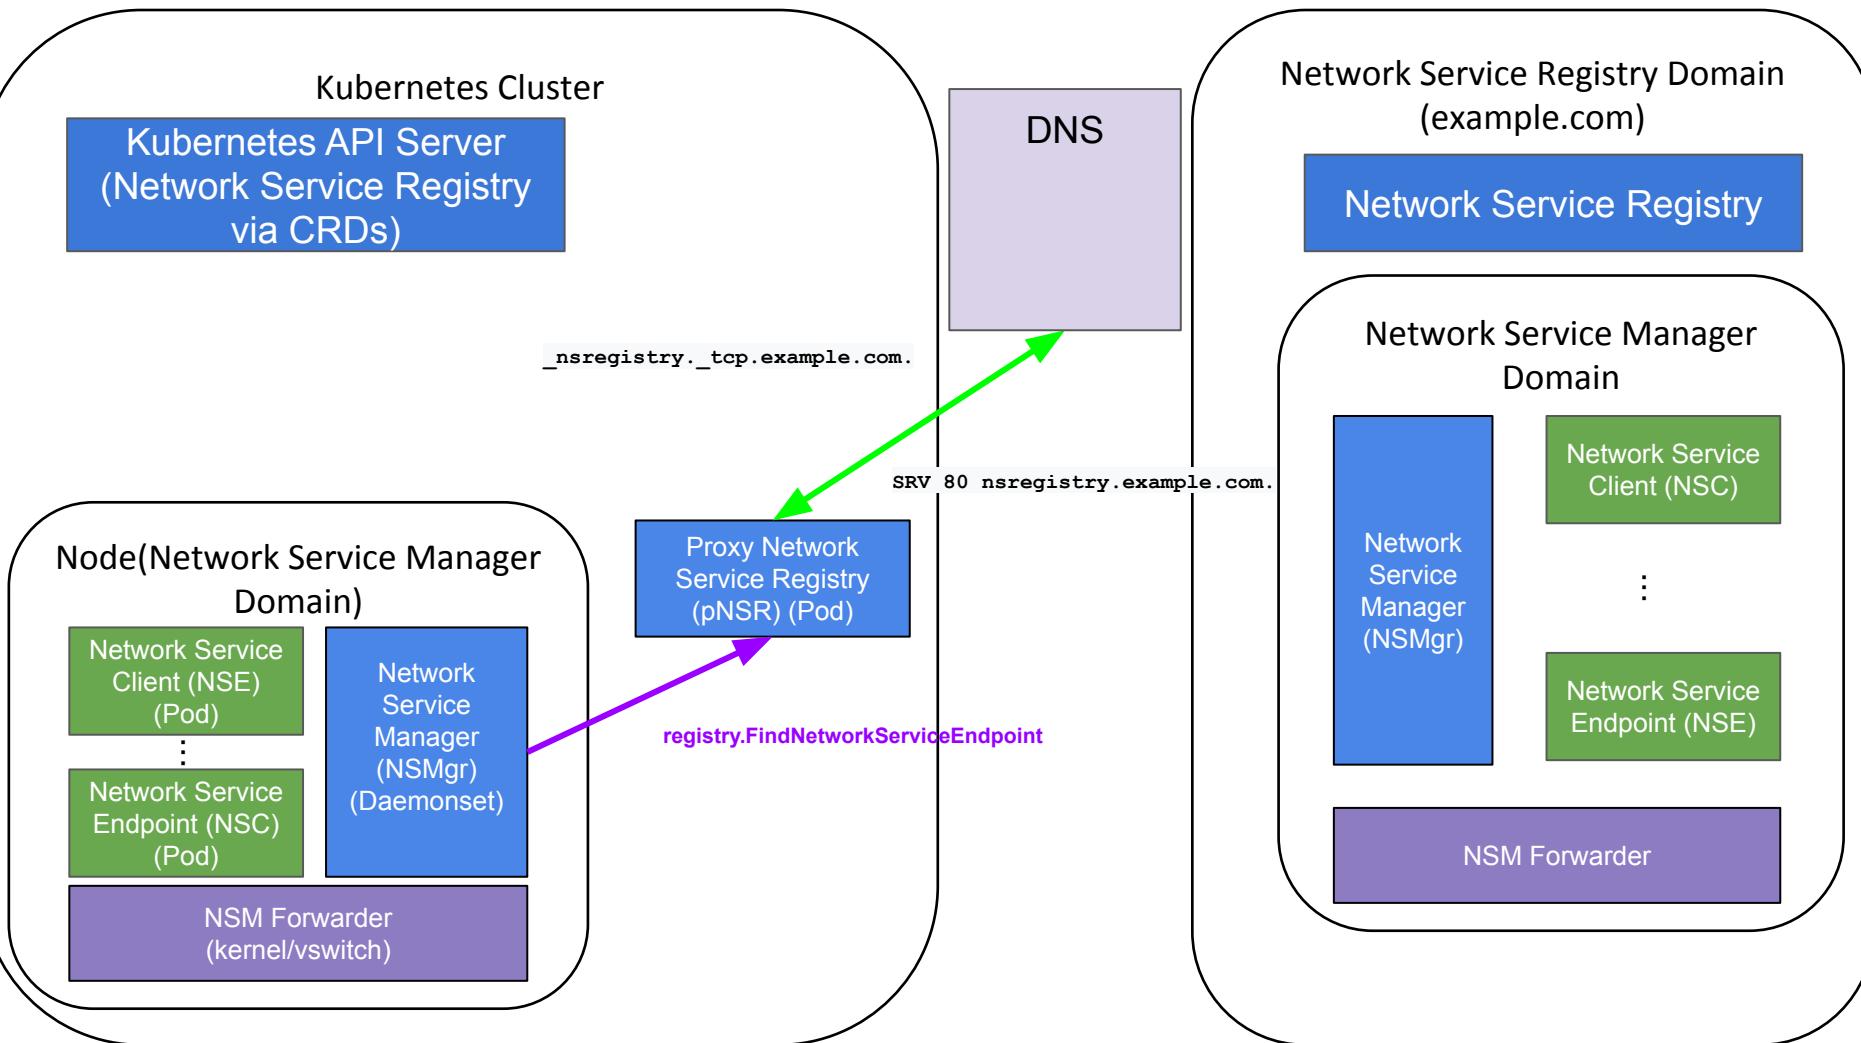
Network Service Manager (NSMgr)

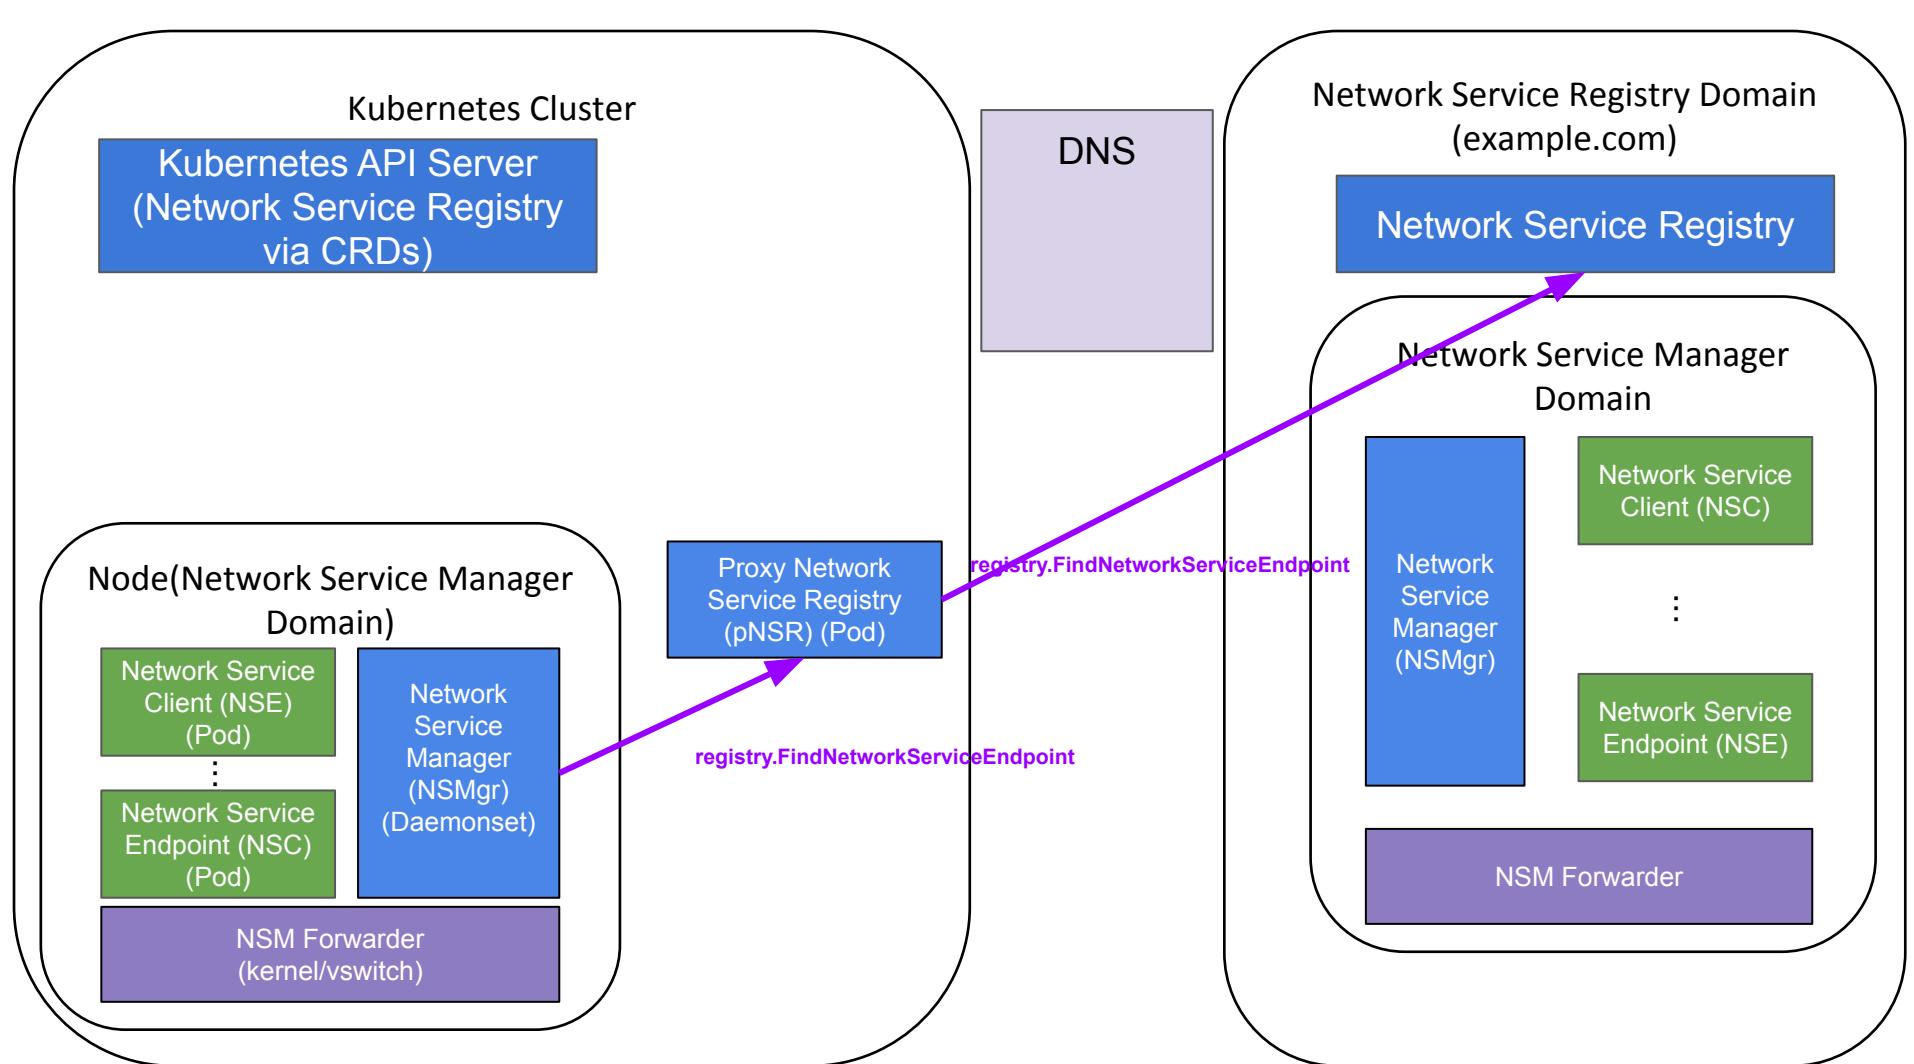
Network Service Client (NSC)

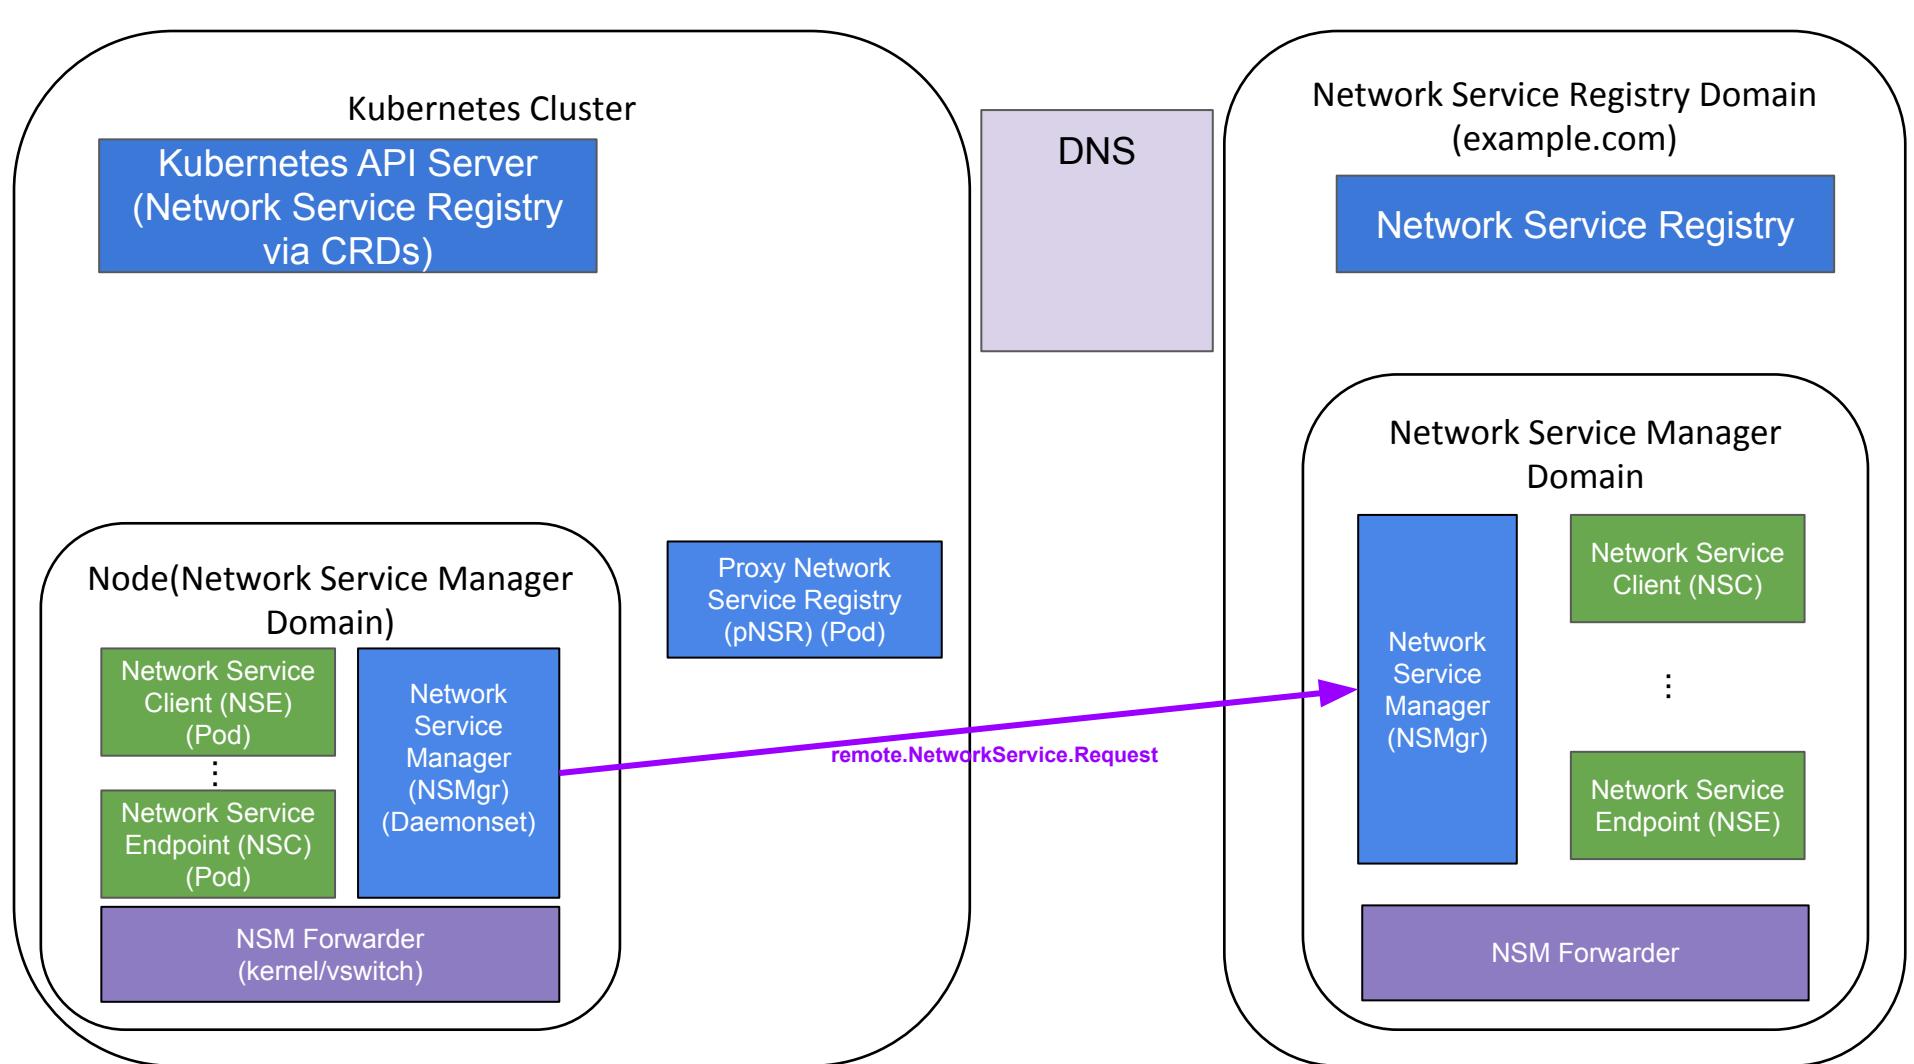
⋮

Network Service Endpoint (NSE)

NSM Forwarder







Kubernetes Cluster

Kubernetes API Server
(Network Service Registry
via CRDs)

Node(Network Service Manager Domain)

Network Service Client (NSE)
(Pod)

⋮
Network Service Endpoint (NSE)
(Pod)

Network Service Manager (NSMgr)
(Daemonset)

NSM Forwarder

Proxy Network Service Registry (pNSR) (Pod)

DNS

Network Service Registry Domain (example.com)

Network Service Registry

Network Service Manager Domain

Network Service Manager (NSMgr)

Network Service Client (NSC)

⋮

Network Service Endpoint (NSE)

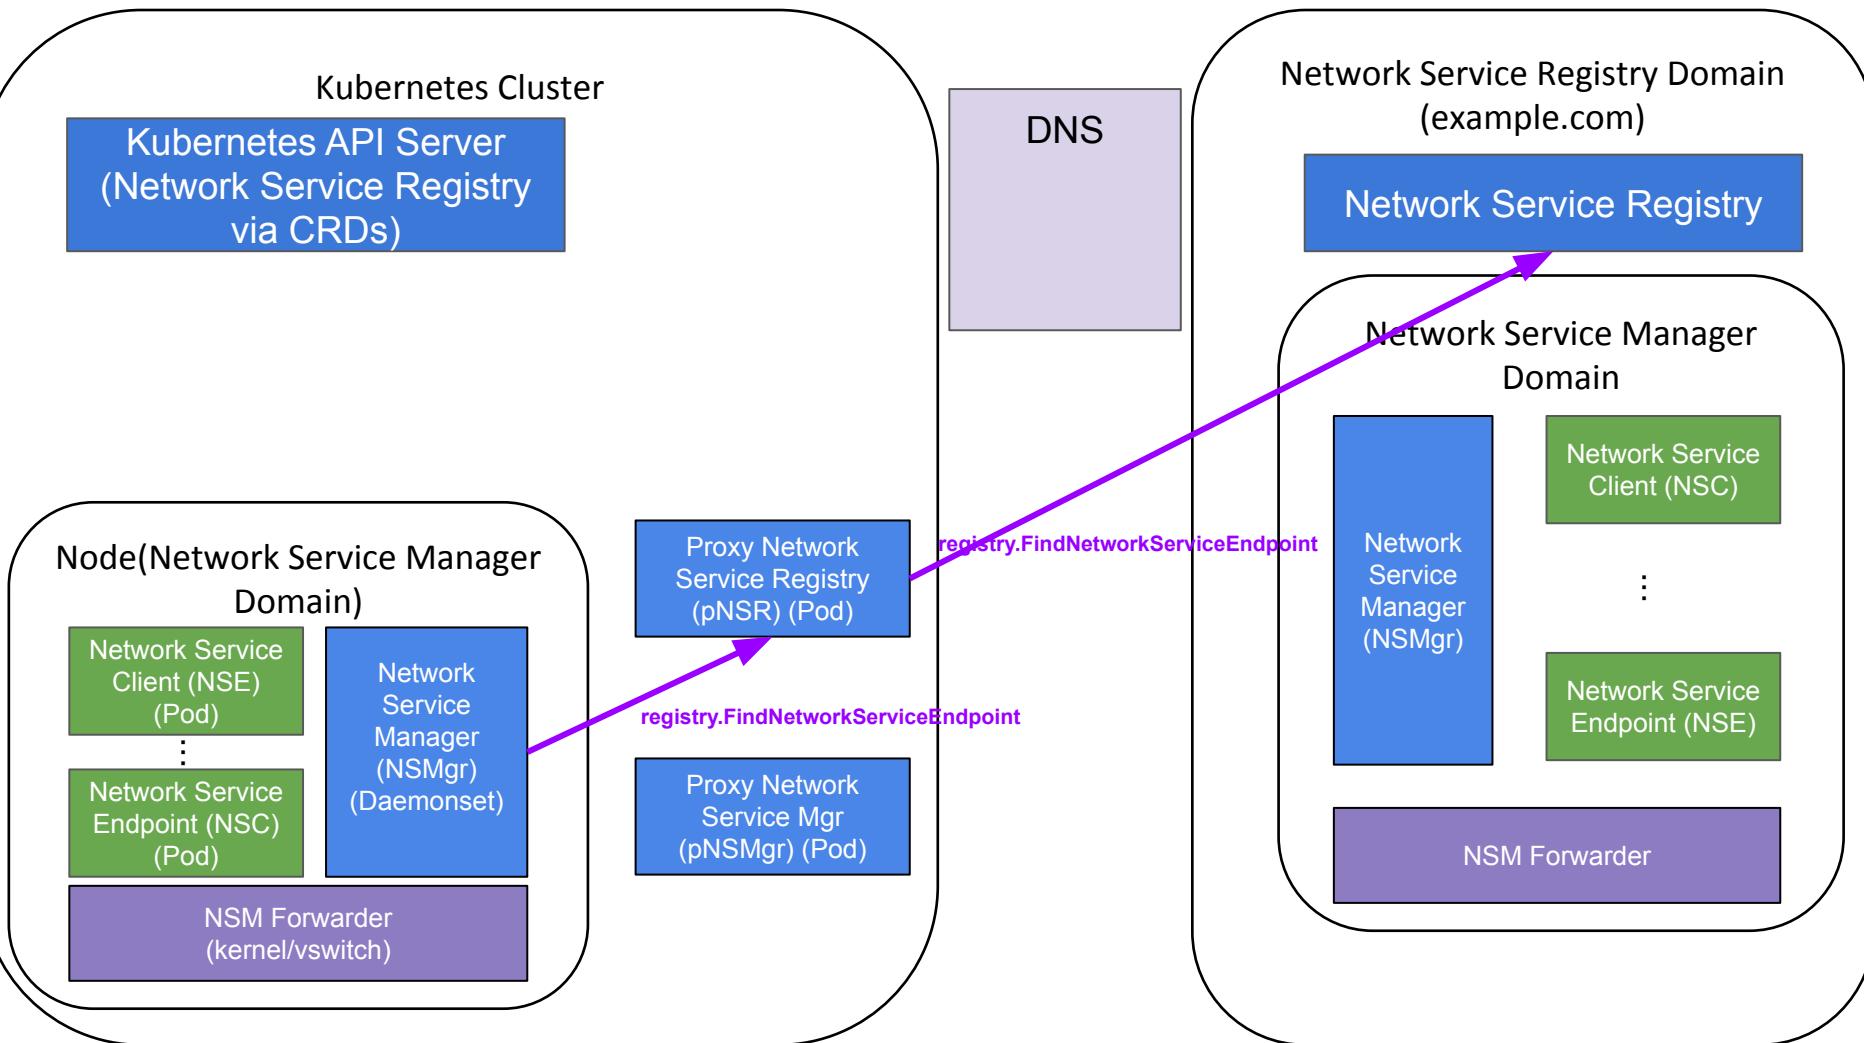
NSM Forwarder

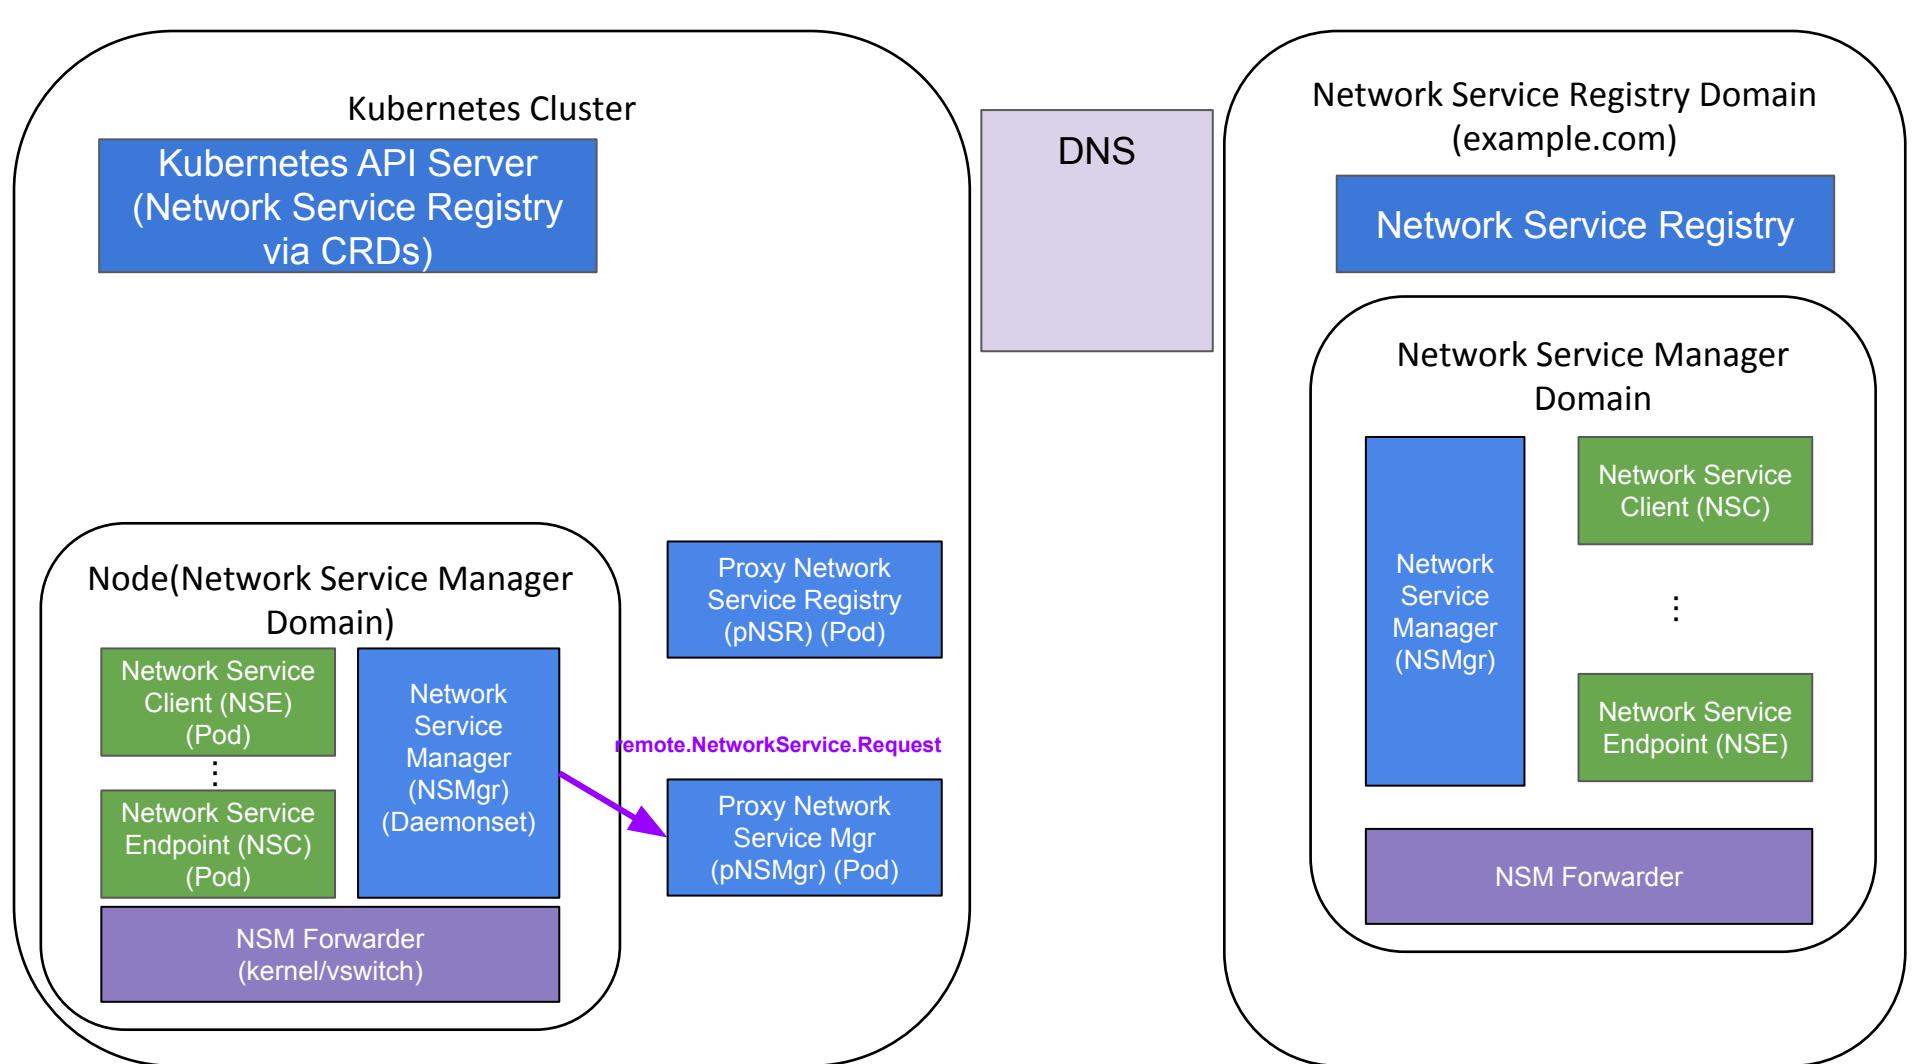
vWire (aka vWire (L2/L3 connection (tunnel)))

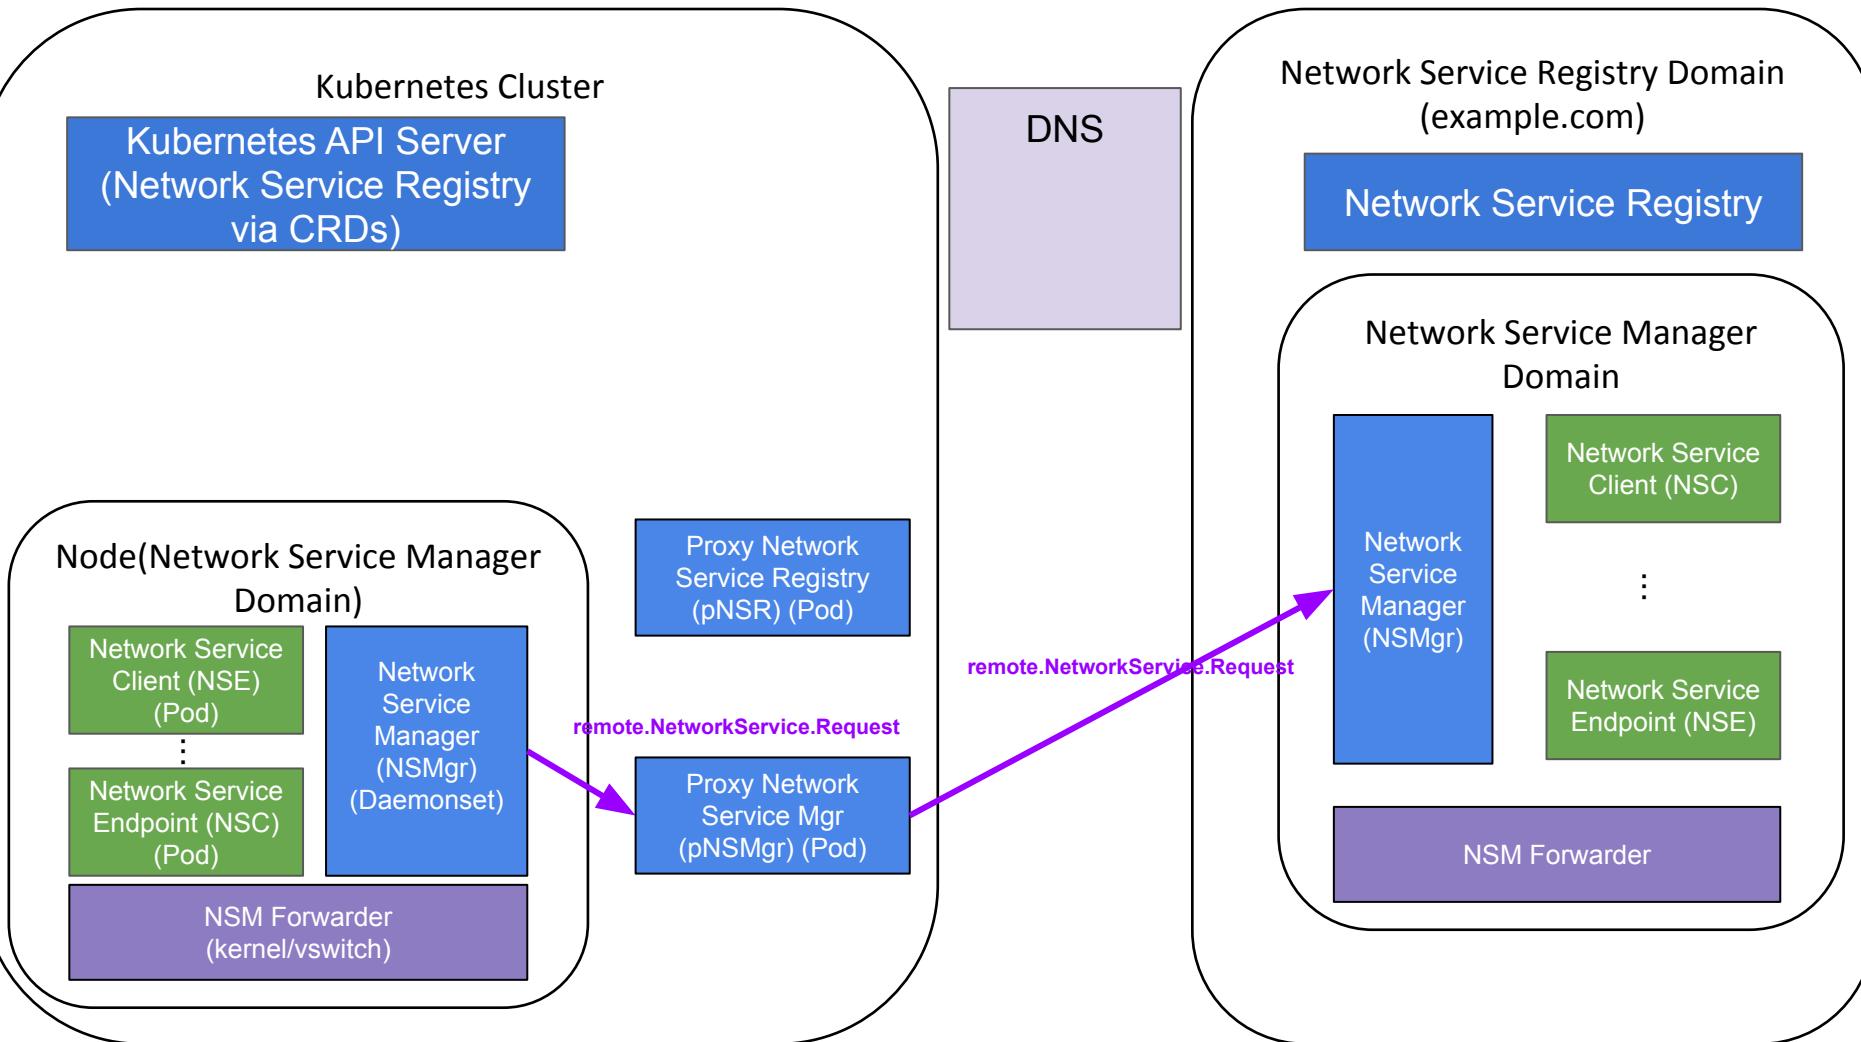


Network Service Mesh

Interdomain w/pNSMgr







Kubernetes Cluster

Kubernetes API Server
(Network Service Registry
via CRDs)

Node(Network Service Manager Domain)

Network Service Client (NSE)
(Pod)

⋮
Network Service Endpoint (NSE)
(Pod)

Network Service Manager (NSMgr)
(Daemonset)

NSM Forwarder

Proxy Network Service Registry
(pNSR) (Pod)

Proxy Network Service Mgr
(pNSMgr) (Pod)

DNS

Network Service Registry Domain
(example.com)

Network Service Registry

Network Service Manager Domain

Network Service Manager (NSMgr)

⋮
Network Service Client (NSC)

Network Service Endpoint (NSE)

NSM Forwarder

vWire (L2/L3 connection (tunnel))



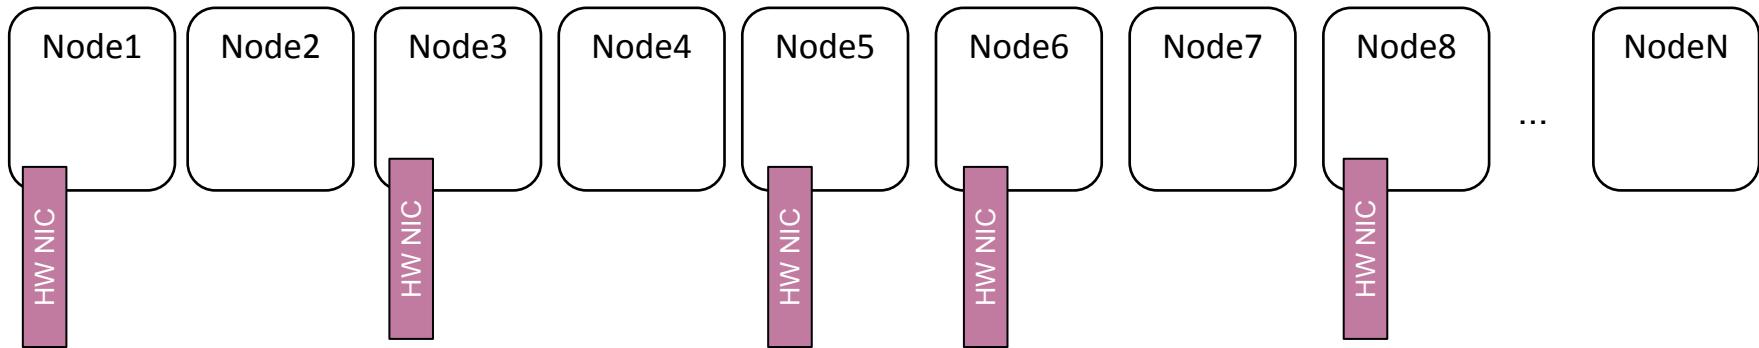
Network Service Mesh

HW NICs



The Problem: NICs

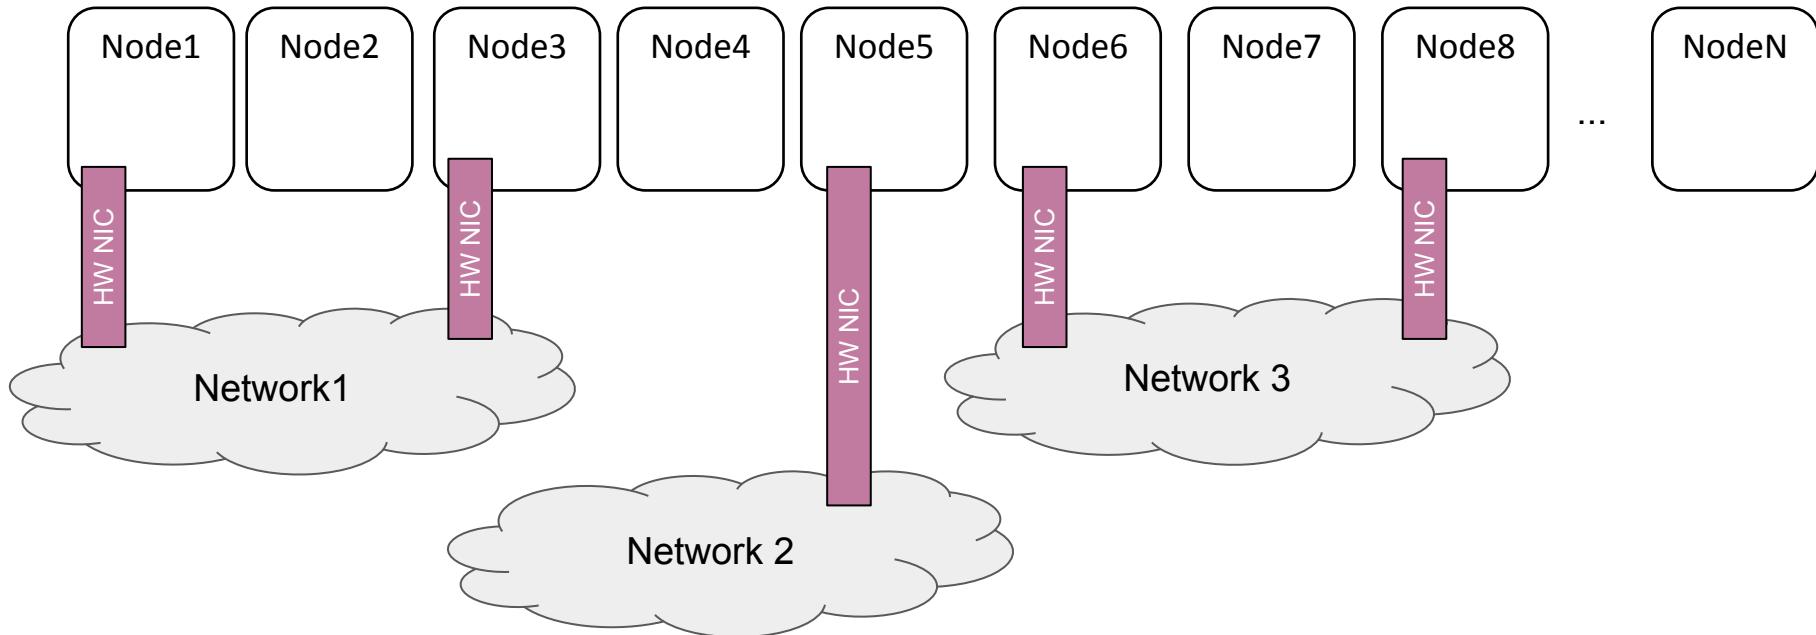
A Kubernetes Cluster may have special NICs in some but not all Nodes:



Not all NICs are on the same Network



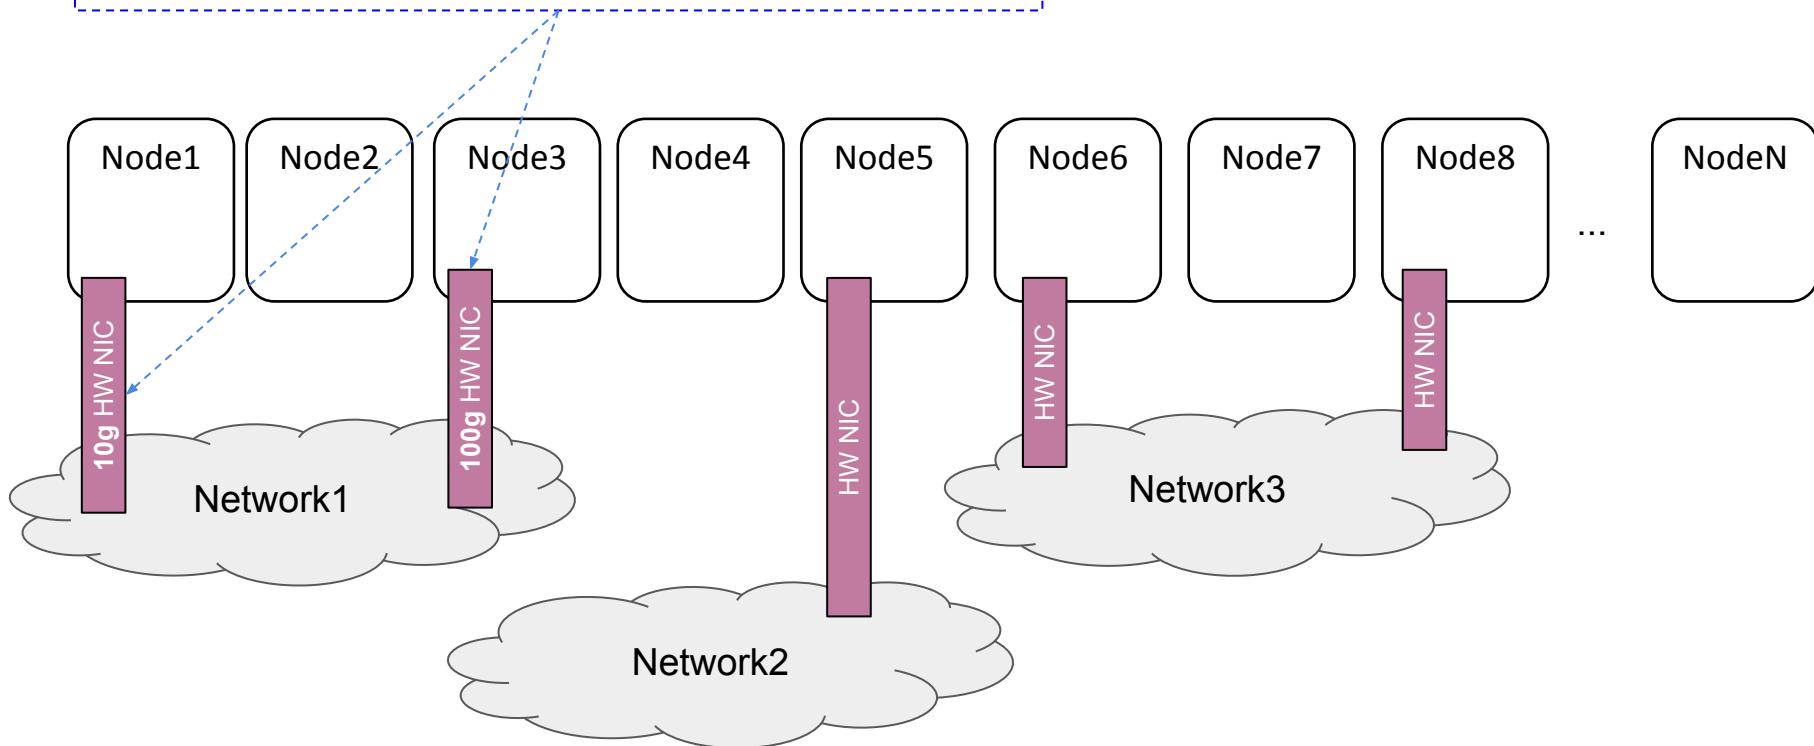
Those NICs may be plugged into a variety of different Networks:



Not all NICs have the same capabilities



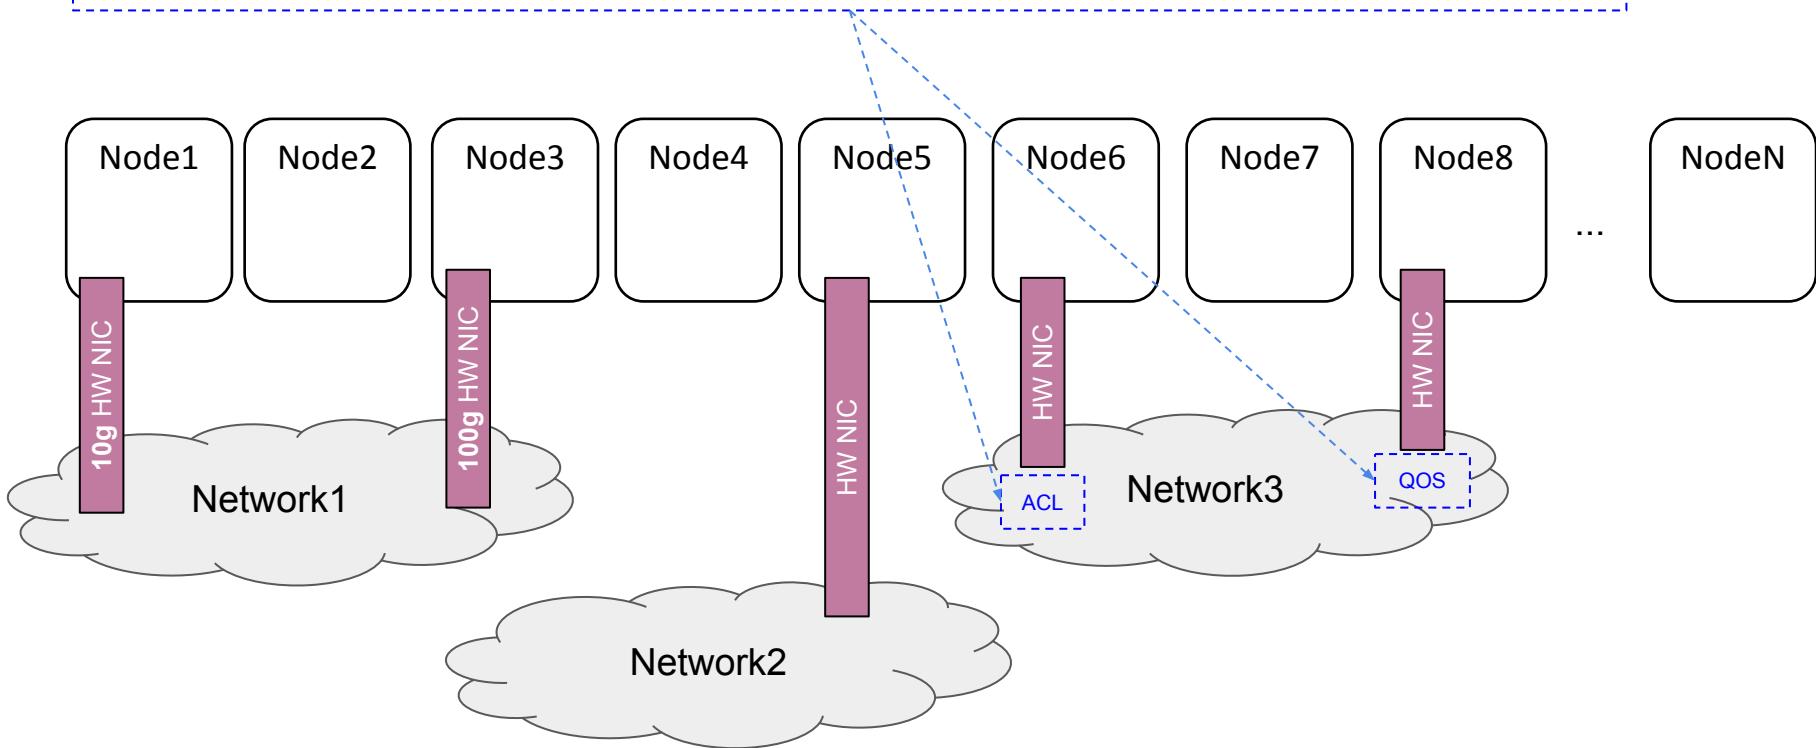
Those NICs may have a variety of capabilities(100G, 10G, etc):





Not all NICs can access the same Network Service

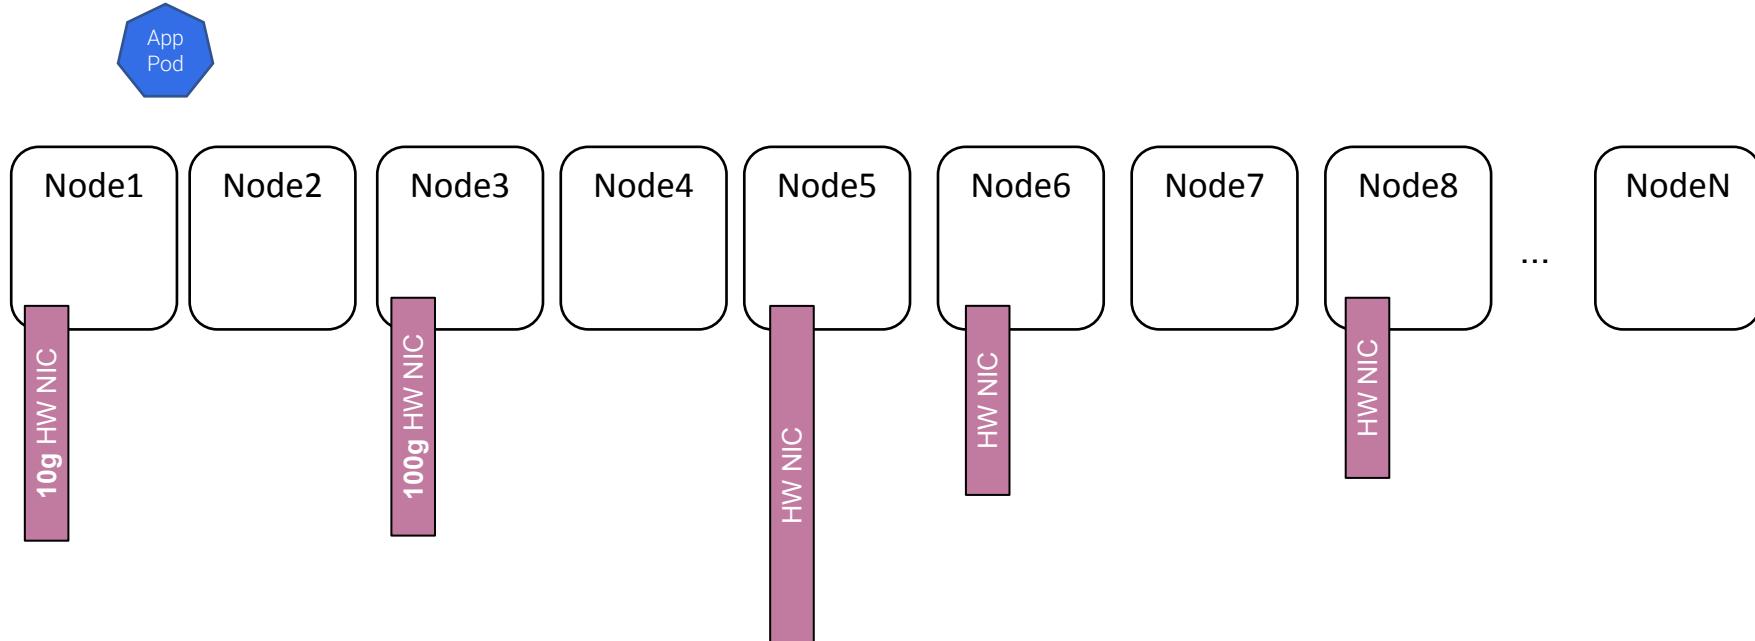
And two NICs plugged into the *same* Network may have different treatment (ACLs, QoS etc) (Network Service):





Scheduling a Pod

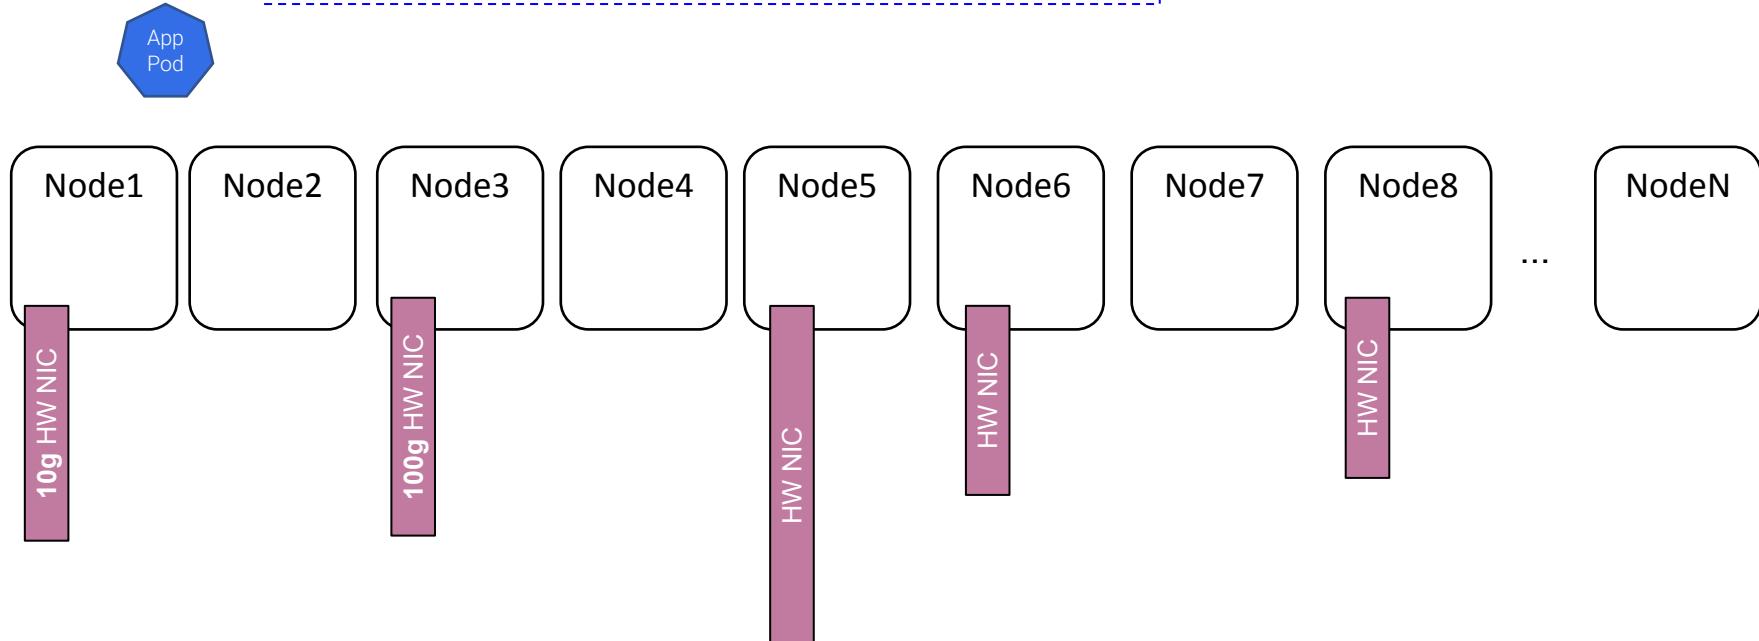
So what do we need when deploying a Pod?





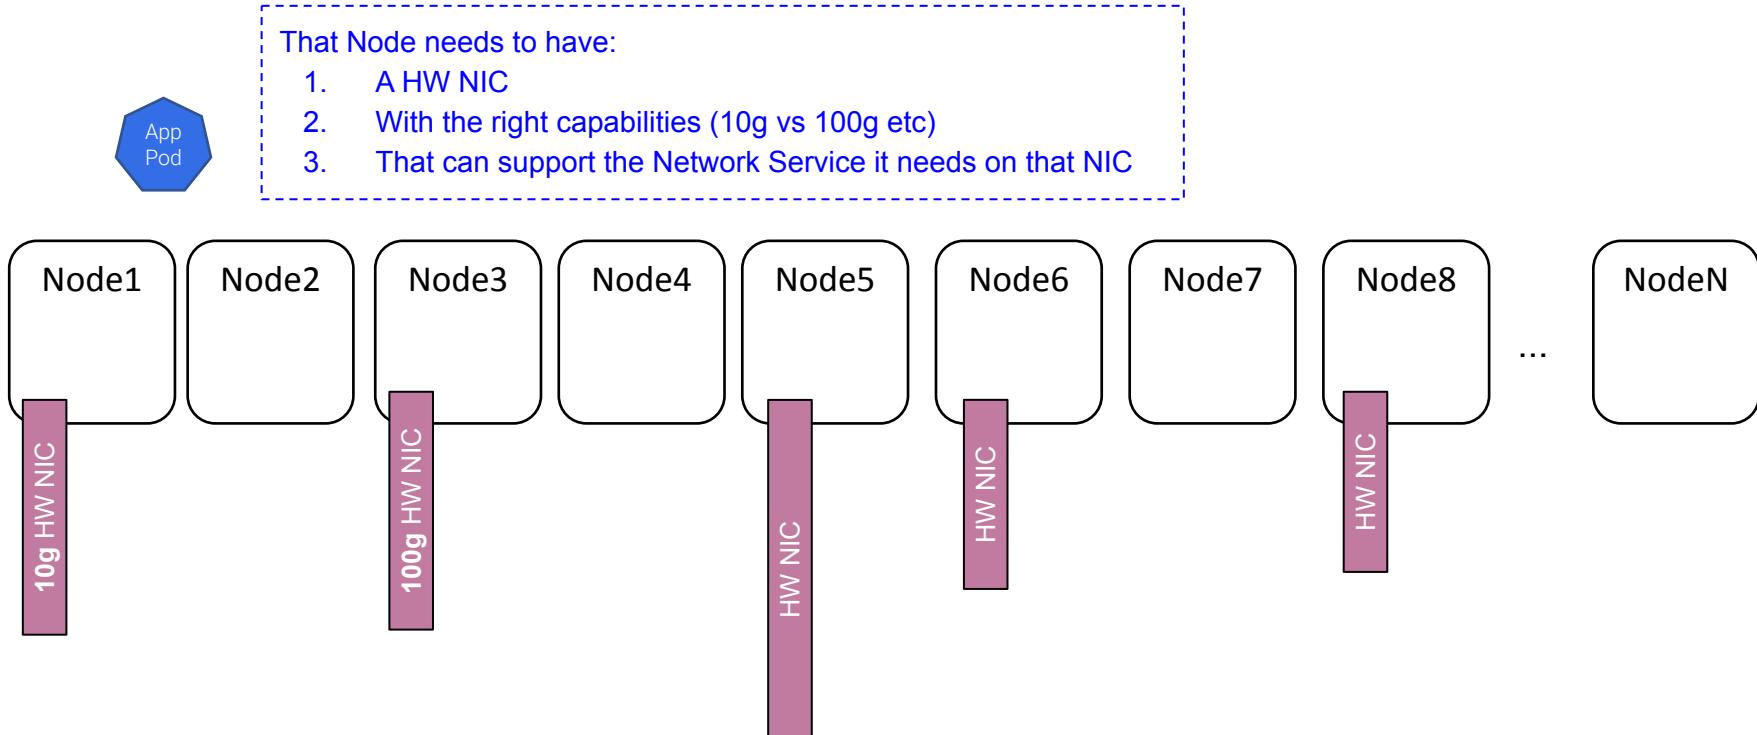
Scheduling a Pod

The K8s Scheduler needs to decide which Node to deploy it to...



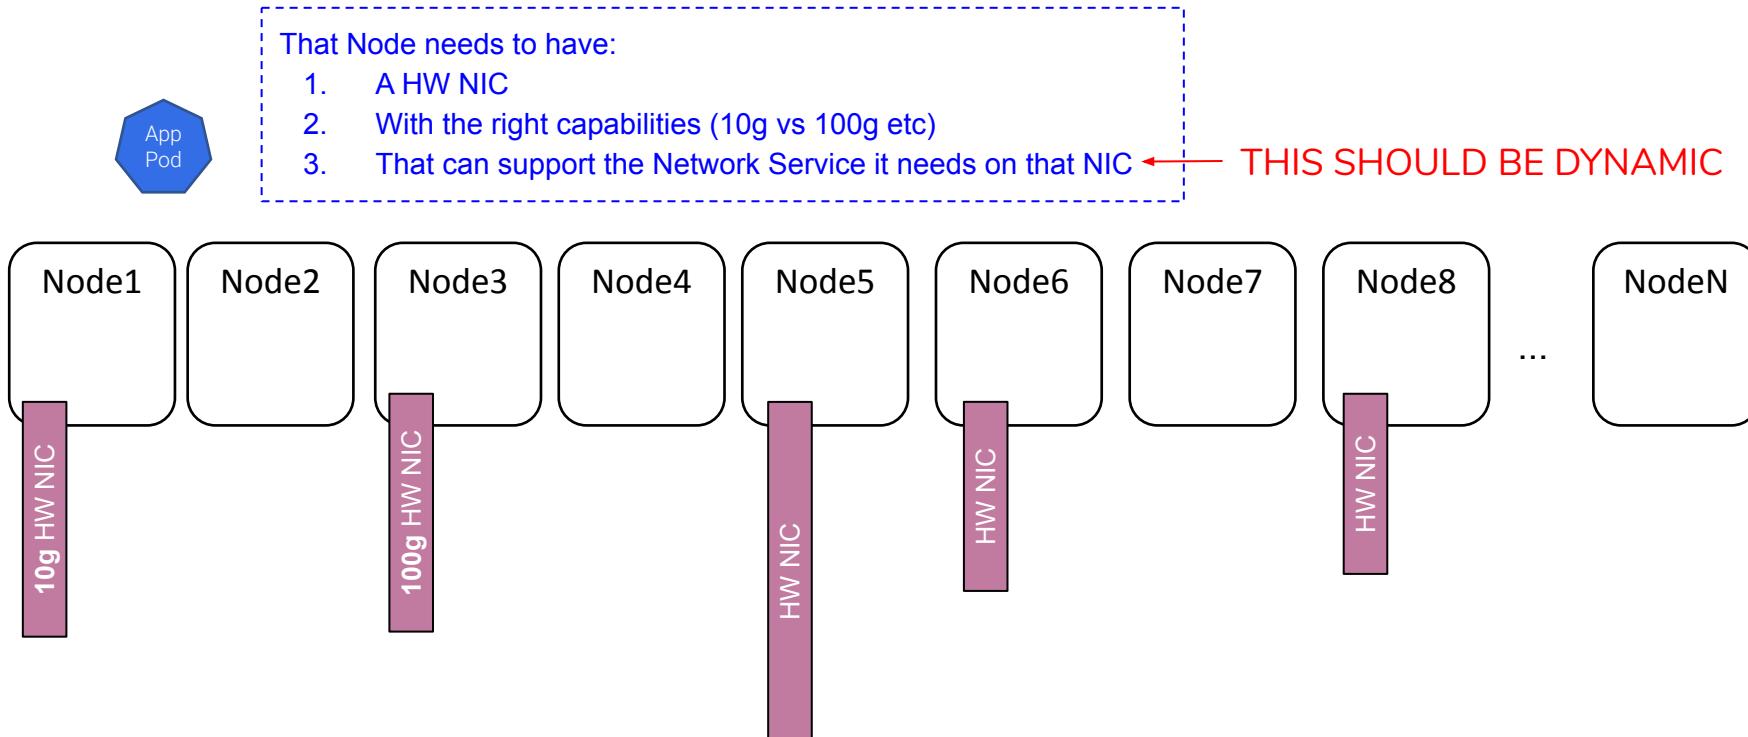


Criteria for scheduling a Pod





Network Service Should be Dynamic





How to ask for it?

```
apiVersion: v1
kind: Pod
metadata:
  name: cnf-1
  annotations:
    ns.networkservicemesh.io: myns.example.com
spec:
  ...
  resources:
    ...
  requires:
    example.com/100g: 1
```

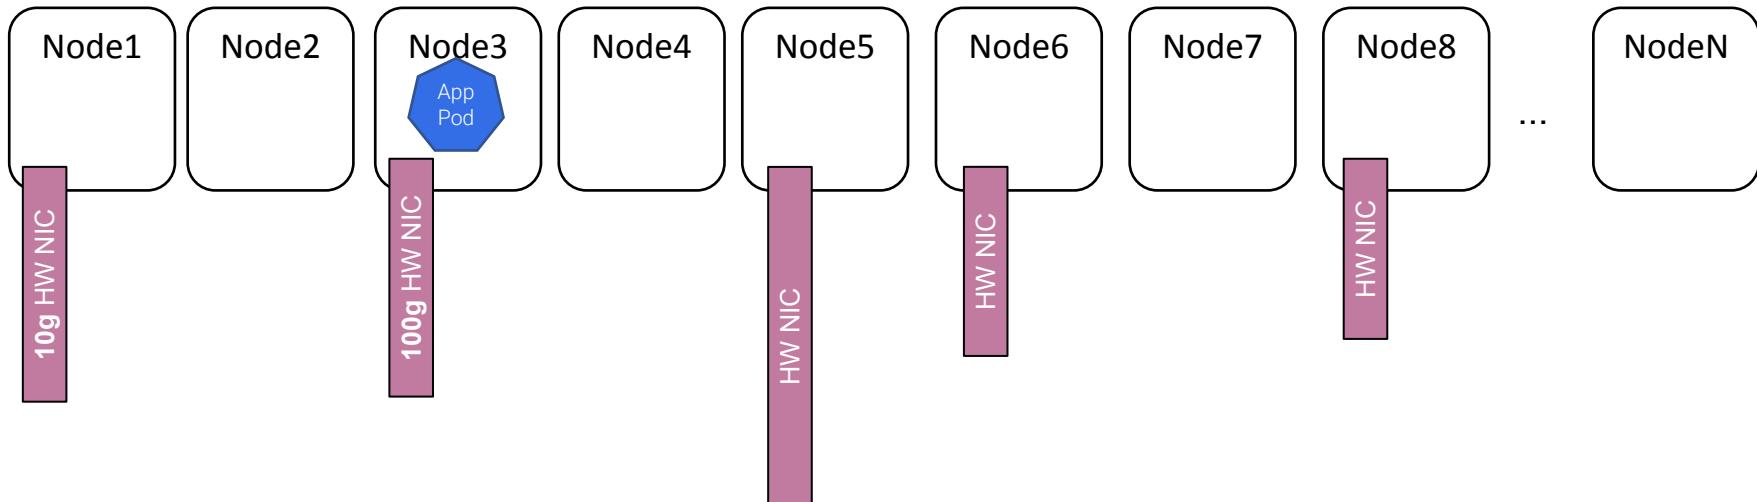
Ask for 'Network Service'

Ask for resource that can provide Network Services from the example.com domain with 100g capability



Scheduling...

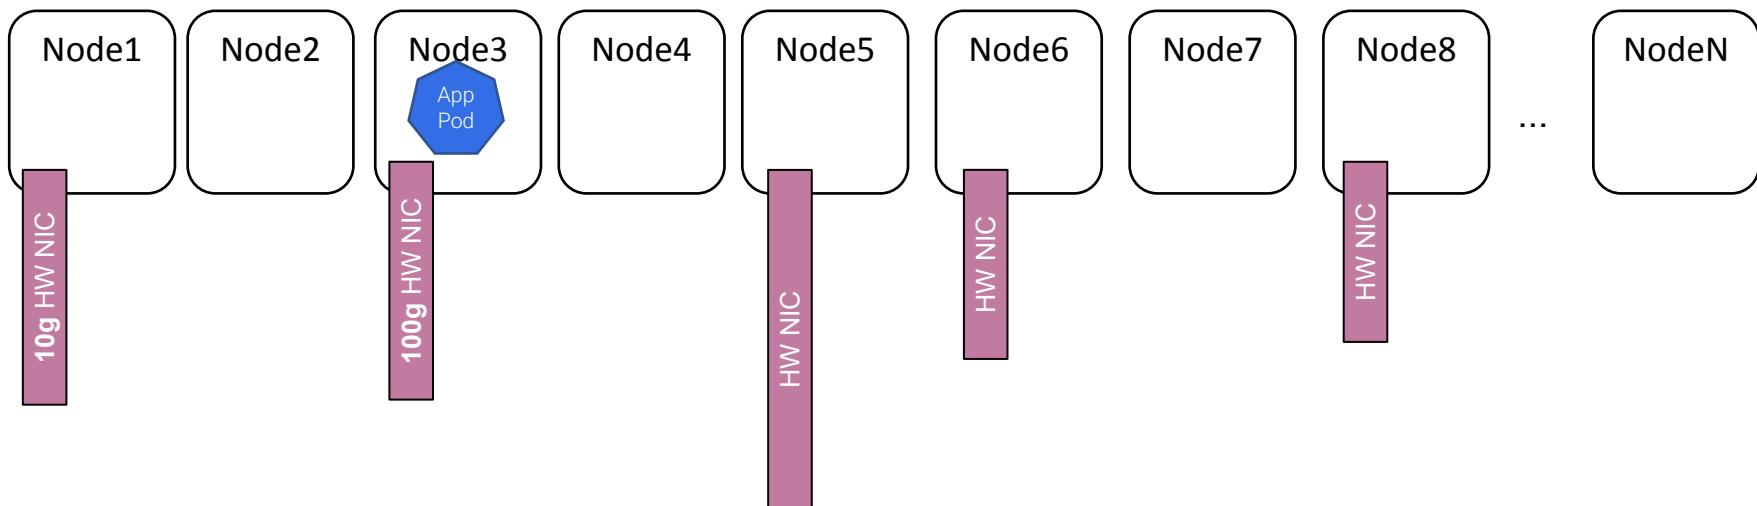
Device Plugin schedules us to Node with an available
'example.com/100g' resource where we can get any Network
Service in the example.com domain with a NIC with 100g
capabilities.





After scheduling...

Once Scheduled... we need to:

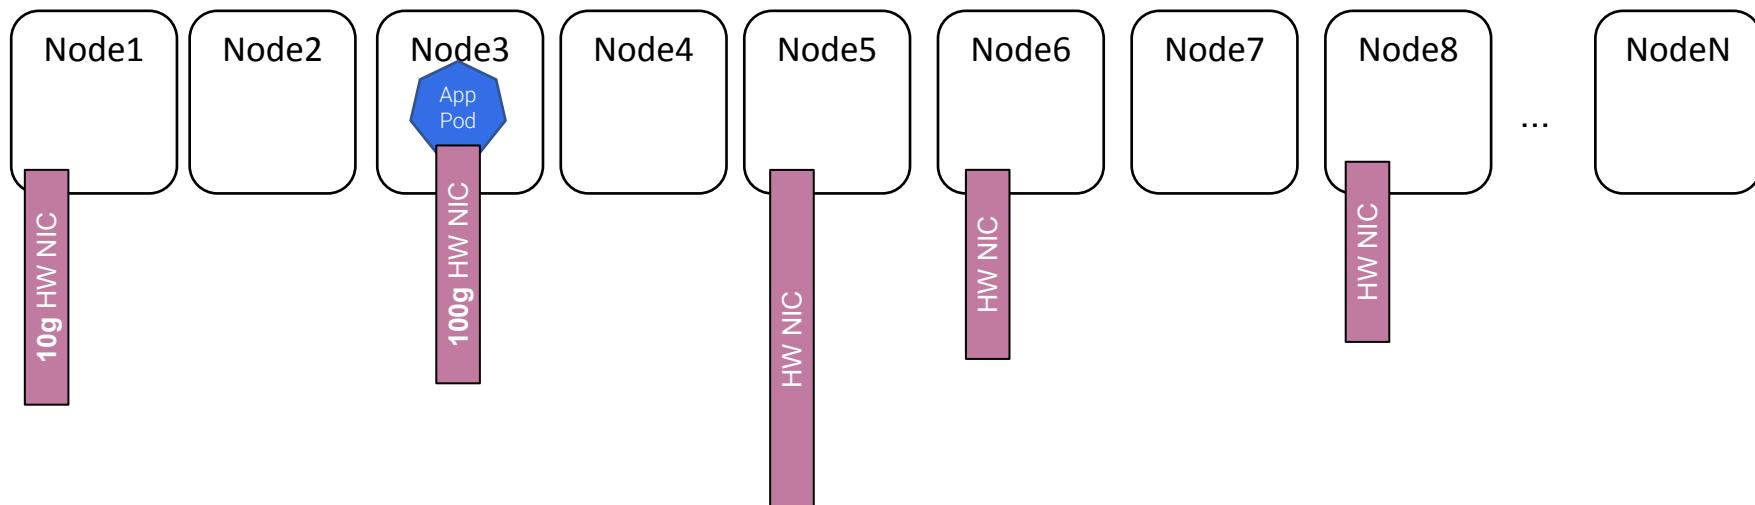




Plug NIC into the Pod

Once Scheduled... we need to:

1. Plug the NIC with the correct capabilities into the Pod

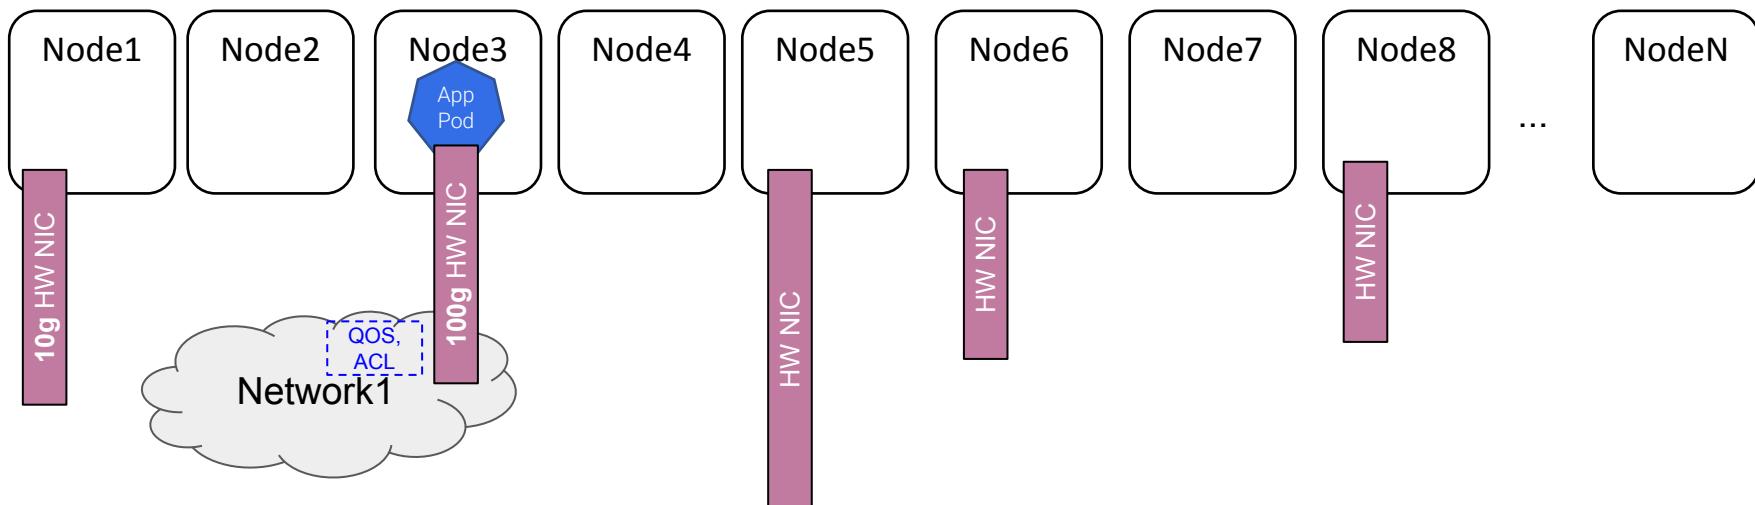




Configure the Network Service

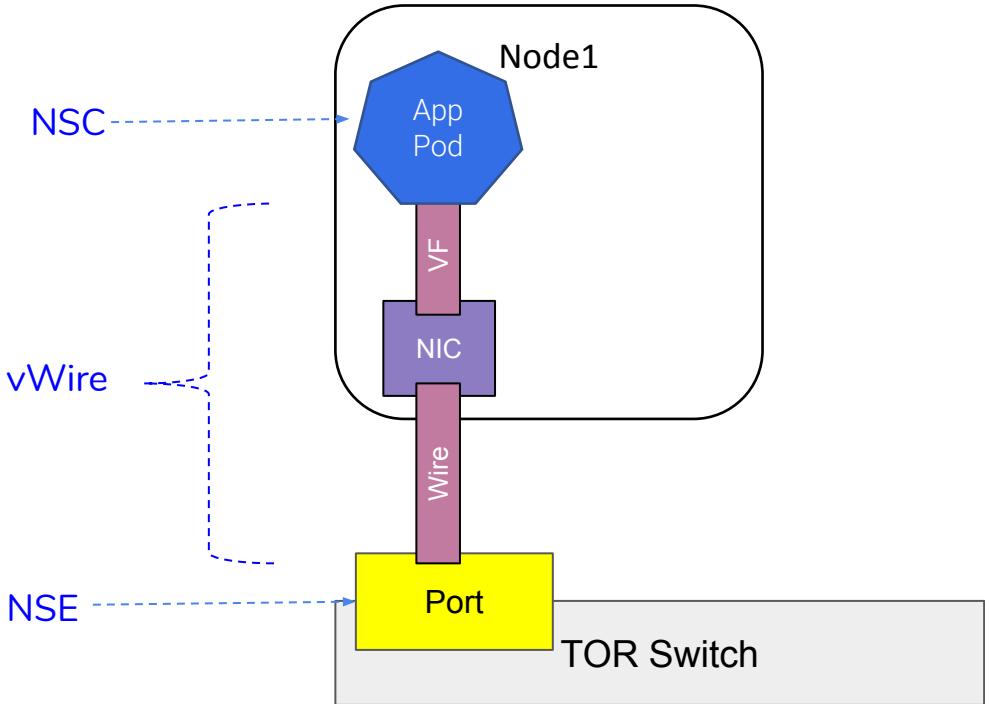
Once Scheduled... we need to:

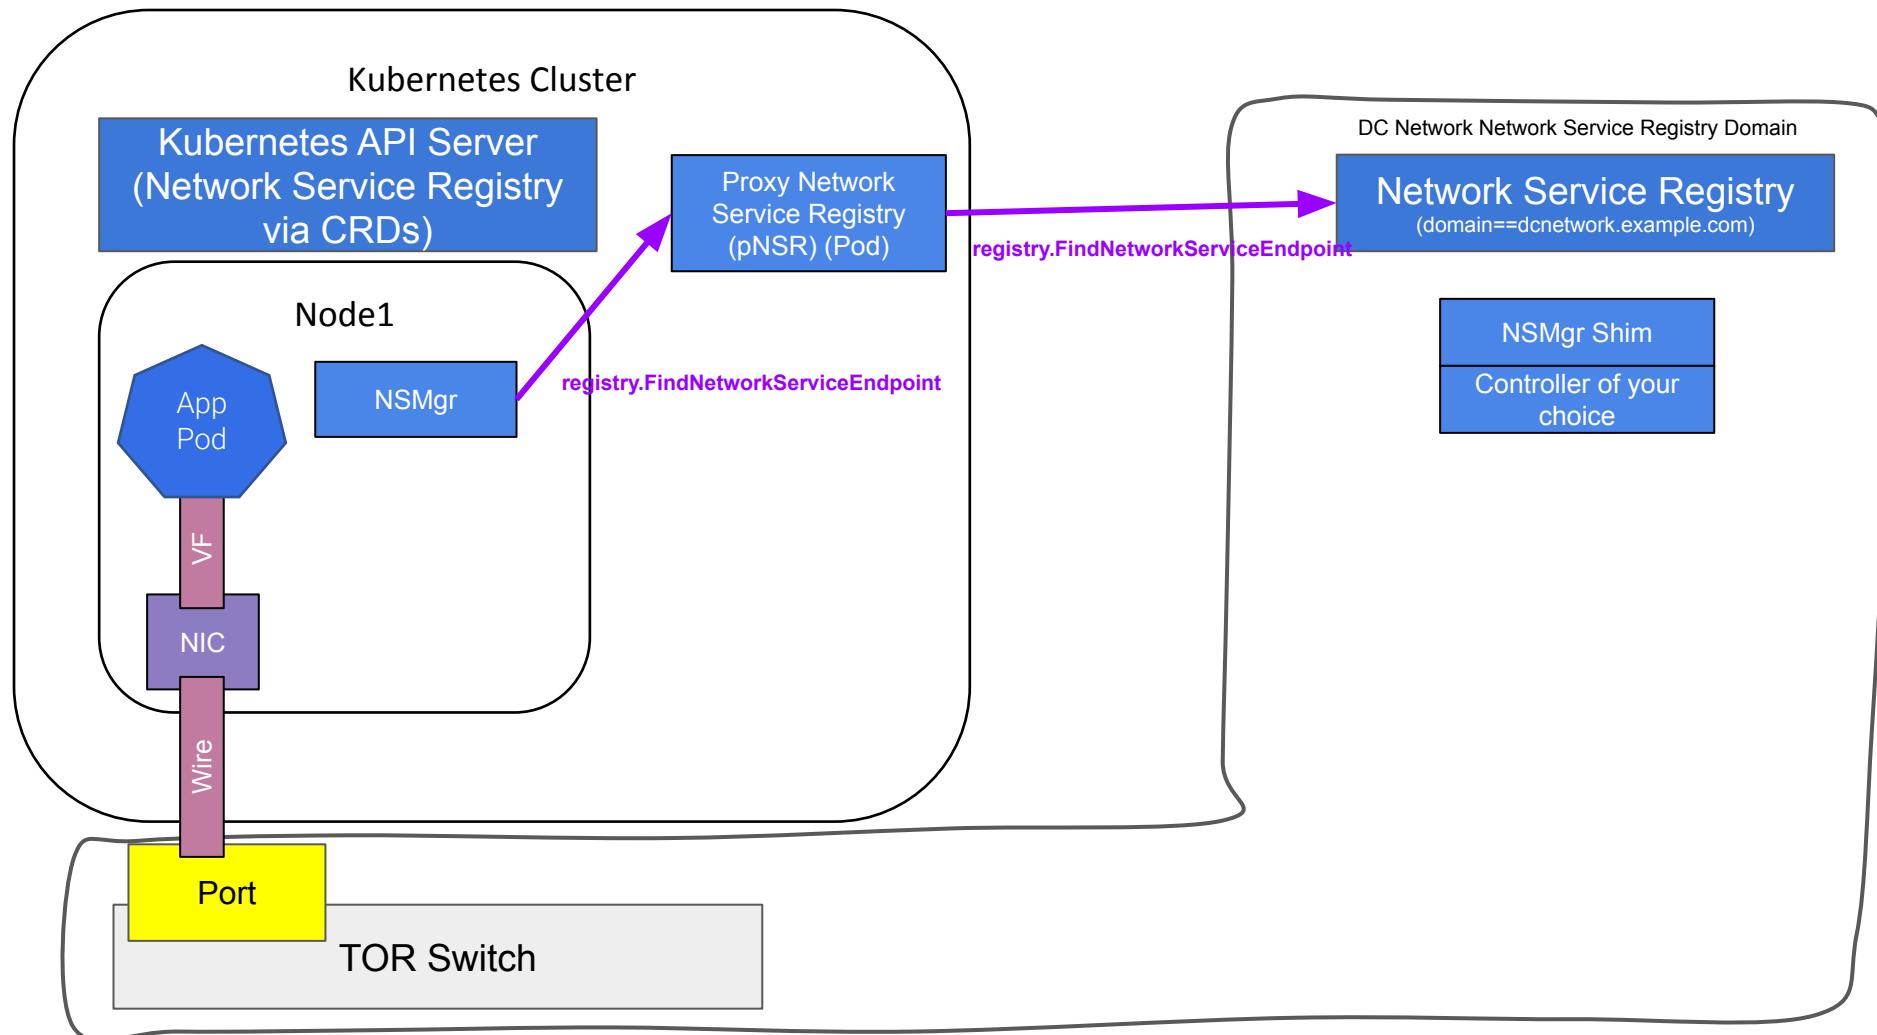
1. Plug the NIC into the Pod
2. Configure the other end of the NIC (TOR port) to provide the Network Service the Pod requested





Relationship to NSM model





Kubernetes Cluster

Kubernetes API Server
(Network Service Registry
via CRDs)

Proxy Network
Service Registry
(pNSR) (Pod)

Node1

App
Pod

NSMgr

VF

NIC

Wire

remote.NetworkService.Request

DC Network Network Service Registry Domain

Network Service Registry
(domain==dcnetwork.example.com)

NSMgr Shim

Controller of your
choice

Controller specific programming of network

Port

TOR Switch



Housekeeping



<https://networkservicemesh.io>



NSMCon

Nov 18, 2019 | San Diego, California
Colocated with Kubecon+CloudNativeCon 2019



These slides