



Everything you need to know about Storage for Kubernetes

Webinar

Alex Chircop

About me



- **Founder & CEO of StorageOS, building a software defined, cloud native storage platform**
- **Co-chair CNCF Storage SIG**
- 25 years of engineering infrastructure platforms before embarking on the startup adventure
- Previously Goldman Sachs and Nomura
- Twitter: **@chira001**



CNCF SIG Storage - Get involved!



The **CNCF Storage SIG** meets on the
2nd and 4th Wednesday of every month at 8AM PT (USA Pacific)

Call	https://zoom.us/my/cncfstoragewg
SIG rep	https://github.com/cncf/sig-storage
Meeting Minutes	http://bit.ly/cncf-storage-sig-minutes
Mail List	https://lists.cncf.io/g/cncf-sig-storage



Why is storage so important?

Why Storage?



There is no such thing
as a stateless
architecture!

All applications store
state somewhere...



Challenge: No Storage Pets!

Cloud Native Storage should be:

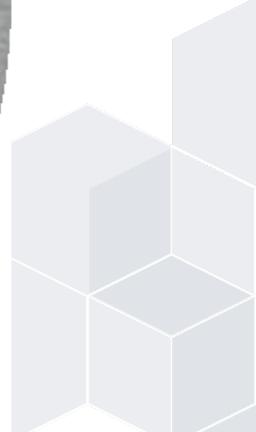
- Declarative & Composable
- API Driven



Challenge: Data needs to follow!

Cloud Native Storage should be:

- **Application Centric**
- **Platform Agnostic**
- **Agile**



Challenge: Humans are fallible - automate everything!

Cloud Native Storage should be:

- **Consistently Available**
- **Performant**
- **Natively Secure**





CNCF Storage Landscape Whitepaper

<http://bit.ly/cncf-storage-whitepaper>

- How is a storage system **deployed**
- Definition of the **attributes** of a storage system
- Definition of the **layers** in a storage solution with a focus on terminology and how they impact the attributes
- Definition of the **data access interfaces** in terms of volumes and application APIs
- Definition of the **management interfaces**



Instantiation and Deployment

Instantiation	Description
Hardware	Deployed as hardware solution in a datacenter. This limits the portability of the application and generally means that such systems cannot be deployed in a public cloud environment
Software	Deployed as software components on commodity hardware, appliances or cloud instances. Software solutions tend to be more platform agnostic and can be installed both on-premises as well as cloud environments. Some software defined storage systems can also be deployed as a container and deployment can be automated by an orchestrator.
Cloud Services	Consumed from public cloud providers. Cloud services provide storage services in cloud environments.

Storage Attributes

Availability	Scalability	Performance	Consistency	Durability
Failover	Clients	Latency	Delay to access correct data after a commit	Data protection
Moving access between nodes	Operations	Operations		Redundancy
Redundancy	Throughput	Throughput	Delay between commit and data being committed to non-volatile store	Bit-Rot
Data protection	Components			

Orchestrator, Host and Operating System

Storage Topology

(centralized, distributed, sharded, hyperconverged)

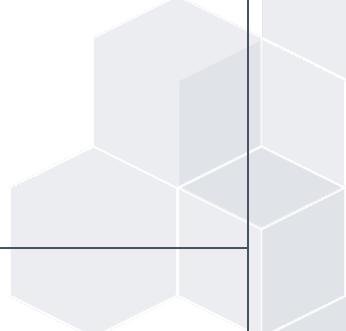
Data Protection

(RAID, Erasure coding, Replicas)

Data Services

(Replication, Snapshots, Clones, etc.)

Physical, Non-Volatile Layer

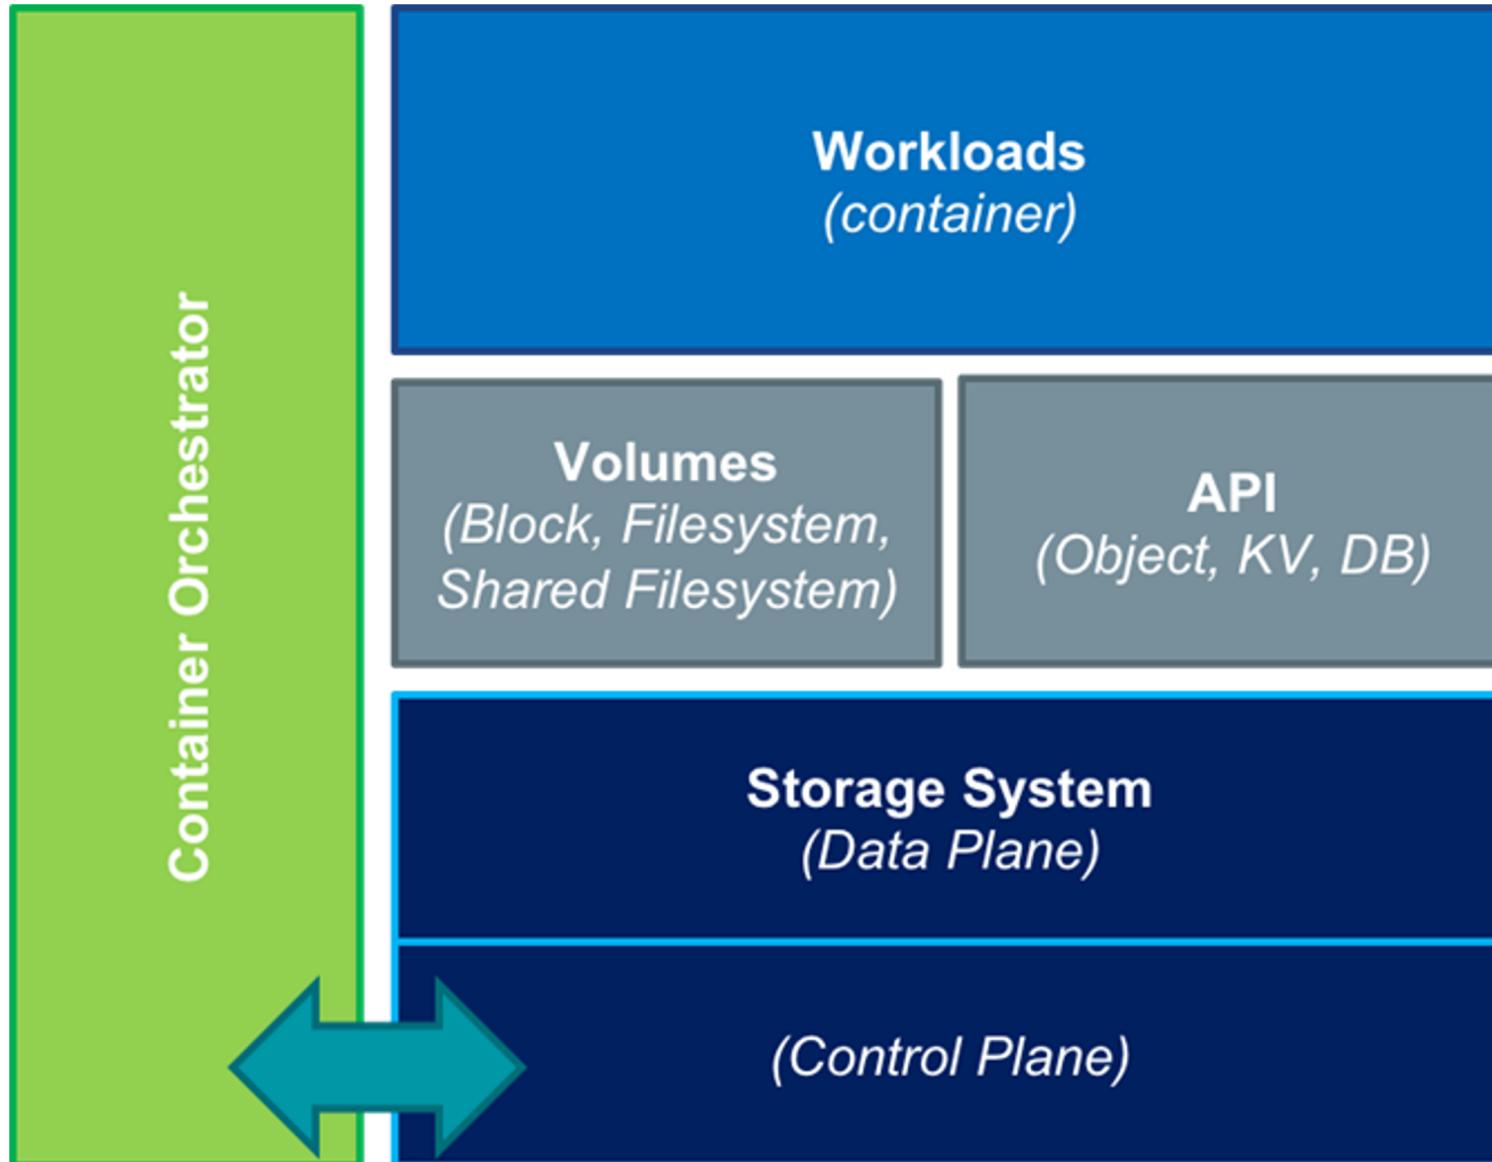


Storage Topology Comparison

	Local	Remote	Distributed
Availability	Limited by failure of components locally and ability to failover. If a node fails, the local storage is isolated to the local node.	May be limited by single points of failure. Workloads can move to another node and reconnect to the remote storage.	Clients may access numerous nodes, and any storage node failures can be mitigated. The additional complexity may also add operational complexity which may affect availability or the ability to recover errors.
Scalability	Limited by local architecture (1 node; typically TB)	Limited by monolithic architecture (2-16 nodes; typically 10s-100s of TB)	Scale by adding additional systems. (3-1000s nodes; often supports PB)
Consistency	Yes (storage system implementation is easy)	Yes (storage system implementation is harder with more nodes)	Yes (storage system implementation is hardest)
Durability	Limited by local components (less)	Limited by monolithic architecture (more)	Scaling out to additional systems increases durability (most)
Performance	Limited by local components, can benefit low-latency applications (100us-5ms, GB/sec)	Similar to local, but additional overhead in network transport (500us-5ms, GB/sec)	Scaling out to additional systems increases performance (500us-5ms, GB/sec)

^{**} The information in this table are generally accepted attributes and measurements among local, remote, and distributed storage systems.

Data Access Interfaces



Storage can be accessed via
Data Access Interfaces:

- **Volumes** – accessed through a more traditional file interface in a **block** or **filesystem** interface
- **API** – other ways to persist data such as **object stores**, **KV stores** or **databases**

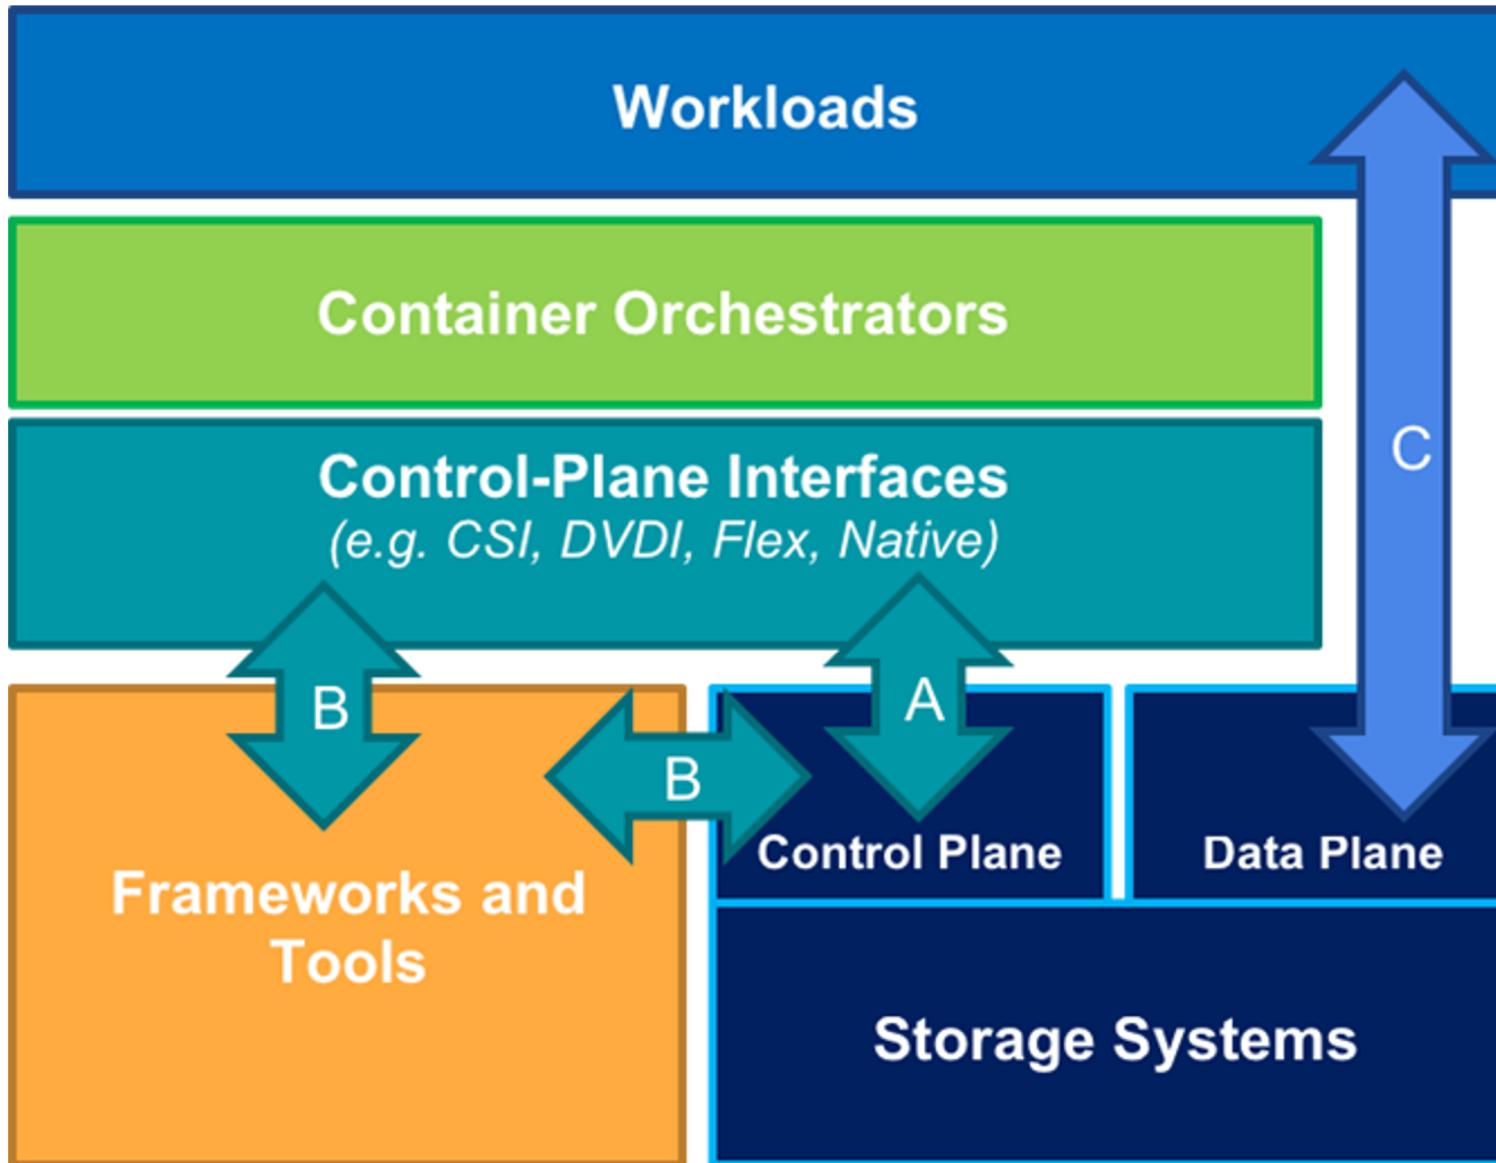


Data Access Interface Comparison

Data Access Interface	Most suited	Least suited
Block	<ul style="list-style-type: none">• Availability• Low latency performance• Good throughput performance for individual workloads	<ul style="list-style-type: none">• Capacity scaling• Sharing data with multiple workloads simultaneously
Filesystem	<ul style="list-style-type: none">• Sharing data with multiple workloads simultaneously• Optimised throughput for aggregated workloads	<ul style="list-style-type: none">• Strong file locking integrity when filesystems are shared
Object Store	<ul style="list-style-type: none">• Availability• Large capacities (PB scale)• Durability• Sharing data with multiple workloads simultaneously• Optimised throughput for parallelised workloads	<ul style="list-style-type: none">• Low Latency performance 

** The information in this table are generally accepted attributes and measurements for data access interfaces.

Orchestration and Management Interfaces



Container Orchestration system (**CO**) uses an interface to interact with a storage system

The storage system can:

- (**A**) support control-plane API directly
- (**B**) interact via an API Framework layer or other Tools

Workloads consume (**C**) storage via a data access interface



Control Plane Interfaces

Container Storage Interface (CSI)

CSI v1.0.0 was released in November 2018. Standard for Kubernetes.

<https://github.com/container-storage-interface/spec>

K8S Native Drivers

Native drivers which are “in-tree” to K8S and provide the interfaces for Persistent Volumes (PV) and Persistent Volume Claims (PVC)

<https://kubernetes.io/docs/concepts/storage/>

Docker Volume Driver Interface

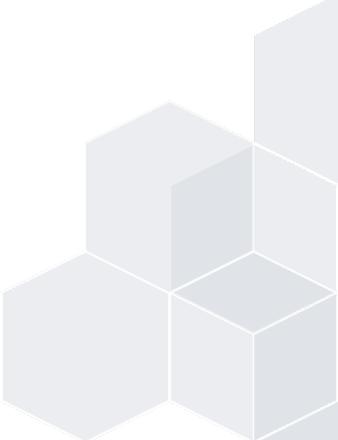
Provides a mechanism for storage vendors to write a volume driver so that storage systems can be used to provide volumes for a docker container.

<https://docs.docker.com/storage/>

K8S Flex Volume

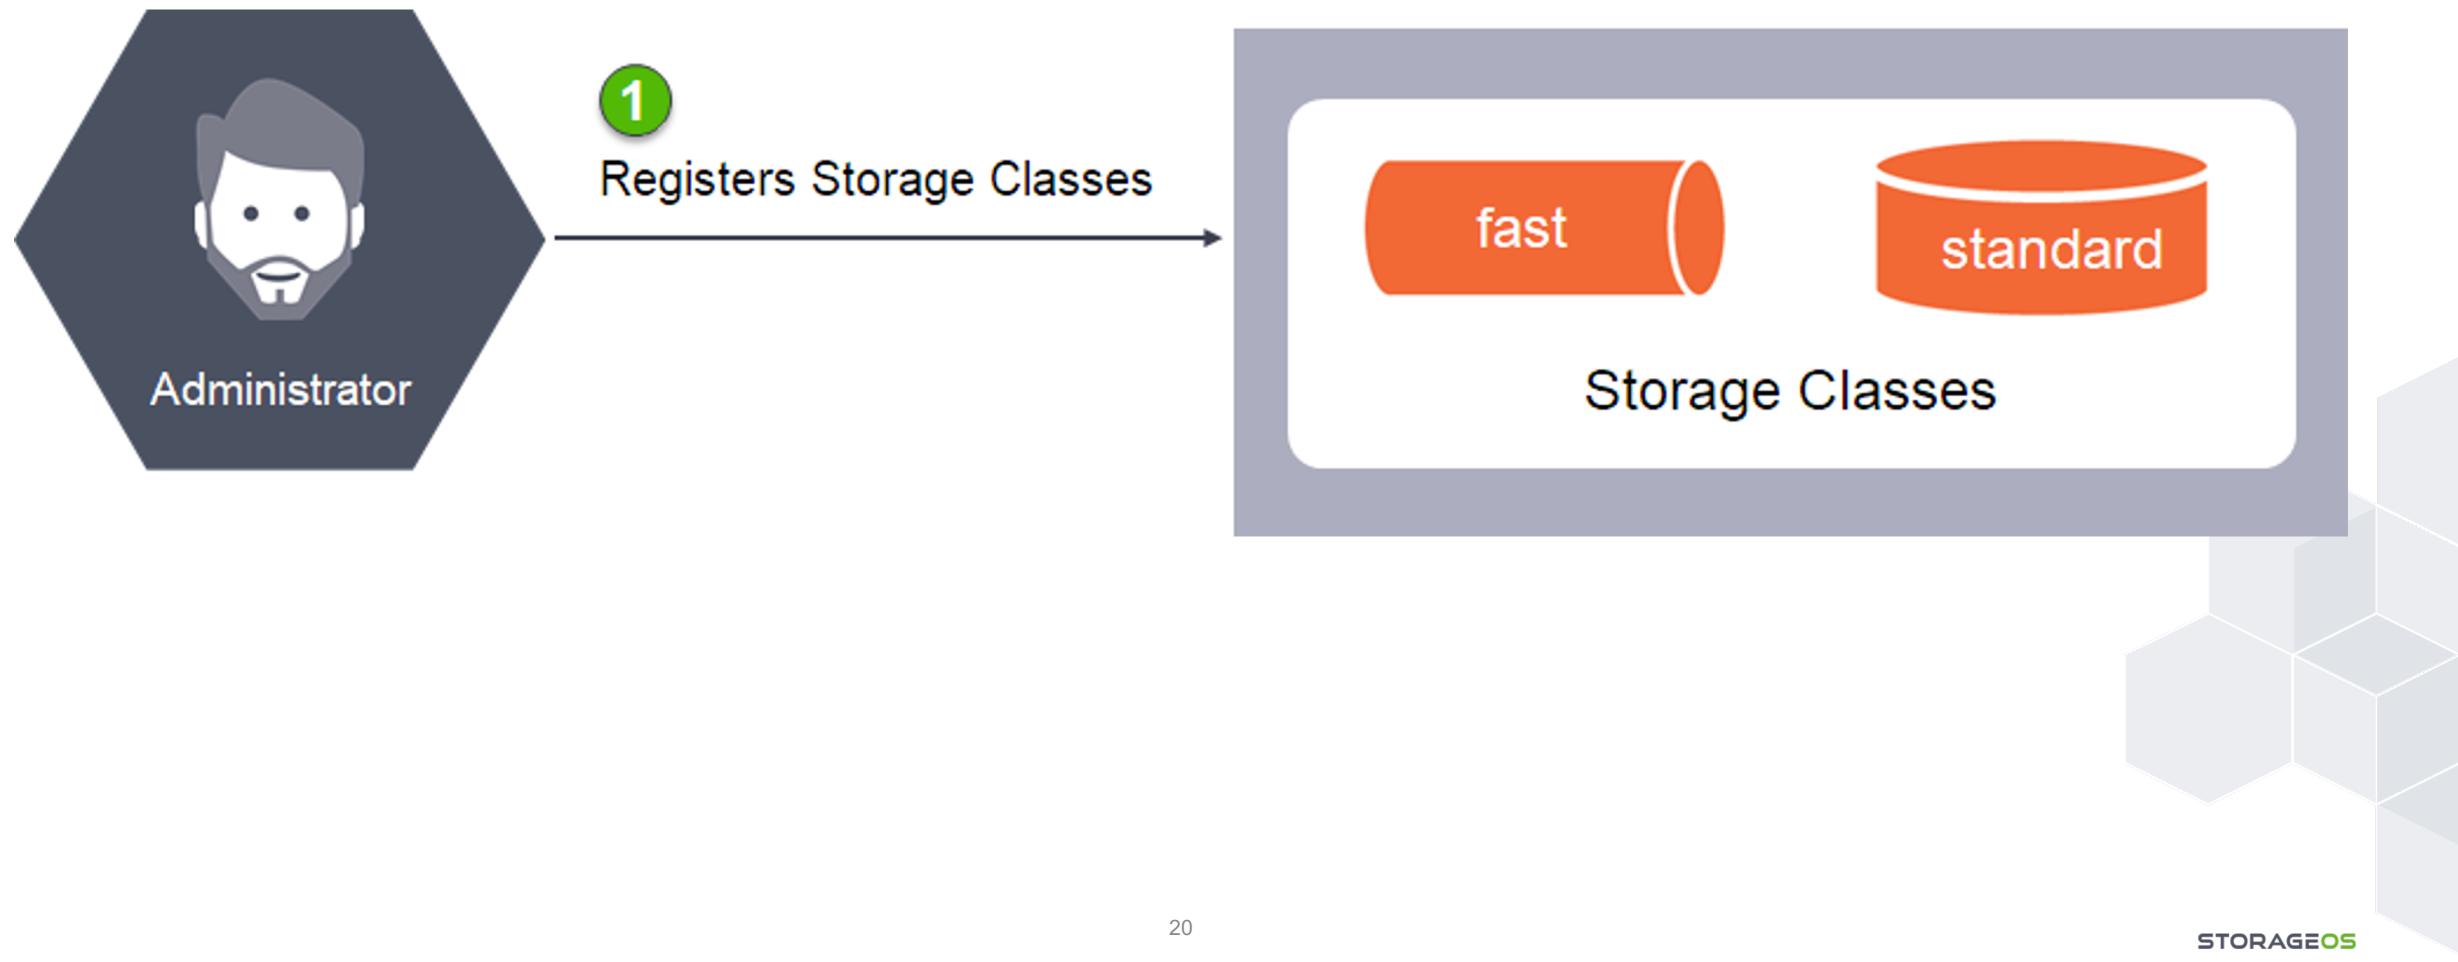
“out of tree” external volume driver for K8S. Preference is now for CSI.

<https://github.com/kubernetes/community/blob/master/contributors/devel/flexvolume.md>



How is storage configured in Kubernetes?

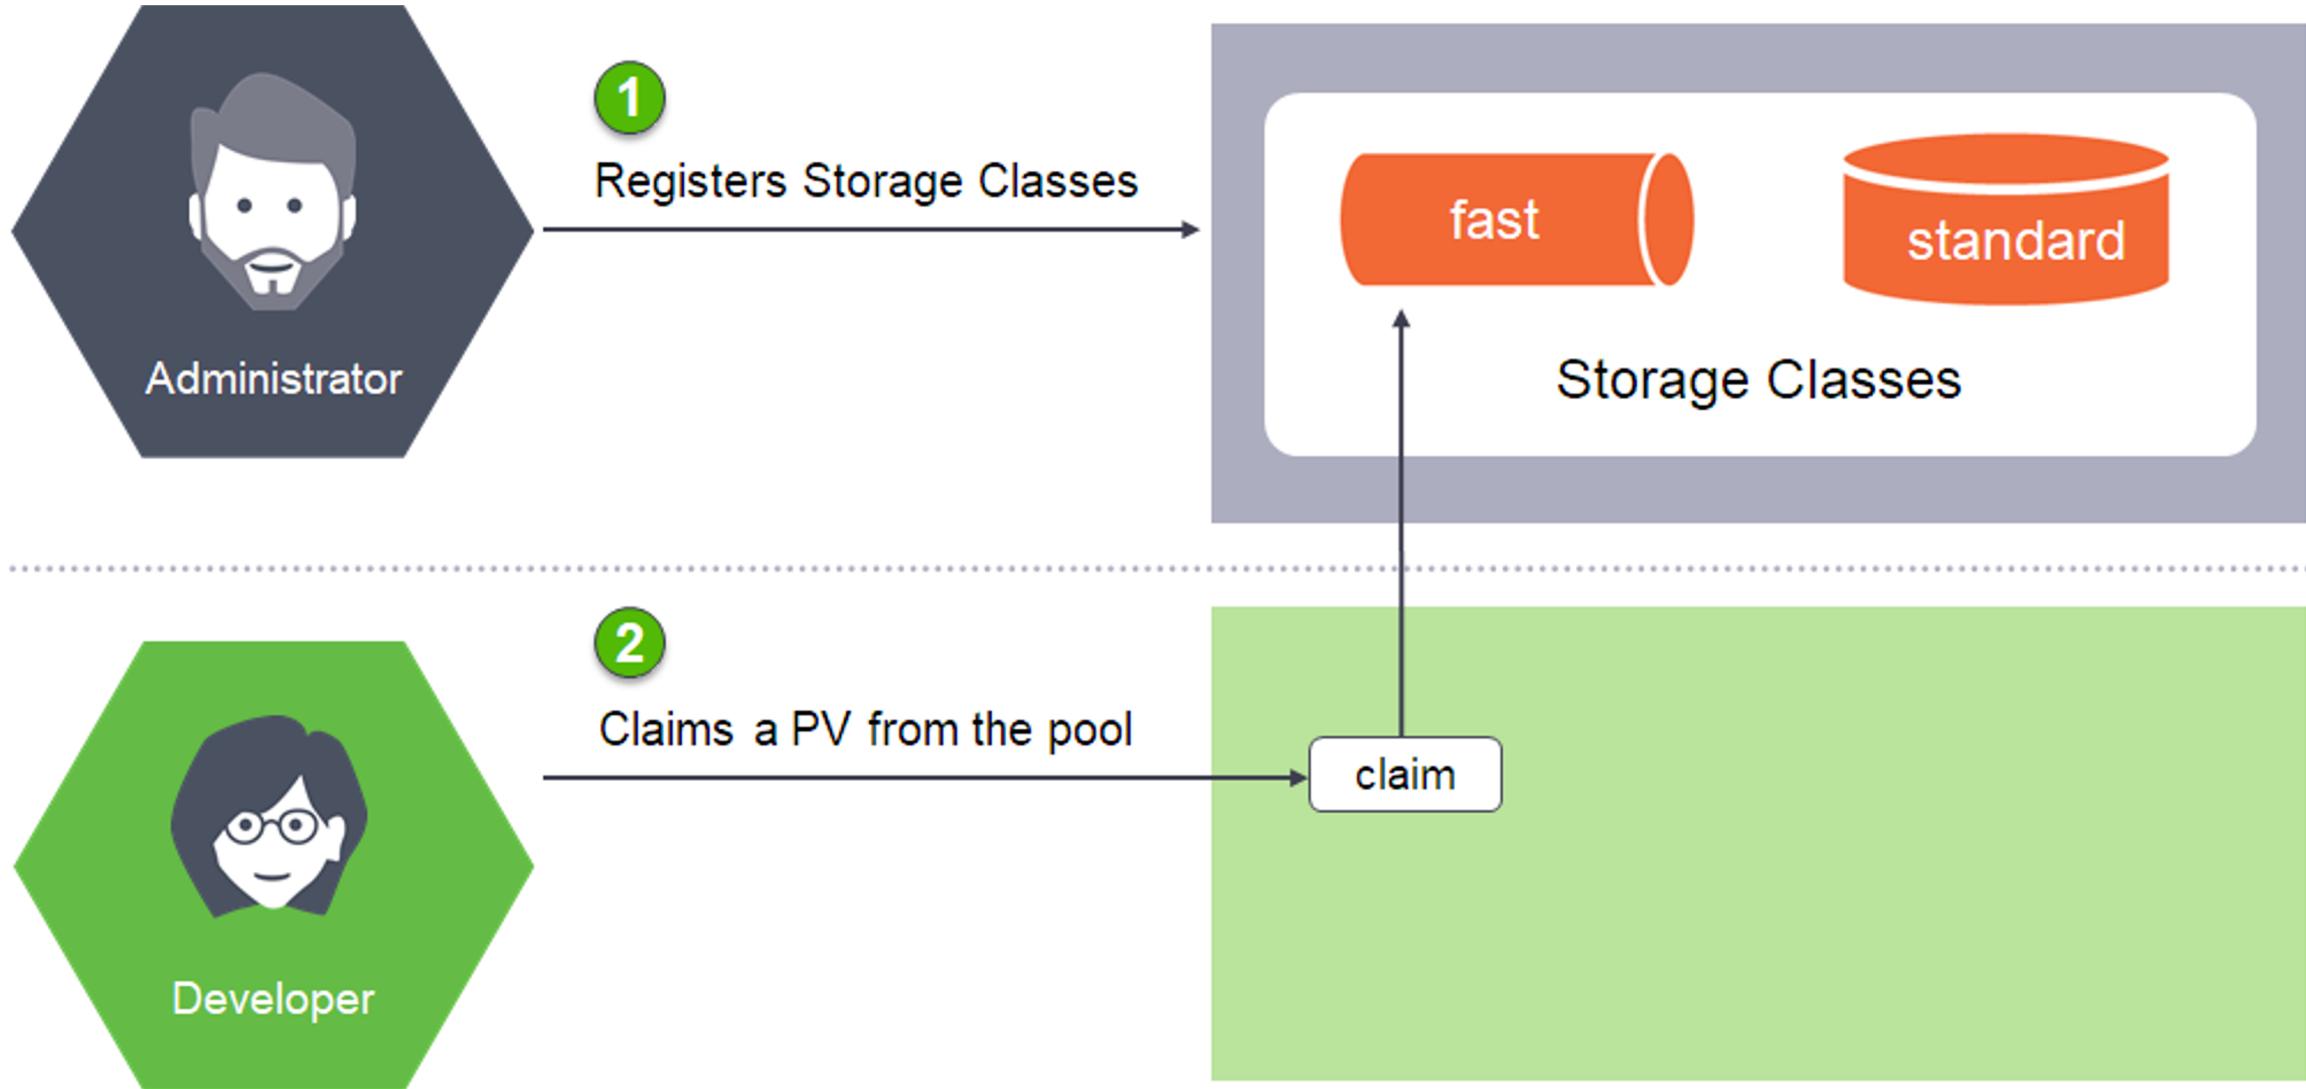
Dynamic Provisioning in Kubernetes: StorageClass



An Example StorageClass

```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: fast
  parameters: fsType: ext4
  pool: default    storageos.com/replicas: "2"
provisioner: storageos
```

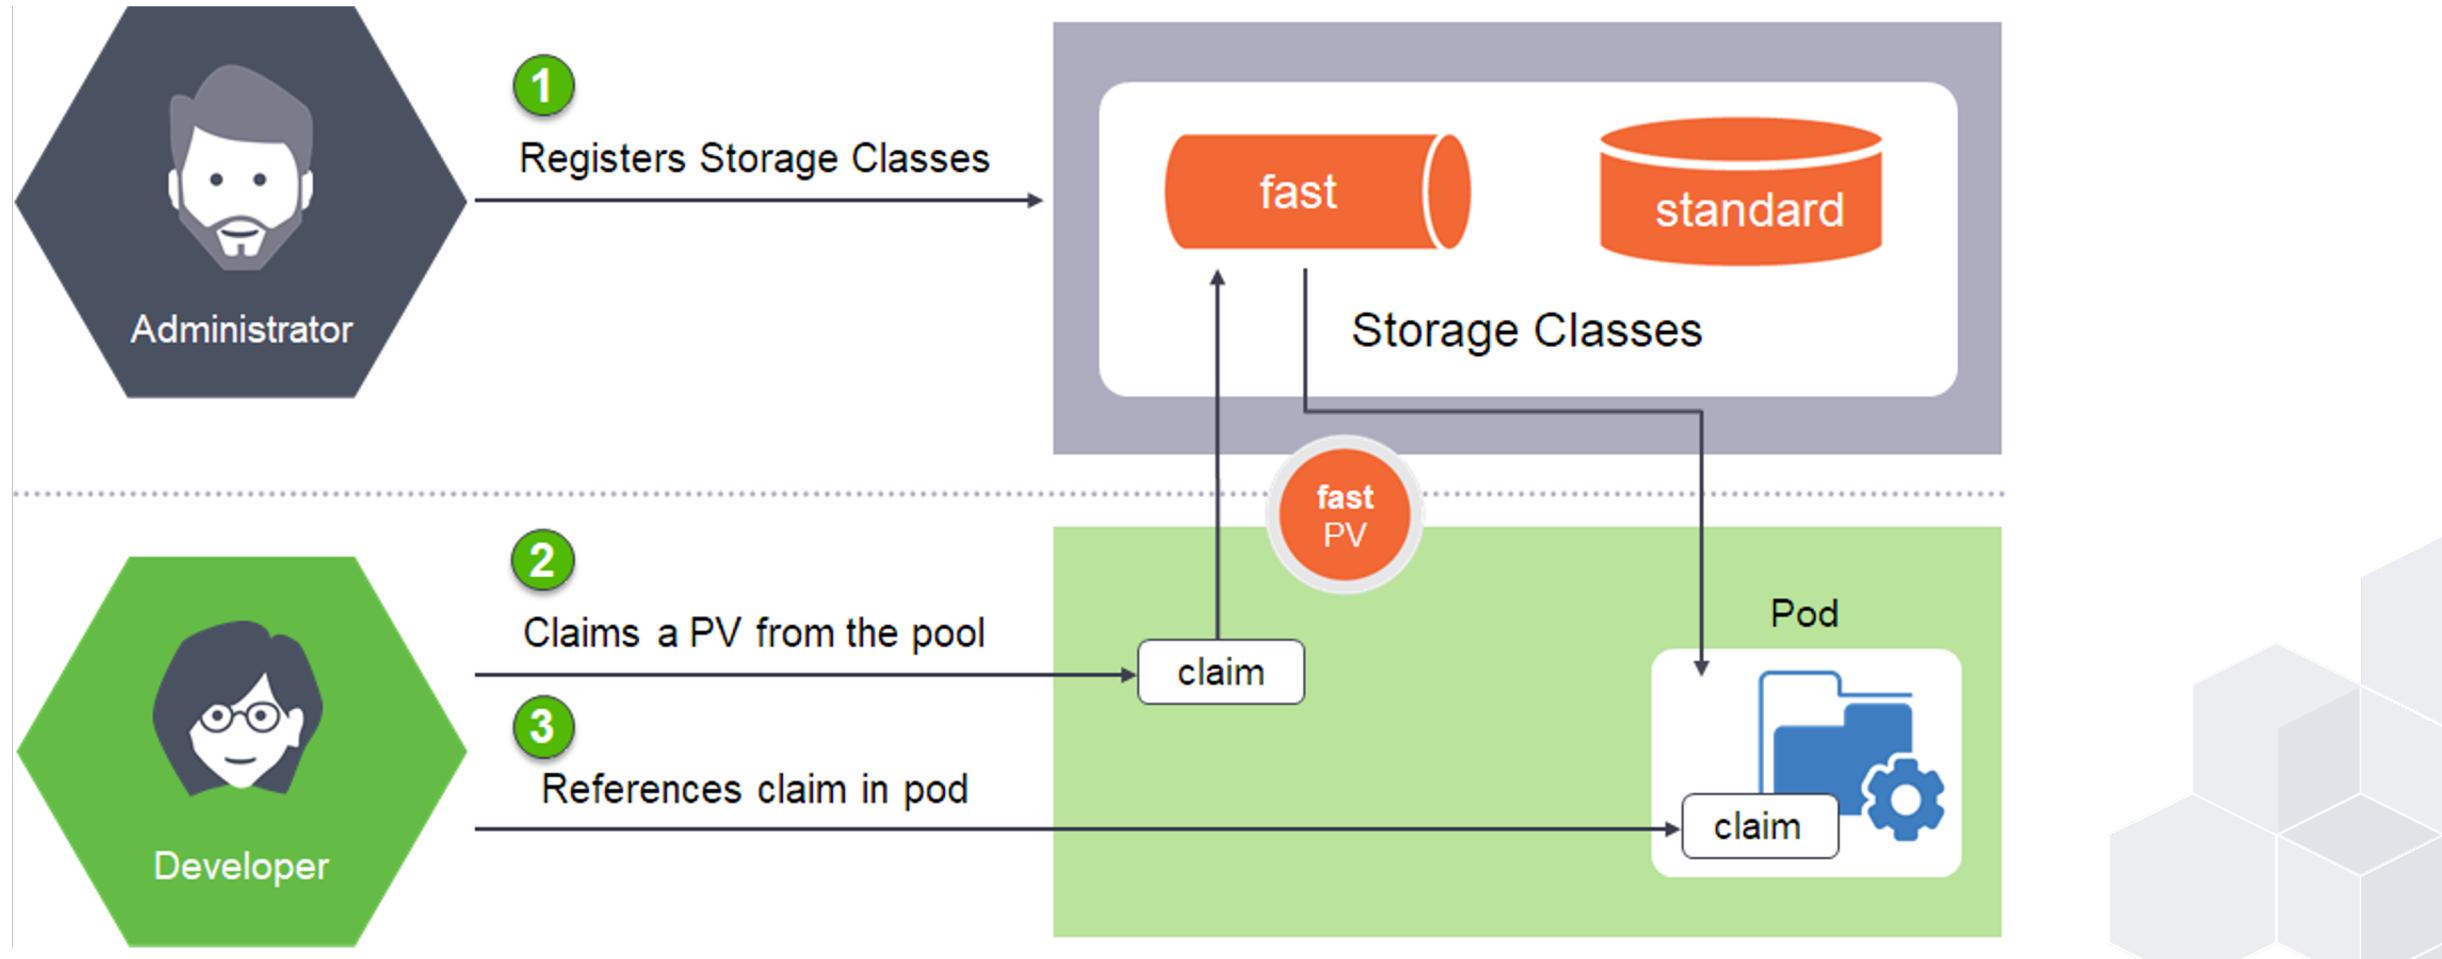
Dynamic Provisioning in Kubernetes: PersistentVolumeClaim



An example PersistentVolumeClaim

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: database-vol1
  annotations:
    volume.beta.kubernetes.io/storage-class: fast
spec:
  accessModes:
    - ReadWriteOnce
resources:
  requests:
    storage: 5Gi
```

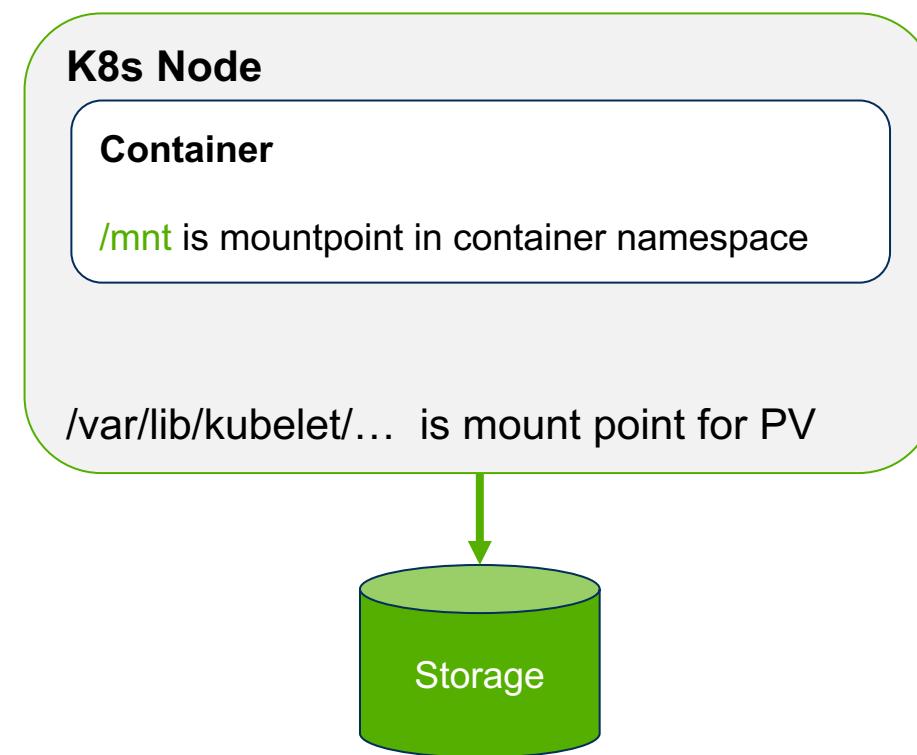
Dynamic Provisioning in Kubernetes: PersistentVolume



Using a volume in a pod

```
apiVersion: v1
kind: Pod
metadata:
  name: test
spec:
  containers:
    - name: debian
      image: debian:9-slim
      command: ["/bin/sleep"]
      args: [ "3600" ]
    volumeMounts:
      - mountPath: /mnt
        name: v1
  volumes:
    - name: v1
      persistentVolumeClaim:
        claimName: pvc-1
```

1. PVC is requested
2. Storage system creates volume (PV)
3. PV is attached to Node and mounted
4. Bind mount of the volume into specified path in container



Volume Access Modes

Access Mode	Description
ReadWriteOnce “RWO”	Volume is mounted and accessed exclusively by only one node.
ReadWriteMany “RWX”	Volume can be mounted on multiple nodes at the same time.



Example of provisioning a stateful workload

Questions?

Thank You

Email: alex.chircop@storageos.com

Twitter: [@chira001](https://twitter.com/@chira001)

Slack: slack.storageos.com

www.storageos.com

