

# The evolution of Ingress through the Gateway API

...

Bowei Du <[bowei@google.com](mailto:bowei@google.com)>

Kaslin Fields <[kaslin@google.com](mailto:kaslin@google.com)>



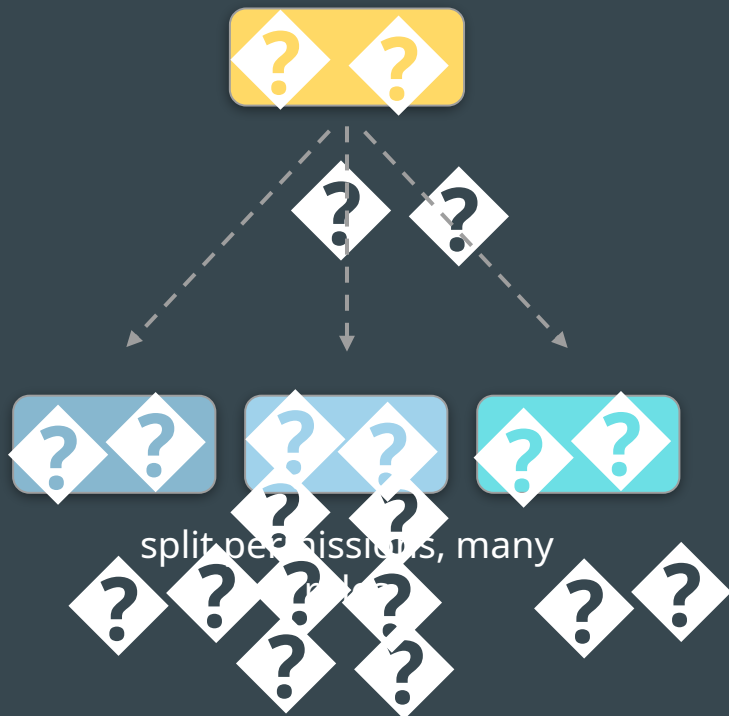
Kaslin Fields <  
[kaslin@google.com](mailto:kaslin@google.com)>



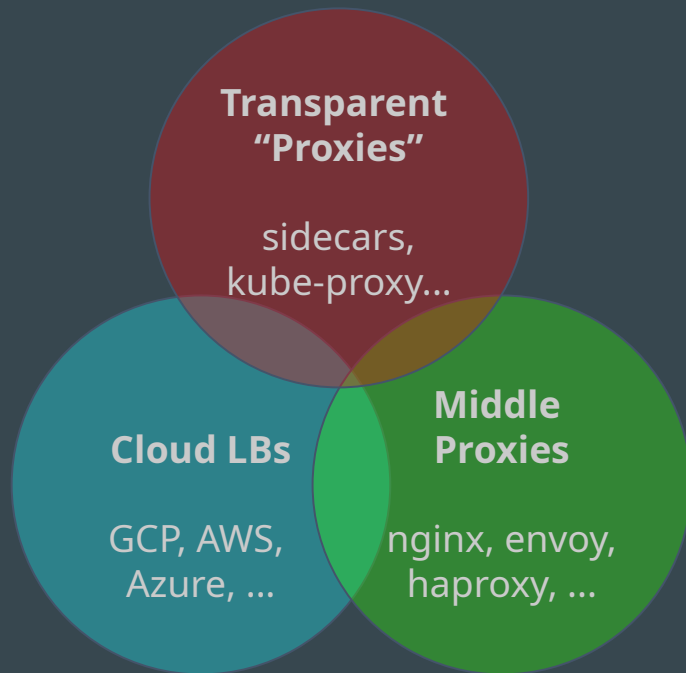
Bowei Du <[bowei@google.com](mailto:bowei@google.com)>

# Evolving Landscape

self-service, single role



split permissions, many



# Goals

Better model the **personas and roles** involved with services and load-balancing.

Support **modern load-balancing features** while maintaining **portability** (or maybe “**predictability**”)

Have **standard mechanisms for extension** for API growth / implementation / vendor-specific behaviors.

## Goals (and how to get there)

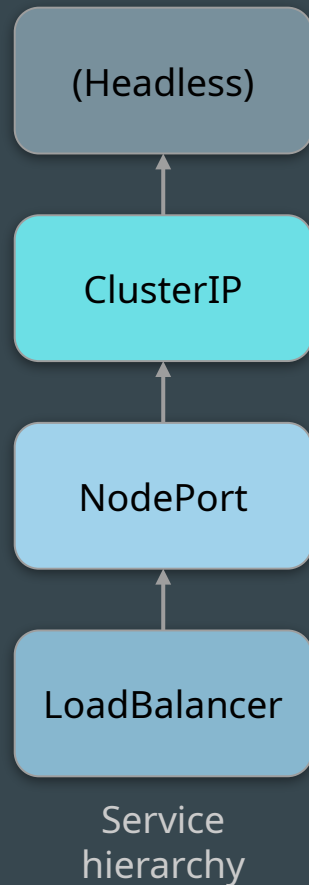
|   |   |  |
|---|---|--|
| Better model the <b>personas and roles</b> involved with services and load-balancing.                                   | → | Resource model + RBAC                            |
| Support <b>modern load-balancing features</b> while maintaining <b>portability</b> (or maybe “ <b>predictability</b> ”) | → | Levels of support, specification and conformance |
| Have <b>standard mechanisms for extension</b> for API growth / implementation / vendor-specific behaviors.              | → | Resource model, polymorphism-like                |

# Modeling: Services

A Service resource describes many things:

- **Method of exposure** (ClusterIP, NodePort, LoadBalancer)
- **Grouping of Backends** (i.e. selector)
- **Attributes** (ExternalTrafficPolicy, SessionAffinity, ...)

Evolving and extending the resource becomes harder and harder due to interactions between fields...



# Personas and Roles

|  |  |  |
|--|--|--|
|  | <b>Infrastructure Provider</b>         | Provides the infrastructure for cluster creation, e.g. cloud provider, internal PaaS team.                                 |
|  | <b>Cluster Operator / NetOps / SRE</b> | Manages the cluster overall once its created.<br>Responsible for overall policies, e.g. which services expose to Internet. |
|  | <b>Application Developer</b>           | Builds the services and applications and defines traffic routing, services.  |

# Designing for RBAC





# Modeling roles: Ingress

Ingress is a self-service model.

IngressClass is created by the infrastructure provider

Application developer manages Ingress + Service;  
Ingress limited to simple L7 descriptions.



# Modeling roles: Gateway

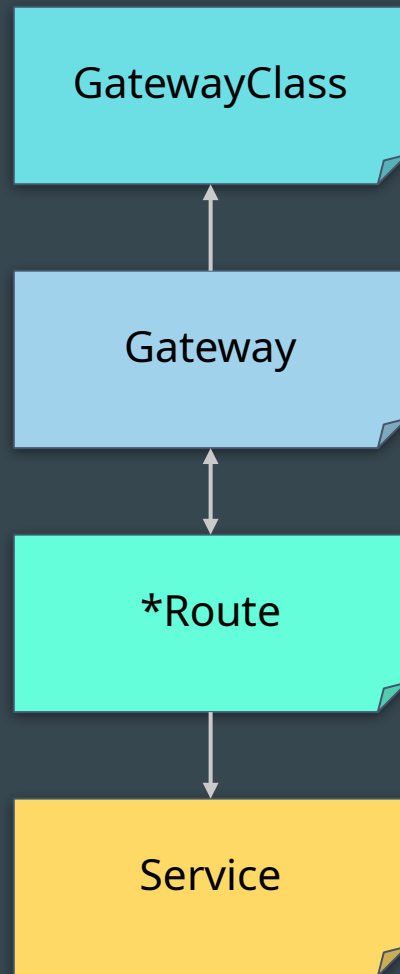
Idea: Decouple along role, concept axes:

Roles:

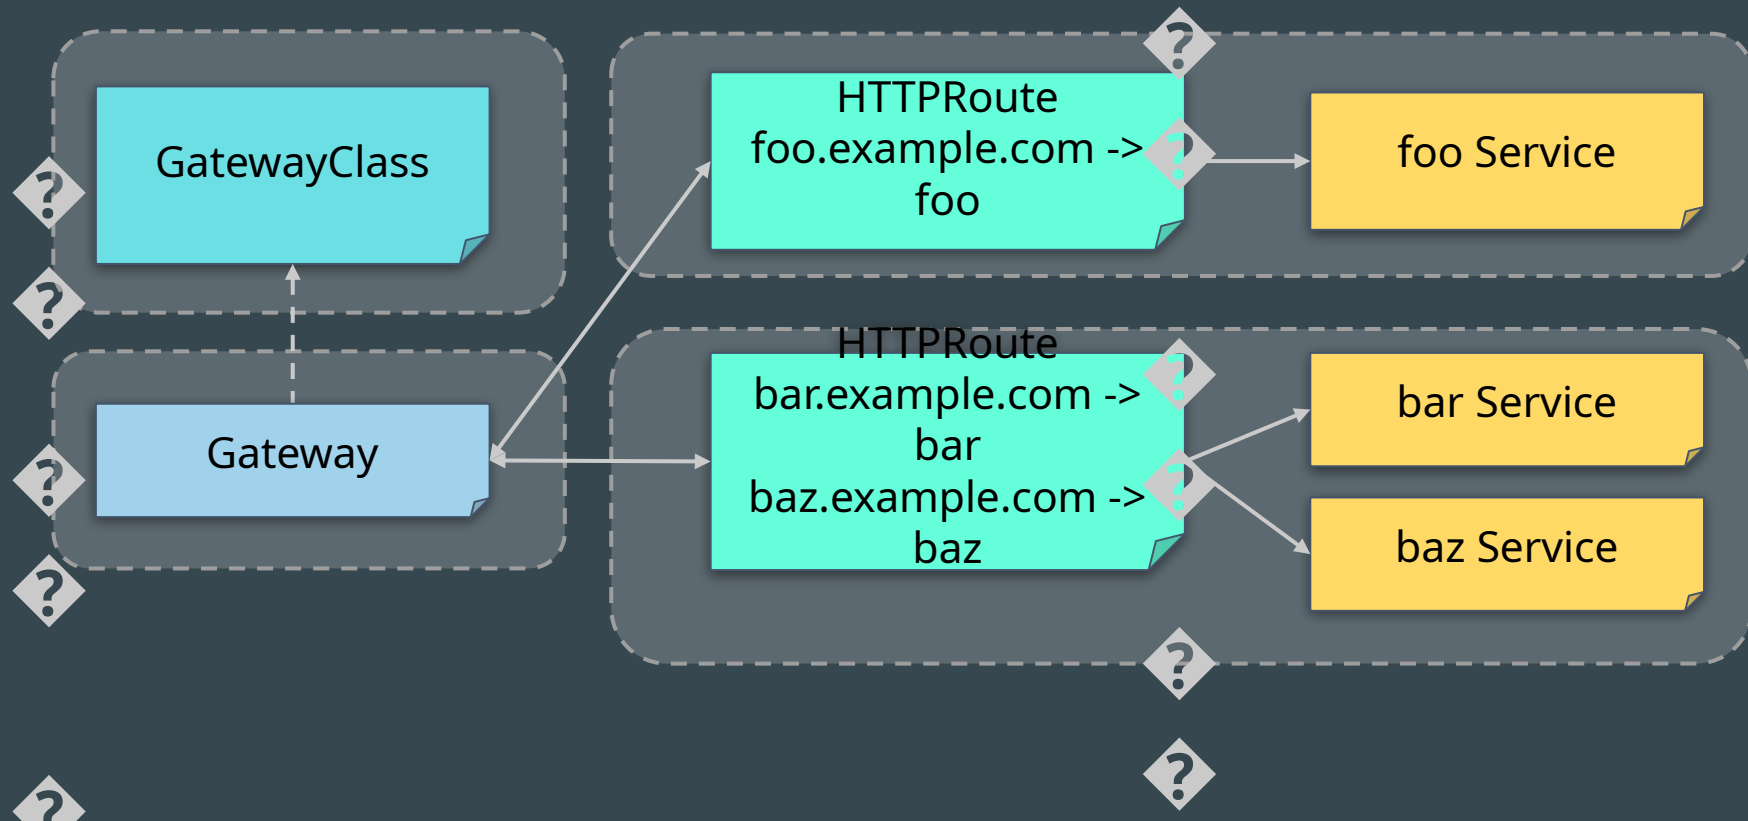
- Infrastructure Provider (GatewayClass)
- Cluster Operator / NetOps (Gateway)
- Application Developer (Routes and Services)

Concepts:

- Exposure and access (Gateway)
- Routing, protocol specific attributes (Routes)
- Grouping, selection (Service)



# API schema



# Portability vs predictability

## Core

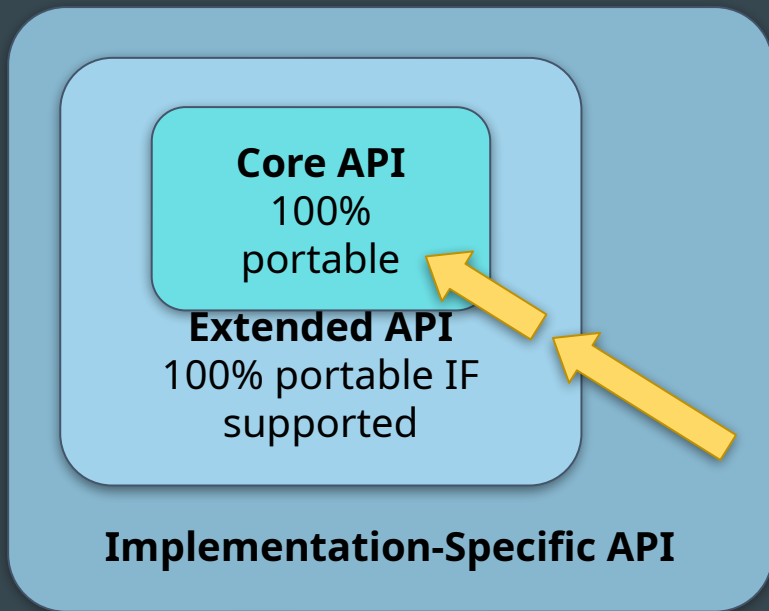
- MUST be supported.

## Extended

- Feature by feature.
- MAYBE supported, but MUST be portable.
- Part of API schema.

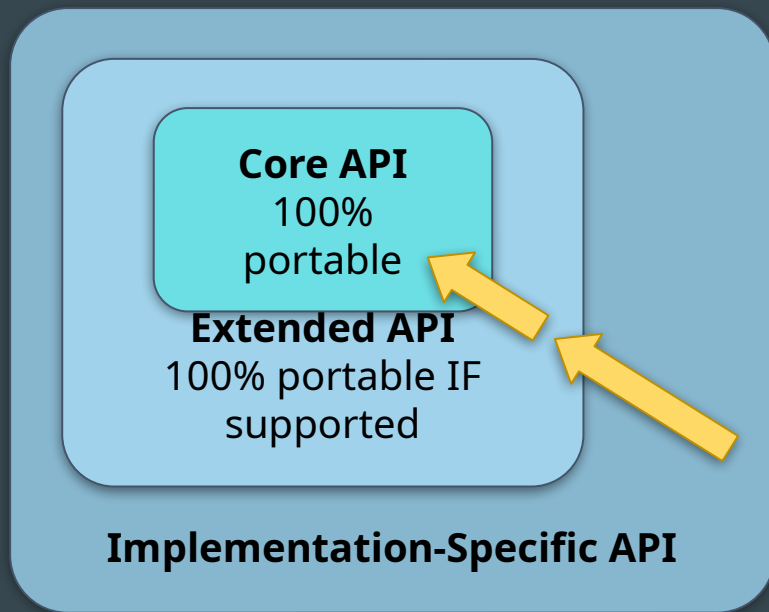
## Implementation-specific

- No guarantee for portability, No k8s API schema.



# Portability vs predictability

- Enforcement by conformance tests
- Extended feature definition **requires** self-contained conformance
- Require all extended features be checkable statically

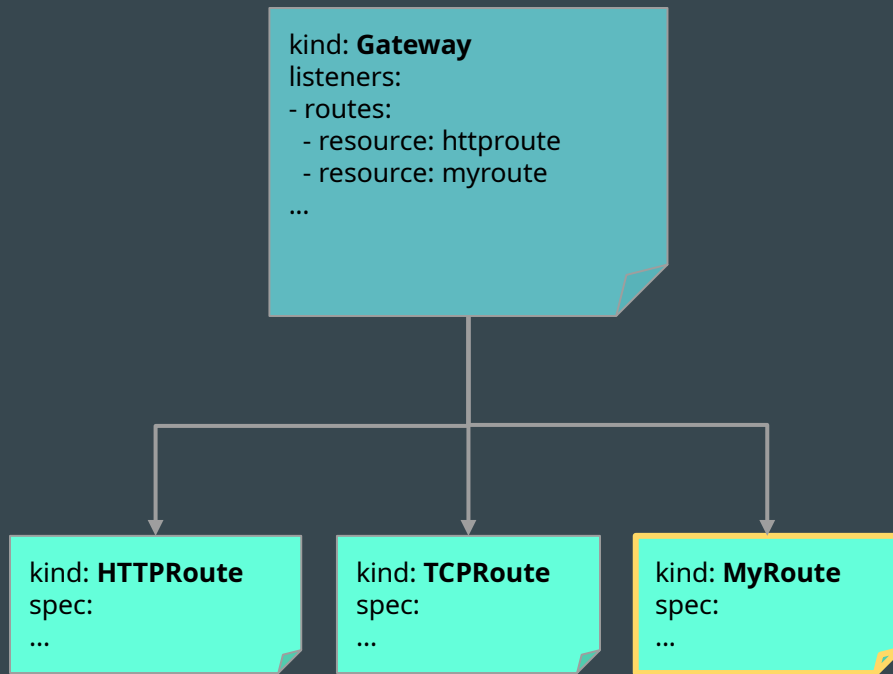


# Extensibility

## Polymorphic object references:

- Gateways can refer to different kinds of Routes: {HTTP, TCP, UDP, SNI} Route
- Route to backends also polymorphic (not just to Services)

Use of Custom Resources (CRs) for extension.



A taste of the API

# User story

**Alice** the IaaS provider

**Bob** the SRE

**Carol** the application developer



# User story: Alice the IaaS

**Alice** the IaaS provider offers two flavors of load-balancers:

- external: Public access to the Internet
- internal: Internal to the VPC

```
kind: GatewayClass
meta:
  name: external
spec:
  controller: alice.io/gw-ctrl
...
```

```
kind: GatewayClass
meta:
  name: internal
spec:
  controller: alice.io/gw-ctrl
...
```

# User story: Bob the SRE

Bob wants:

- Only certain namespaces can deploy external LBs (namespaces with label `internet:external`)
- Anyone can deploy an internal LB
- Anyone can deploy a in-cluster proxy for testing (Bob installs an `acme.io/proxy`)

```
kind: GatewayClass
meta:
  name: external
spec:
  controller: alice.io/gw-ctrl
allowedGatewayNamespaces:
matchLabels:
  internet: external
```

```
kind: GatewayClass
meta:
  name: internal
spec:
  controller: alice.io/gw-ctrl
# Default allows all ns
```

```
kind: GatewayClass
meta:
  name: test
spec:
  controller: acme.io/proxy
# Default allows all ns
```

# User story: Carol the dev

Creates Routes and Services to describe her applications: "store", "checkout"

```
kind: HTTPRoute
meta:
  name: store
spec:
  hostnames: ["store.acme.io"]
  rules:
    - match:
        - path: { value: "/" }
      forwardTo:
        targetRef:
          name: store
    - match:
        - path: { value: "/search" }
      forwardTo:
        targetRef:
          name: search
```

```
kind: Service
meta:
  name: store
...
```

```
kind: Service
meta:
  name: search
```

```
kind: HTTPRoute
meta:
  name: checkout
spec:
  hostnames: ["checkout.acme.io"]
  rules:
    - match:
        - headers:
            type: Exact
            values: {"canary": "y"}
      forwardTo:
        targetRef:
          name: checkout-canary
    ...
```

```
kind: Service
meta:
  name: checkout
...
```

namespace: carol

# User story: Carol the dev

To test out her applications, she creates a Gateway of class “test” in her namespace:

```
kind: Gateway
meta:
  name: test-gw
spec:
  gatewayClassName: test
  listeners:
    - protocol: HTTP # take defaults
  routes:
    resource: httproutes
    routeNamespaces:
      onlySameNamespace: true
```

namespace: carol

```
kind: HTTPRoute
meta:
  name: checkout
```

```
kind: Service
meta:
  name: checkout
```

```
kind: HTTPRoute
meta:
  name: store
```

```
kind: Service
meta:
  name: store
```

```
kind: Service
meta:
  name: search
```

# User story: Carol the dev

A Gateway is a “request” for an LB. It may be underspecified and the controller for the class will fill in the blanks:

```
kind: Gateway
meta:
  name: test-gw
spec:
  gatewayClassName: test
  listeners:
    - protocol: HTTP # take defaults
  routes:
    resource: httproutes
    routeNamespaces:
      onlySameNamespace: true
```

namespace: carol

```
kind: HTTPRoute
meta:
  name: checkout
```

```
kind: Service
meta:
  name: checkout
```

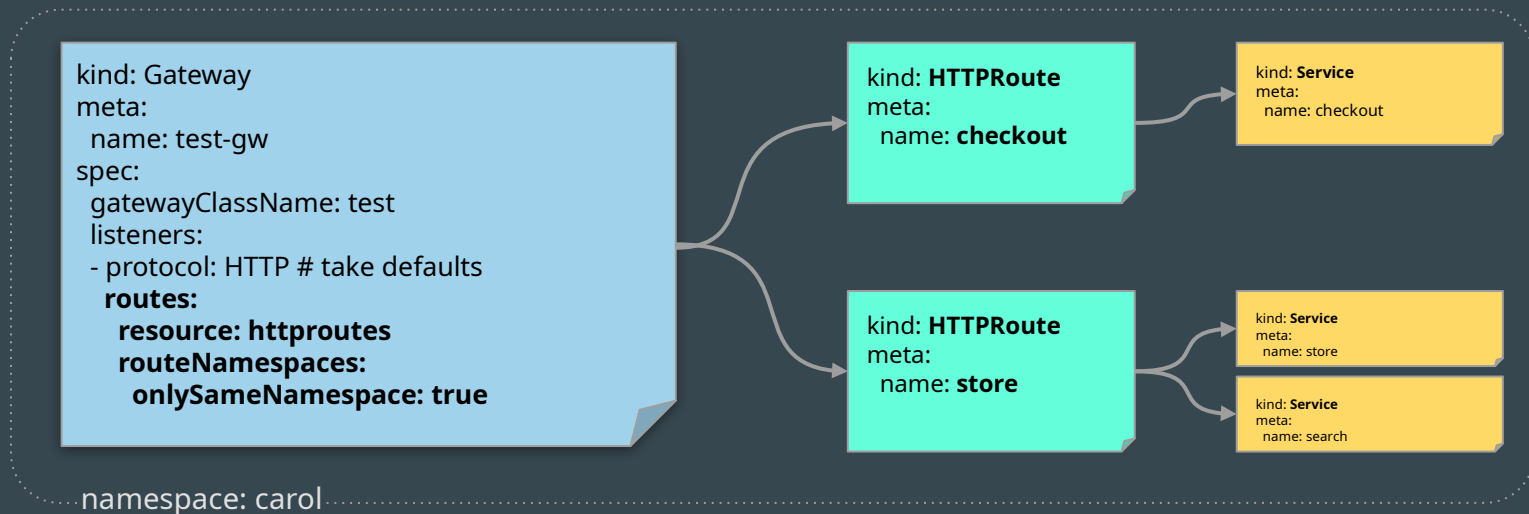
```
kind: HTTPRoute
meta:
  name: store
```

```
kind: Service
meta:
  name: store
```

```
kind: Service
meta:
  name: search
```

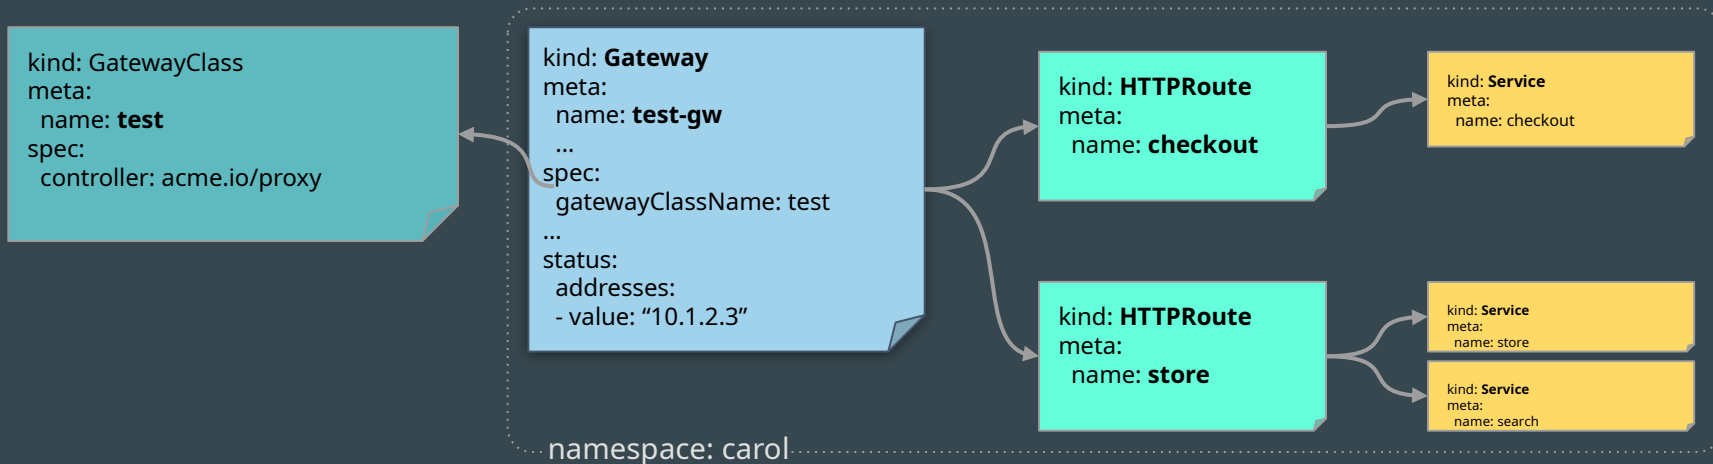
# User story: Carol the dev

A Gateway serves Routes. Carol uses the simplest configuration which is to pick up routes in her own namespace:



# User story: Carol the dev

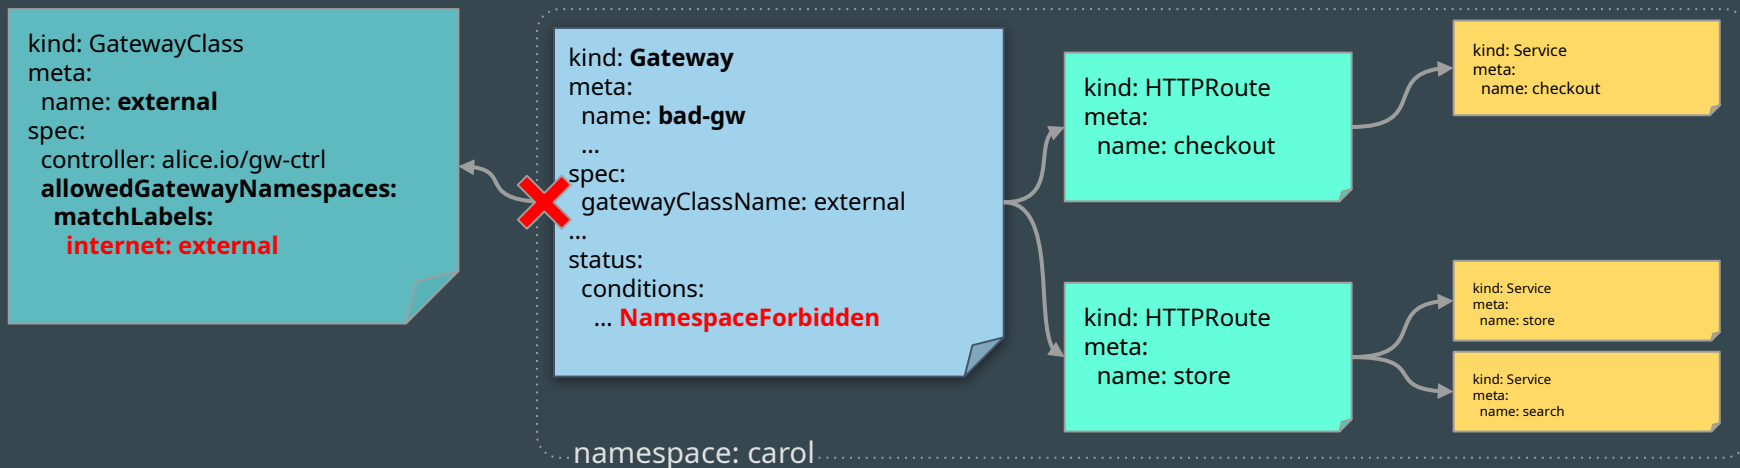
Creating the Gateway wires up the app(s) and the acme.io/proxy starts serving traffic:



```
carol@acme$ curl -H store.acme.io http://10.1.2.3
Hello, world! ....
```

# User story: Carol the dev

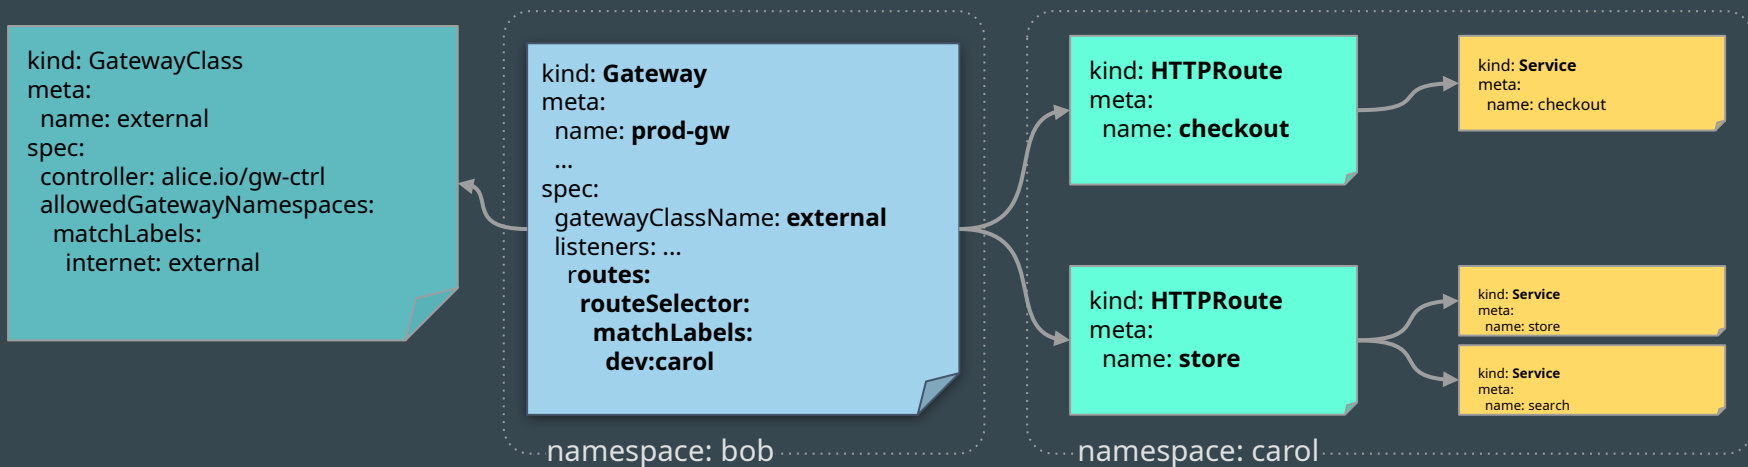
Carol is cannot use the external class and serve production traffic - she isn't allowed:





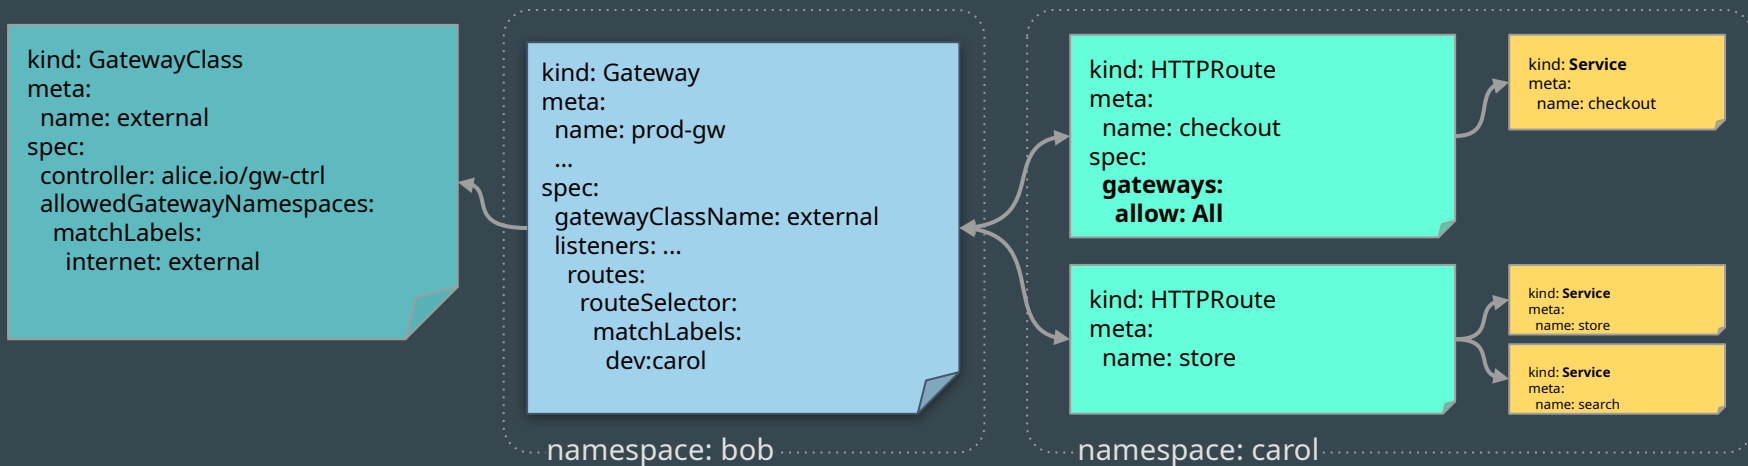
# User story: Bob the SRE

Bob, who manages external Gateways for the cluster, can reference Carol's Routes and put them into production:



# User story: Carol the dev

Carol also needs to allow her Route to the prod Gateway - otherwise anyone with a Gateway can start serving traffic (default is to allow in the same namespace):



```
carol@acme$ curl -H http://store.acme.io
Hello, world! ...
```

## User story: recap

GatewayClass: Supports multiple classes of load-balancing, reflecting capabilities of the IaaS or deployed infrastructure

Gateway: defines how your app is **exposed** (e.g. VIP:port, what kind of proxy used)

Route: defines the **routing** of your app for a given **protocol**

Service: defines your Backends (**grouping**)

## User story: recap

Supports multiple classes of load-balancing, reflecting capabilities of the IaaS or deployed infrastructure.

Access to GatewayClasses is controlled.

Gateways can reference and aggregate cross namespace Routes.  
Access is controlled using handshake between Gateway and Route

# Conflicts

Merging Routes into Gateways can result in conflicts (same host, same path)...

# Conflicts principles

## Do no harm

- Don't break things that are working
- Drop as little traffic as possible

## Be consistent

- Provide a consistent behavior when conflicts occur ( $A \text{ then } B \equiv B \text{ then } A$ )
- Prefer more specific matches to less specific ones
- Decide using stable properties: creation timestamp, canonical order `<namespace, name>`

## Be clear

- Make it clear which configuration has been chosen
- Communicate conflicts via object status

# Extension points

GatewayClass parameters for overall LB configuration.

Gateway.Listeners have an ExtensionRef to customize Listener properties

Routes

- Custom filters via ExtensionRef
- Backends can be more than Services

Feedback needed!

# What in the Alpha?

Basic applications, data types:

- GatewayClass, Gateway
- {HTTP, TCP, UDP, SNI} Route
- TLS (HTTPS, TLS)

Implementations:

- Merging style (multiple Gateways hosted on in-cluster proxy)
- Provisioning/Cloud (Gateways mapped to externally managed resources)

What remains to be done:

- Feedback (users, users, users)!
- Conformance tests
- Delegation within Routes

→ Gain confidence towards Beta



Demo:

<https://bit.ly/2RY9Xzr>

# Thanks for coming!

## Kubernetes SIG-NETWORK subproject

### Project homepage

[github/kubernetes-sigs/service-apis](https://github.com/kubernetes-sigs/service-apis)


### How to contribute

[kubernetes-sigs.github.io/service-apis/community/](https://kubernetes-sigs.github.io/service-apis/community/)

### Meetings

Wednesday, Thursdays

Alternating times AM/PM Pacific, check calendar,  
meeting code: "77777"

 Kubernetes Service API Evolution

## How to contribute

This page contains links to all of the meeting notes, design docs and related discussions around the APIs.

## Communications

Major discussions and notifications will be sent on the [SIG-NETWORK mailing list](#).



We also have a [Slack channel \(sig-network-service-apis\)](#) on k8s.io for day-to-day questions, discussions.


## Meetings

Meetings discussing the evolution of the service APIs will alternate times to accommodate participants from various time zones. This calendar includes all Service APIs meetings as well as any other SIG-Network meetings.

**k8s-sig-network**build-label: calendar\_200921.08\_p0, build-cl: 332867598, Megastore, io.gammas-us.web.21

Today

  September 2020

 Print **Week** Month Agenda

| Sun | Mon             | Tue   | Wed                            | Thu | Fri | Sat |
|-----|-----------------|-------|--------------------------------|-----|-----|-----|
| 30  | 31              | Sep 1 | 2                              | 3   | 4   | 5   |
|     | 1pm Network Pol |       | 3pm [SIG-NETW( 10:30am [SIG-NE |     |     |     |

Fin

# References

- <https://www.youtube.com/watch?v=Ne9UJL6irXY> (very old)
- [Kubecon San Diego 2019 - Evolving the Kubernetes Ingress APIs to GA and Beyond \[PUBLIC\]](#)
- [Virtual KubeCon EU 2020: SIG-Network Intro and DeepDive](#)
- <https://github.com/kubernetes-sigs/service-apis>