



Introducing Jaeger 1.0

Yuri Shkuro (Uber Technologies)

CNCF Webinar Series, Jan-16-2018

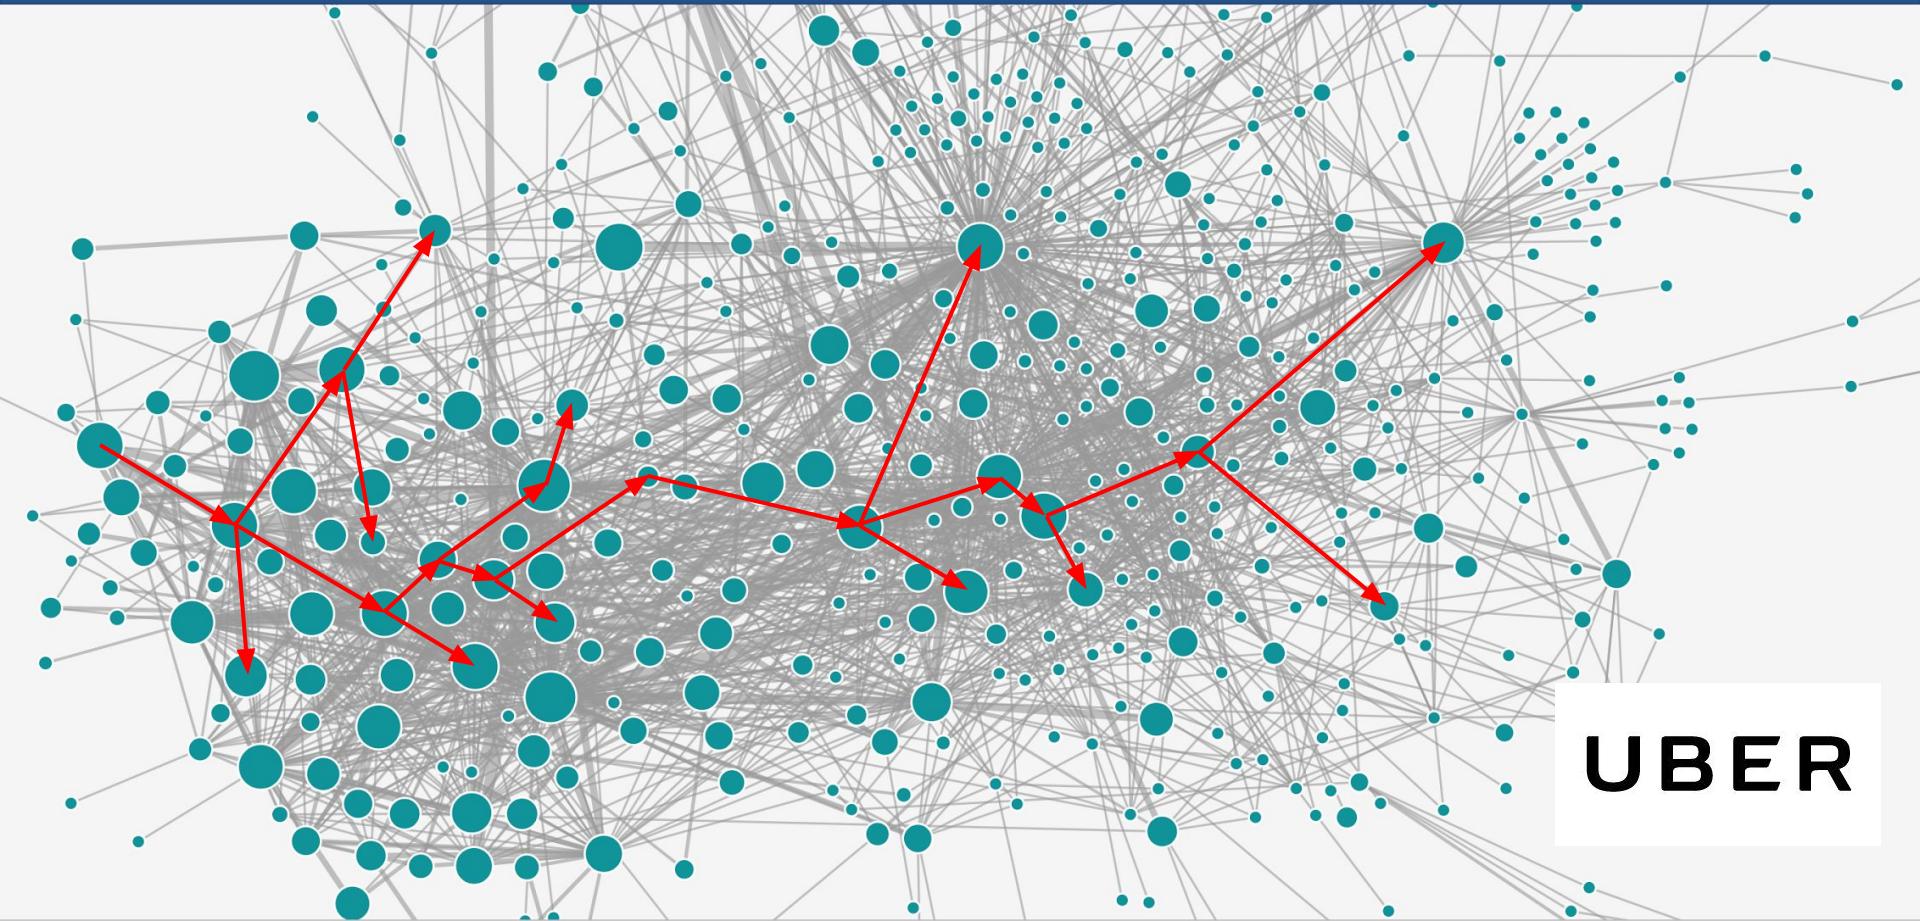
Agenda

- What is distributed tracing
- Jaeger in a HotROD
- Jaeger under the hood
- Jaeger v1.0
- Roadmap
- Project governance, public meetings, contributions
- Q & A

About

- Software engineer at Uber
 - NYC Observability team
- Founder of Jaeger
- Co-author of OpenTracing Specification

BILLIONS times a day!



How do we know what's going on?

We use MONITORING tools

Metrics / Stats

- Counters, timers, gauges, histograms
- Four golden signals
 - utilization
 - saturation
 - throughput
 - errors
- Prometheus, Grafana

Logging

- Application events
- Errors, stack traces
- ELK, Splunk, Sentry

Monitoring tools must “tell stories” about your system

How do you debug this?

```
2017/12/04 21:30:37 scanning error: bufio.Scanner: token too long
```

WHAT IS THE CONTEXT?

Metrics and logs don't cut it anymore!

Metrics and logs are

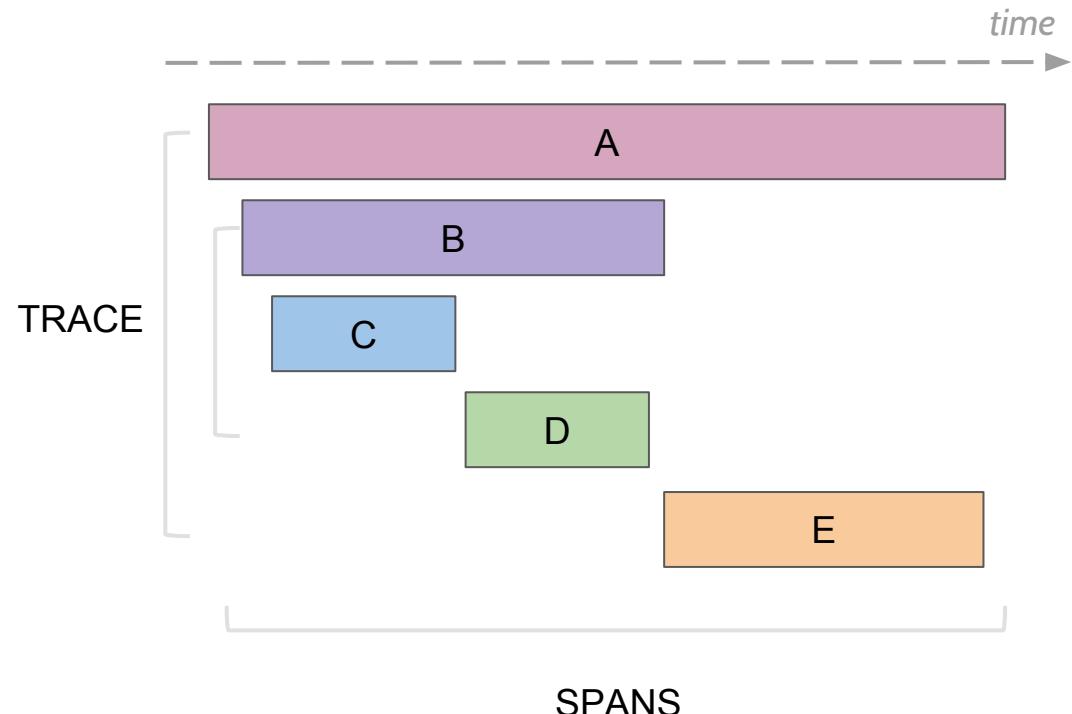
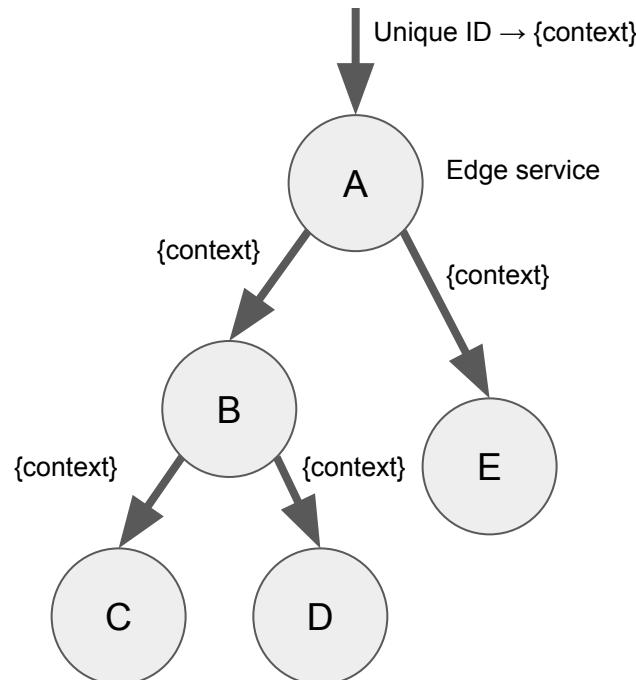
- per-instance
- missing the context

It's like debugging
without a stack trace

We need to monitor
distributed transactions



Distributed Tracing In A Nutshell





Let's look at some traces

demo time: <http://bit.do/jaeger-hotrod>



Distributed Tracing Systems

distributed
transaction
monitoring

performance
and latency
optimization

root cause
analysis

service
dependency
analysis

distributed context propagation



CLOUD NATIVE
COMPUTING FOUNDATION

Jaeger under the hood
Architecture, etc.



Jaeger - /'yāgər/, noun: hunter

- Inspired by Google's Dapper and OpenZipkin
- Started at Uber in August 2015
- Open sourced in April 2017
- Official CNCF project since Sep 2017
- Built-in OpenTracing support
- <http://jaegertracing.io>



UBER

Community

- 10 full time engineers at Uber and Red Hat
- 80+ contributors on GitHub
- Already used by many organizations
 - including Uber, Symantec, Red Hat, Base CRM, Massachusetts Open Cloud, Nets, FarmersEdge, GrafanaLabs, Northwestern Mutual, Zenly

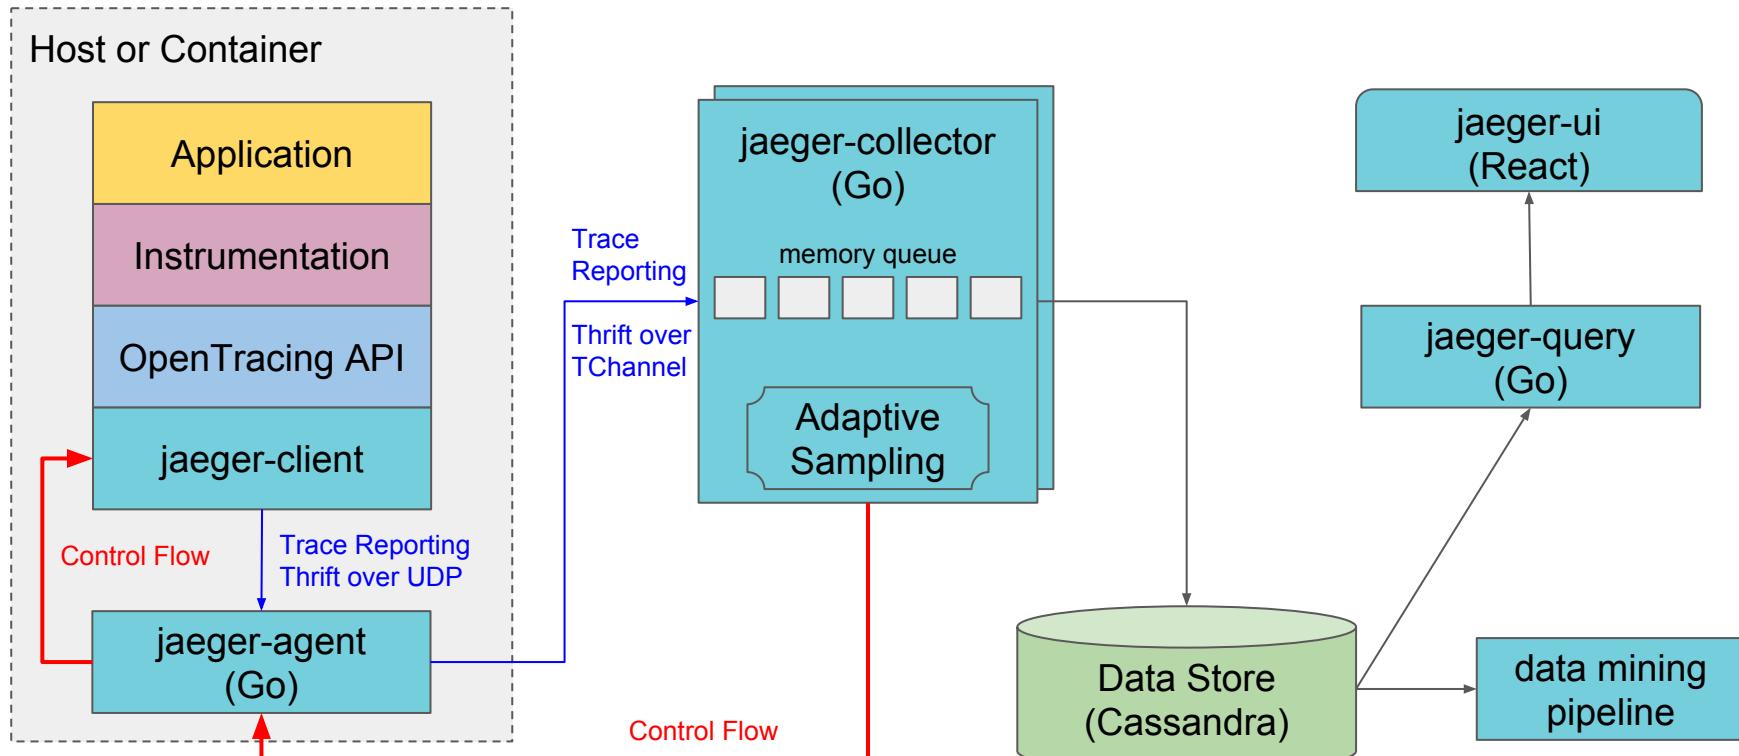


Technology Stack

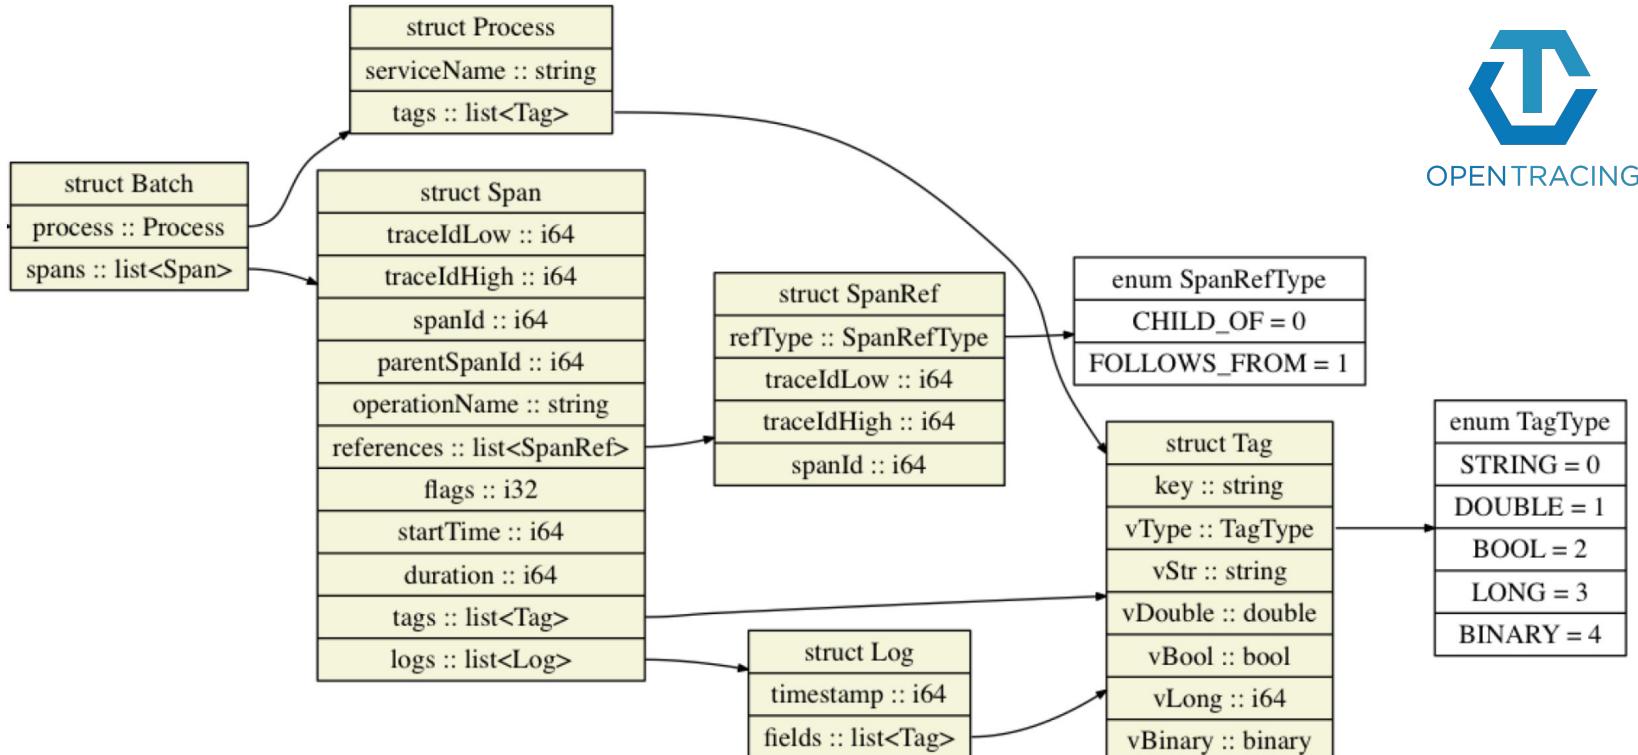
- Backend components in Go
- Pluggable storage
 - Cassandra, Elasticsearch, memory, ...
- Web UI in React/Javascript
- OpenTracing instrumentation libraries



Architecture



Data model



Understanding Sampling

Tracing data can exceed business traffic.

Most tracing systems sample transactions:

- **Head-based sampling:** the sampling decision is made just before the trace is started, and it is respected by all nodes in the graph
- **Tail-based sampling:** the sampling decision is made after the trace is completed / collected



Jaeger 1.0

Released 06-Dec-2017



Jaeger 1.0 Highlights

Announcement: <http://bit.do/jaeger-v1>

- Multiple storage backends
- Various UI improvements
- Prometheus metrics by default
- Templates for Kubernetes deployment
 - Also a Helm chart
- Instrumentation libraries
- Backwards compatibility with Zipkin

Multiple storage backends

Official

- Cassandra 3.4+
- Elasticsearch 5.x, 6.x
- Memory storage

Experimental (by community)

- InfluxDB, ScyllaDB, AWS DynamoDB, ...
- <https://github.com/jaegertracing/jaeger/issues/638>

Jaeger UI

- Improved performance in all screens
- Viewing large traces (e.g. 80,000 spans)
- Keyboard navigation
- Minimap navigation, zooming in & out
- Top menu customization

Zipkin drop-in replacement

Collector can accept Zipkin spans:

- JSON v1/v2 and Thrift over HTTP
- Kafka transport not supported yet

Clients:

- B3 propagation
- Jaeger clients in Zipkin environment

Monitoring

- Metrics
 - `--metrics-backend`
 - prometheus (default), expvar
 - `--metrics-http-route`
 - `/metrics` (default)
- Scraping Endpoints
 - Query service - API port 16686
 - Collector - HTTP API port 14268
 - Agent - sampler port 5778



CLOUD NATIVE
COMPUTING FOUNDATION

Roadmap

Things we are working on



Adaptive Sampling

- APIs have endpoints with different QPS
- Service owners do not know the full impact of sampling probability

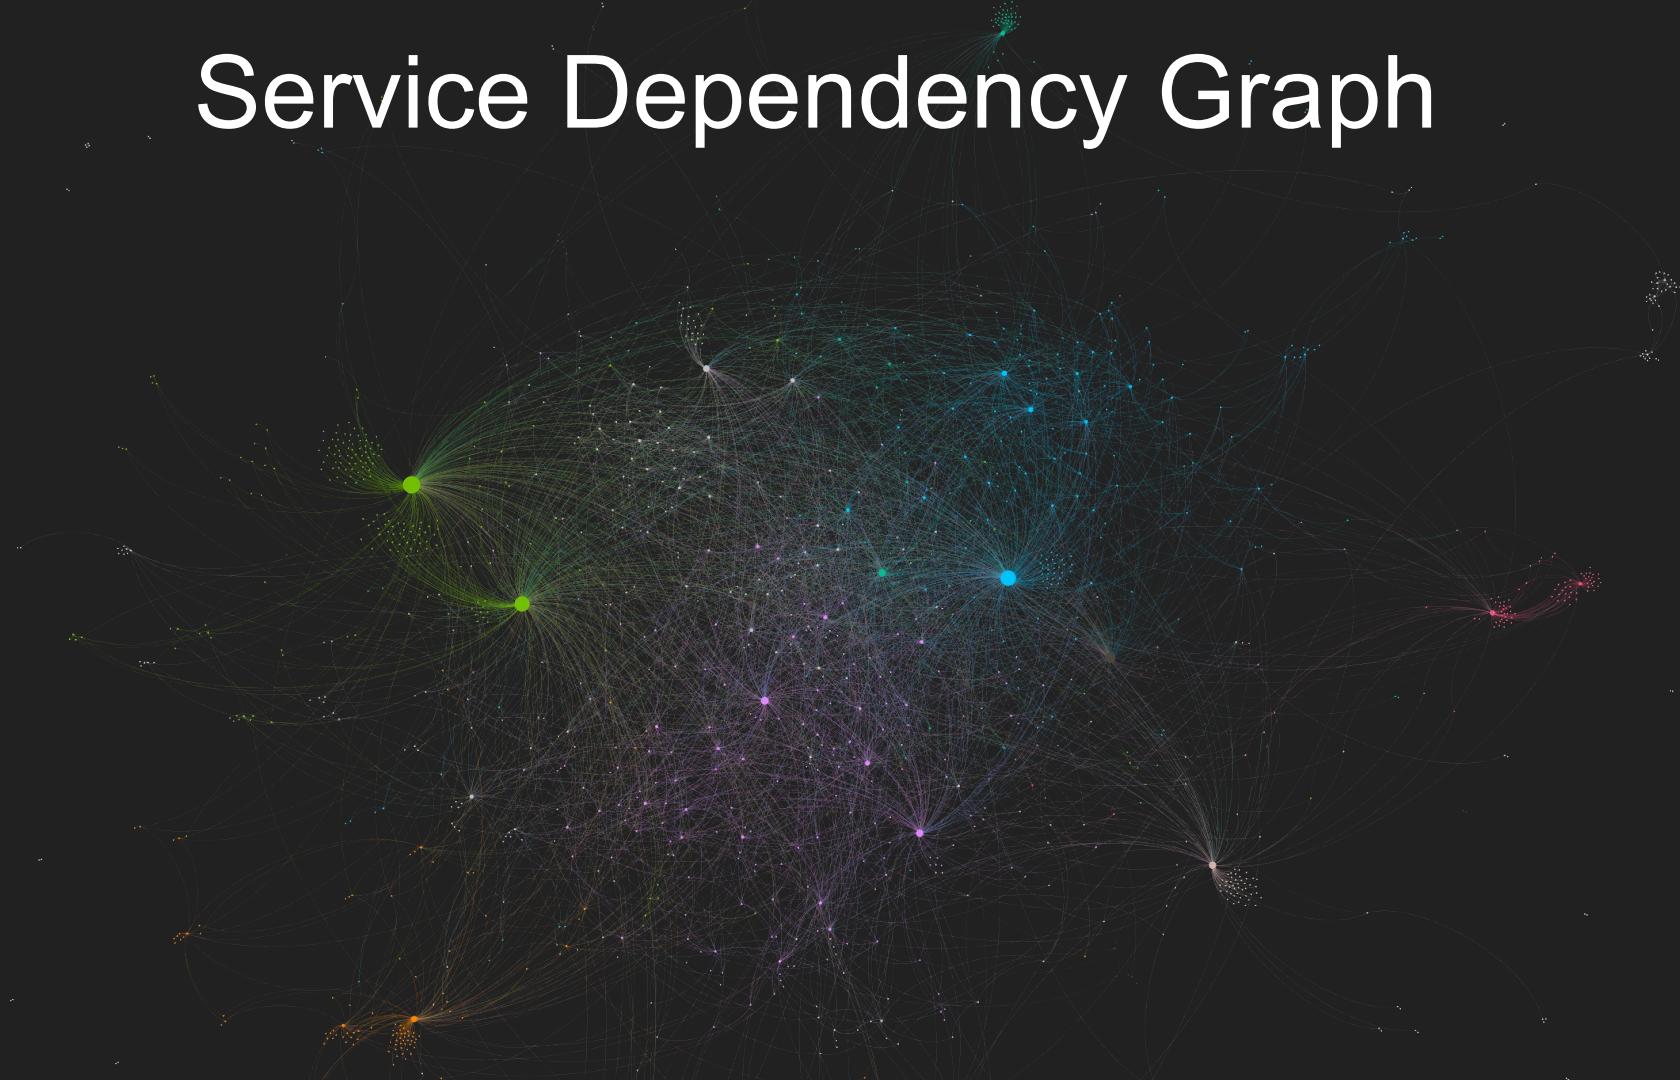
Adaptive Sampling is per service + endpoint,
decided by Jaeger backend based on traffic

Data Pipeline

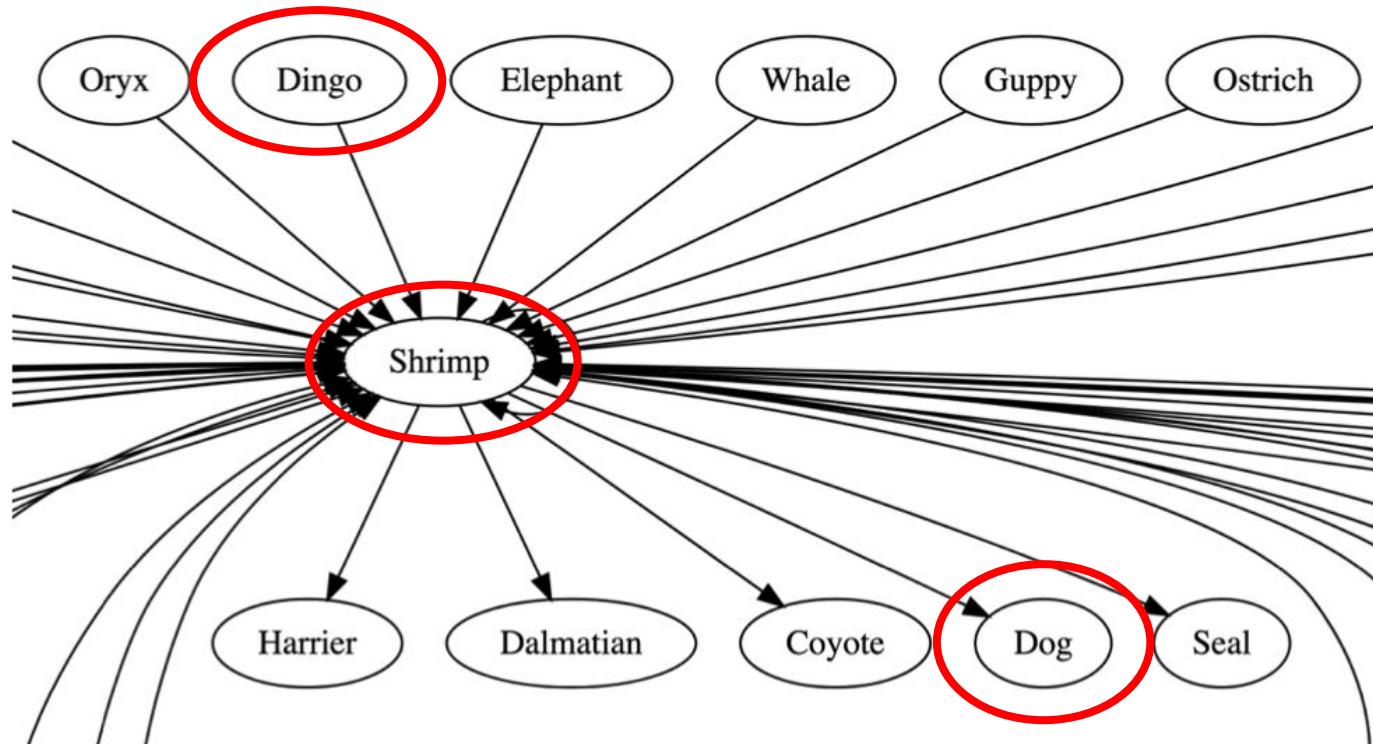
- Based on Kafka and Apache Flink
- Support aggregations and data mining
- Examples:
 - Pairwise dependency graph
 - Path-based, per endpoint dependency graph
 - Latency histograms by upstream caller



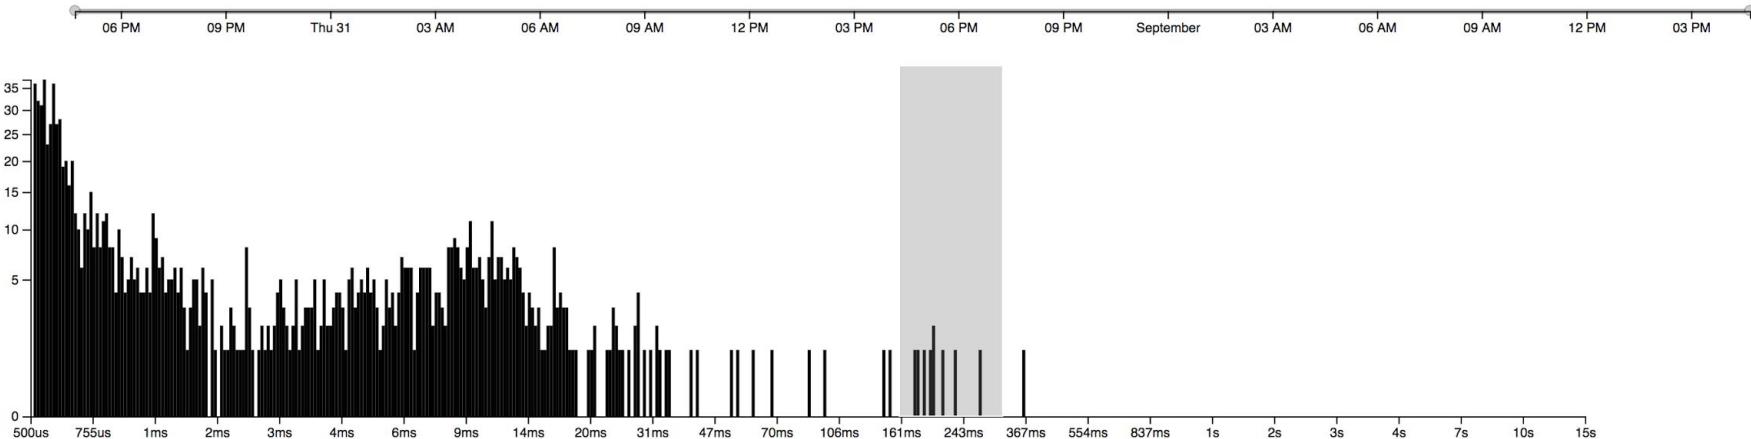
Service Dependency Graph



Does Dingo Depend on Dog?



Latency Histogram



Representative Jaeger Traces: 6 Traces Found!

Estimated Call Count by Endpoint

4	Postmaster::testEmail
4	GET
1	POST
0	Meta::health

Estimated Call Count by Upstream

5	[postmaster]Postmaster::testEmail
3	[cli-user1]
1	[cli-user2]
0	[cfgtools]
0	[tcheck]
0	[health-agent]



Project & Community

Contributors are welcome



Contributing

- ! **iOS client library** help wanted outreachy
#576 opened on Dec 2, 2017 by yurishkuro
- ! **Flaky test TestPeerListManager_updatePeers** bug help wanted
#574 opened on Nov 30, 2017 by yurishkuro
- ! **Unify metric name and format in client libraries** help wanted
#572 opened on Nov 30, 2017 by NeoCN
- ! **Document external interfaces** documentation help wanted outreachy
#456 opened on Oct 6, 2017 by jpkrohling  Core Infra Best ...

Contributing

- Agree to the Certificate of Origin
- Sign all commits (`git commit -s`)
- Test coverage cannot go ↓ (backend - 100%)
- Plenty of work to go around
 - Backend
 - Client libraries
 - Kubernetes templates
 - Documentation

References

- GitHub: <https://github.com/jaegertracing>
- Chat: <https://gitter.im/jaegertracing/>
- Mailing List - jaeger-tracing@googlegroups.com
- Blog: <https://medium.com/jaegertracing>
- Twitter: <https://twitter.com/JaegerTracing>
- Bi-Weekly Online Community Meetings



Q & A

Open Discussion