



The Open Source Observability Playbook

Hen Peretz - Head of Solutions Engineering @epsagon

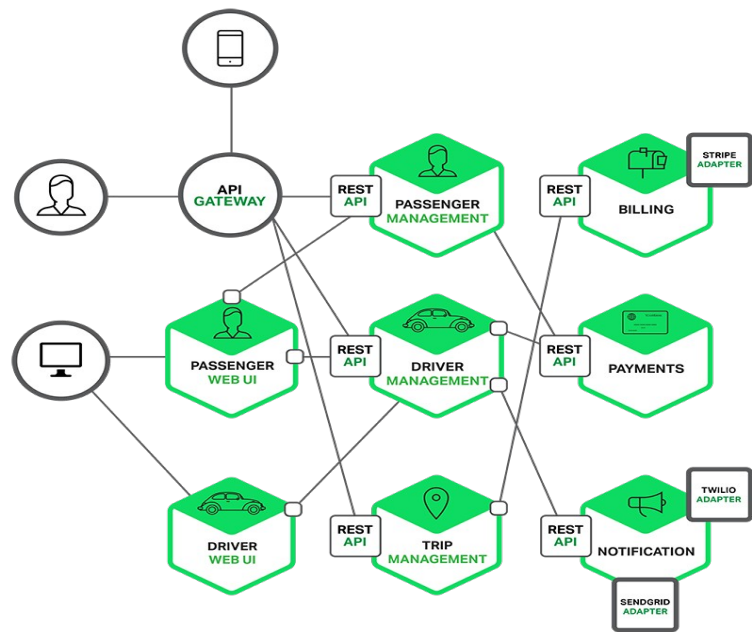
What We'll Discuss Today

Old-School Observability methods

How to achieve observability

Logging, monitoring best practices

Open-source monitoring landscape



Why Monitoring?



Make sure our business works

What Should We Monitor?

4 golden signals from [SRE book](#) by Google:

Latency

Traffic

Errors

Saturation

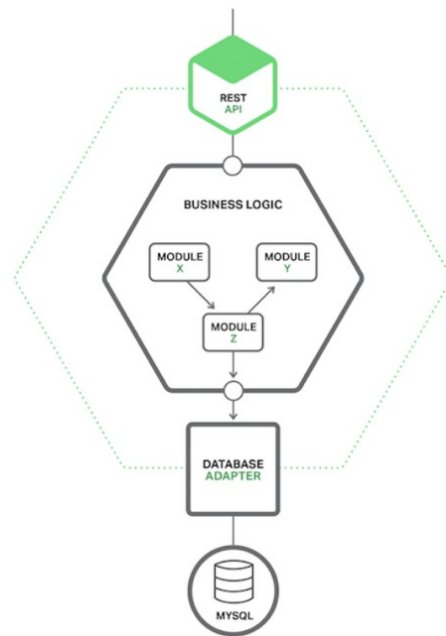


Old-School Monitoring

Agent based

Collects only host data

Collects only metrics



Old-School Troubleshooting


We need more debug data → logs

Old-School Logging

Agent based

Dumps locally or remotely

Collects only logged data






Fast Forward into the Future

New Era for Applications: Cloud + Microservices


Annual revenue of Amazon Web Services from 2013 to 2018 (in million U.S. dollars)




Source:
Amazon
© Statista 2019

Additional information:
Worldwide; Amazon; 2013 to 2018

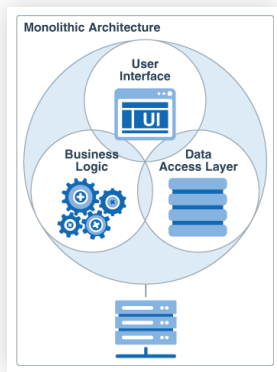
AWS Lambda adoption



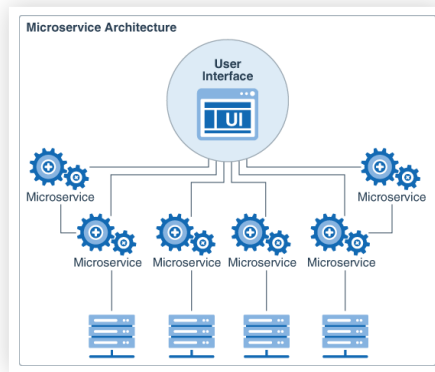
Containerization timeline



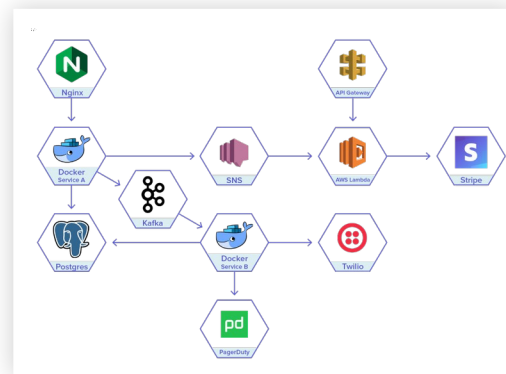
The Rise of Microservices in the Cloud



Host-based
Monolithic



Host-based
Distributed




Abstracted-host
Highly Distributed

**Extremely hard to monitor and
troubleshoot**

Challenges for Engineering and DevOps

- **Troubleshooting**
basic logs and metrics the right tool for highly distributed applications?
- **Monitoring**
my application working properly”?
- **Development**
not sure what’s currently running in production. How can I build new services?



epsagon

The Three Pillars of Observability

Combining metrics, logs, and traces for observability is the **only** way to understand complex environments

Metrics tell us the “what”

Logs tell us the “why”

Traces tell us the “where”



epsagon

Logging Best Practices

Print out JSONed logs with metadata
(service name, stage, etc.)

Automate the process of logging

Index the fields you're using

Logging Best Practices - Setup





Logging Best Practices

Demo

Monitoring Best Practices

Aggregate all metrics into a unified dashboard

Define your critical metrics (thresholds)

Use custom business metrics

Monitoring Best Practices

Monitor application metrics:

Avg. duration of calls to an HTTP API

Minimum number of calls to a message queue

Number of 500/400 errors



Monitoring Best Practices

Demo

Something Is Still Missing

How do we correlate between metrics and logs

How do we correlate between data in different services



epsagon



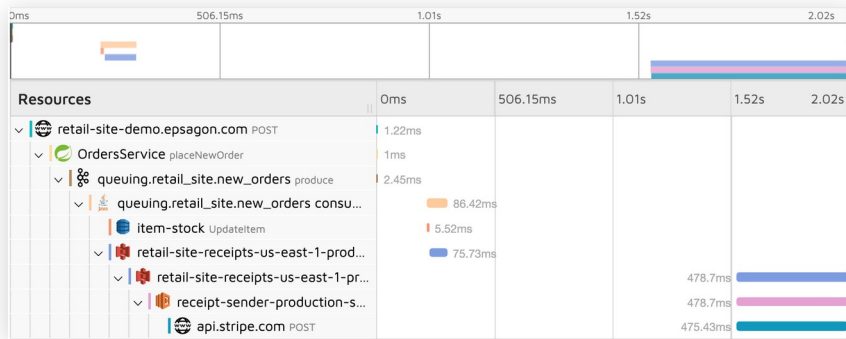
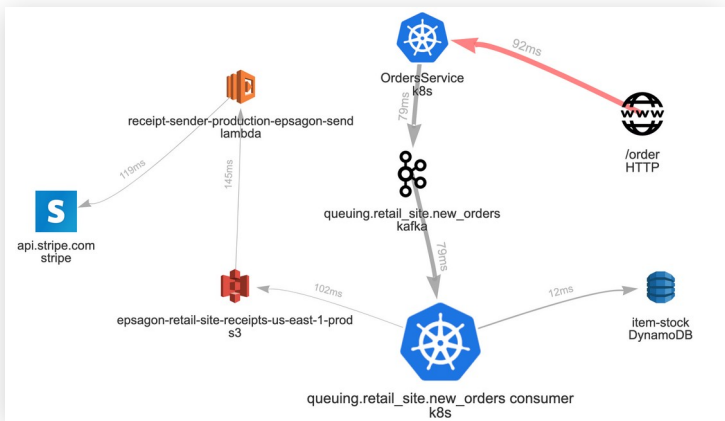
Distributed Tracing DIY

Why Distributed Tracing Is Critical Today

The Only Way to Understand Cloud-Native Workloads

Due to high complexity and the need for manual instrumentation, distributed tracing remained an approach viable only for very **tech savvy** companies

“A trace tells the story of a transaction or workflow as it propagates through a distributed system”



Distributed Tracing Landscape



+



OpenCensus

=



Generating Traces

Ingestion and visualization



JAEGER



CLOUD NATIVE
COMPUTING FOUNDATION



epsagon

Distributed Tracing DIY - Generating Traces

Instrument every call

(AWS-SDK, http, postgres, Flask, ...)

A **span** for every request and response

Add **context** to every span

Inject and **Extract** IDs in relevant calls

```
def handle_request(request):
    span = before_request(request, opentracing.global_tracer())
    # store span in some request-local storage using Tracer.scope_manager,
    # using the returned 'Scope' as Context Manager to ensure
    # 'Span' will be cleared and (in this case) 'Span.finish()' be called.
    with tracer.scope_manager.activate(span, True) as scope:
        # actual business logic
        handle_request_for_real(request)

def before_request(request, tracer):
    span_context = tracer.extract(
        format=Format.HTTP_HEADERS,
        carrier=request.headers,
    )
    span = tracer.start_span(
        operation_name=request.operation,
        child_of(span_context))
    span.set_tag('http.url', request.full_url)

    remote_ip = request.remote_ip
    if remote_ip:
        span.set_tag(tags.PEER_HOST_IPV4, remote_ip)

    caller_name = request.caller_name
    if caller_name:
        span.set_tag(tags.PEER_SERVICE, caller_name)

    remote_port = request.remote_port
    if remote_port:
        span.set_tag(tags.PEER_PORT, remote_port)

    return span
```



epsagon



Distributed Tracing DIY

Demo

Distributed Tracing DIY - Ingestion & Visualization

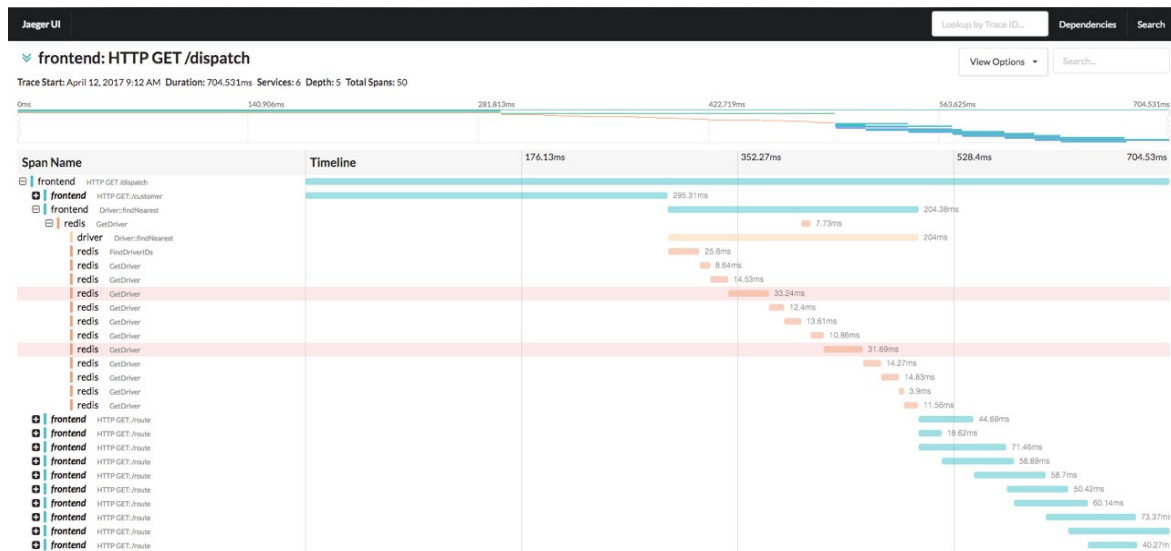
Ingestion according to our scale (millions? billions?)

Index context and tags for easy search

Visualize traces (timeline, graph)

Set alerts

Distributed Tracing DIY - Ingestion & Visualization



epsagon

Tagging Traces

Adding tags for search and aggregations:

Identifiers – user ID, customer ID, device ID

Flow control – event type, business logic

Business metrics – items in cart, minutes watched

Tracing with payload

Search an event according to:

HTTP headers or body - IDs

Key or query in SQL/NoSQL

Response payload from HTTP call



Tracing As a Glue

Trace \rightarrow Logs

Trace \leftrightarrow Environment



Best Practices for Observability

Best Practices for Observability

Automated setup and zero maintenance

Support any environment

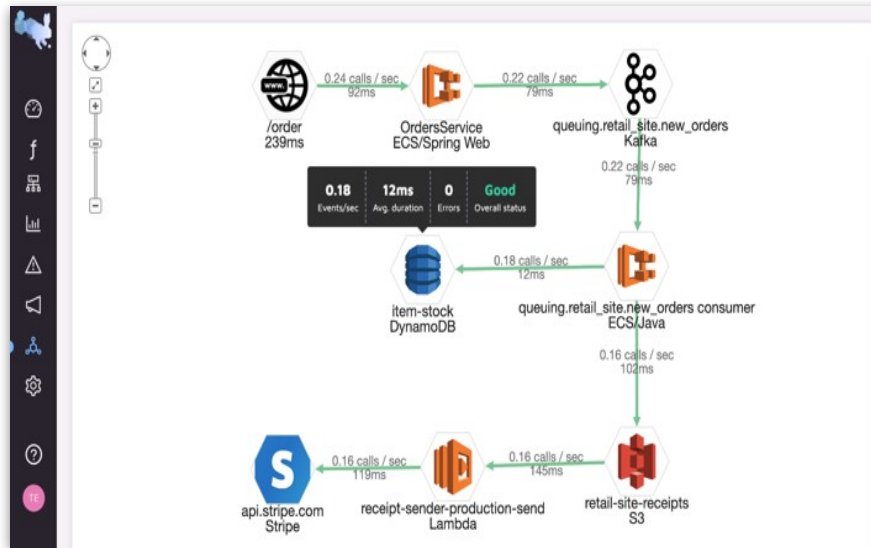
(K8s, Serverless, ECS, AKS, On Prem)

Connects **every request**

in a transaction

Search and analyze your **data**

Helps to quickly pinpoint problems



Summary

Modern applications requires more than just monitoring

Distributed tracing becomes a crucial component in such environments

Stop implementing your own solutions unless needed



Q&A
