

Kubernetes native two-level resource management for AI workloads

Diana Arroyo and Alaa Youssef
IBM Research

October 2020





Agenda

Why the Kubernetes scheduler is not enough

Desired capabilities

MCAD - open source project

How it works

Demo

Call to action

Why the Kubernetes scheduler is not enough

Desired capabilities

MCAD - open source project

How it works

Demo

Call to action

BLOG

Why the Kubernetes Scheduler Is Not Enough for Your AI Workloads

CNCF Member Blog Post

 Posted on August 10, 2020

KubeCon + CloudNativeCon sponsor guest post from Alaa Youssef, manager of the Container Cloud Platform at IBM Research



Recent CNCF Blog Post

<https://www.cncf.io/blog/2020/08/10/why-the-kubernetes-scheduler-is-not-enough-for-your-ai-workloads/>

Some Characteristics of AI Workloads

- Multiple concurrent learners or executors
- Co-location/affinity preferences
- Massively parallel; big number of short running tasks
- Resource hungry
- Elastic jobs
- Elapsed time may vary while total consumed resources remain the same
- Increasing trend of interactive use cases



TensorFlow

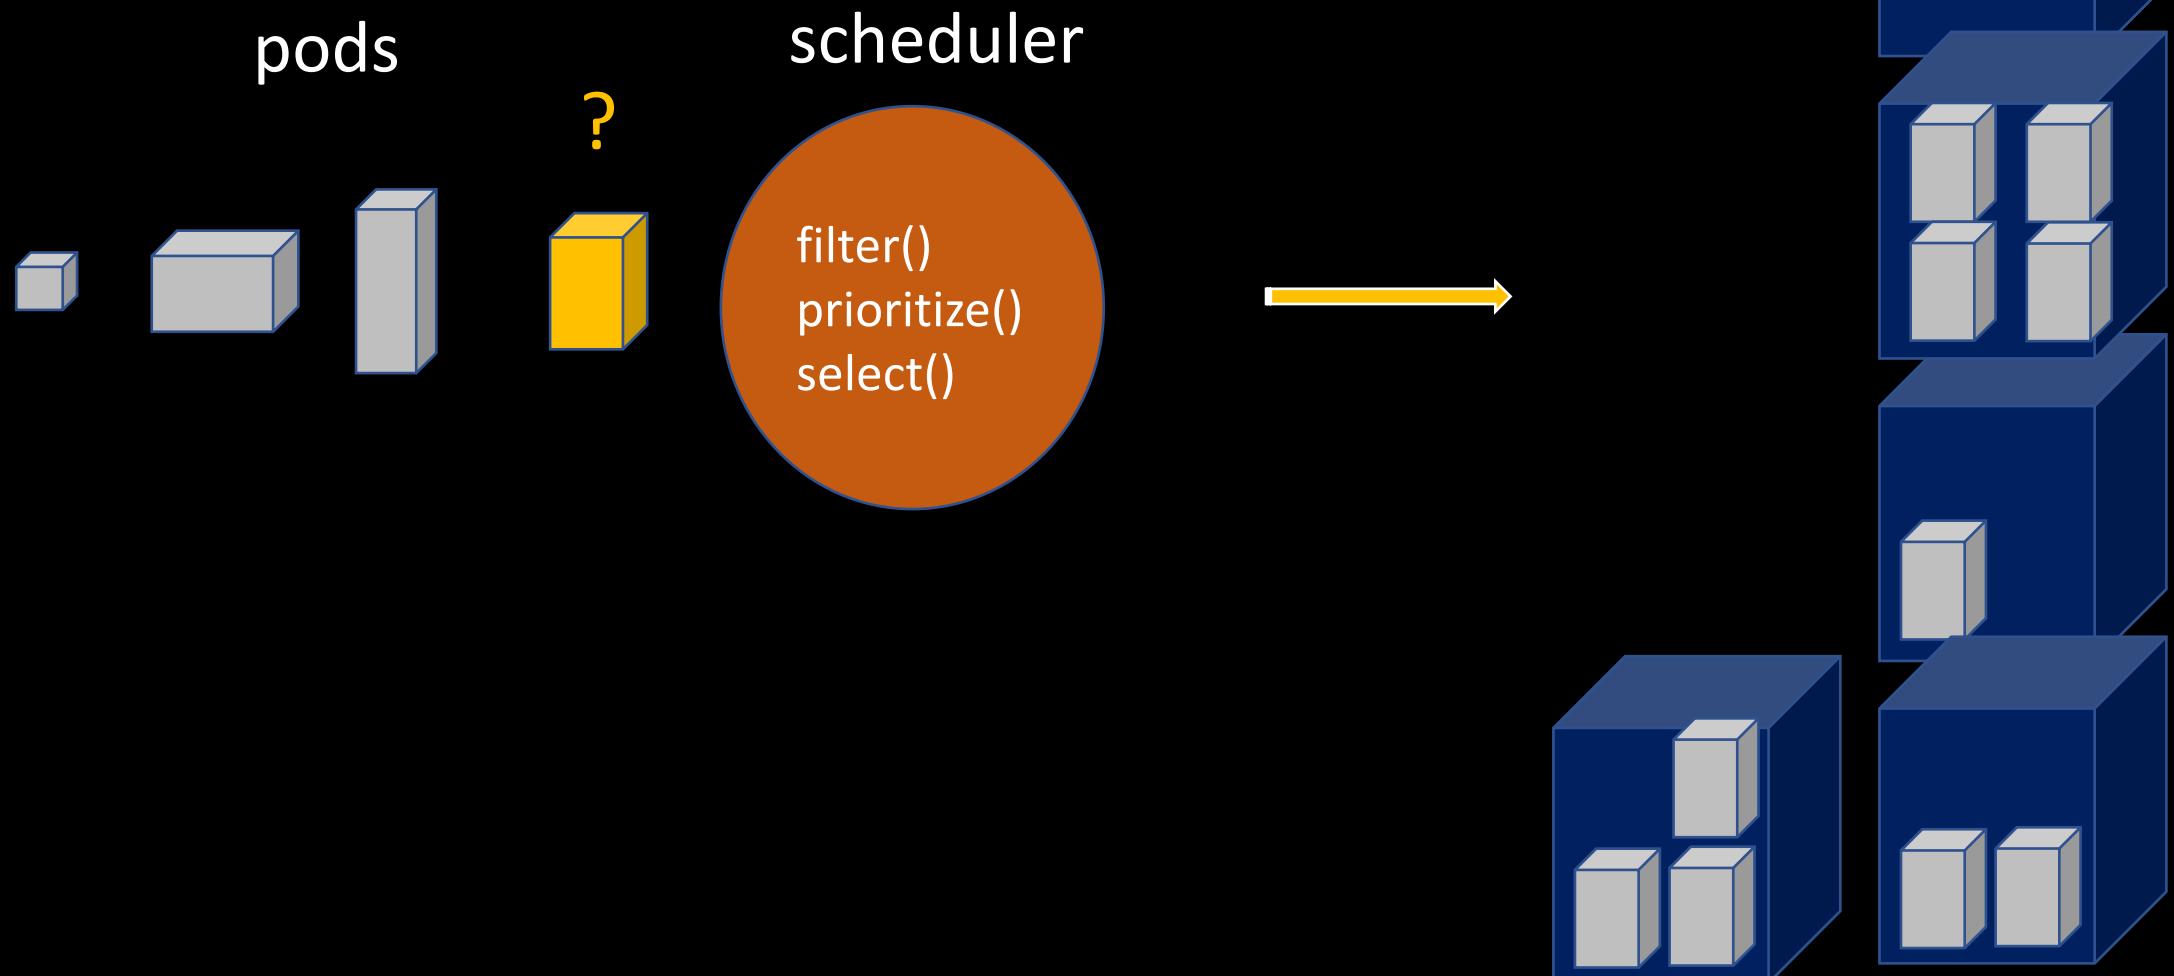


Kubeflow

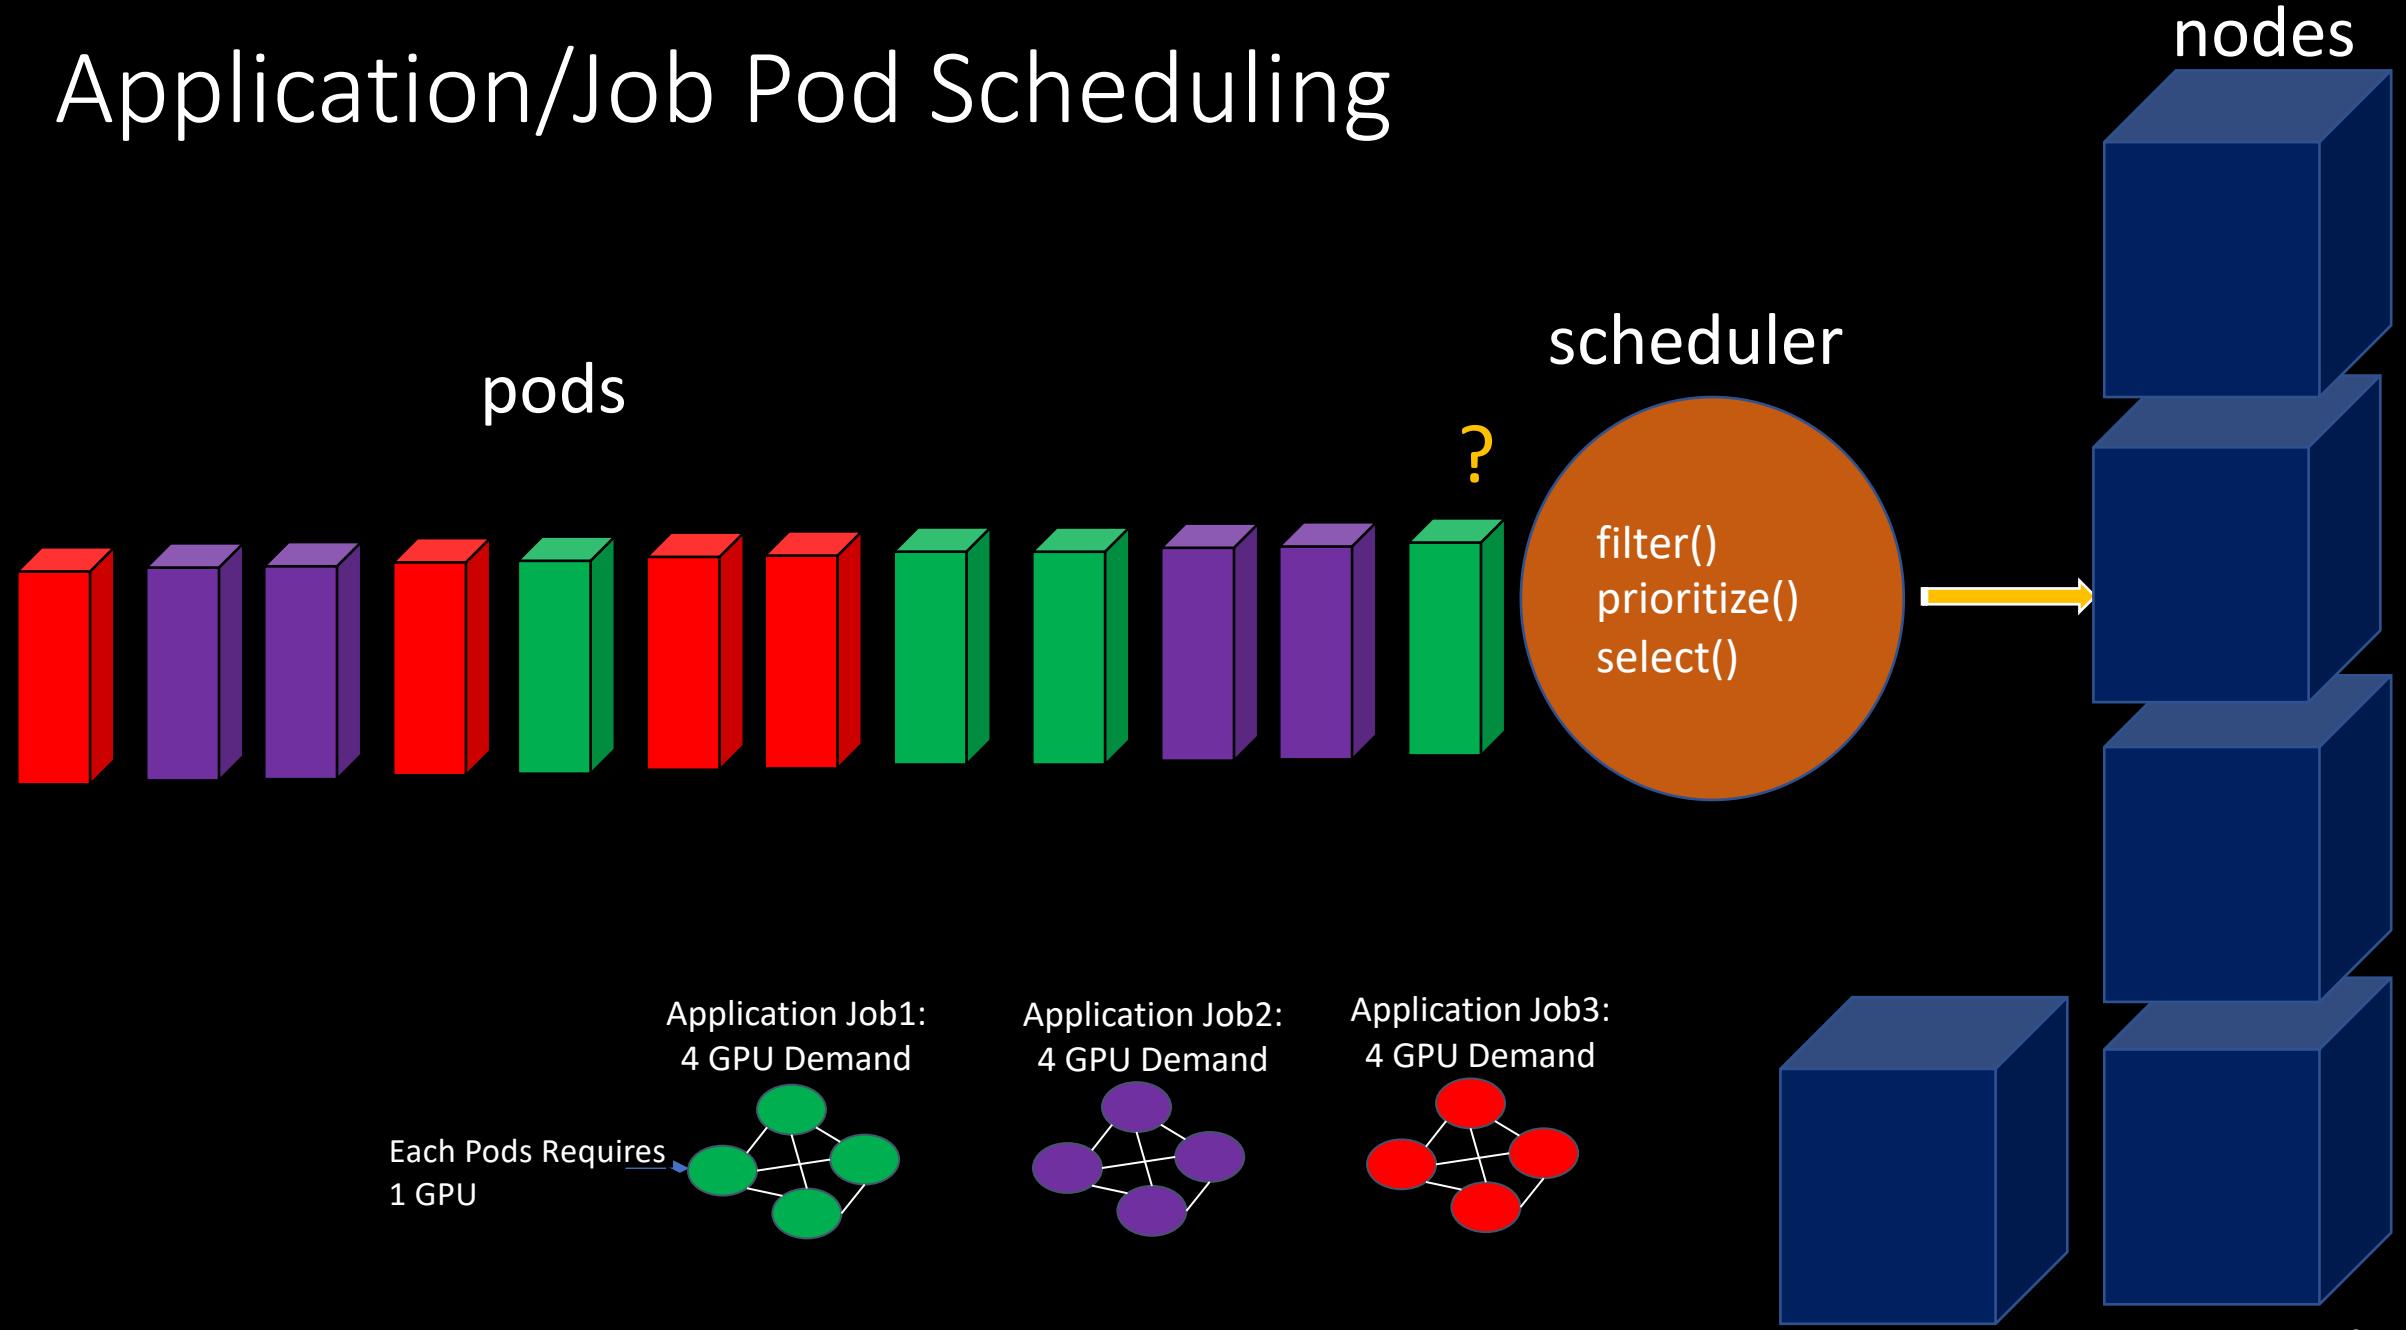
Managing Jobs vs Tasks

- Tasks map to PODs
- Jobs are composed of many tasks
- At which level do you set...
 - Priorities
 - Classes of service
 - Quotas
- Task, job, user, org, service, ...?
- At which level do you queue, allocate resources, preempt, ...
- What happens to your scheduler when 1000s of PODs are pending?
... Overwhelming!!

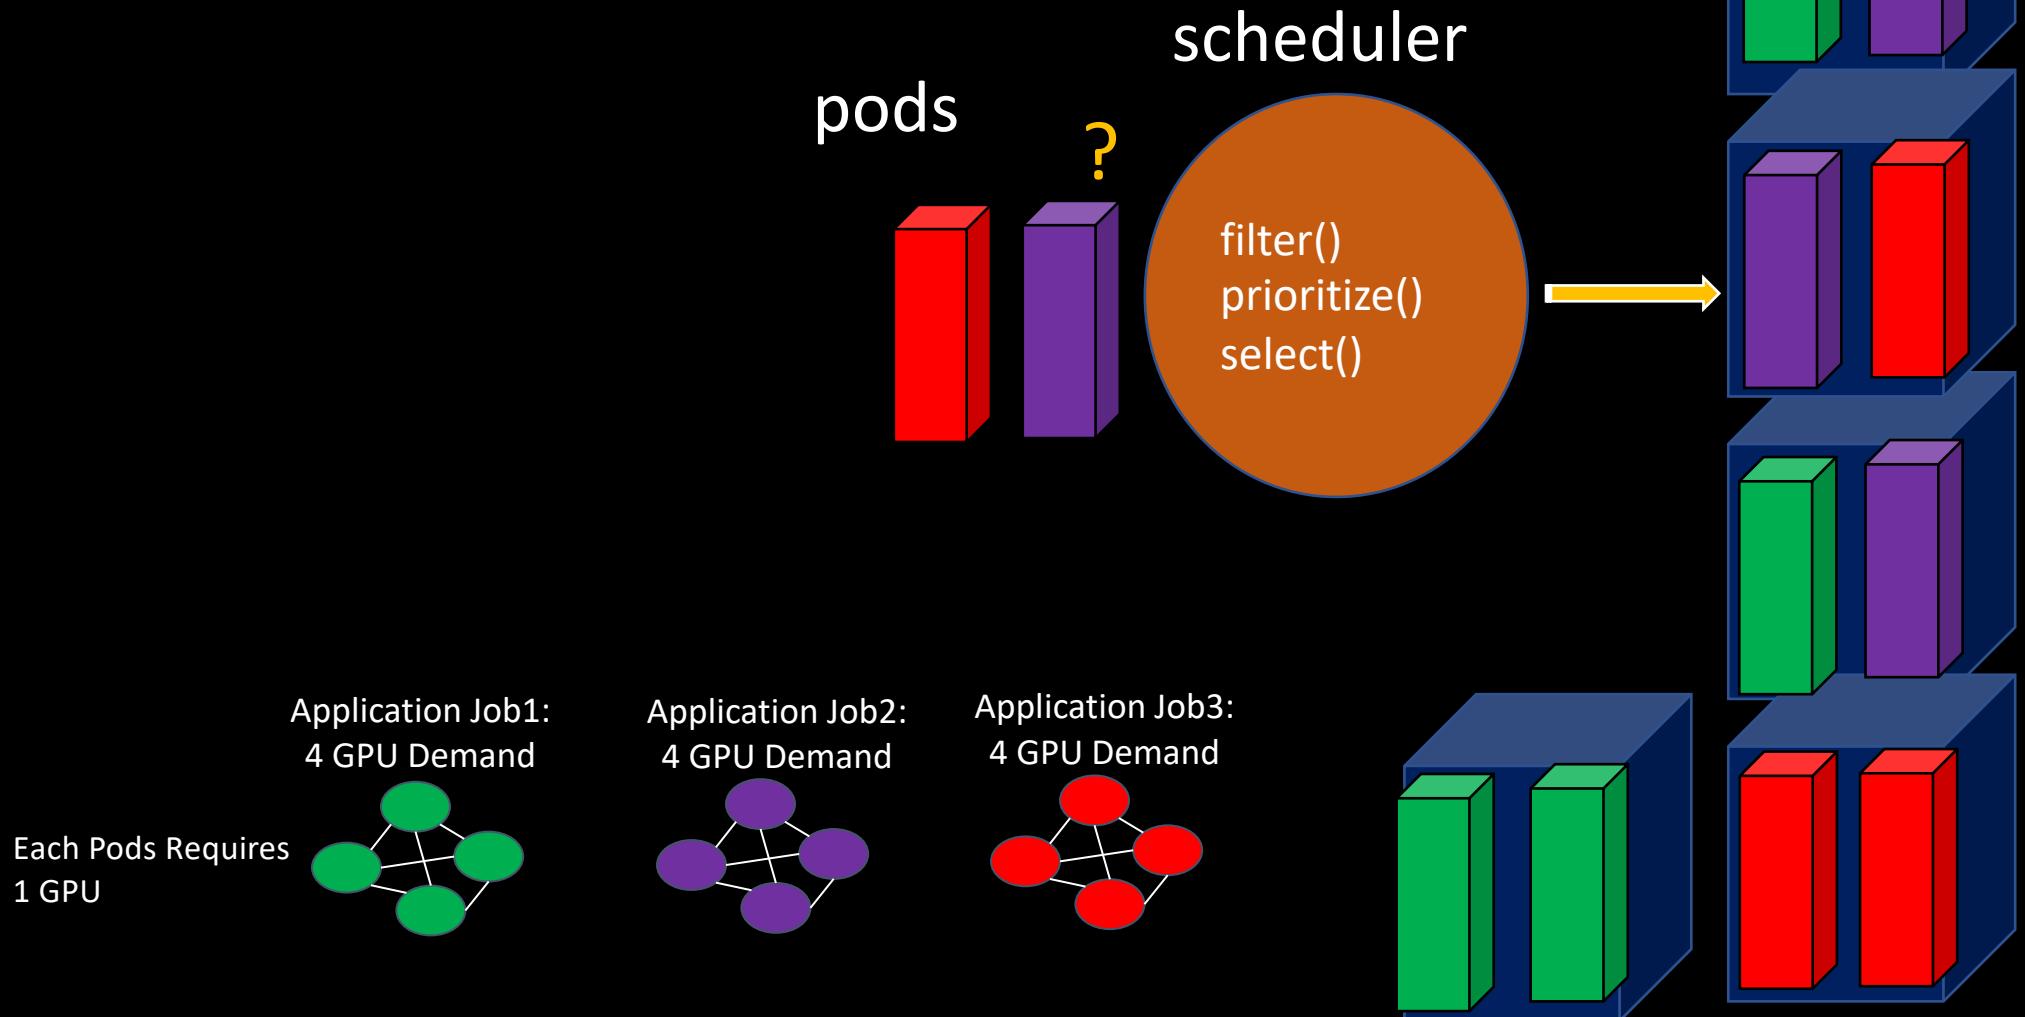
Kubernetes Pod Based Scheduling



Application/Job Pod Scheduling



Application/Job Pod Scheduling Problem



Cluster Resources Are Limited

- Scaling up the cluster by adding nodes takes time
- These jobs are resource hungry and are willing to consume any resources you provide
- Practically, there is a limit on available resources, at every moment, even if not permanent

“The sky is the limit”, but this is only a cloud!

Multi-cluster, Hybrid Cloud, and Edge

- Organizations own tens of clusters
- Smaller clusters are easier to manage
- Static assignment of users or apps to clusters is not efficient
- Bursting from on-prem to public is needed more often and less predictable for AI workloads
- The rise of edge computing paradigm introduces more clusters to manage and choose from for running a job

Why the Kubernetes scheduler is not enough

Desired capabilities

MCAD - open source project

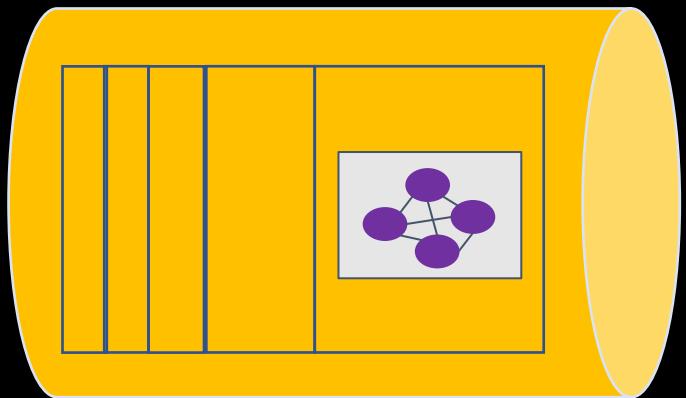
How it works

Demo

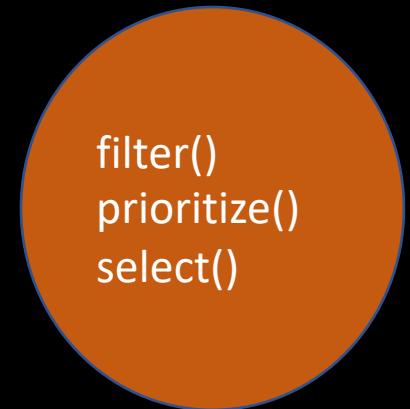
Call to action

Holistic Application/Job Queuing/Dispatching

Queue Controller



scheduler



Application Job1:
4 GPU Demand
Each Pod Requires
1 GPU

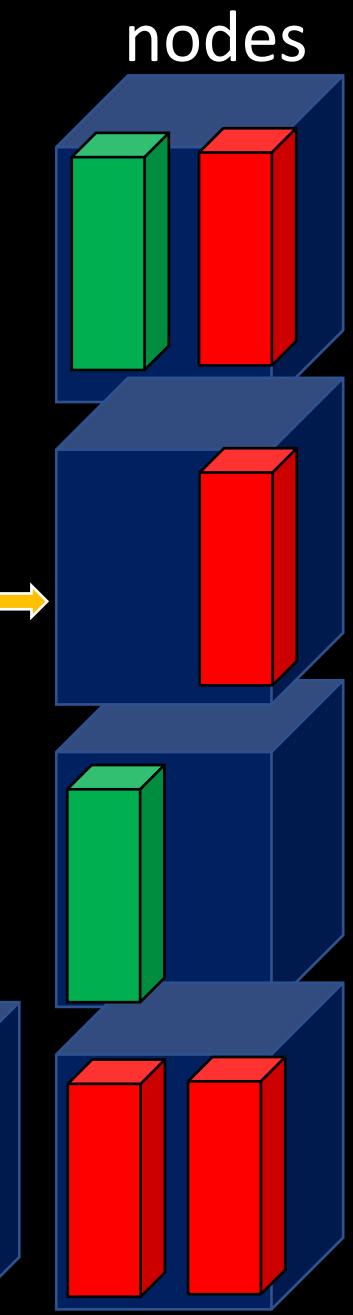
A green network graph icon representing a pod. It consists of four green circles connected by lines, with one circle highlighted in red.

Application Job2:
4 GPU Demand

A purple network graph icon representing a pod. It consists of four purple circles connected by lines, with one circle highlighted in red.

Application Job3:
4 GPU Demand

A red network graph icon representing a pod. It consists of four red circles connected by lines, with one circle highlighted in green.



More Desired Capabilities

- Multi-cluster dispatching
- Priorities and classes of service (Gold/Silver/Bronze).
- Hierarchical quota management – soft and hard quotas, multiple resources
- Preemption
- Unified view of jobs belonging to multiple AI/ML frameworks – AppWrapper

Why the Kubernetes scheduler is not enough

Desired capabilities

MCAD - open source project

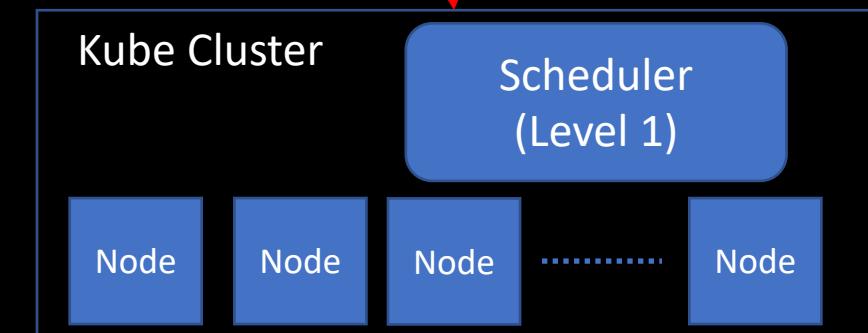
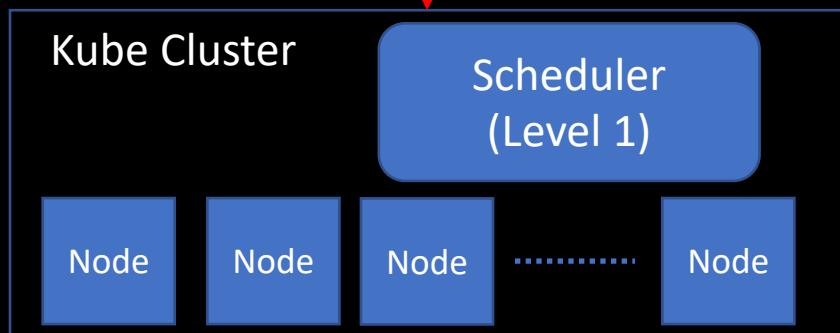
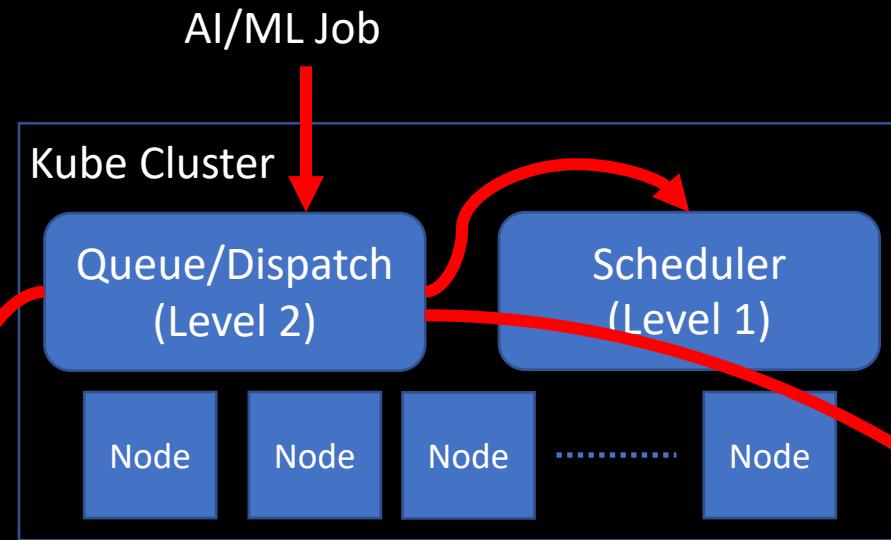
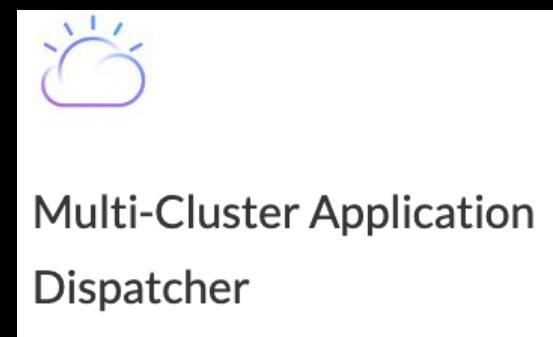
How it works

Demo

Call to action

What is MCAD?

OperatorHub.io



Why the Kubernetes scheduler is not enough

Desired capabilities

MCAD - open source project

How it works

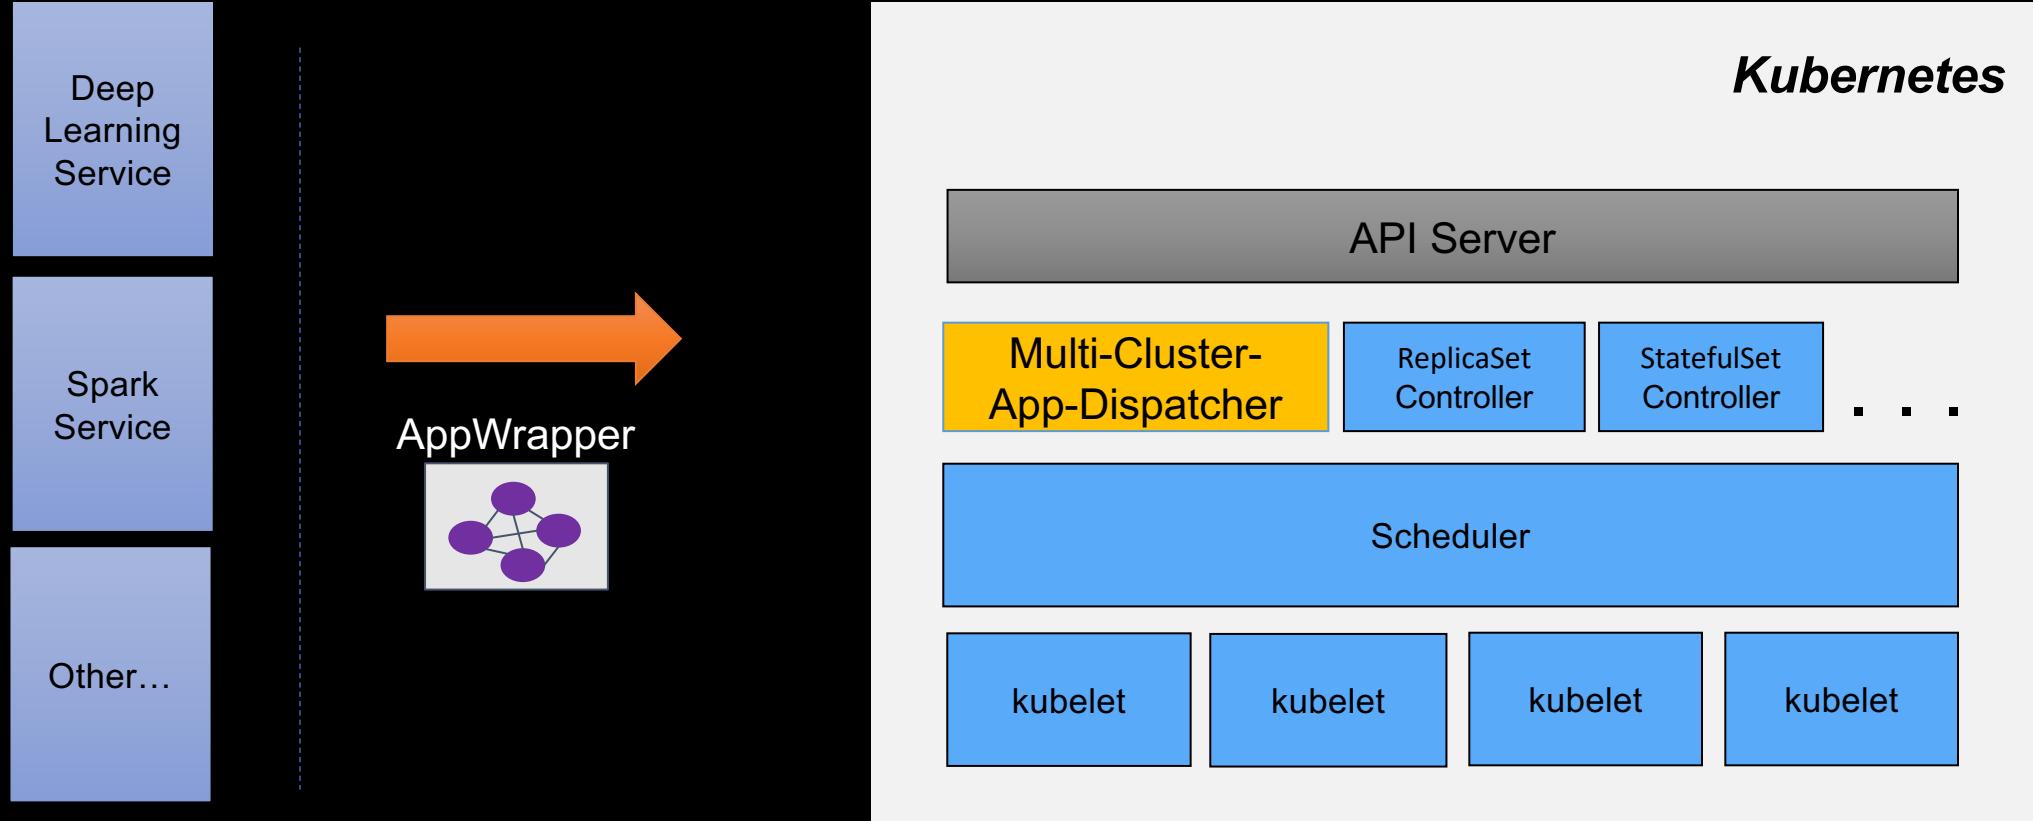
Demo

Call to action

MCAD Controller

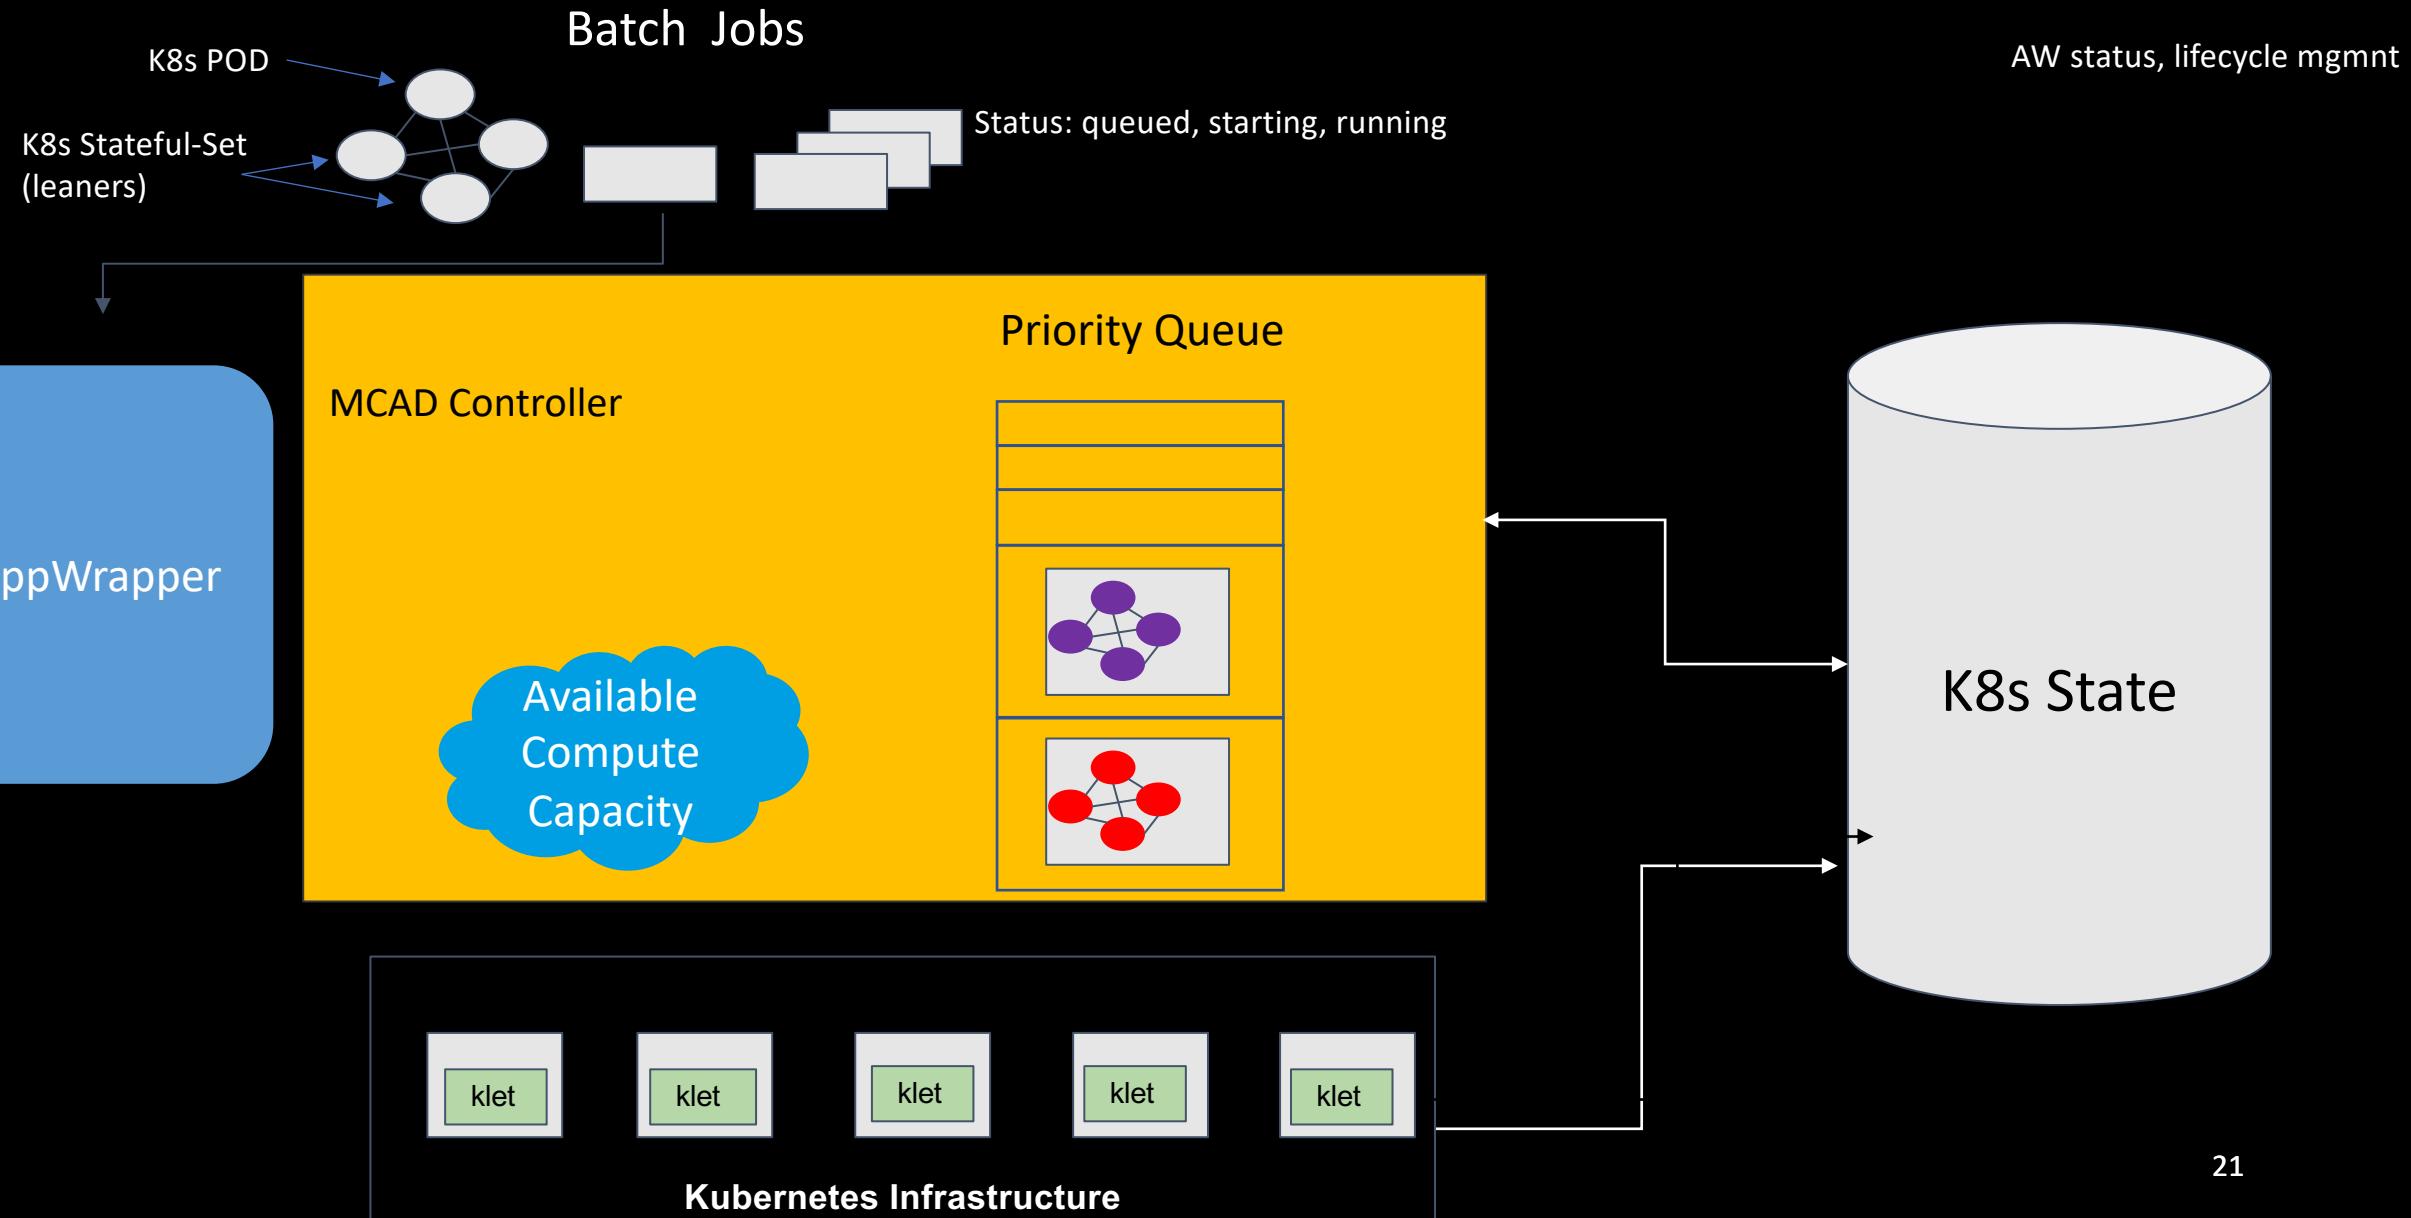
- AppWrapper Custom Resource Definition
 - Wraps All Compute Resource Consuming Kubernetes Objects for a Job
 - E.g. Deployments, Pods, Statefulsets
 - Wraps All Non-Compute Resource Consuming Kubernetes Objects for a Job
 - E.g. Services, Namespaces
- MCAD Custom Resource Controller
 - Determines Runnable Job for All Kubernetes Objects within an AppWrapper Holistically
 - Unwraps and Dispatches All Kubernetes Objects within an AppWrapper when Job is Runnable
 - Queues Non-Runnable Jobs
 - Supports Preemption/Requeuing

Multi-Cluster Application Dispatcher

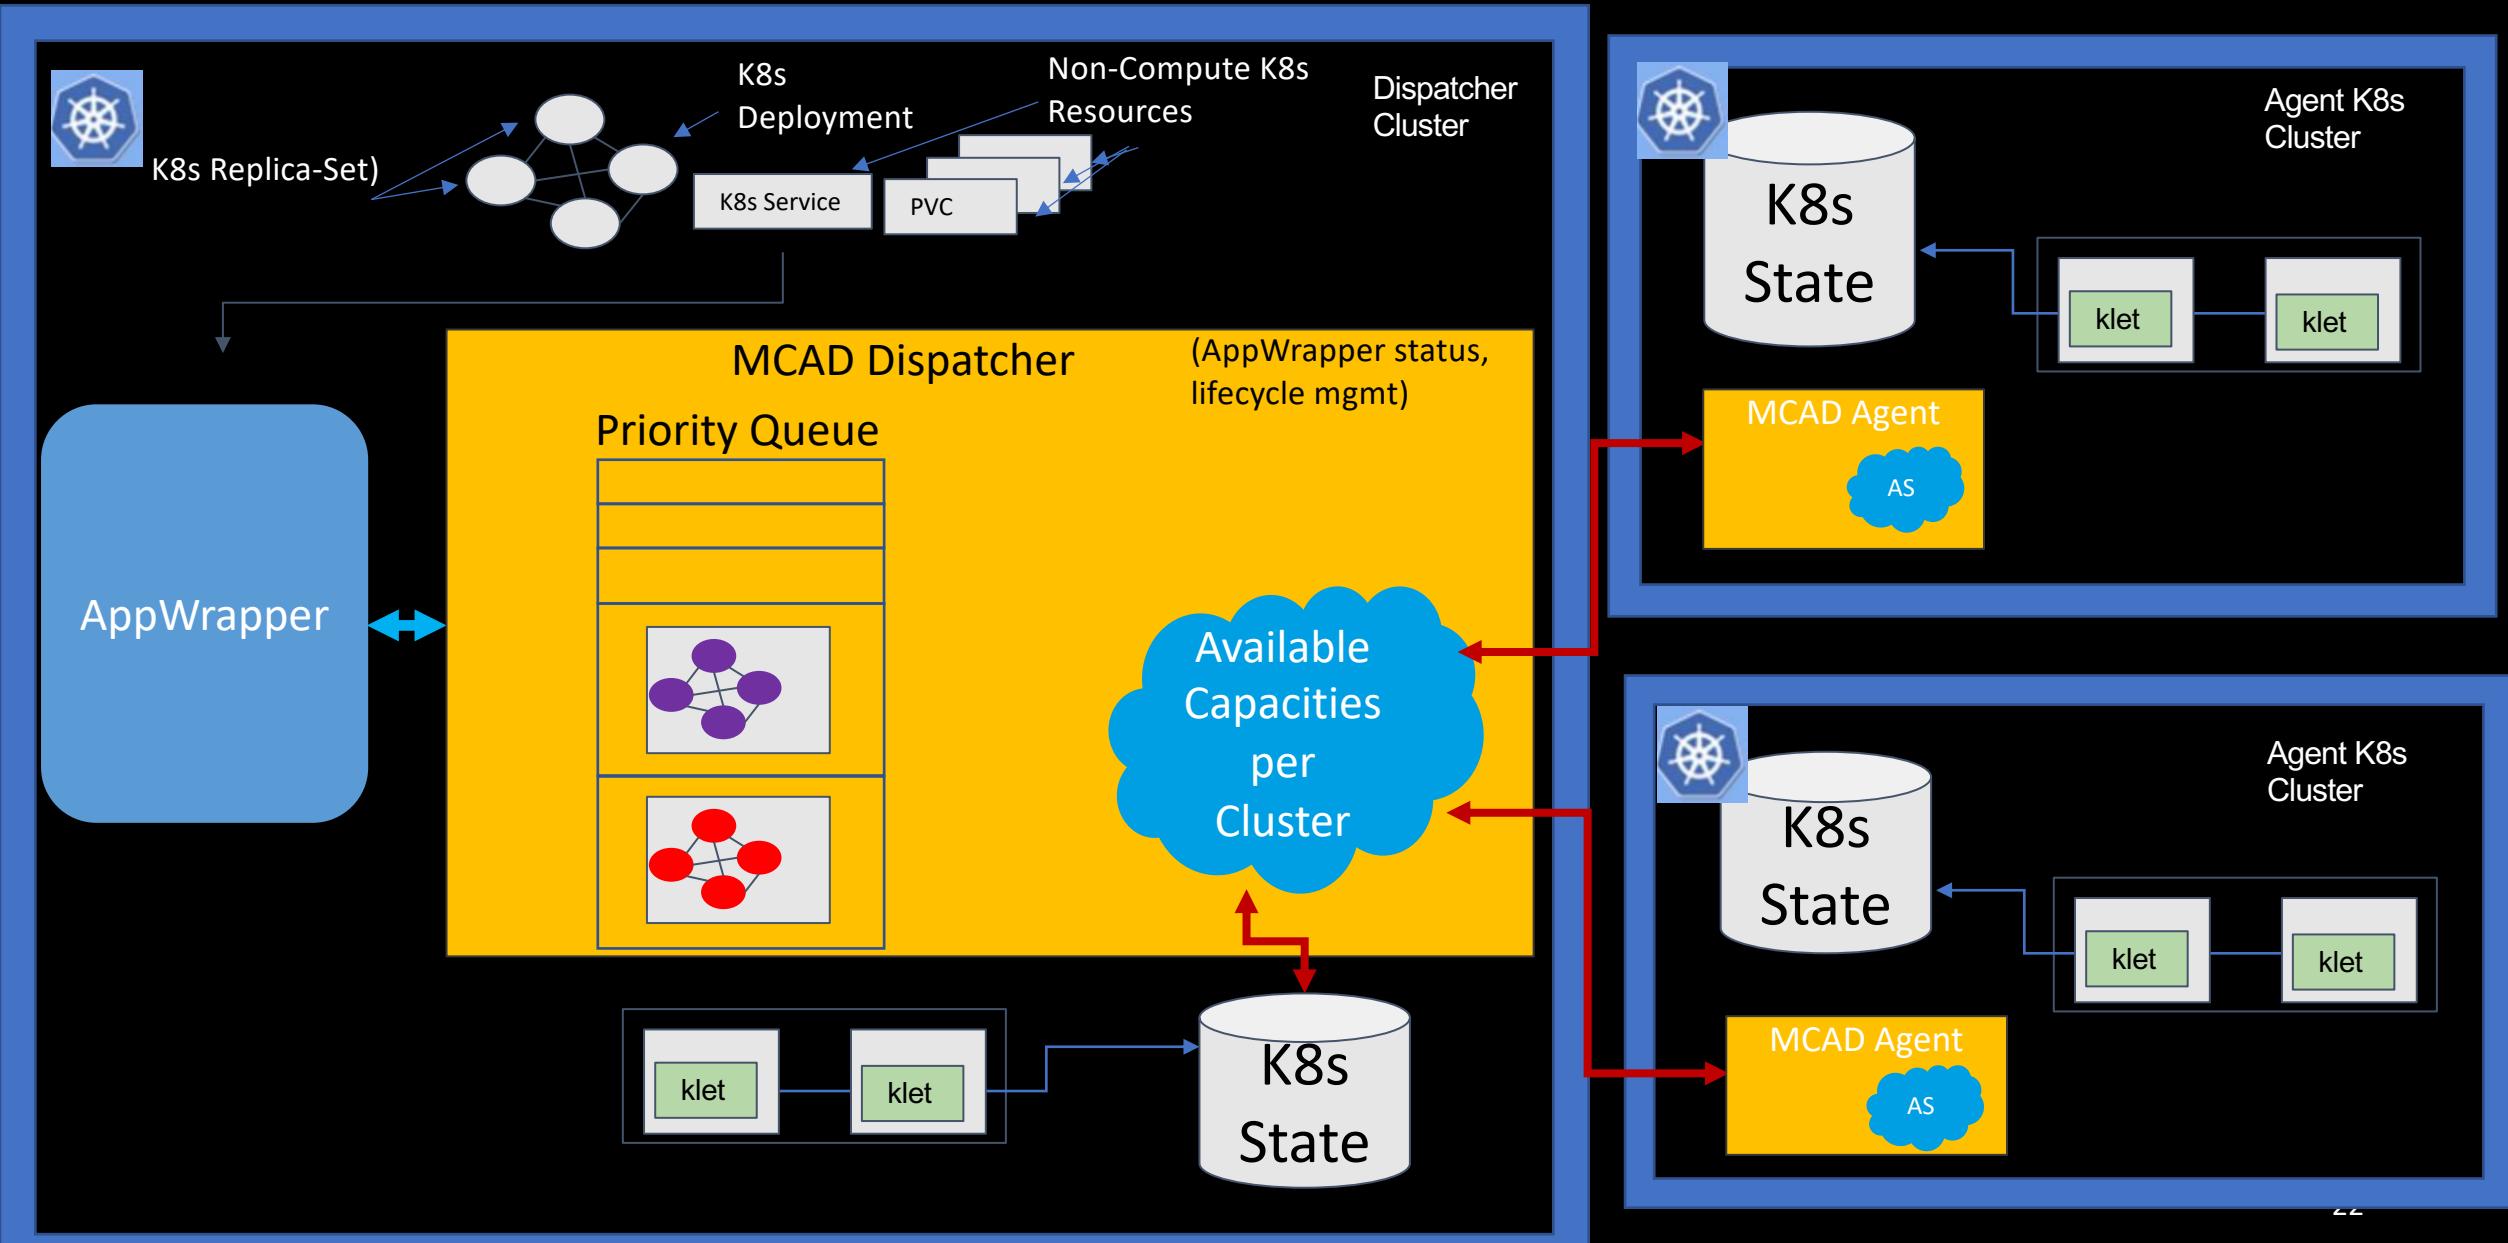


- 2 Runtime Configurations
 - Standalone
 - Dispatcher/Agent

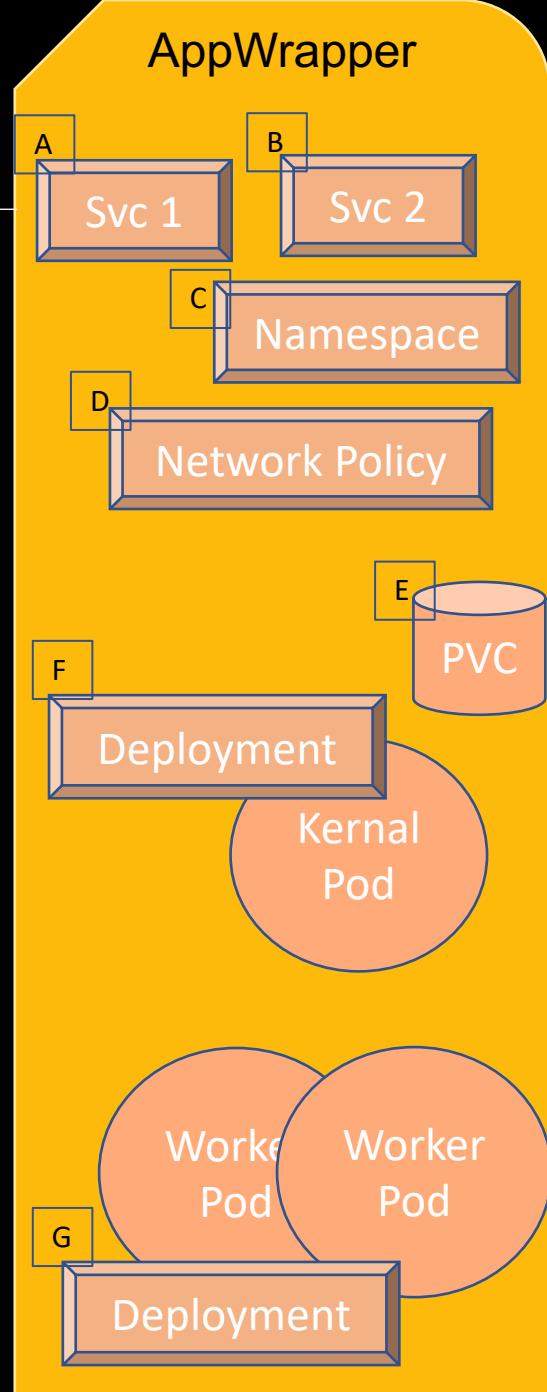
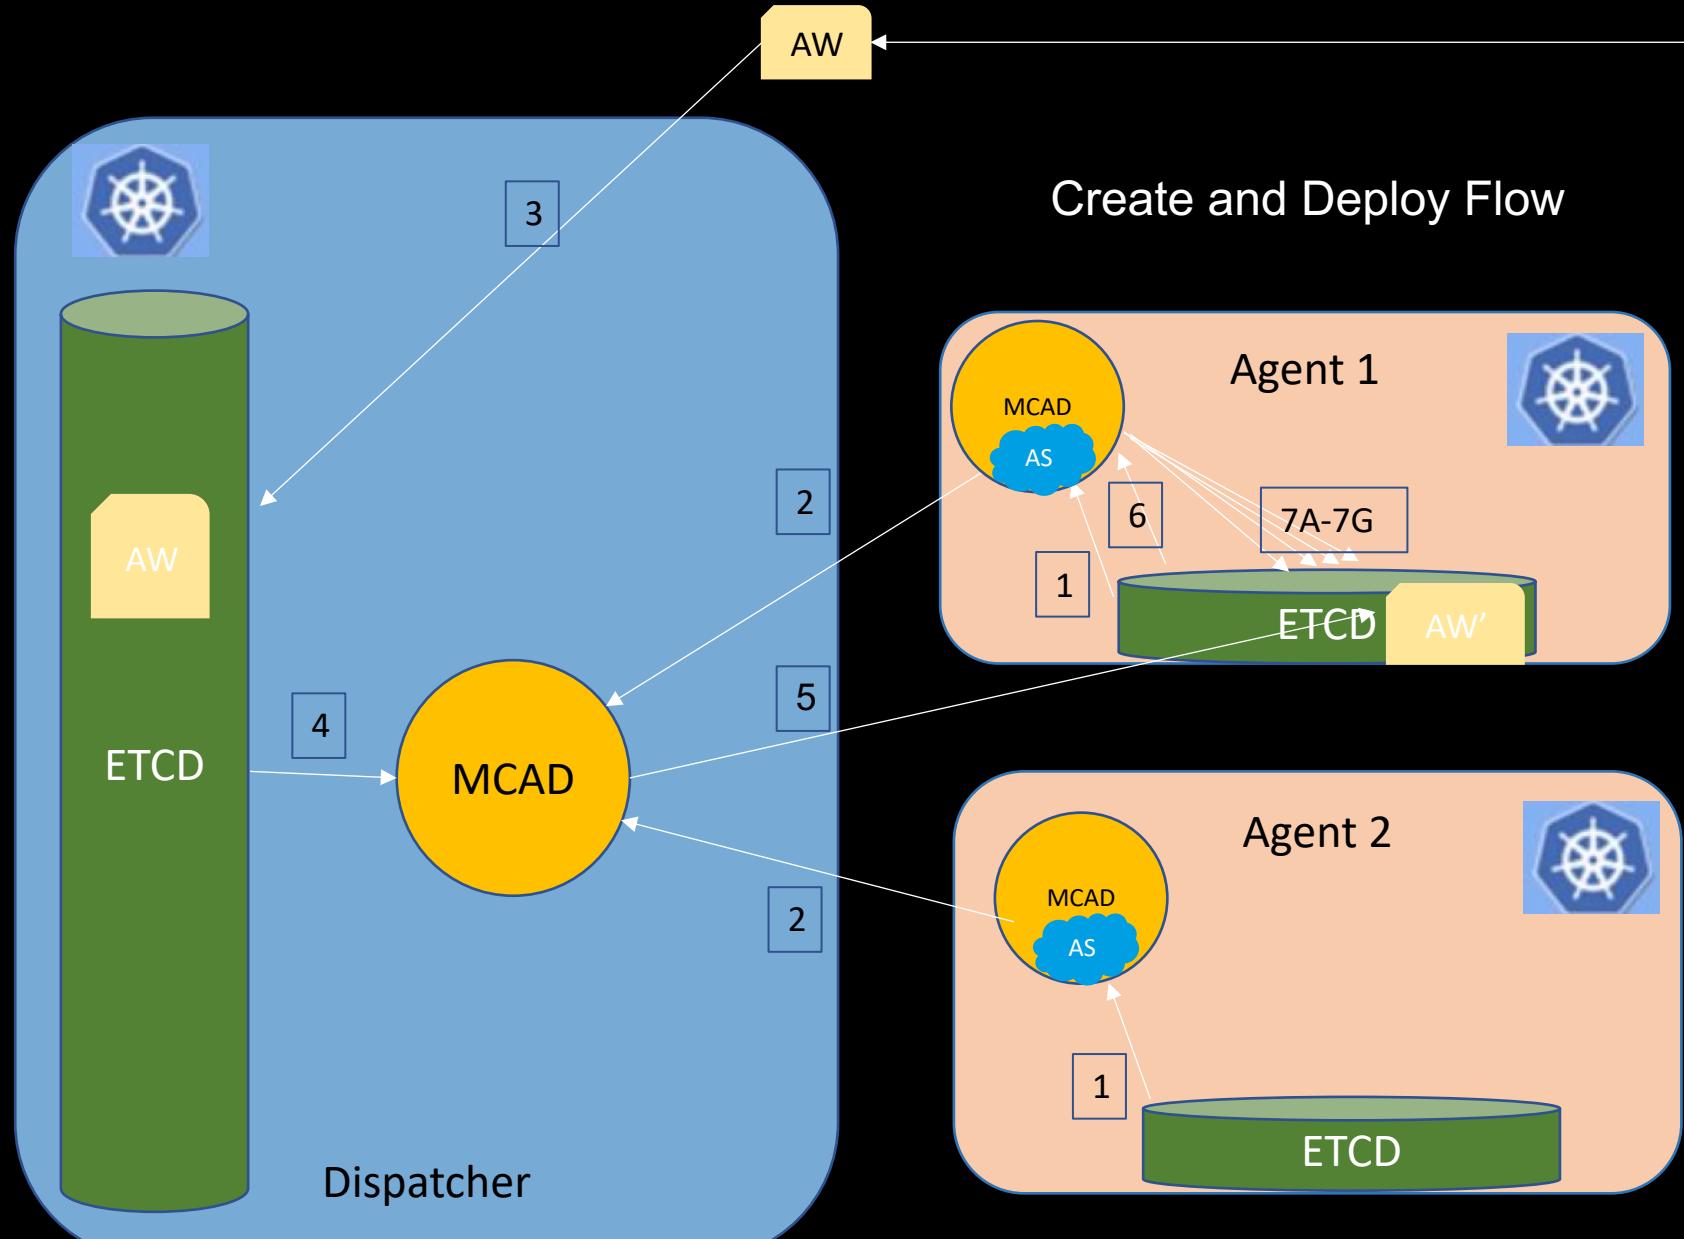
Standalone: Multi-Cluster Application Dispatcher



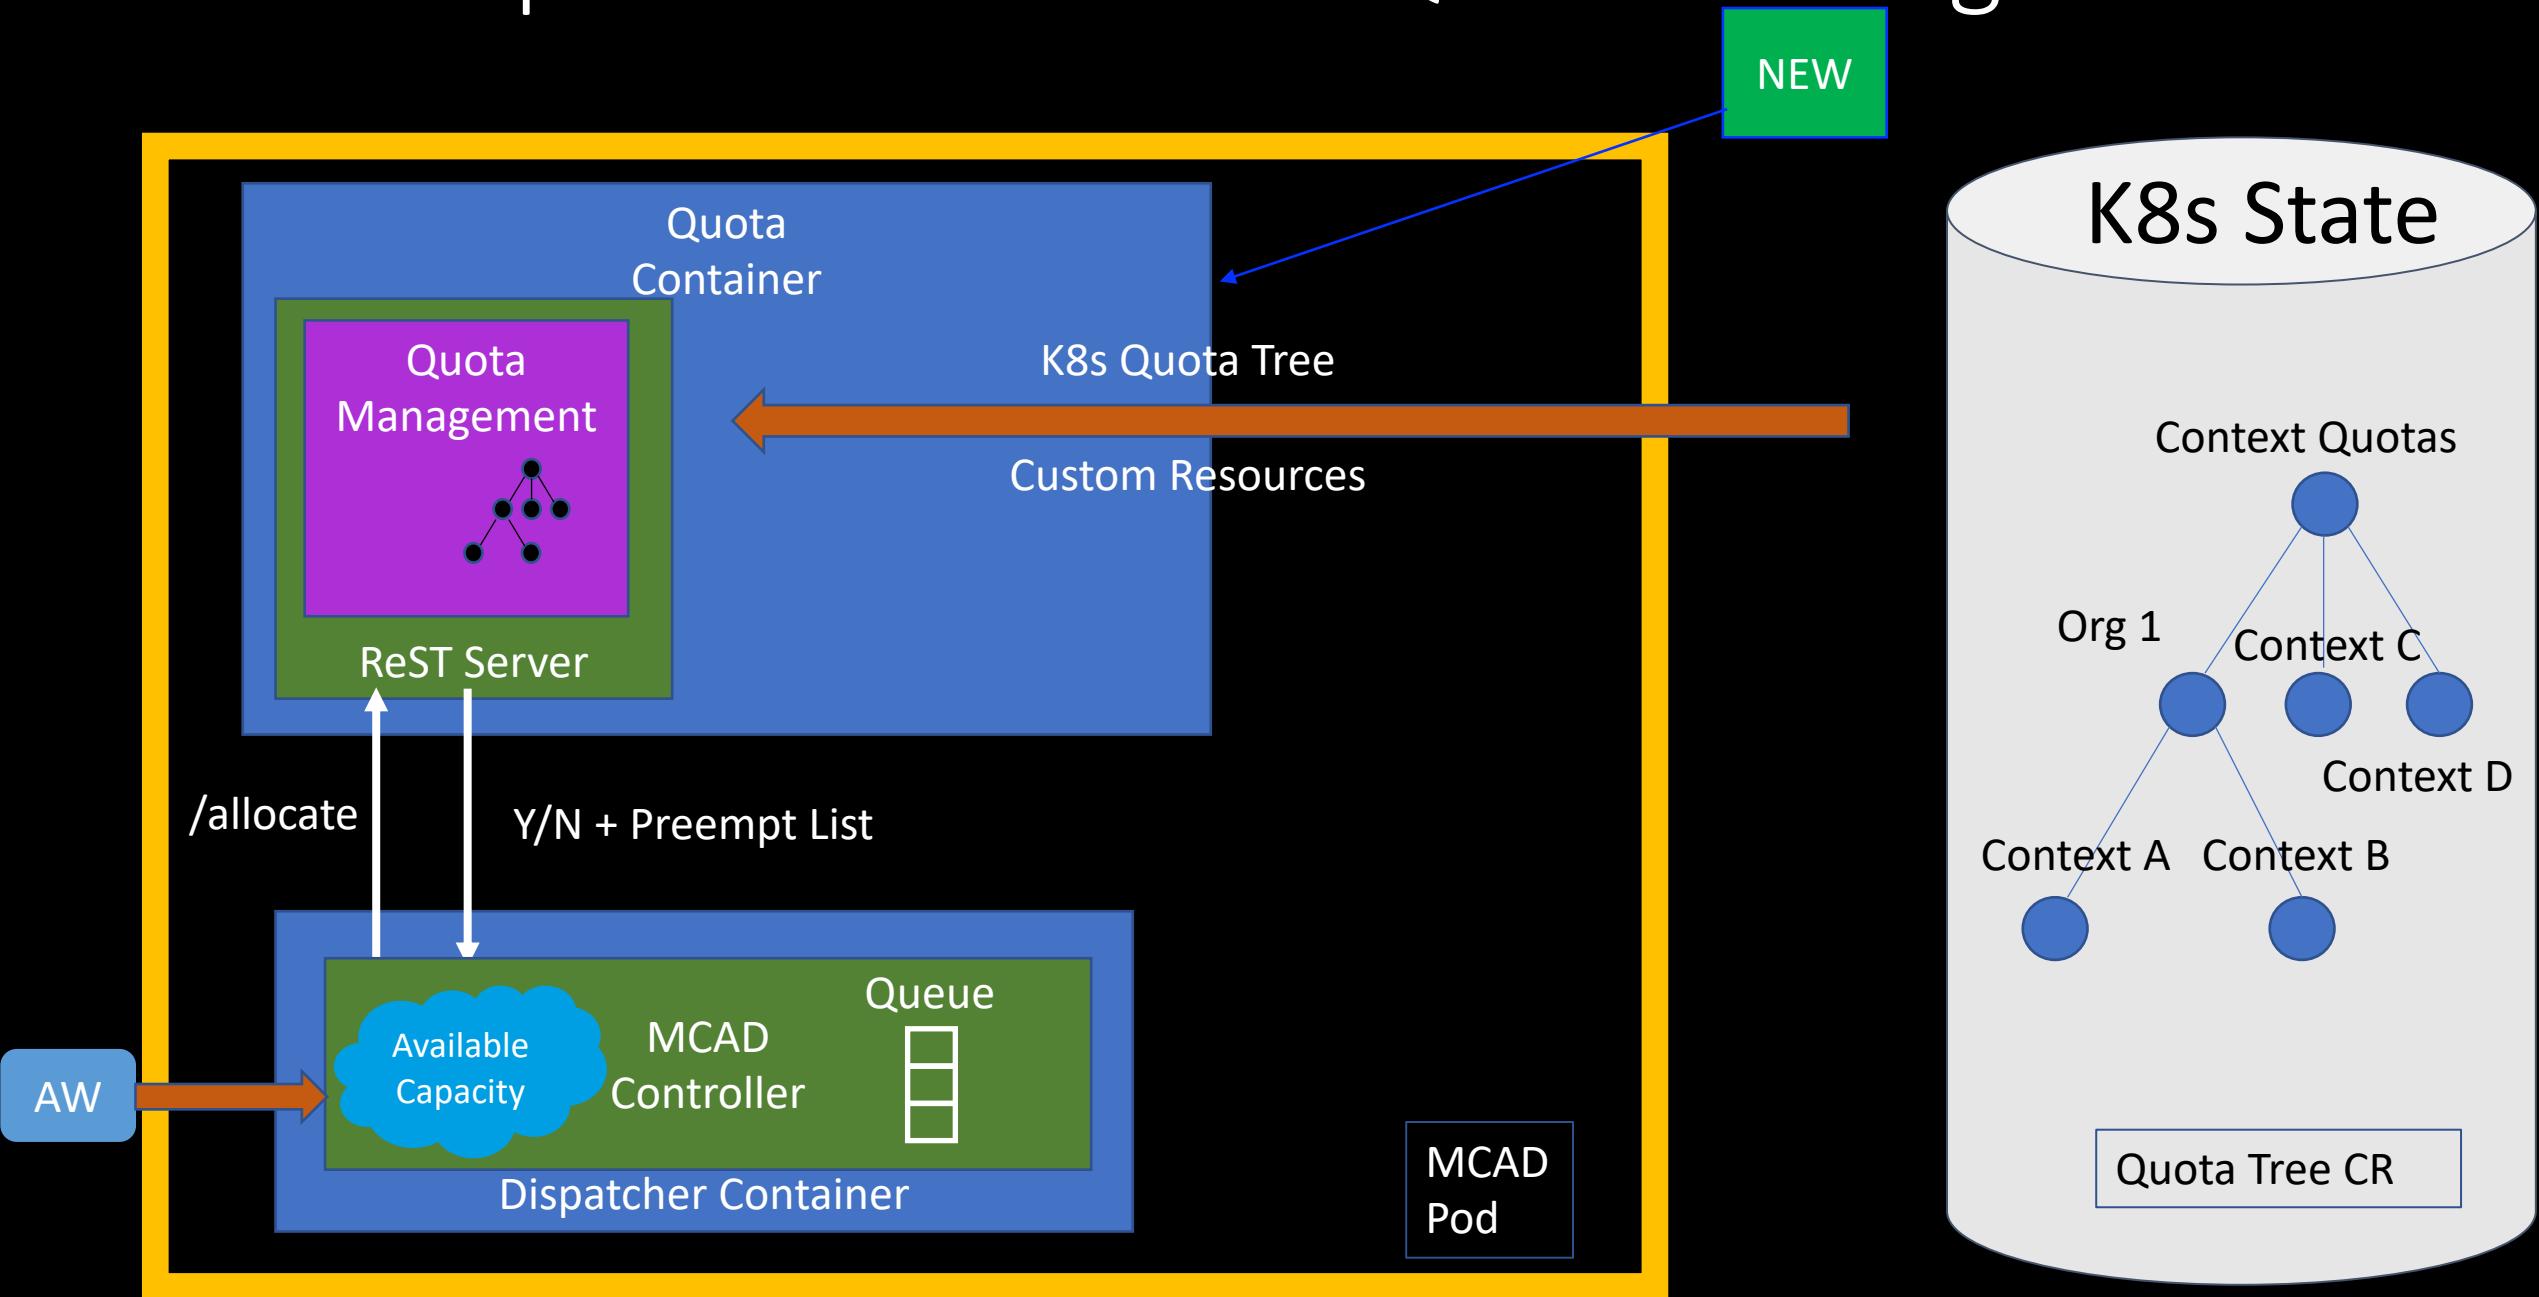
Multi-Cluster: Dispatcher/Agent



Example Use Case



Roadmap: Current Work - Quota Management



AppWrapper Yaml

```
kind: AppWrapper
metadata:
  name: job1-context2-1replica
  labels:
    quota_context: Context-2
  ...
spec:
  priority: 1000
  resources:
    Items:
    - replicas: 1
      type: StatefulSet
      template:
        apiVersion: apps/v1
        kind: StatefulSet
        metadata:
          name: job1-context2-1replica
          labels:
            app: job1-context2-1replica
```

AppWrapper
Custom
Resource

(Roadmap)
Quota
Management

(Optional)
Dispatch
Priority

List of
Kubernetes
Resources

Why the Kubernetes scheduler is not enough

Desired capabilities

MCAD - open source project

How it works

Demo

Call to action

Why the Kubernetes scheduler is not enough

Desired capabilities

MCAD - open source project

How it works

Demo

Call to action

We Need Your Help!!

OperatorHub.io



Multi-Cluster Application
Dispatcher



[https://github.com/IBM/
multi-cluster-app-
dispatcher](https://github.com/IBM/multi-cluster-app-dispatcher)

Try it out!
Give feedback!
Contribute!