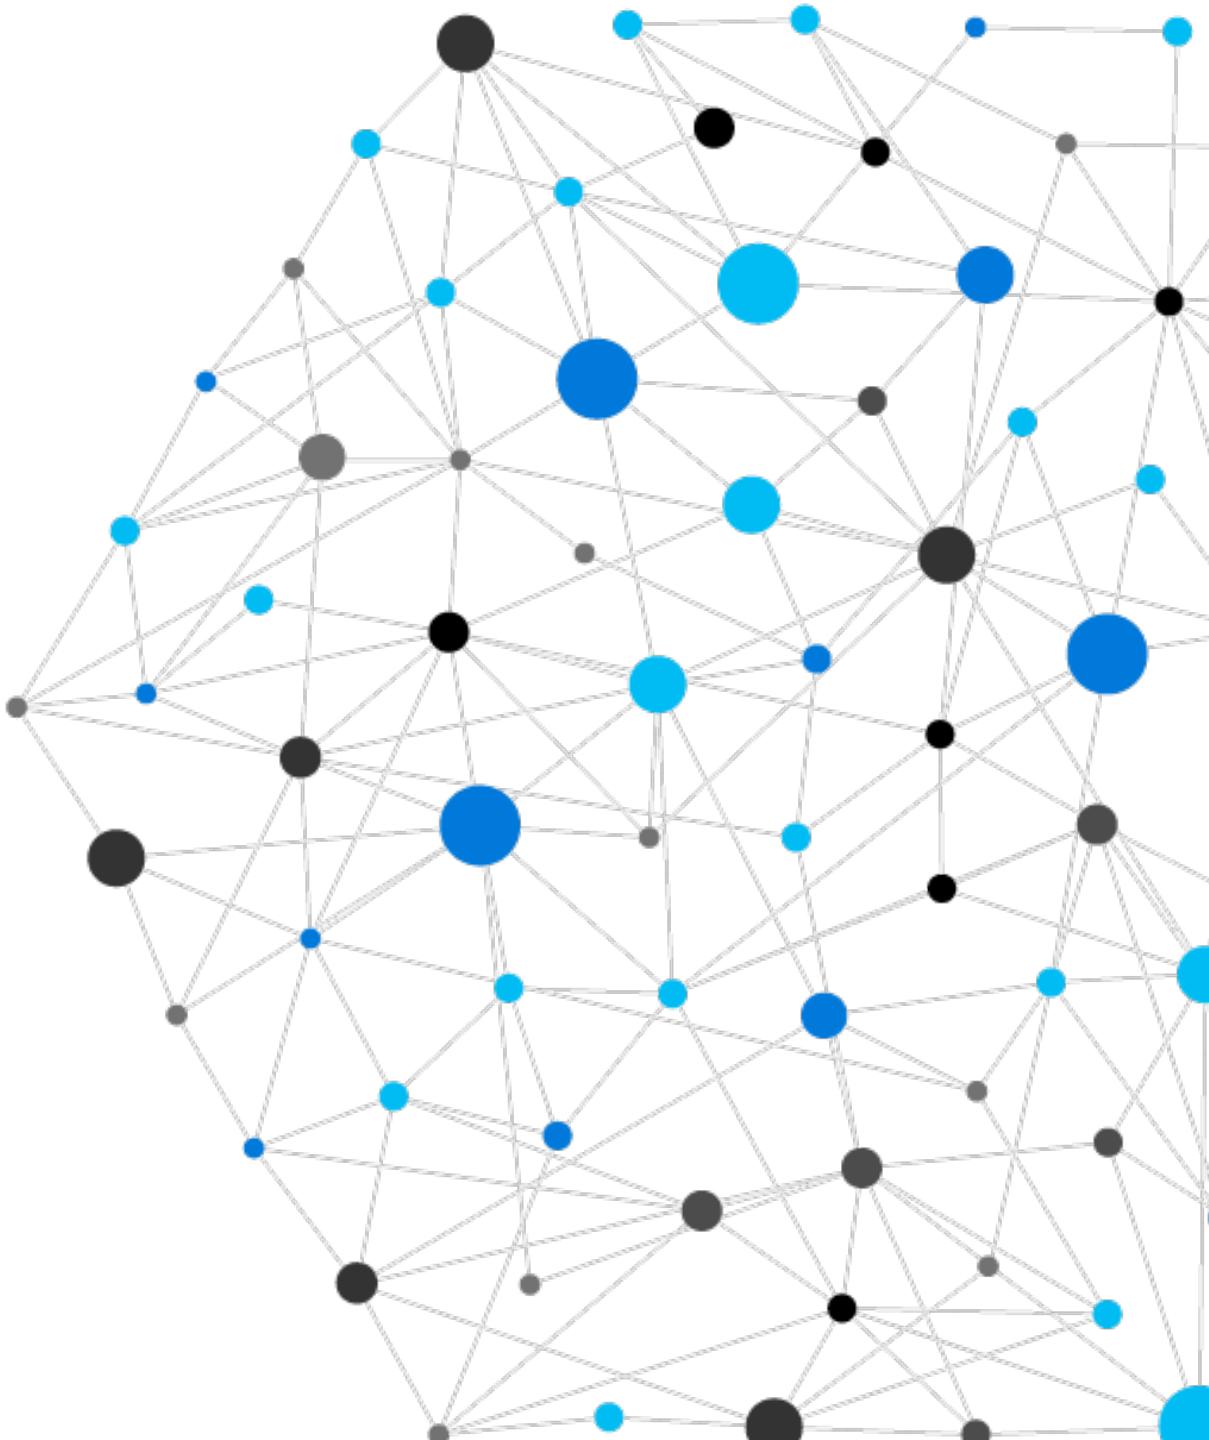


# Optimize your Kubernetes Clusters on Azure with Built-in Best Practices



# Your speaker today

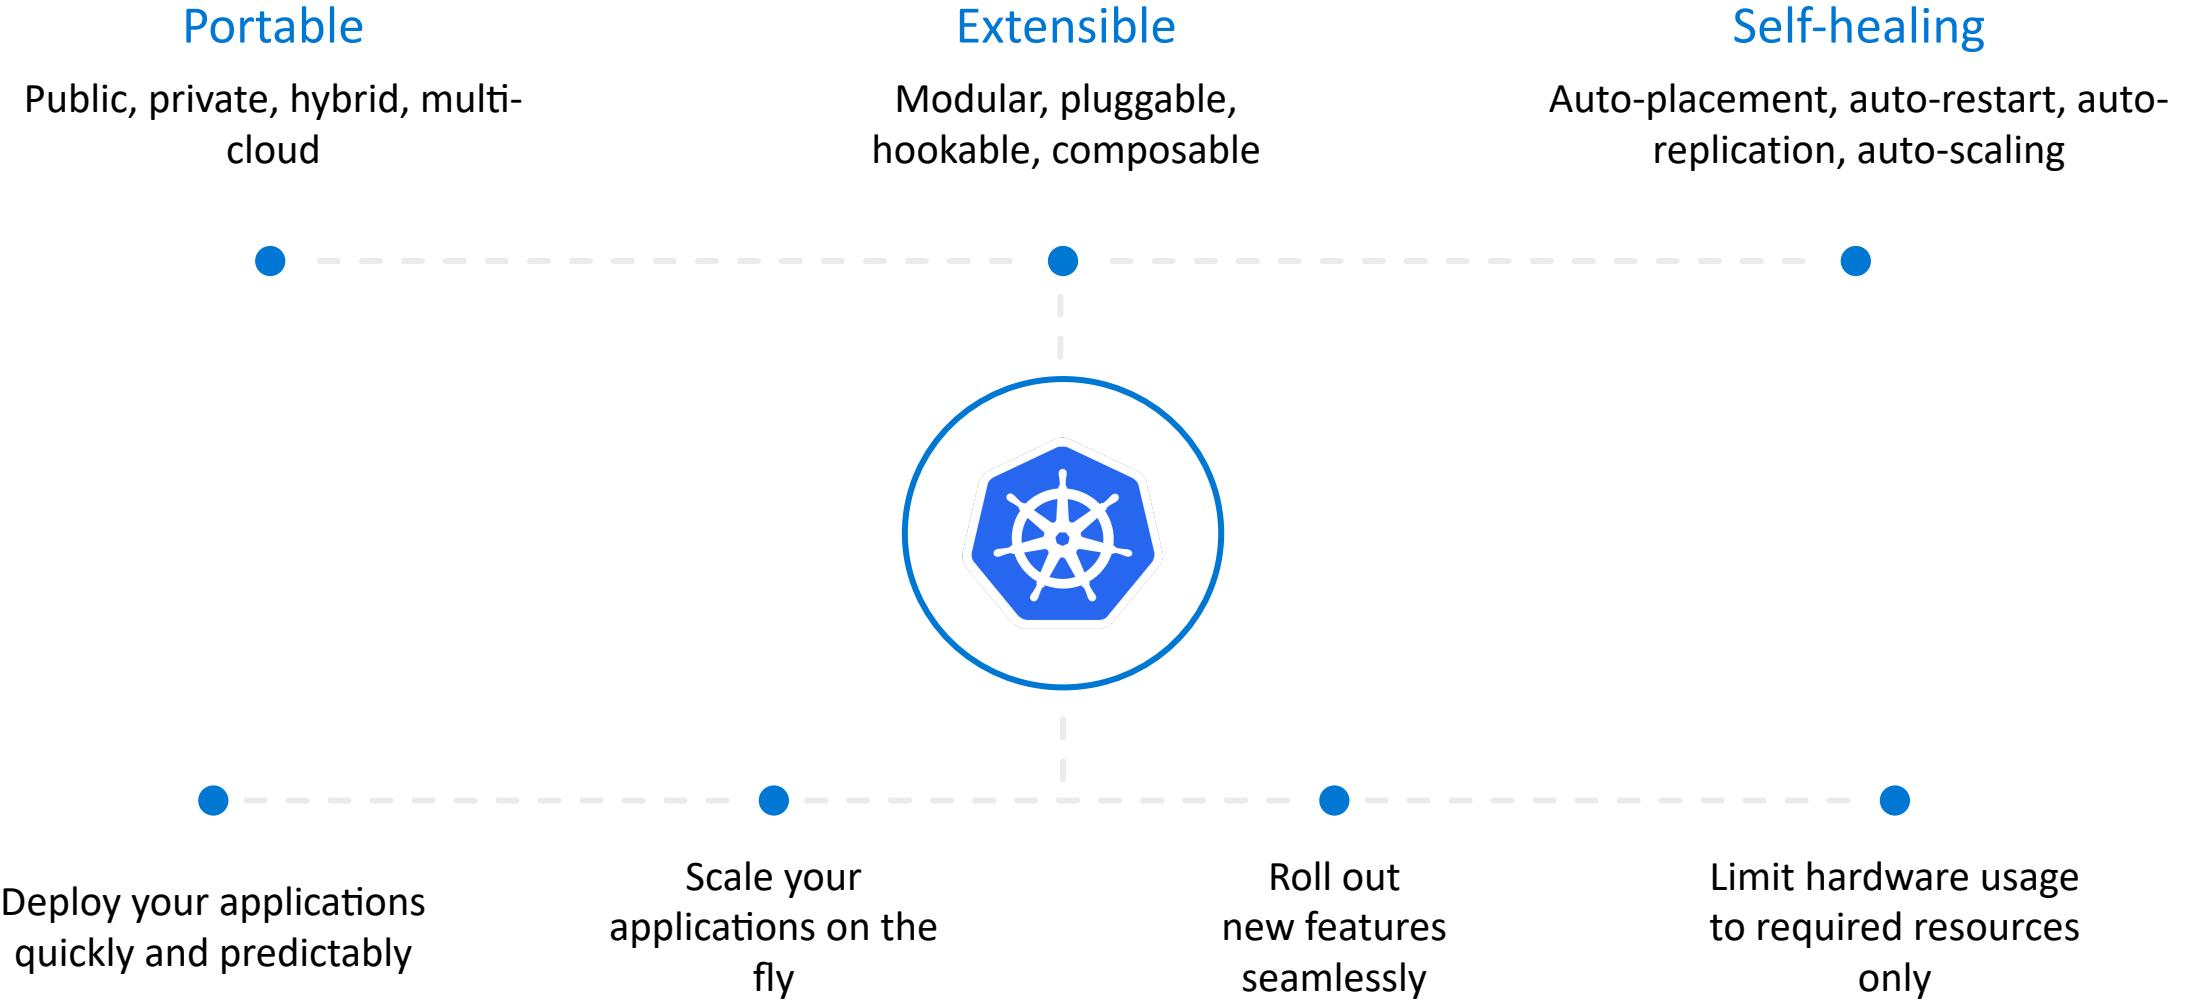


Jorge Palma

Senior Program Manager - Azure Container Compute

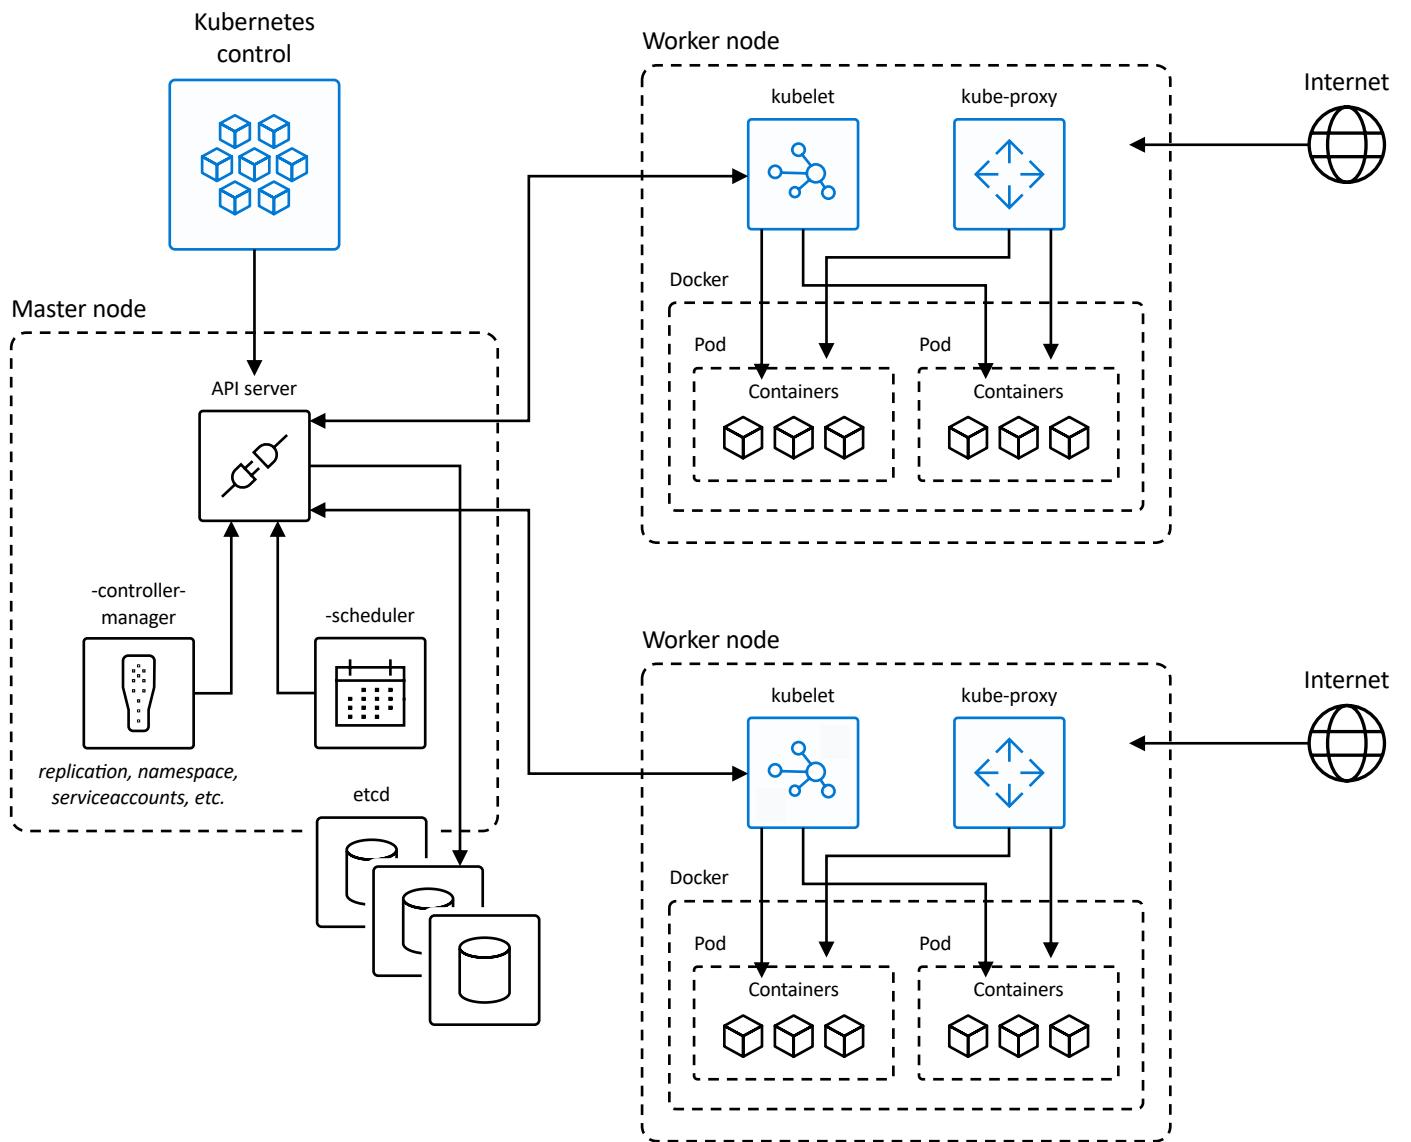
 @jorgefpalma

# Kubernetes: empowering you to do more



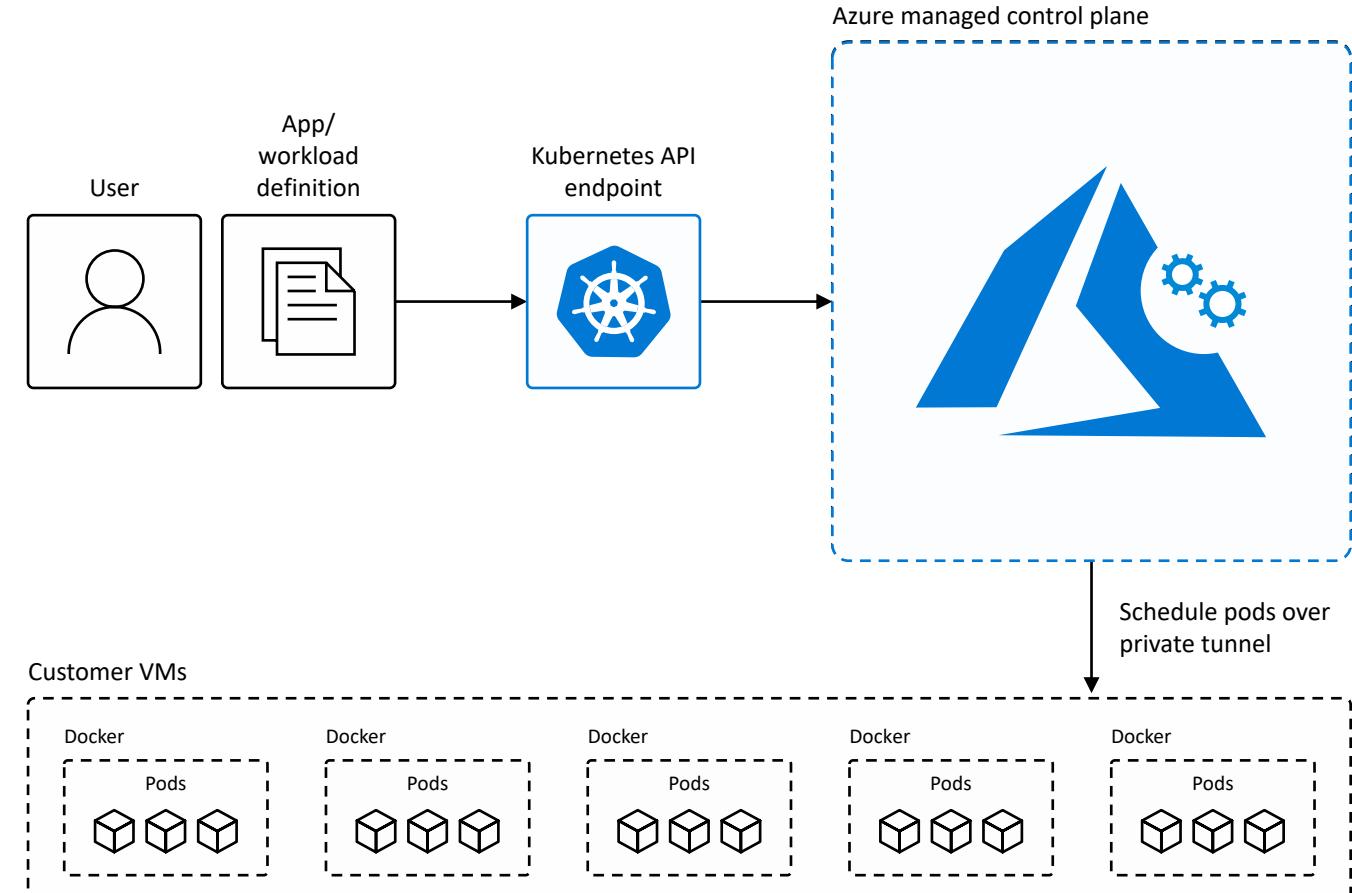
# How Kubernetes works

1. Kubernetes users communicate with API server and apply desired state
2. Master nodes actively enforce desired state on worker nodes
3. Worker nodes support communication between containers
4. Worker nodes support communication from the Internet



# Ok, so managed kubernetes it is

- Automated upgrades, patches
- High reliability, availability
- Easy, secure cluster scaling
- Self-healing
- API server monitoring
- At no charge

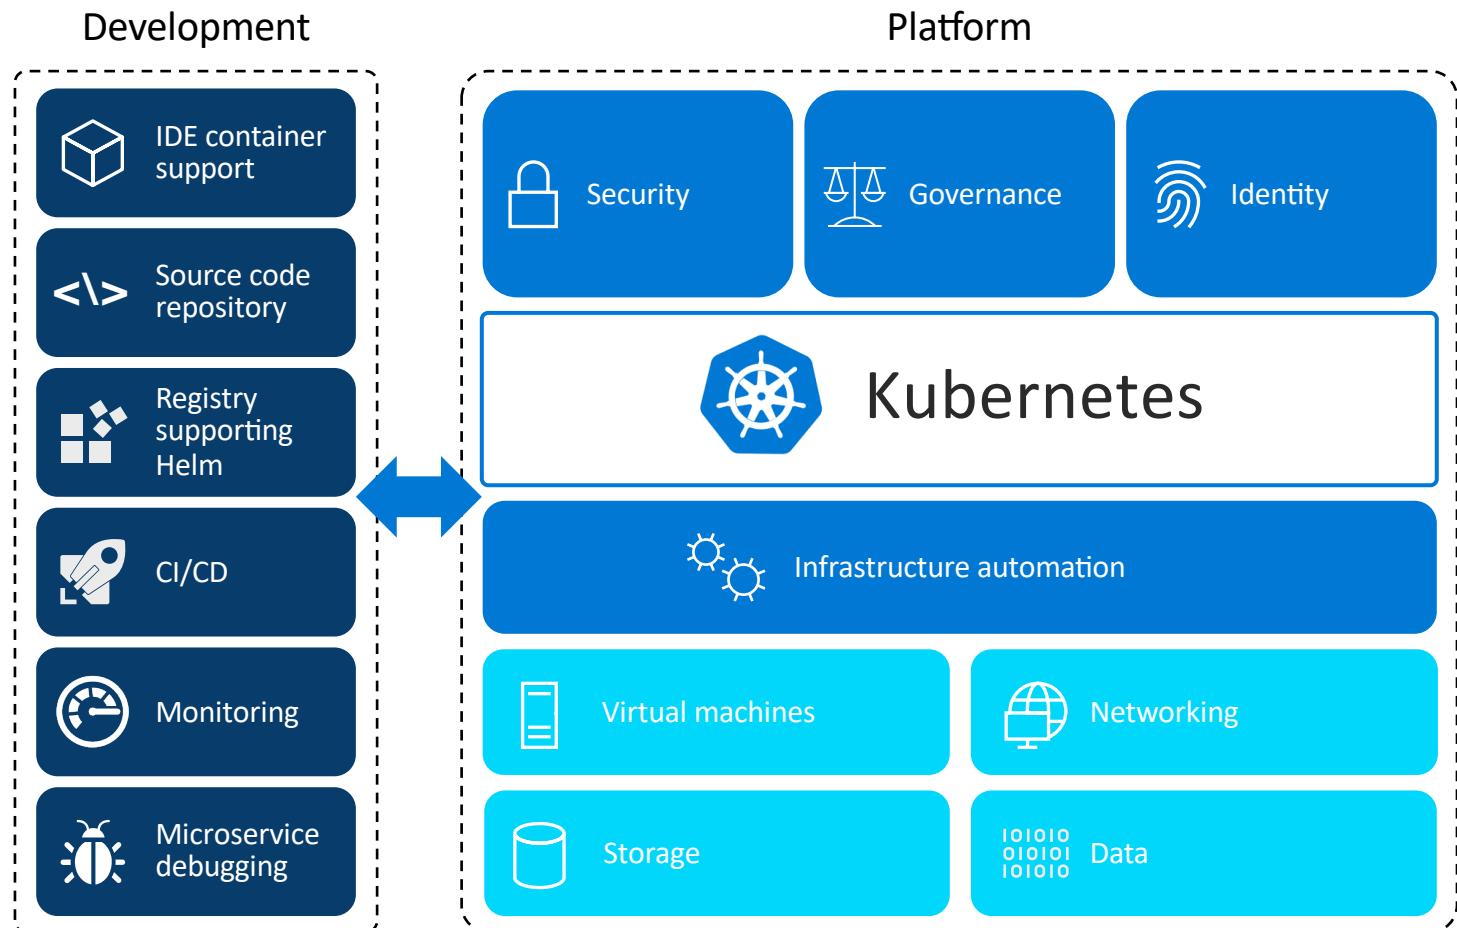


# Is that enough?

Well, depends

Unlock the agility for containerized applications using:

- Infrastructure automation that simplifies provisioning, patching, and upgrading
- Tools for containerized app development and CI/CD workflows
- Services that support security, governance, and identity and access management



# What is missing?

## Common issues

1. Incorrect Kubernetes configurations
2. Permissive or insecure clusters and lack of pod security
3. Lack of governance policies
4. Lack of central authentication at scale
5. Lack of fine-grained access control at scale
6. “Voluntary” disruptions and unplanned downtime
7. Inefficiency

# What can we do?

## Common needs

1. Validating Kubernetes Configurations
2. Manage and harden cluster and pod security, especially at scale
3. Create and curate default secure configurations, and we need to be able to enforce them via Policies
4. Leverage central authentication directories
5. Apply fine-grained access controls at scale
6. Prevent human errors that cause down time, harden against disruption
7. How can we be more efficient
  - a) Autoscalling?
  - b) Different agent node groups/pools?
  - c) Spot instances?

# Kubernetes Configurations

## Sample configuration checks

1. Configure readiness and health probes
2. Configure Resource quotas
3. Do not use “latest” image tags
4. Restrict the registries your images can come from
5. Configure pod security
6. Cluster’s Kubernetes version up to date
7. Cluster agent nodes up to date
8. Administrative access minimized, audited and API access controlled.

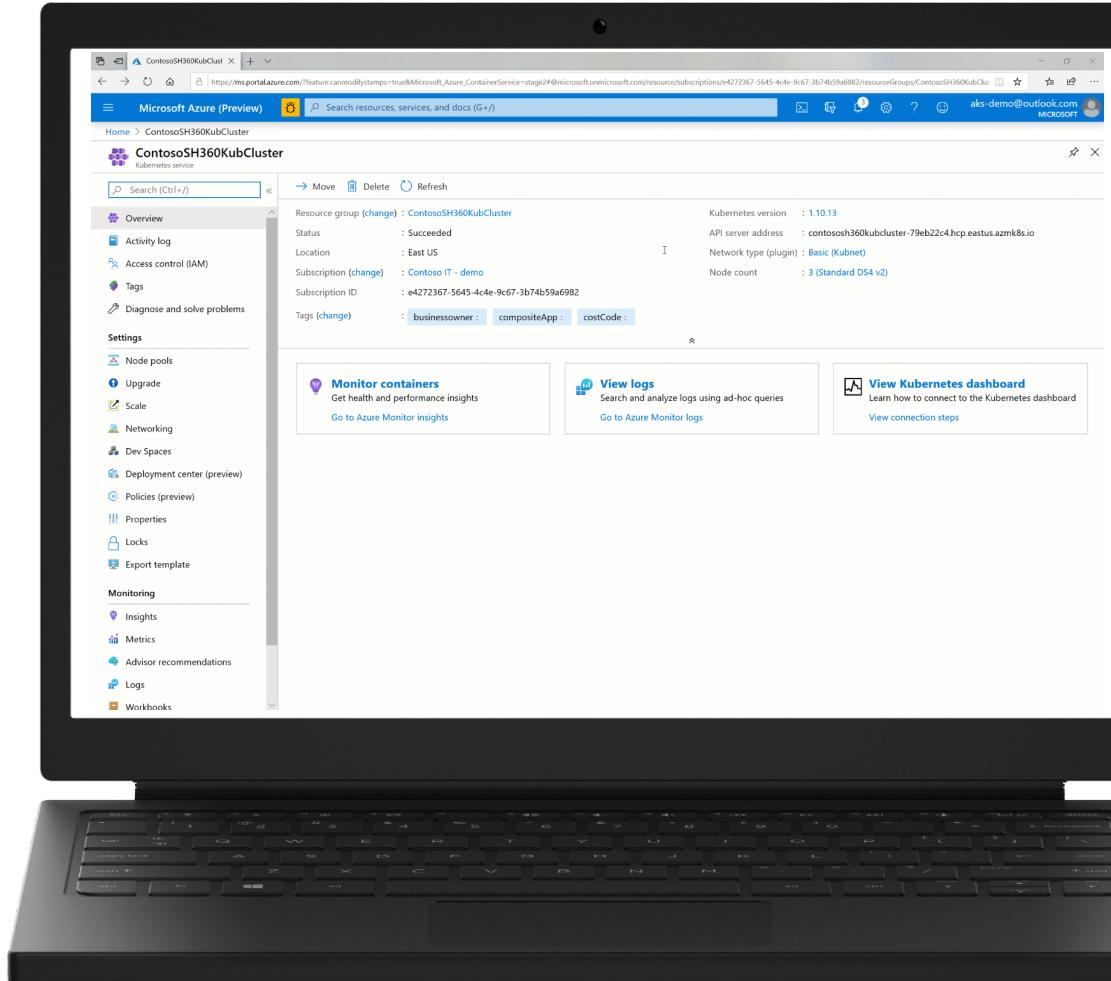
# Kubernetes Configurations

## Configuration recommendation and validation tools

- [Polaris](#)
- [Kube-scan](#)
- [Kube-advisor](#)
- [Secure DevOps Kit](#)

# Kubernetes Built-in best practices with Azure Advisor

- Proactive and actionable recommendations from Azure Advisor based on your configuration and usage telemetry
- Grounded with knowledge from thousands of customer engagements
- Improve the performance, availability, efficiency and security of your cluster before there's a problem



# Tips for getting started with Advisor alerts

1. Start simple with a few high impact recommendations and email
2. Notify the person or team who can remediate the recommendations
3. Explore automation scenarios using webhooks and/or runbooks

# Cloud Workload Protection for Containers



Now available in standard Azure Security Center with the new Container Service add-on

## Protecting Container hosts (IaaS)

- CIS Docker Benchmark assessment
- Node Threat Protection

## Protecting AKS

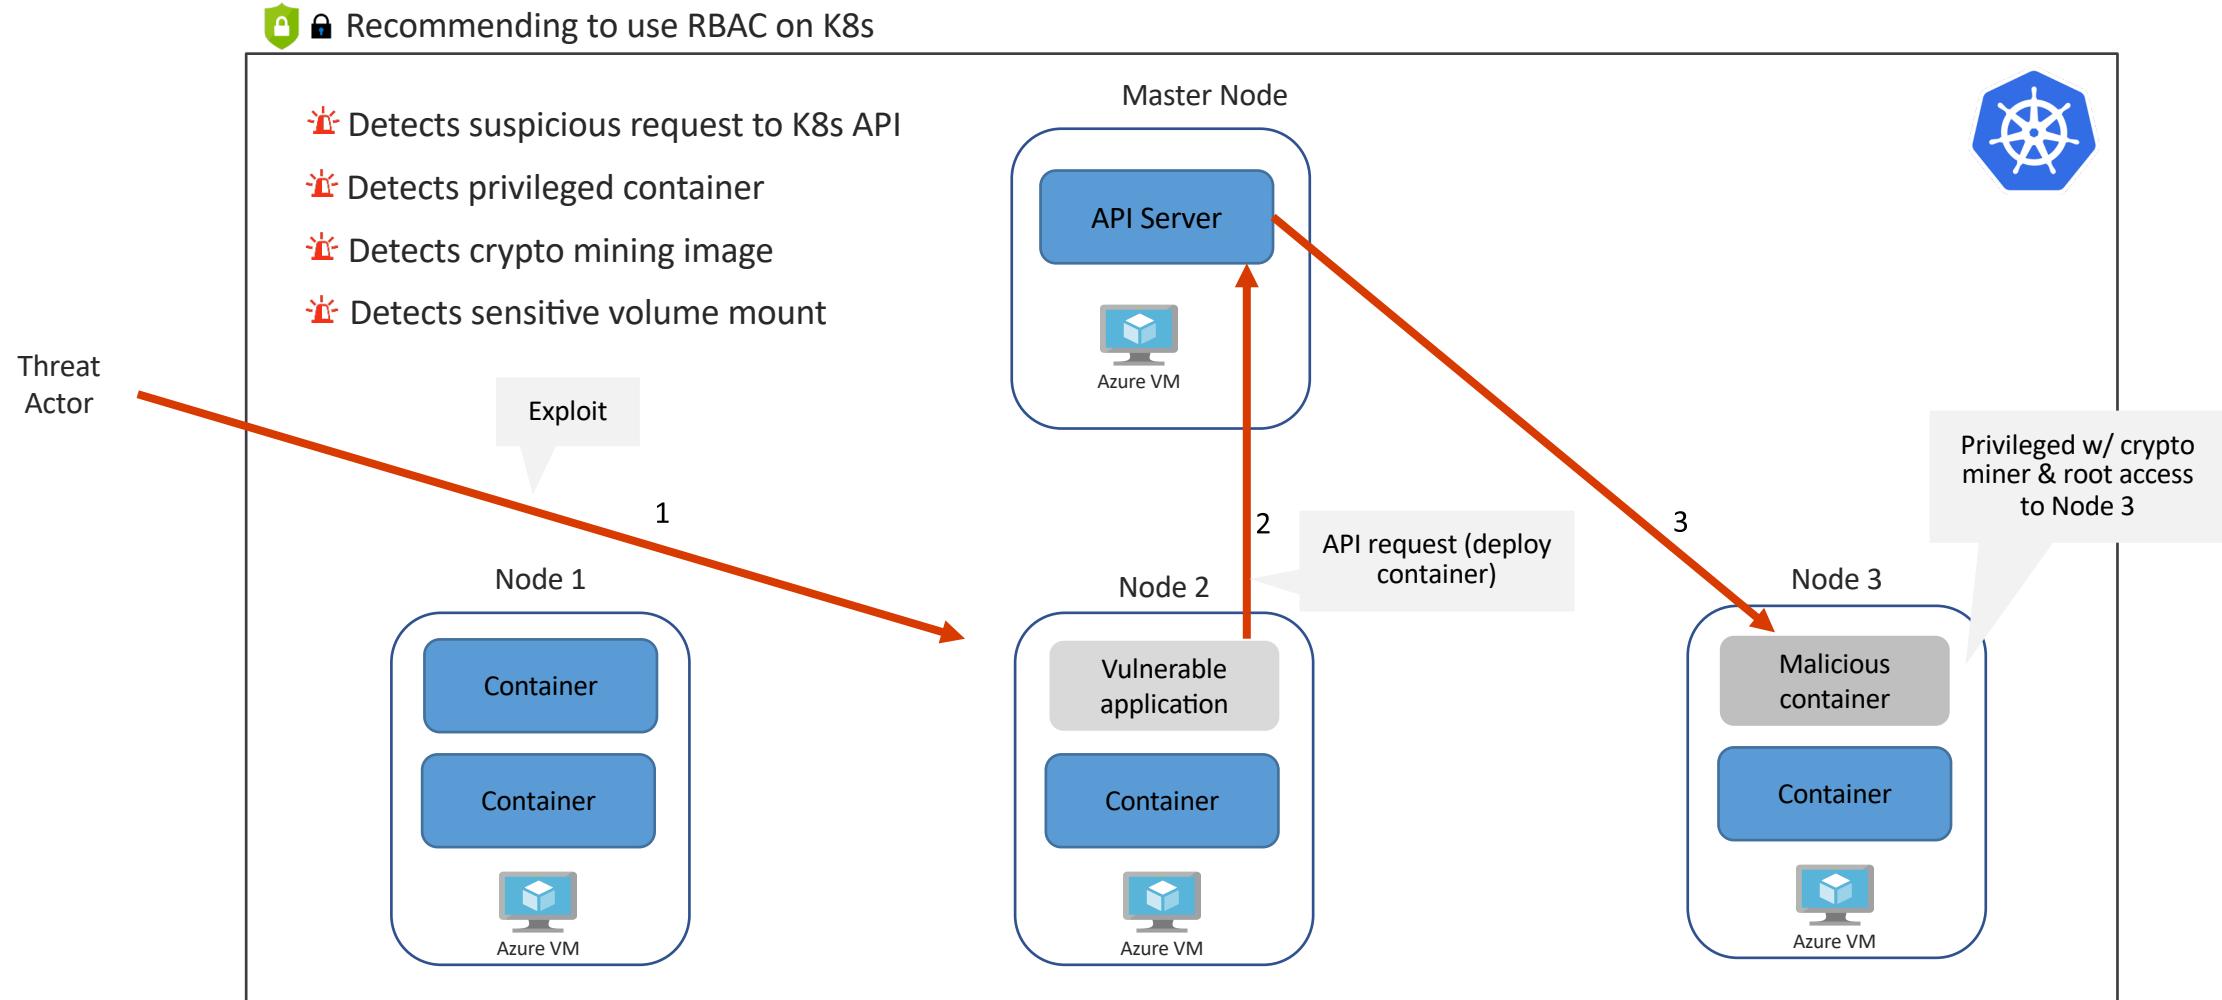
- Actionable recommendations based on AKS best practices
- Cluster and Node Threat detection based on AKS audit log and Node Auditd

The screenshot displays the Microsoft Azure Security Center interface for Compute & apps. On the left, a sidebar lists various security categories: Overview, Getting started, Pricing & settings, POLICY & COMPLIANCE (Coverage, Secure score, Security policy, Regulatory compliance), RESOURCE SECURITY HYGIENE (Recommendations, Compute & apps, Networking, IoT Hubs & resources, Data & storage, Identity & access, Security solutions), ADVANCED CLOUD DEFENSE (Adaptive application controls, Just in time VM access, Adaptive network hardening, File Integrity Monitoring), and THREAT PROTECTION (Security alerts, Custom alert rules (Preview), Security alerts map (Preview)).

The main content area shows a summary of security alerts. At the top right, there is a chart showing the count of alerts by severity: High severity (18), Medium severity (38), and Low severity (47). Below the chart, a table lists the detected security alerts:

DESCRIPTION	COUNT	DETECTED BY	ENVIRONMENT
PREVIEW - Role binding to the cluster-admin role detected	1	Microsoft	Azure
PREVIEW - Privileged container detected	1	Microsoft	Azure
PREVIEW - New high privileges role detected	1	Microsoft	Azure
PREVIEW - New container in the kube-system namespace detected	1	Microsoft	Azure
PREVIEW - Container with a sensitive volume mount detected	1	Microsoft	Azure
Detected suspicious network activity	2	Microsoft	Azure
Detected suspicious network activity	2	Microsoft	Azure
Detected suspicious network activity	3	Microsoft	Azure

# Protecting against advanced cloud threats



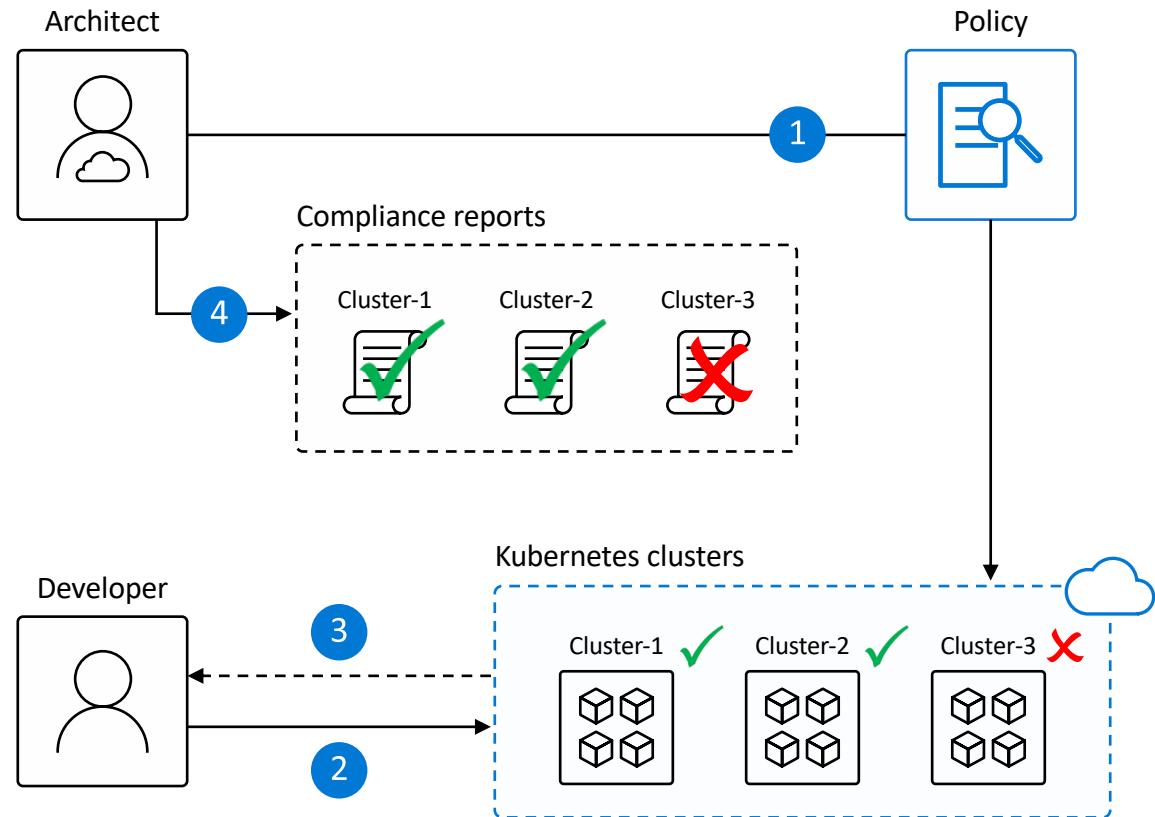
# Policy Management

## Policy management/enforcement tools and projects

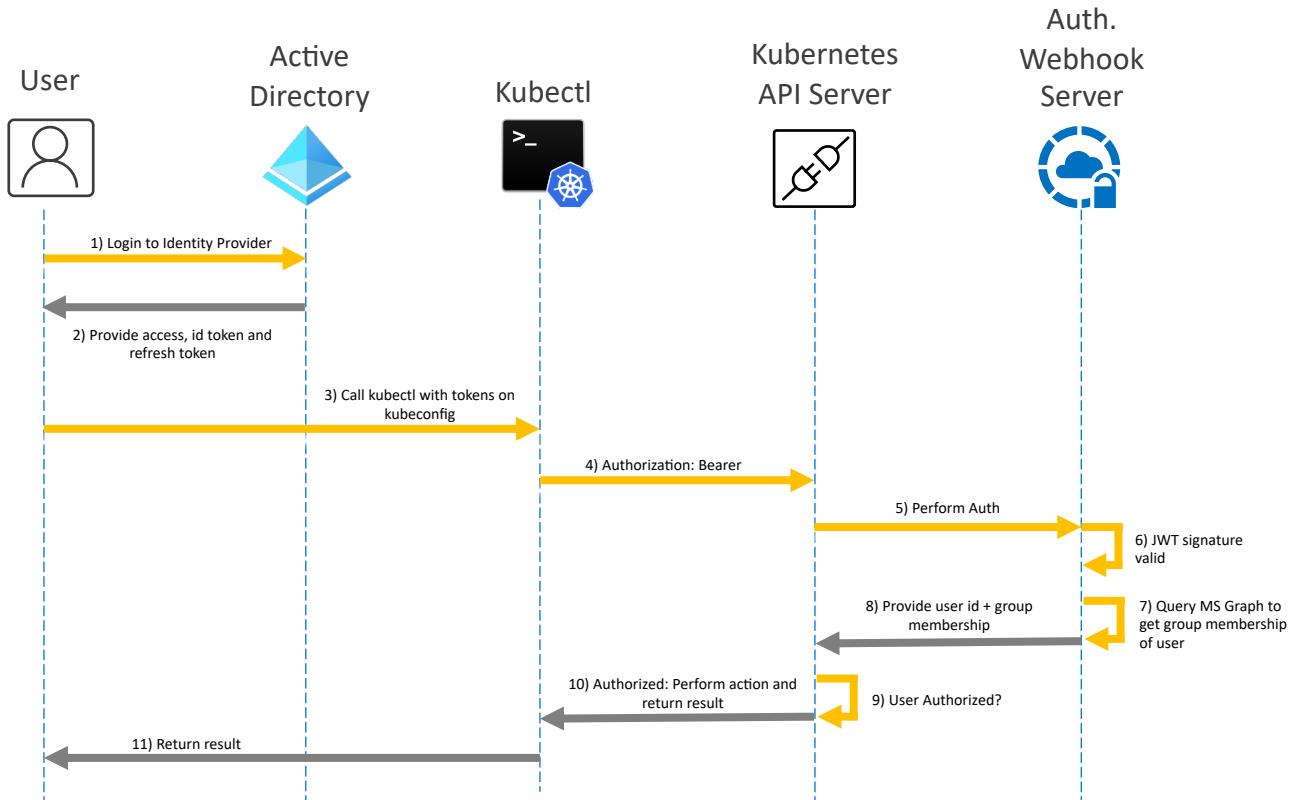
- [OPA Gatekeeper](#)
- [Kyverno](#)
- [K-rail](#)

# Policy management at scale

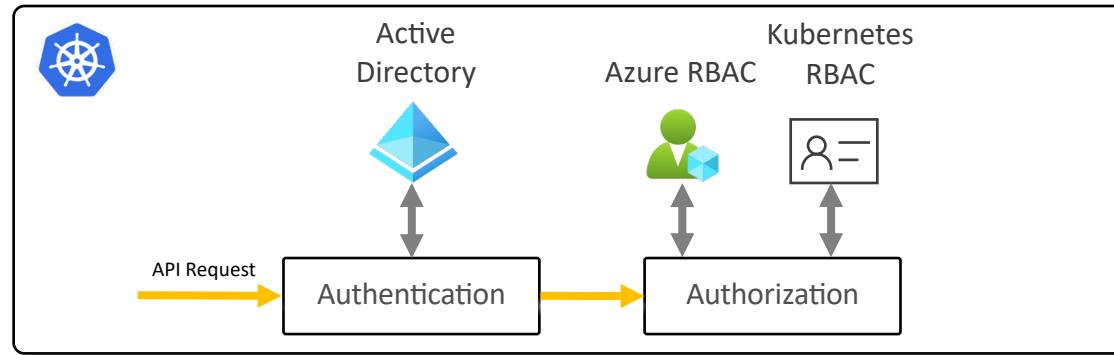
1. Architect assigns a deployment policy across cluster(s)
2. Developer uses standard Kubernetes API to deploy to the cluster
3. Real-time deployment enforcement (acceptance/denial) provided to developer based on policy
4. Cloud architect obtains compliance report for the entire environment and can drill down to individual pod level



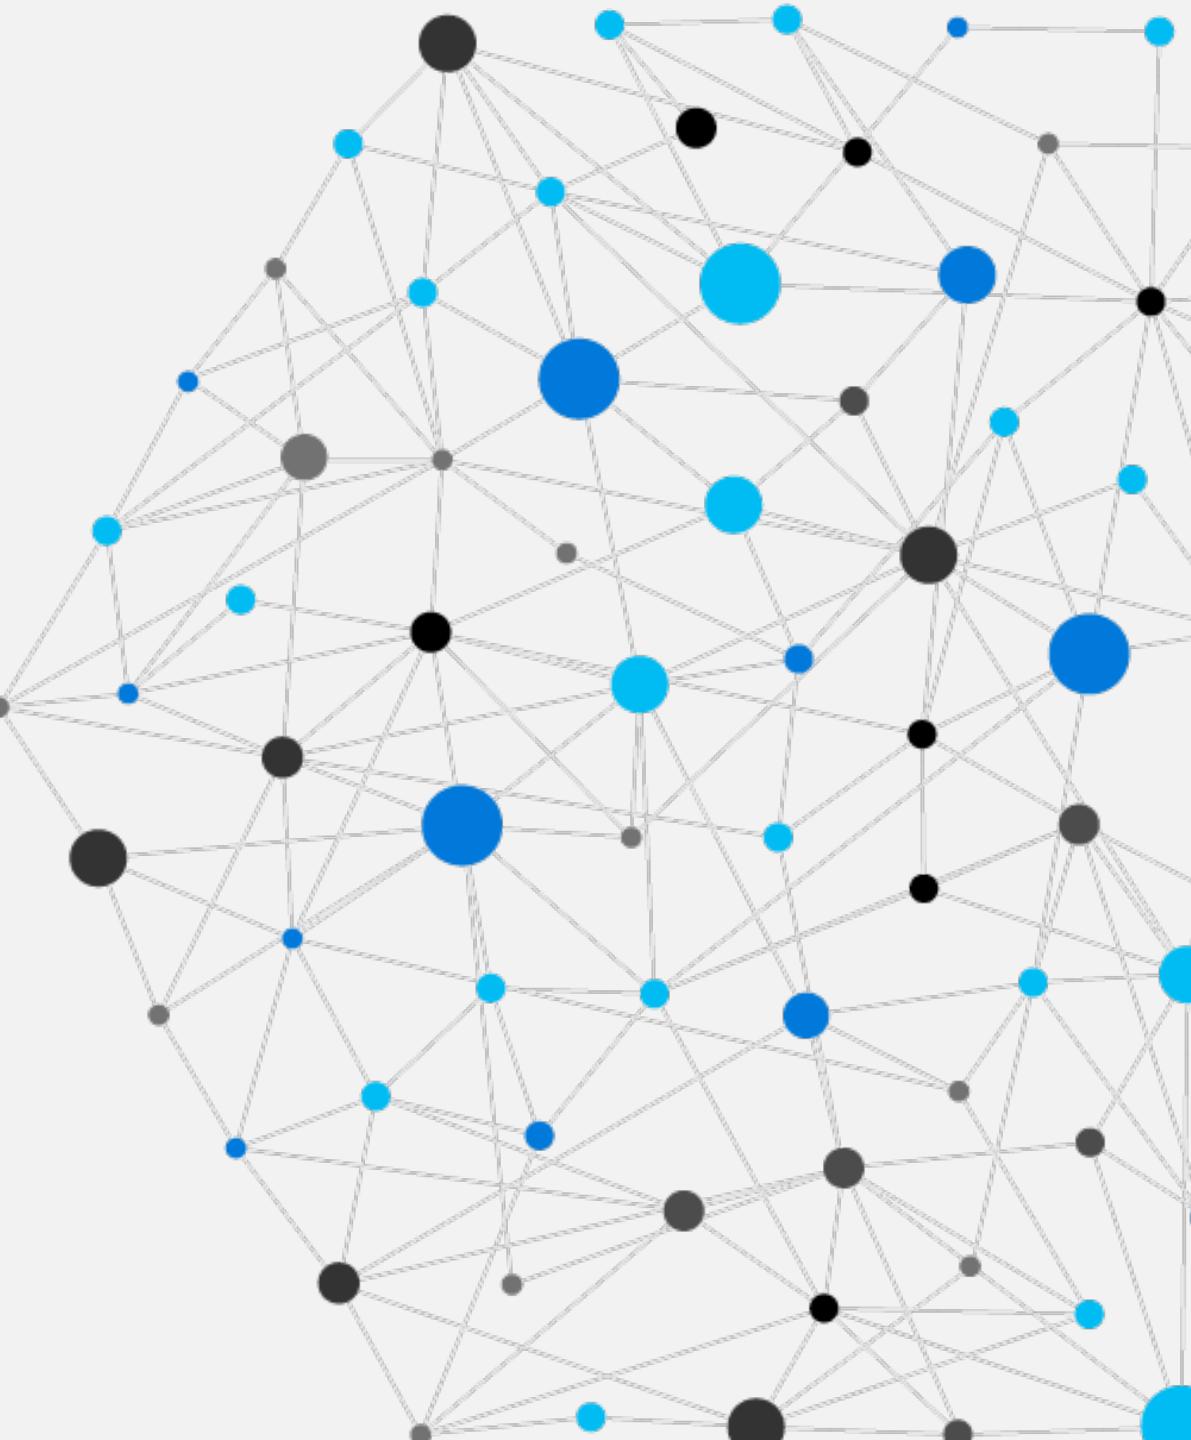
# Authentication at scale



# Authorization at scale



# Demo



# Diagnostics and troubleshooting

## Diagnostics and troubleshooting tools

- [Sysdig](#)
- [Weave Scope](#)
- [Periscope + AKS diagnostics](#)

# Built-in diagnostics and troubleshooting solutions with AKS Diagnostics

Available in the Azure Portal and via VS Code extension

- Quickly pinpoint cluster issues by analyzing backend telemetry on common issues such as networking misconfiguration, CRUD operation failures, and more
- Use red, orange, and green status icons to identify whether there is a pertinent issue
- Get recommendations for next steps and potential solutions
- Run a basic health check on your network configuration using the VS code extension or visit the portal for the deep dive

The screenshot shows a mobile device displaying the Azure Portal. The main header reads "Network and Connectivity". Below it, a sub-header says "Checks for potential network related issues for a given cluster." Two buttons are visible: "Send Feedback" and "Copy Report". A table titled "Network Configuration" lists various cluster properties and their values:

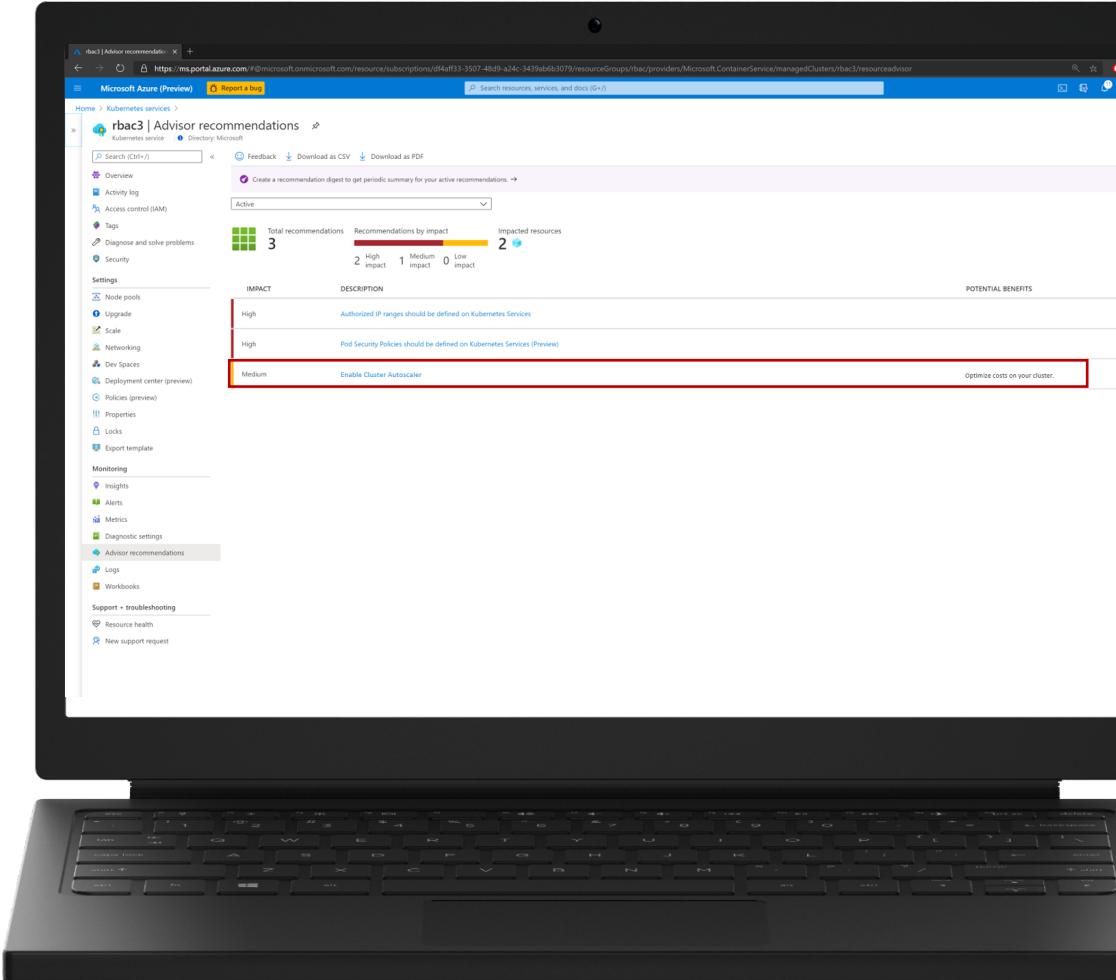
Property	Value
Region	westus
Network Plugin	Kubenet
Cluster Subnet Address Range	10.123.0.0/16
VNet Name	aks-vnet-1234567
VNet Address Ranges	0.0.0.0/8
NSG Name	aks-agentpool-1234567-nsg
Route Table Name	aks-agentpool-1234567-routetable
Allocated Outbound Ports	0 (default)

At the bottom, two notifications are shown:

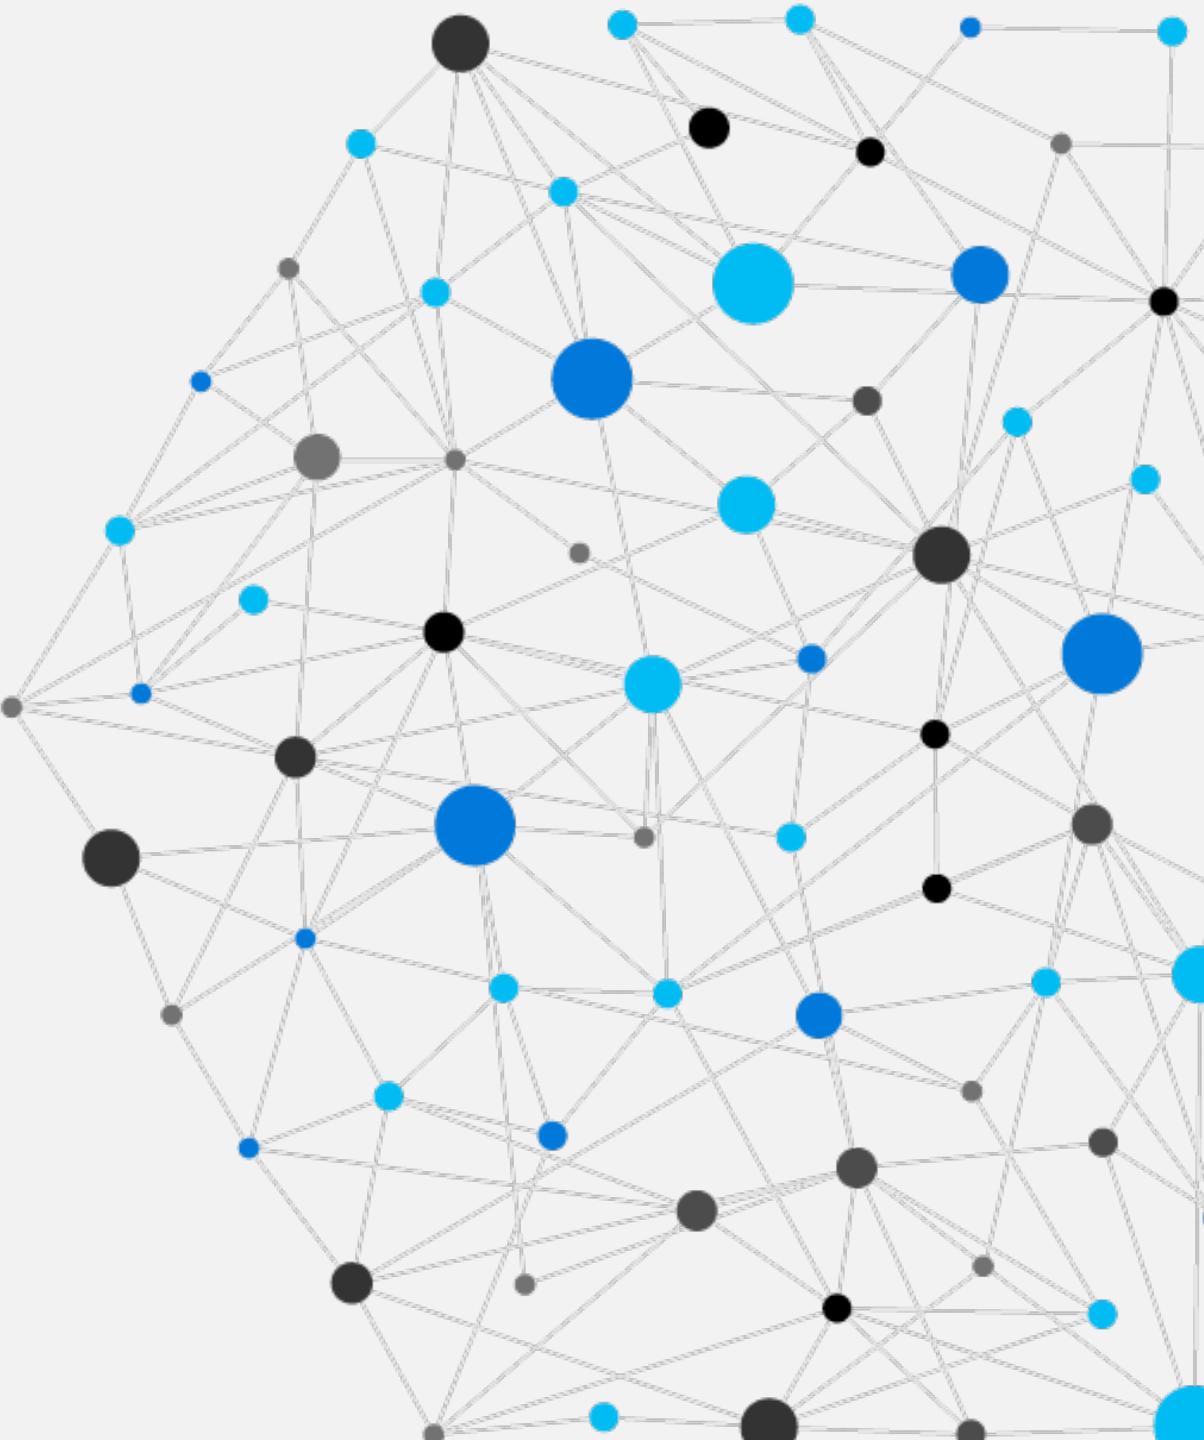
- A blue info icon with a checkmark: "Allocated Outbound Ports is set to the default value" followed by a thumbs up and thumbs down icon.
- A green checkmark icon with a checkmark: "Our analysis did not find any issues with the Kubenet subnet configuration" followed by a thumbs up and thumbs down icon.

# Efficiency and cost savings

- Use Node Groups/Pools with VMs right sized to the workload of that pool(s).
- Leverage Cluster Autoscaler and Horizontal Pod Autoscaler to maximize your cluster infrastructure.
- Use projects like Virtual Kubelet to quickly respond to temporary bursts
- Use Descheduler to respond to unbalanced or unoptimized clusters.
- Use more fit for purpose scaling projects like KEDA to optimize scaling based on your workload architecture and needs.



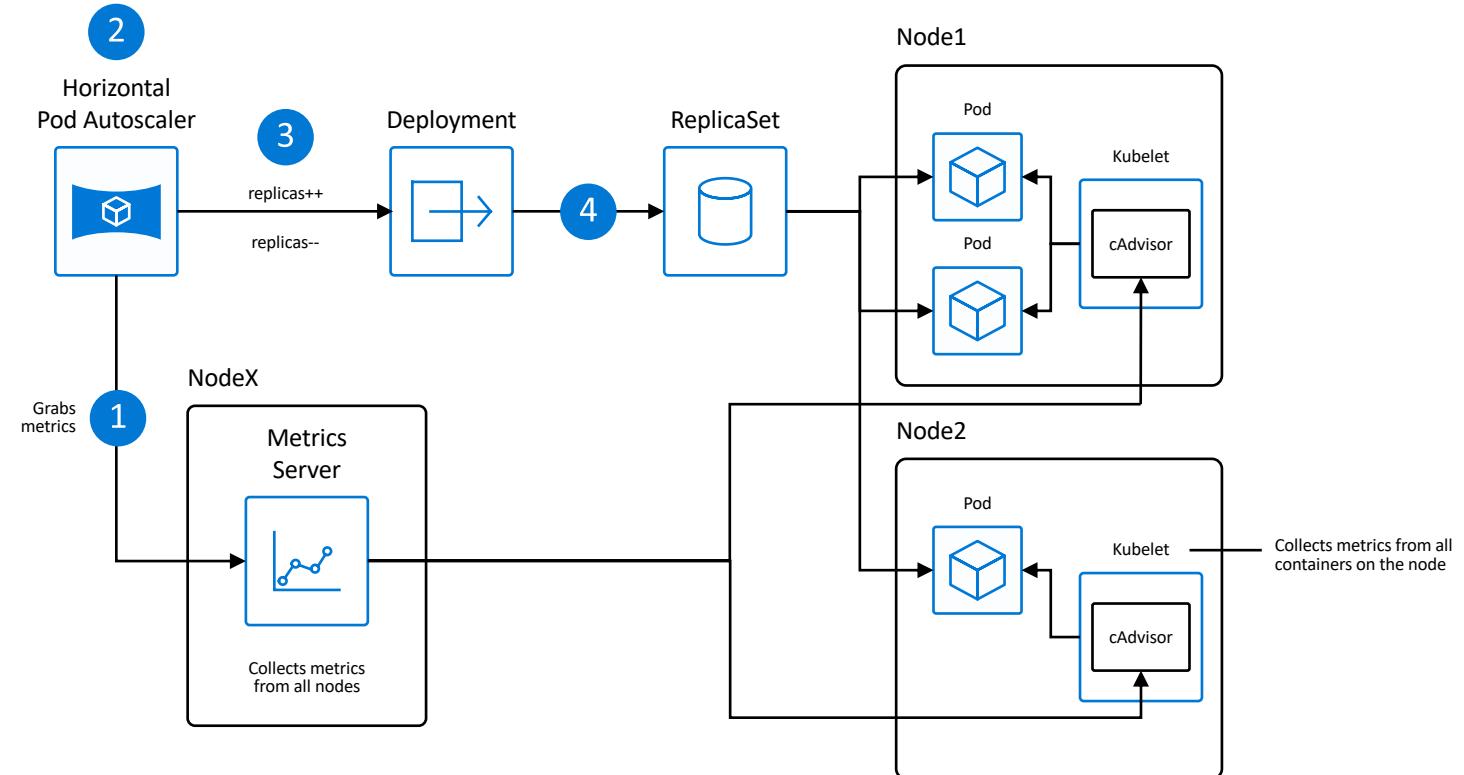
# Efficiency and cost savings



# Horizontal Pod Autoscaler

The horizontal pod autoscaler (HPA) uses the Metrics Server in a Kubernetes cluster to monitor the resource demand of pods. If a service needs more resources, the number of pods is automatically increased to meet the demand.

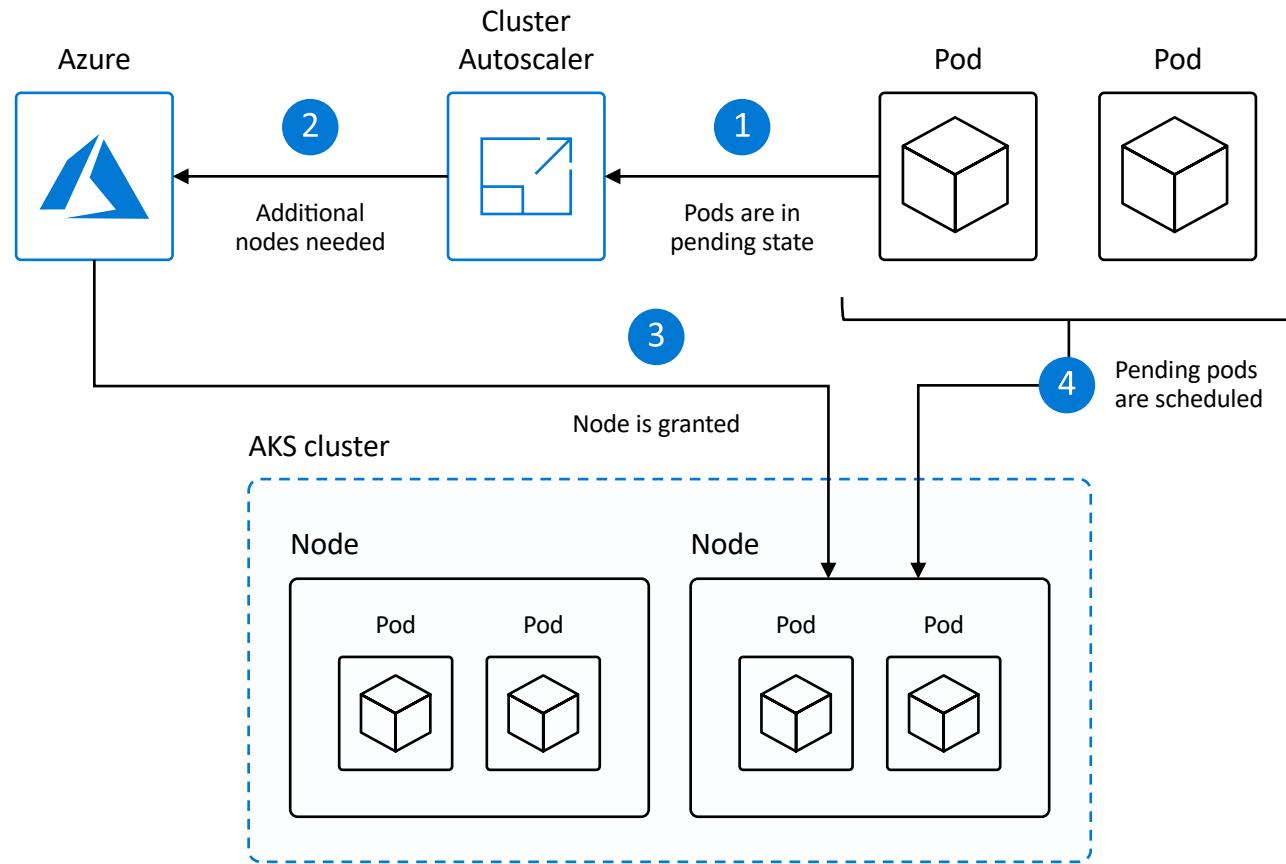
1. HPA obtains resource metrics and compares them to user-specified threshold
2. HPA evaluates whether user specified threshold is met or not
3. HPA increases/decreases the replicas based on the specified threshold
4. The Deployment controller adjusts the deployment based on increase/decrease in replicas



# Cluster Autoscaler

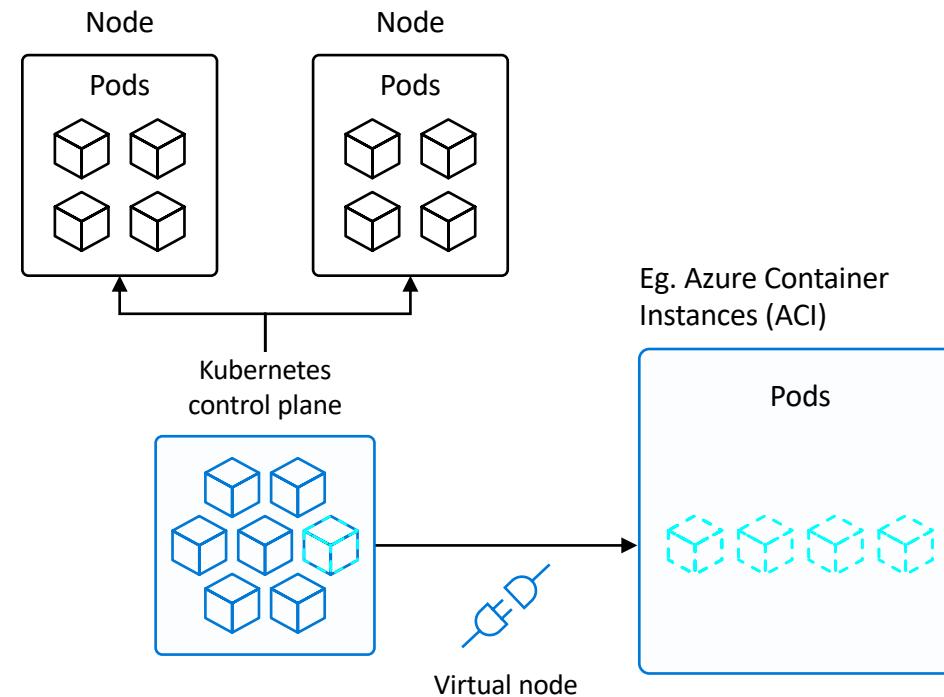
The cluster autoscaler watches for pods that can't be scheduled on nodes because of resource constraints. The cluster then automatically increases the number of nodes.

1. HPA obtains resource metrics and compares them to user-specified threshold
2. HPA evaluates whether user specified threshold is met or not
3. HPA increases/decreases the replicas based on the specified threshold
4. The Deployment controller adjusts the deployment based on increase/decrease in replicas



# Nodeless Kubernetes using Virtual Kubelet

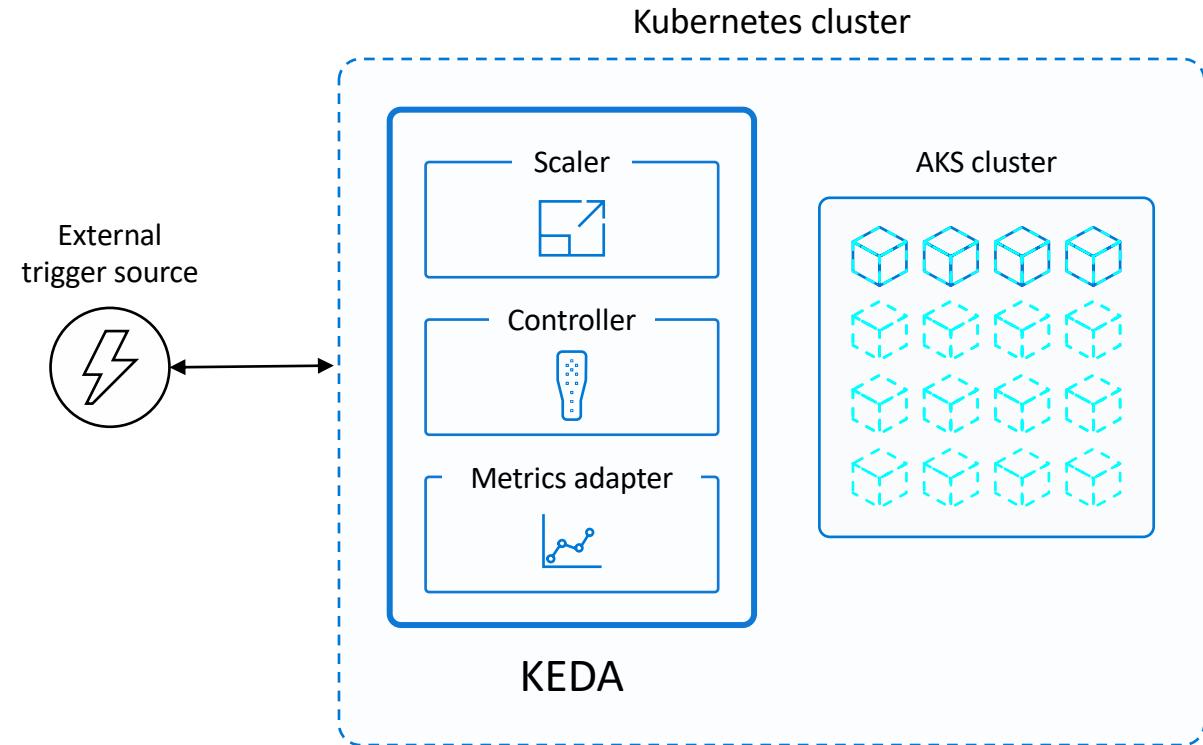
- Elastically provision compute capacity in seconds
- No infrastructure to manage
- Built on open sourced Virtual Kubelet technology, Cloud Native Computing Foundation (CNCF) sandbox project
- Burst faster than what would normally require infrastructure to provision



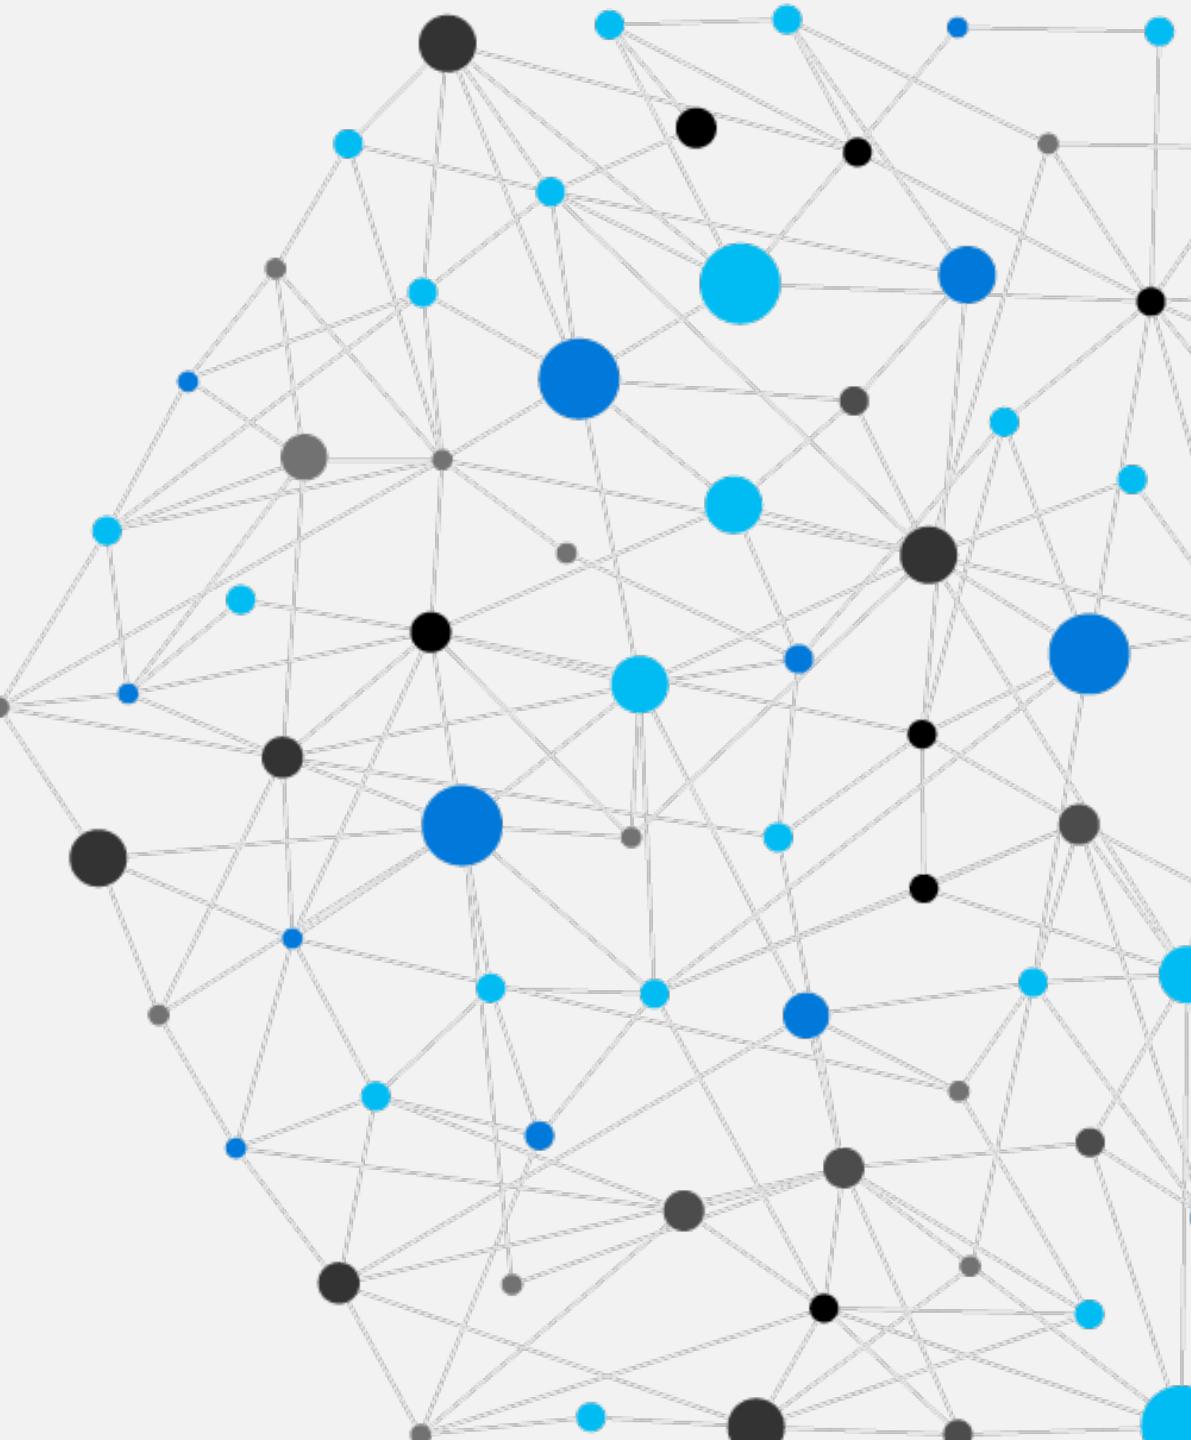
# Kubernetes-based event-driven auto-scaling (KEDA)

Open-source component jointly built by Microsoft and RedHat

- Event-driven container creation & scaling  
Allows containers to “scale to zero” until an event comes in, which will then create the container and process the event, resulting in more efficient utilization and reduced costs
- Native triggers support  
Containers can consume events directly from the event source, instead of routing events through HTTP
- Can be used in any Kubernetes service  
This includes in the cloud (e.g., AKS, EKS, GKE, etc.) or on-premises with OpenShift—any Kubernetes workload that requires scaling by events instead of traditional CPU or memory scaling can leverage this component.

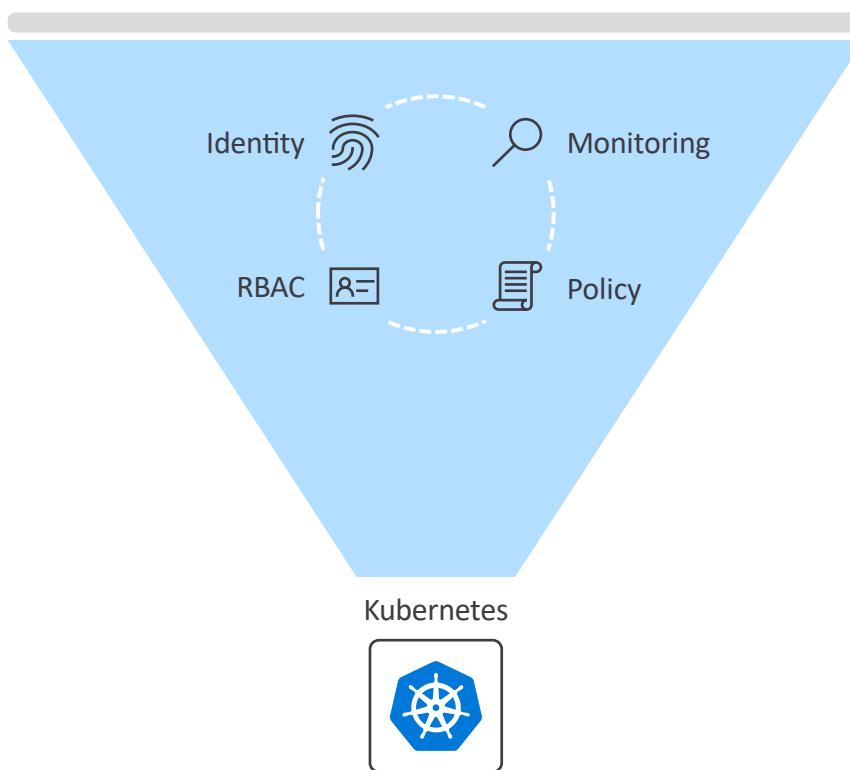


# Demo



# Takeaways

- Consistently apply policies, role-based-access-controls (RBAC) for at-scale governance
- Deploy Kubernetes configuration to all clusters using a GitOps-based workflow
- Enforcing enterprise policies via Policy management tools like OPA Gatekeeper, Kyverno, etc
- Directory integration for centralized identity management at scale
- Workload-specific tuned autoscaling features and fit for purpose projects like Virtual Kubelet, Descheduler and KEDA to maximize the efficiency of clusters and resources



# Resources

Kubernetes on Azure

<https://azure.microsoft.com/en-us/overview/kubernetes-on-azure/>

Kubernetes learnings and training

<https://azure.microsoft.com/en-us/resources/kubernetes-learning-and-training/>

Best Practices for Azure Kubernetes Services (AKS)

<https://aka.ms/aks/bestpractices>

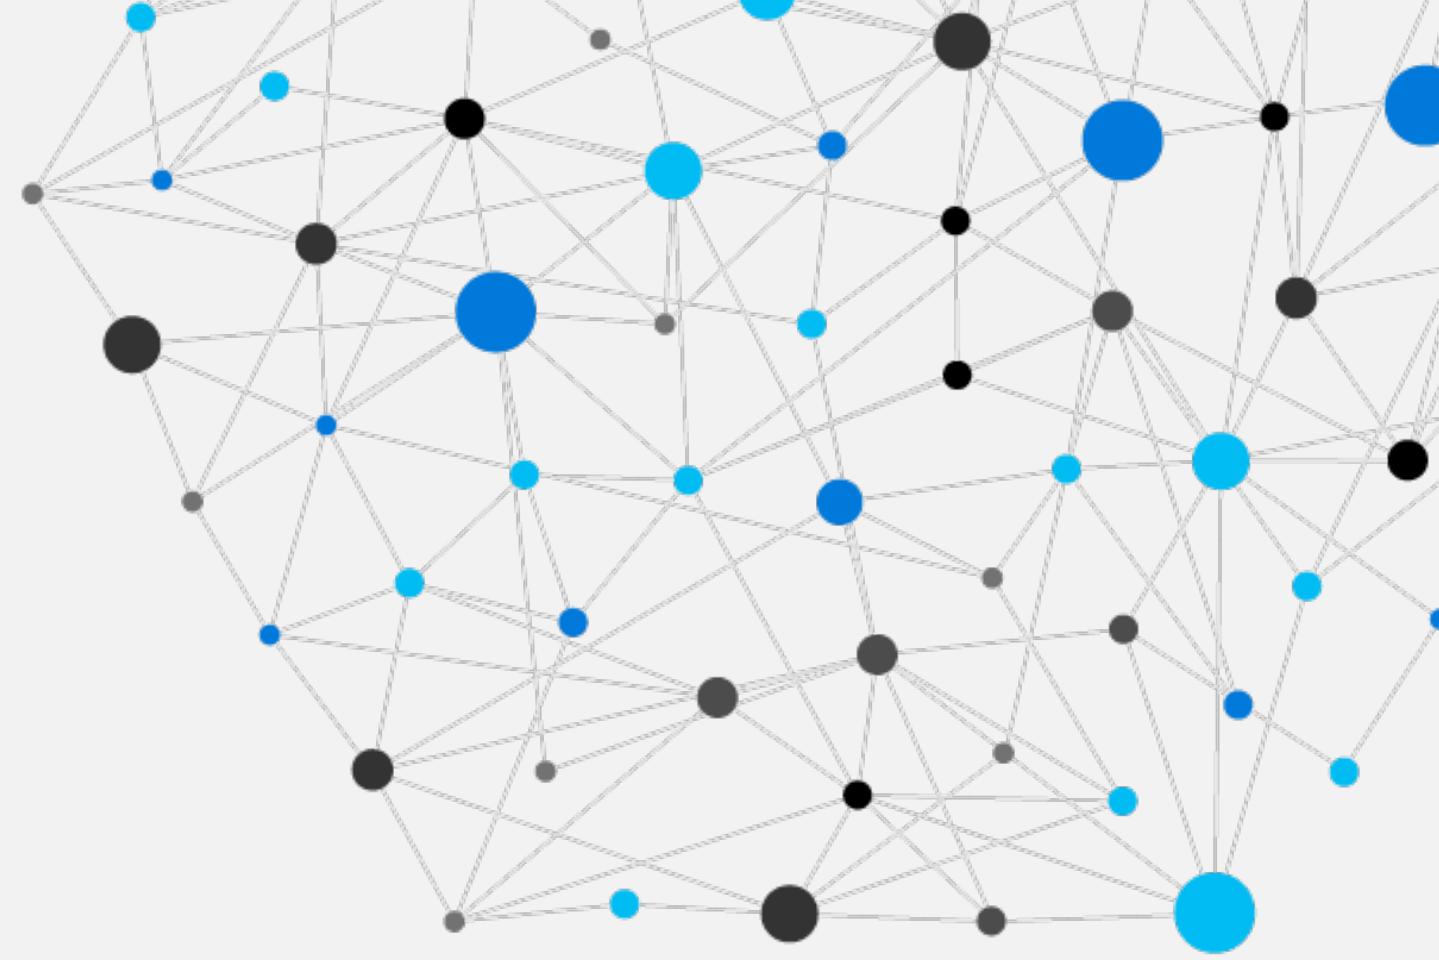
Azure Kubernetes Services (AKS) Diagnostics

<https://docs.microsoft.com/en-us/azure/aks/concepts-diagnostics>

Azure Security Center and Azure Kubernetes Service

<https://docs.microsoft.com/en-us/azure/security-center/azure-kubernetes-service-integration>

# Q&A





Thank you for joining us.