

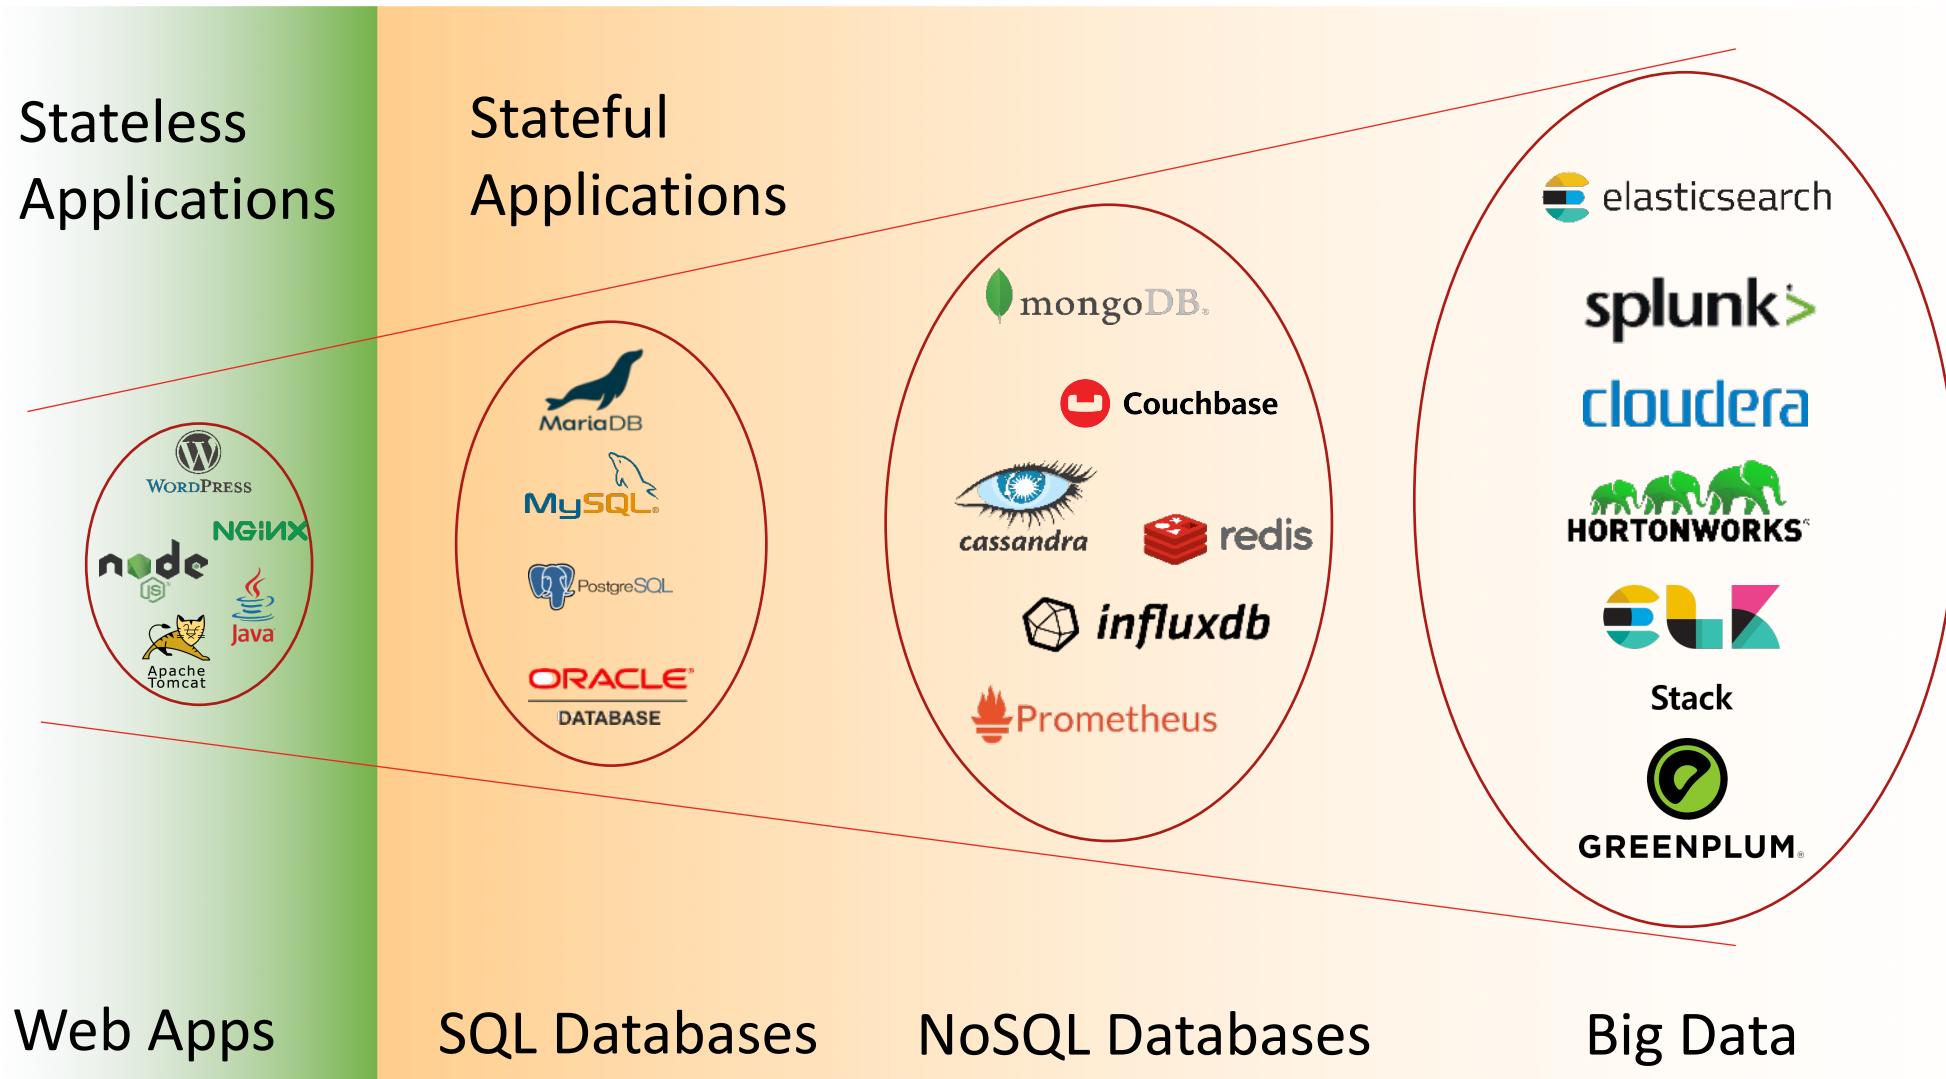
# Application Modernization And Portability Across Clouds

Ravikumar Alluboyina

Senior Product Architect, Robin.io  
<https://get.robin.io>



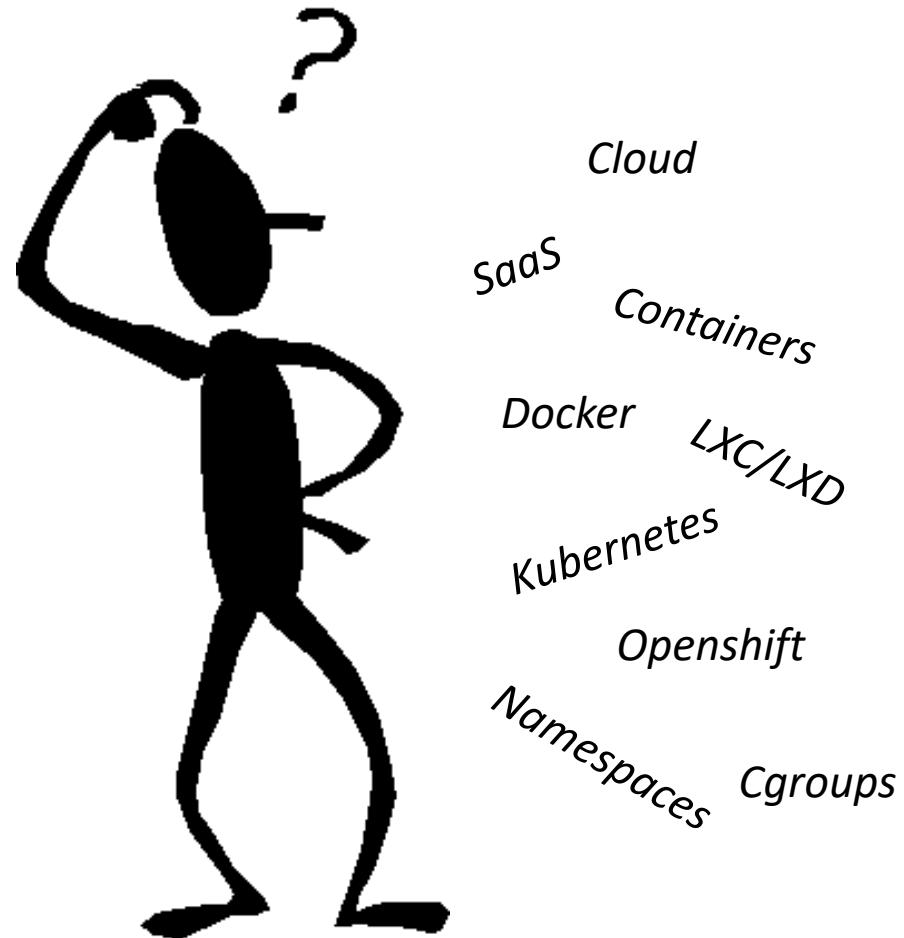
# Spectrum of Applications



# Application Modernization

# Modernization.. The trend

- › **What** is Modernization?
  - › Define modernization?
  - › Is there a standard?
- › **Why** would I want to modernize my application?
  - › What will I gain by modernizing
- › **How** can I modernize my application?
  - › What are the tools available?
  - › What processes should I adopt?



**Q.** Virtualization did not force me to modernize my applications, why now? What changed?

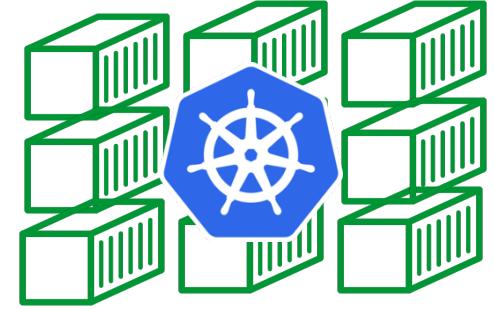
# Modern Application .. The ‘What’

**Application that can adapt to any environment and perform equally well**

- › Bare metal, Virtual machines, Cloud, Containers
- › Mode of delivery (rpm, deb, tar vs runnable formats vmdk, docker container etc)
- › Can integrate easily with pluggable infrastructure component.
  - › Security, Performance monitoring, Logging, backup etc.,
- › Is elastic, can grow and shrink depending on the load and usage
- › Is agile / portable

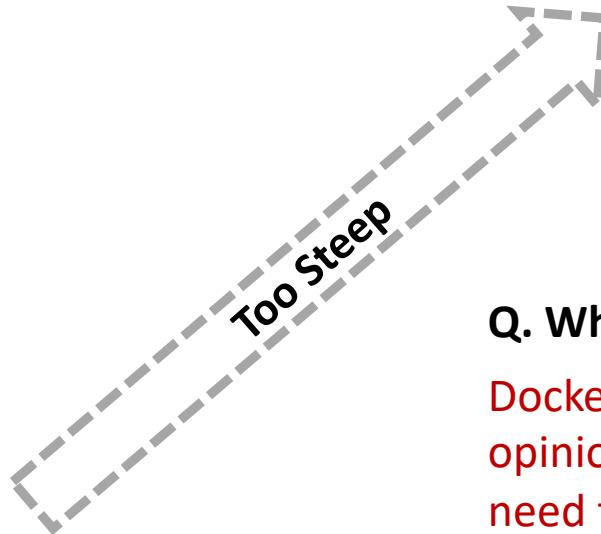
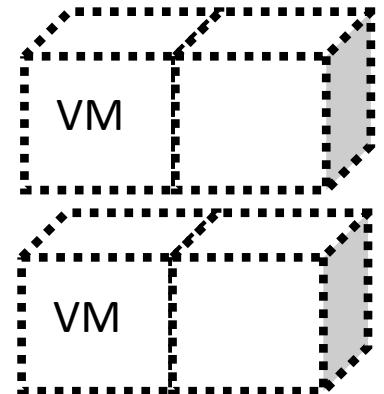
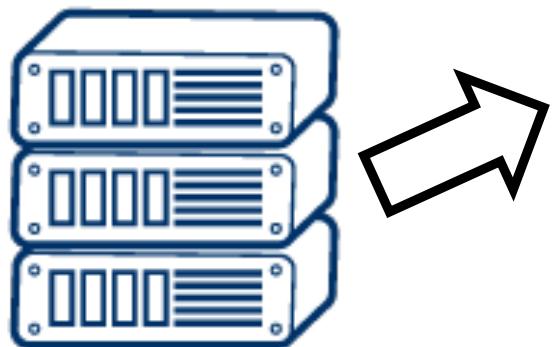
# Why now, what changed?

- › The infrastructure landscape and tooling has changed drastically
  - › We have Cloud, Virtual Machines, Containers, Dockers
  - › Security, Logging, Monitoring frameworks have evolved



## Q. Why is this a smooth transition for applications?

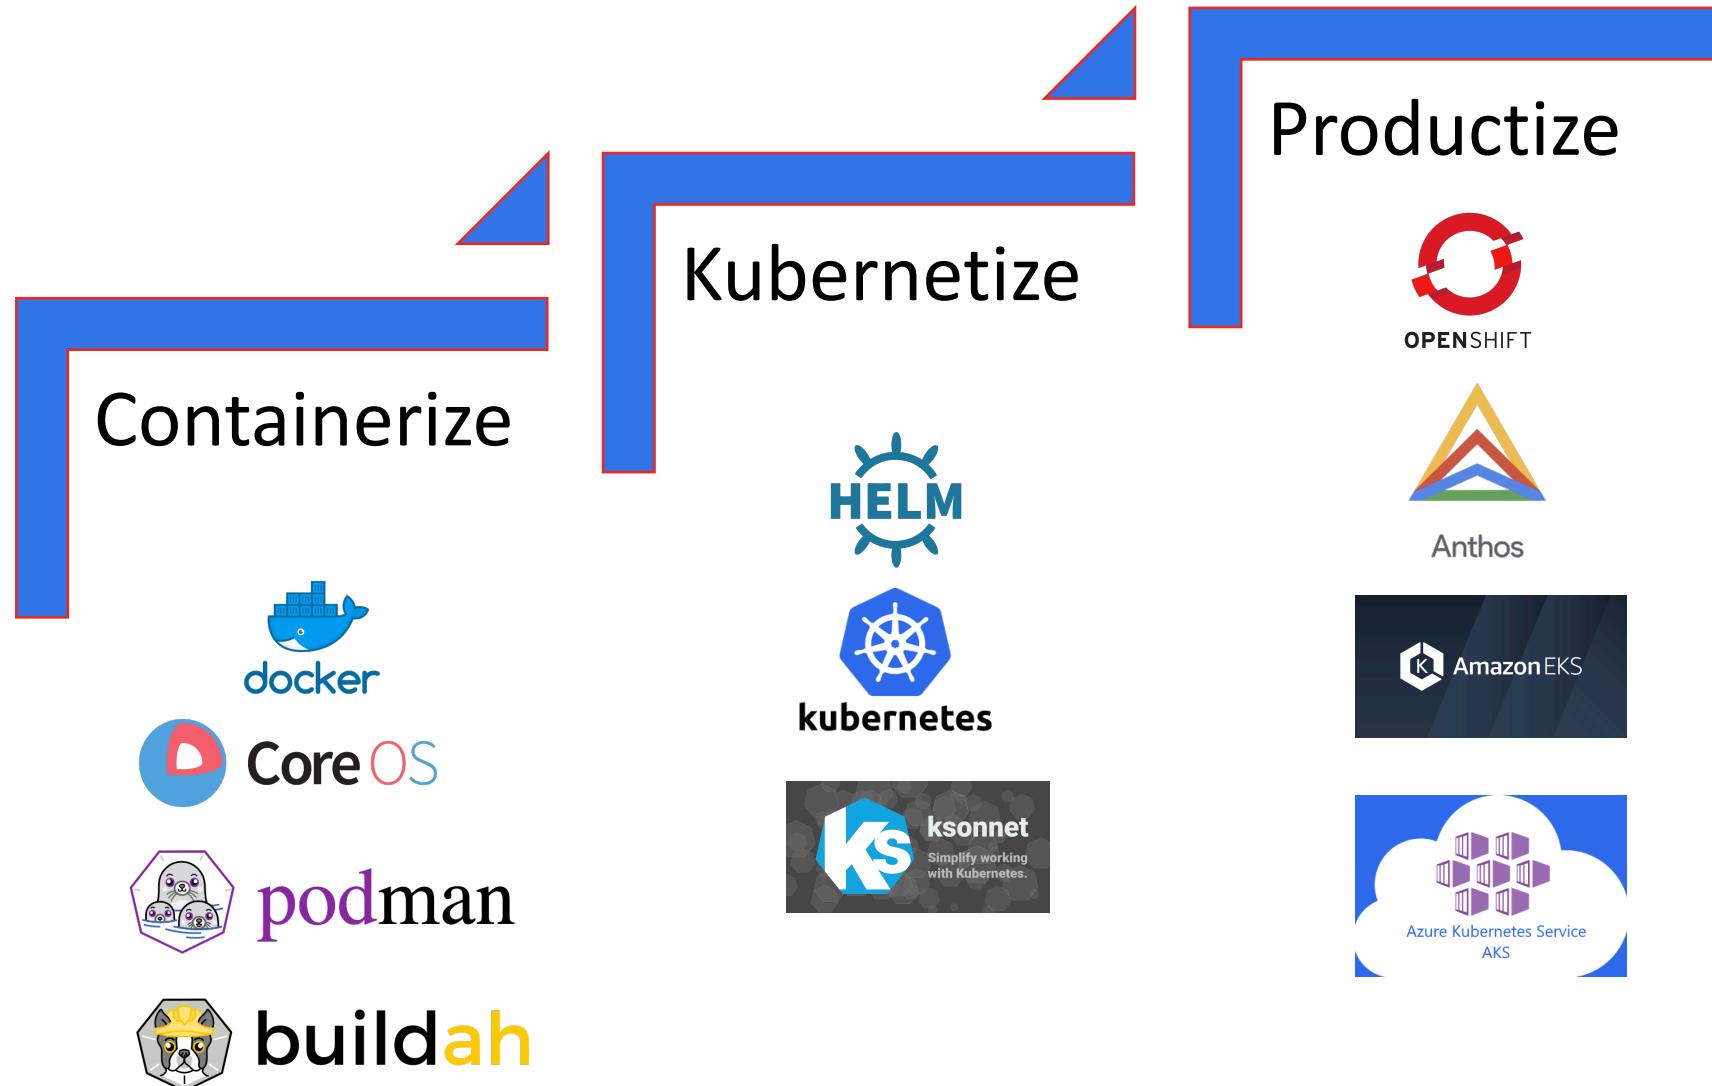
Application didn't need to change. Infrastructure emulated the baremetal



## Q. Why is this a huge leap?

Docker / CRI-O, Kubernetes are all very opinionated frameworks. Applications need to be re-designed for these platforms.

# Modernization .. The ‘How’



# Containerization

# Containerization: The process

- › Process separation
  - › Separate the application into multiple processes (if possible and if it makes sense)
  - › For example:
    - › Separate the nodemanager and datanode in Hadoop,
    - › (or) query router and config server in MongoDB
- › Configuration separation
  - › Find the configuration paths (/etc, /var, etc) and persist it on a different volume (on the host)
  - › This is by far the most difficult process for legacy applications
- › Data separation
  - › Attach volumes to the container so data can be preserved across restarts
- › Entrypoint
  - › Add enough logic in entrypoint to handle first-start, subsequent starts and upgrades.

Did you really think this was going to be enough to Dockerize?

# Docker: The real issues

**Incomplete cgroups virtualization causes many Big Data and Databases to misbehave**

## CPU

- › Contiguous core IDs, CPU ID mapping (Kudu), accurate threads: cores mapping (DB)
- › NUMA aware assignment (HANA)

## Memory:

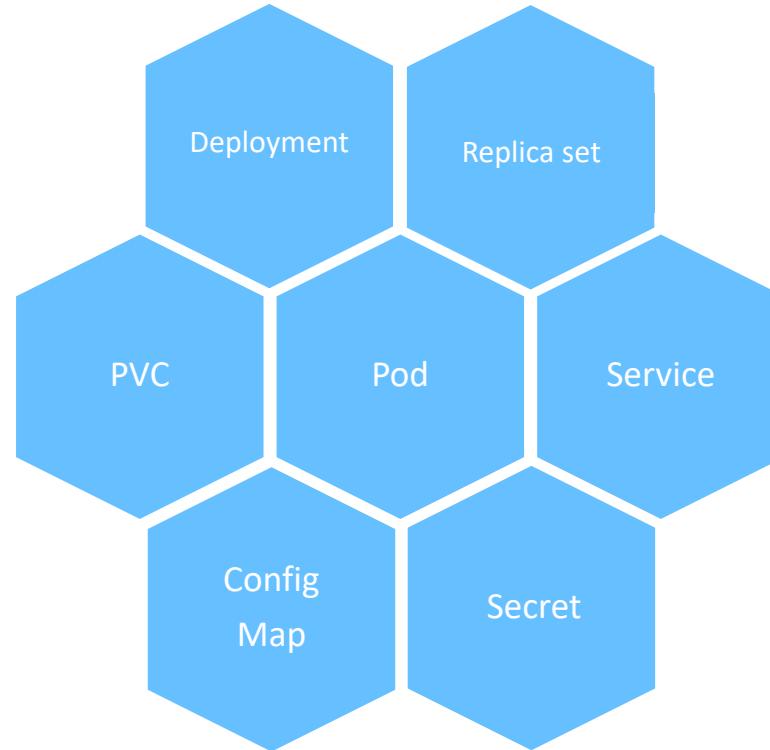
- › JVM sees entire host memory even if you cap the memory for container (Any JVM app)
- › Memory allocation inconsistencies (hugepages, shared page cache) (Oracle)

## Storage

- › Apps that need raw block devices need correct WWNs management (e.g., Oracle, MapR)
- › blkio cgroups setting is useless to avoid noisy neighbor problems (All apps)

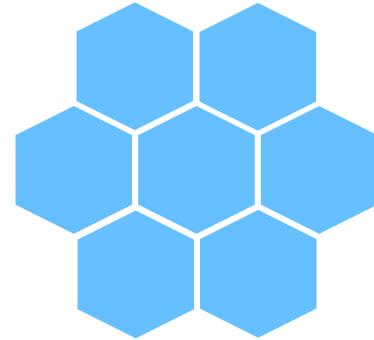
# Kubernetes

# Application Composition

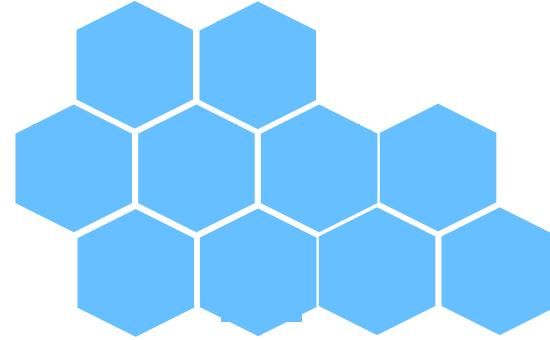


<https://github.com/helm/charts/tree/master/stable/mysql>

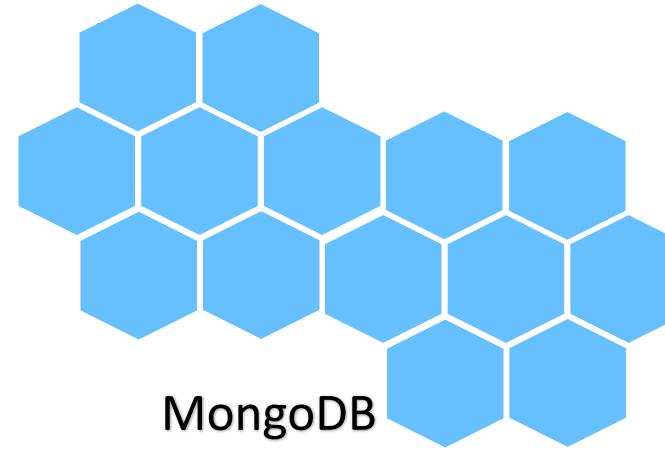
# Application Composition .. The complexity



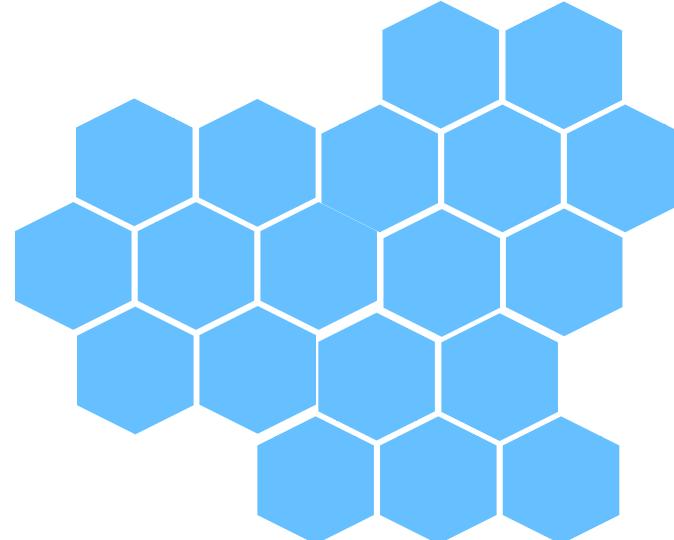
MySQL



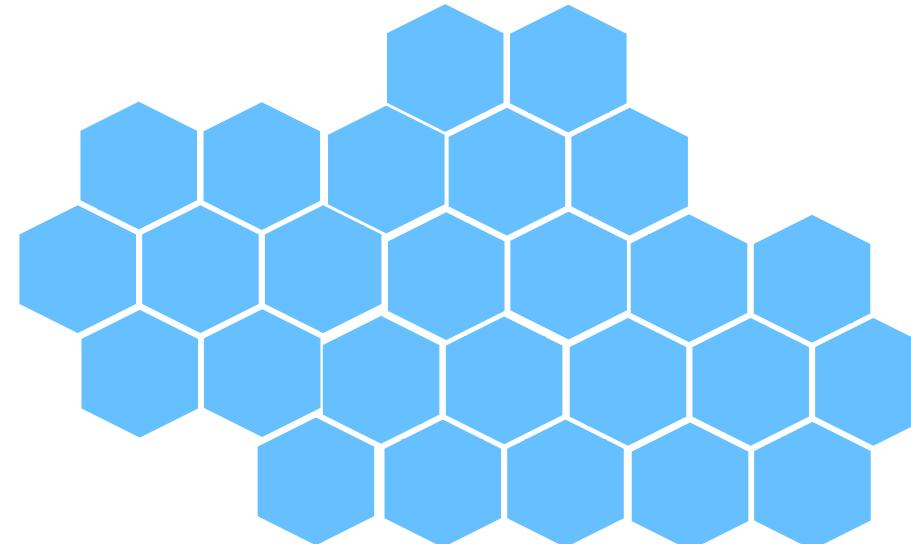
MariaDB



MongoDB



ElasticSearch

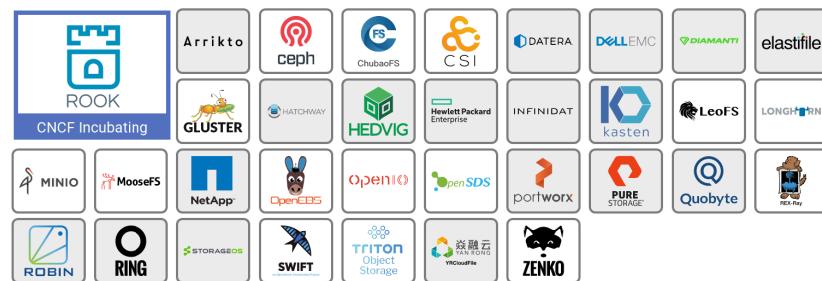


ELK Stack

# Productize

# State of Kubernetes eco-system

- There are **over 30 Storage vendors** and **20 Network vendors** providing Storage & Networking solutions for containers and Kubernetes<sup>1</sup>



Storage vendors



Network vendors

- CNCF<sup>2</sup>: 30% say Storage is a big challenge, 30% say Networking is a challenge in Kubernetes

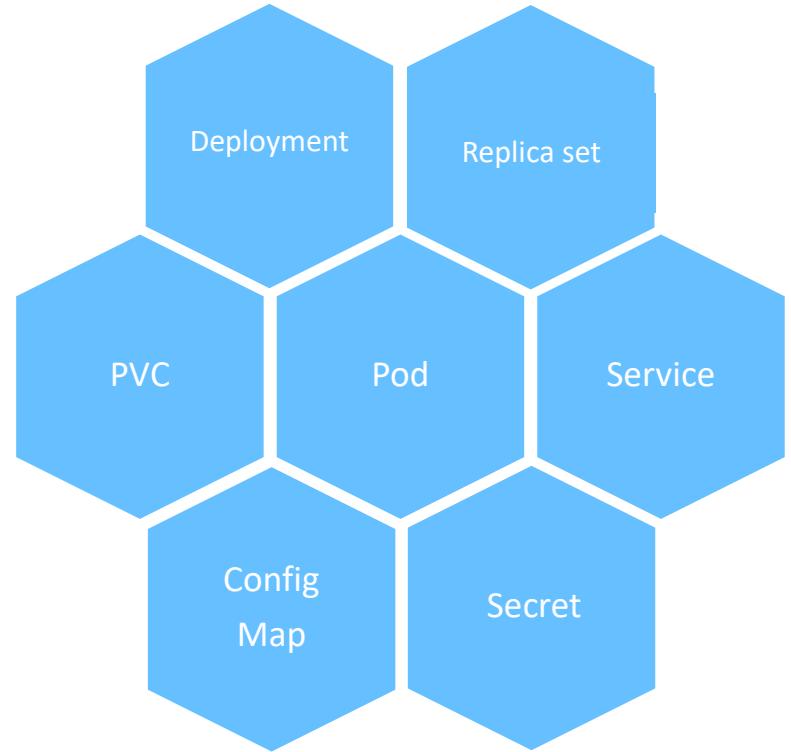
Despite so many vendor solutions, why is it still a challenge for so many people?

1 <https://github.com/cncf/landscape>

2 <https://www.cncf.io/blog/2018/08/29/cncf-survey-use-of-cloud-native-technologies-in-production-has-grown-over-200-percent/>



# MySQL Application Composition



Wait .....

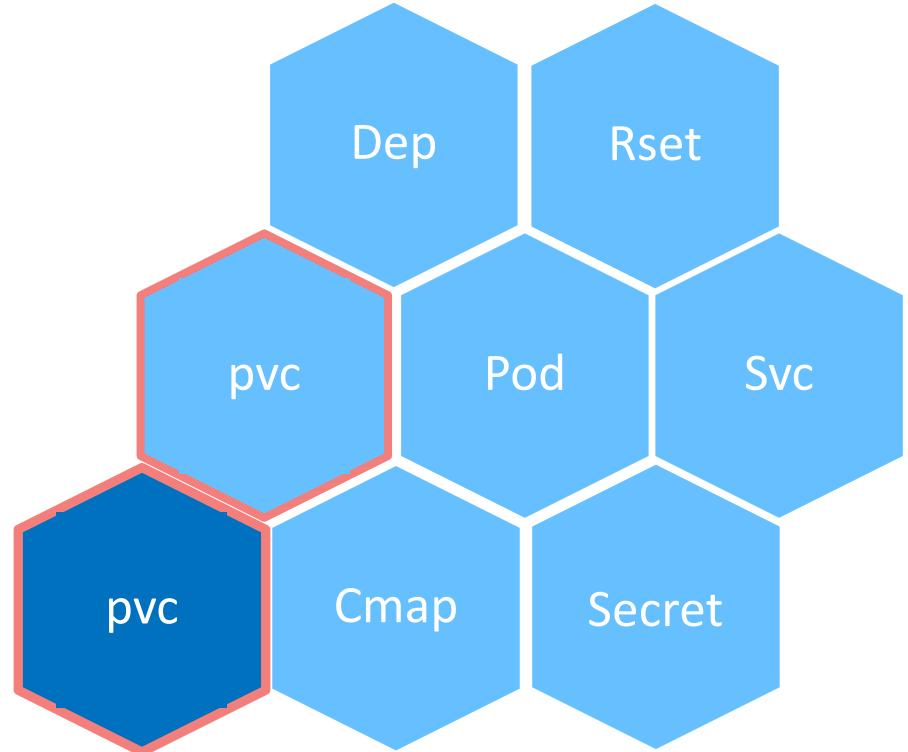
Performance issues !!!!

# What about ..... InnoDB storage recommendation

- › Non-rotational storage generally provides better performance for random I/O operations; and rotational storage for sequential I/O operations. When distributing data and log files across rotational and non-rotational storage devices, consider the type of I/O operations that are predominantly performed on each file.
- › Random I/O-oriented files typically include [file-per-table](#) and [general tablespace](#) data files, [undo tablespace](#) files, and [temporary tablespace](#) files. Sequential I/O-oriented files include InnoDB [system tablespace](#) files (due to [doublewrite buffering](#) and [change buffering](#)) and log files such as [binary log](#) files and [redo log](#) files.

Let us be a good MySQL Admin and  
follow recommendations

# MySQL Application Composition



Wait .....

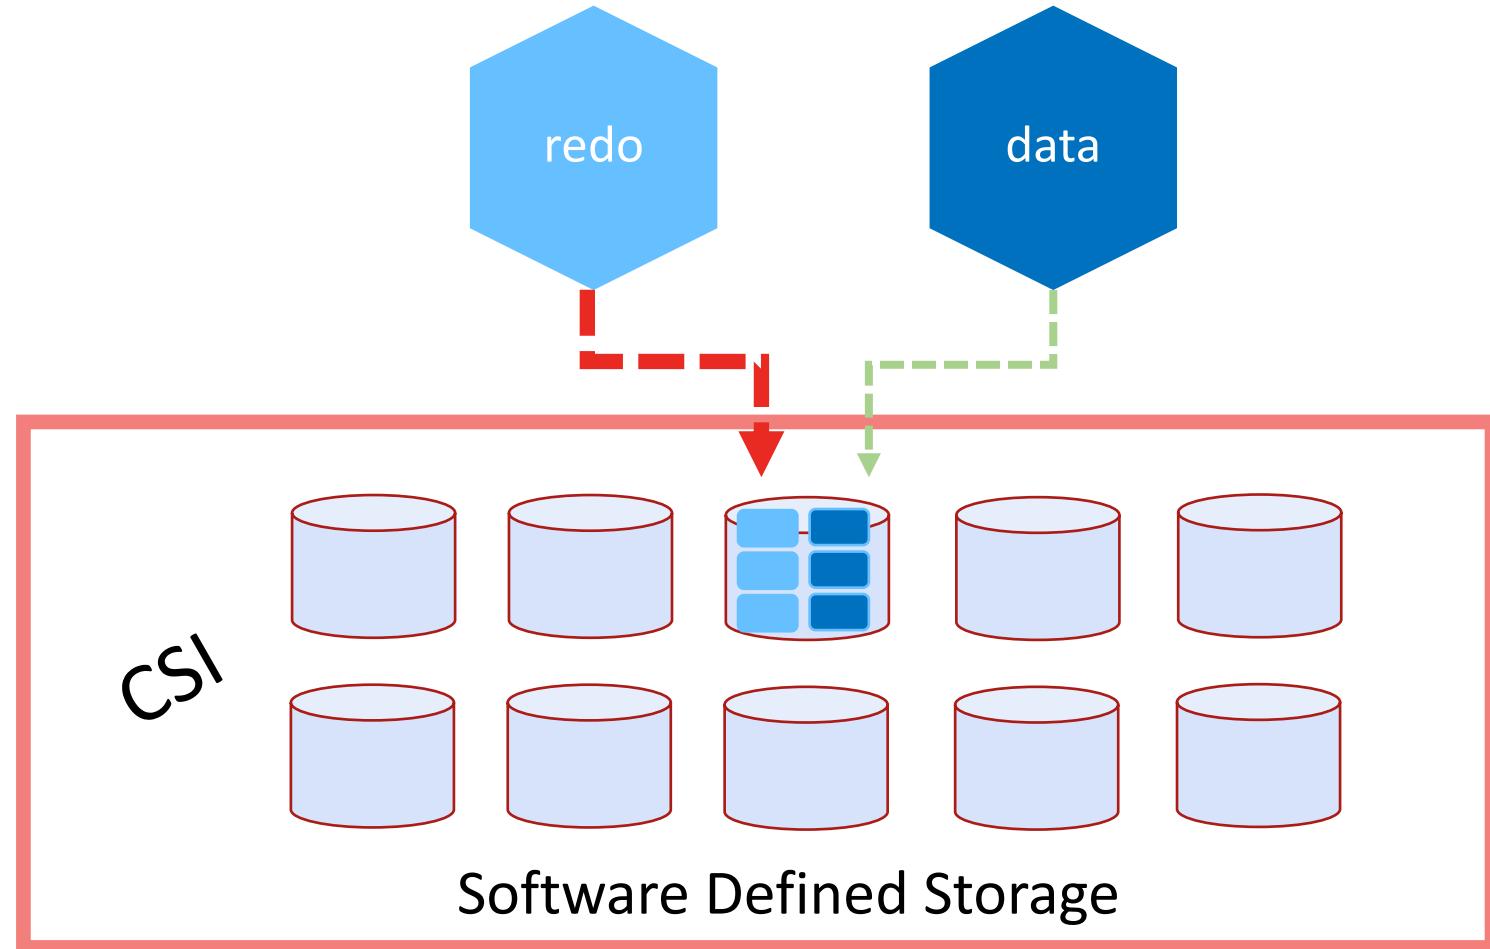
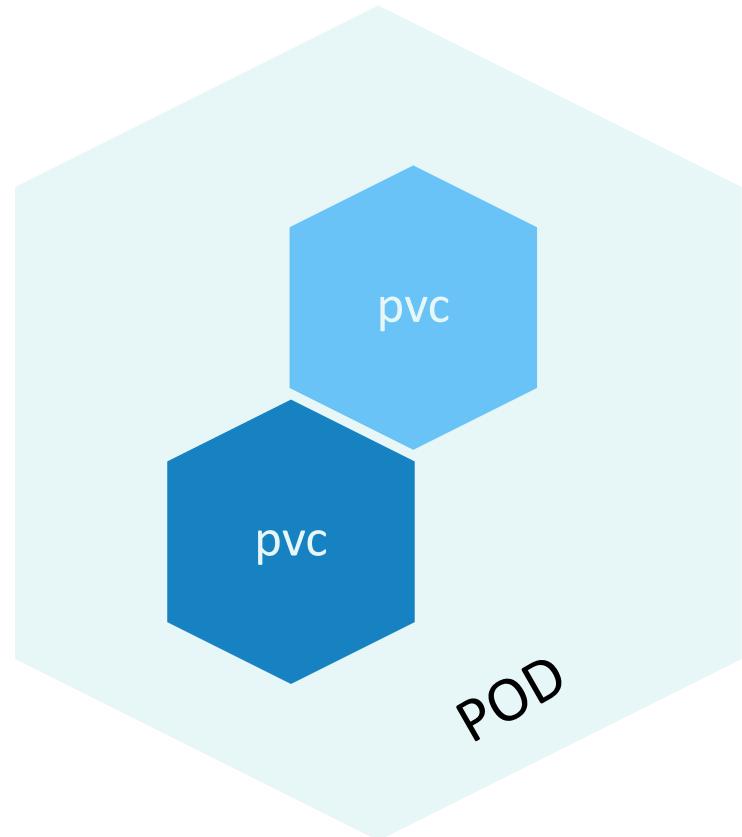
Performance issues !!!!



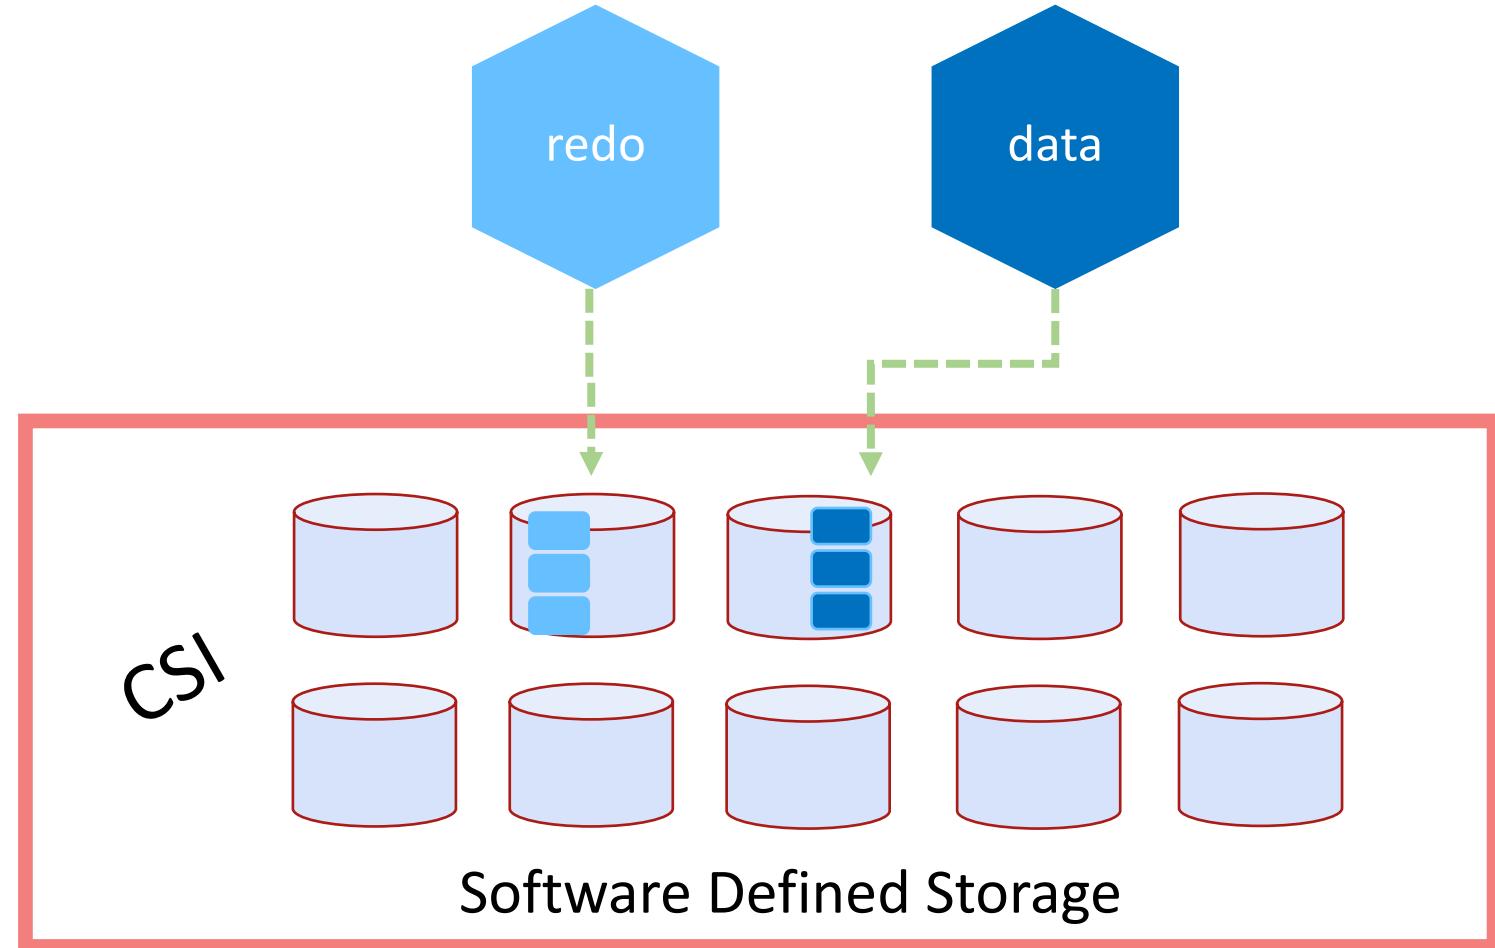
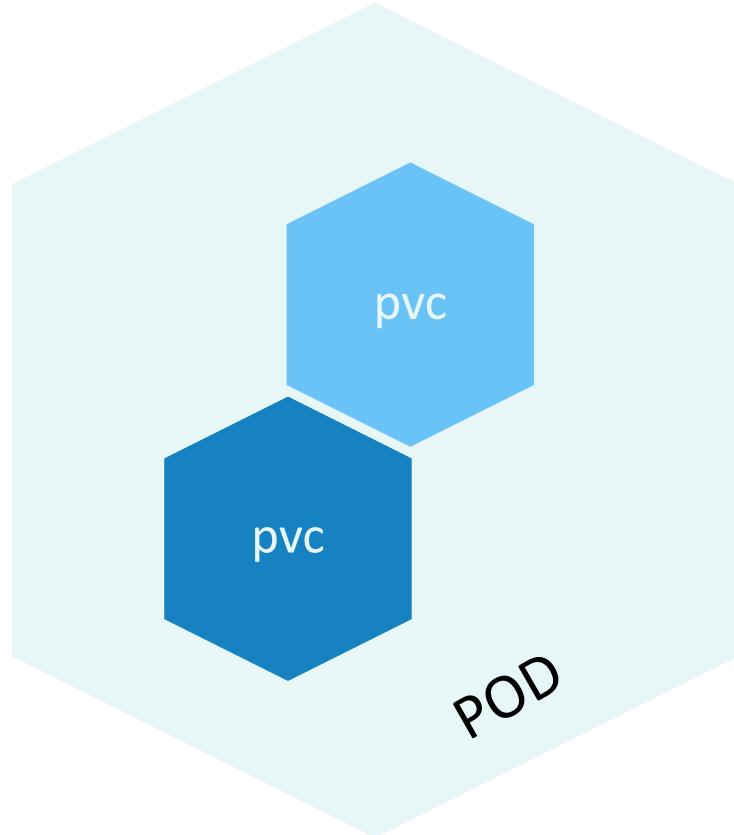
Any guesses ??

Let us look at it more closely

# Troubleshooting .....



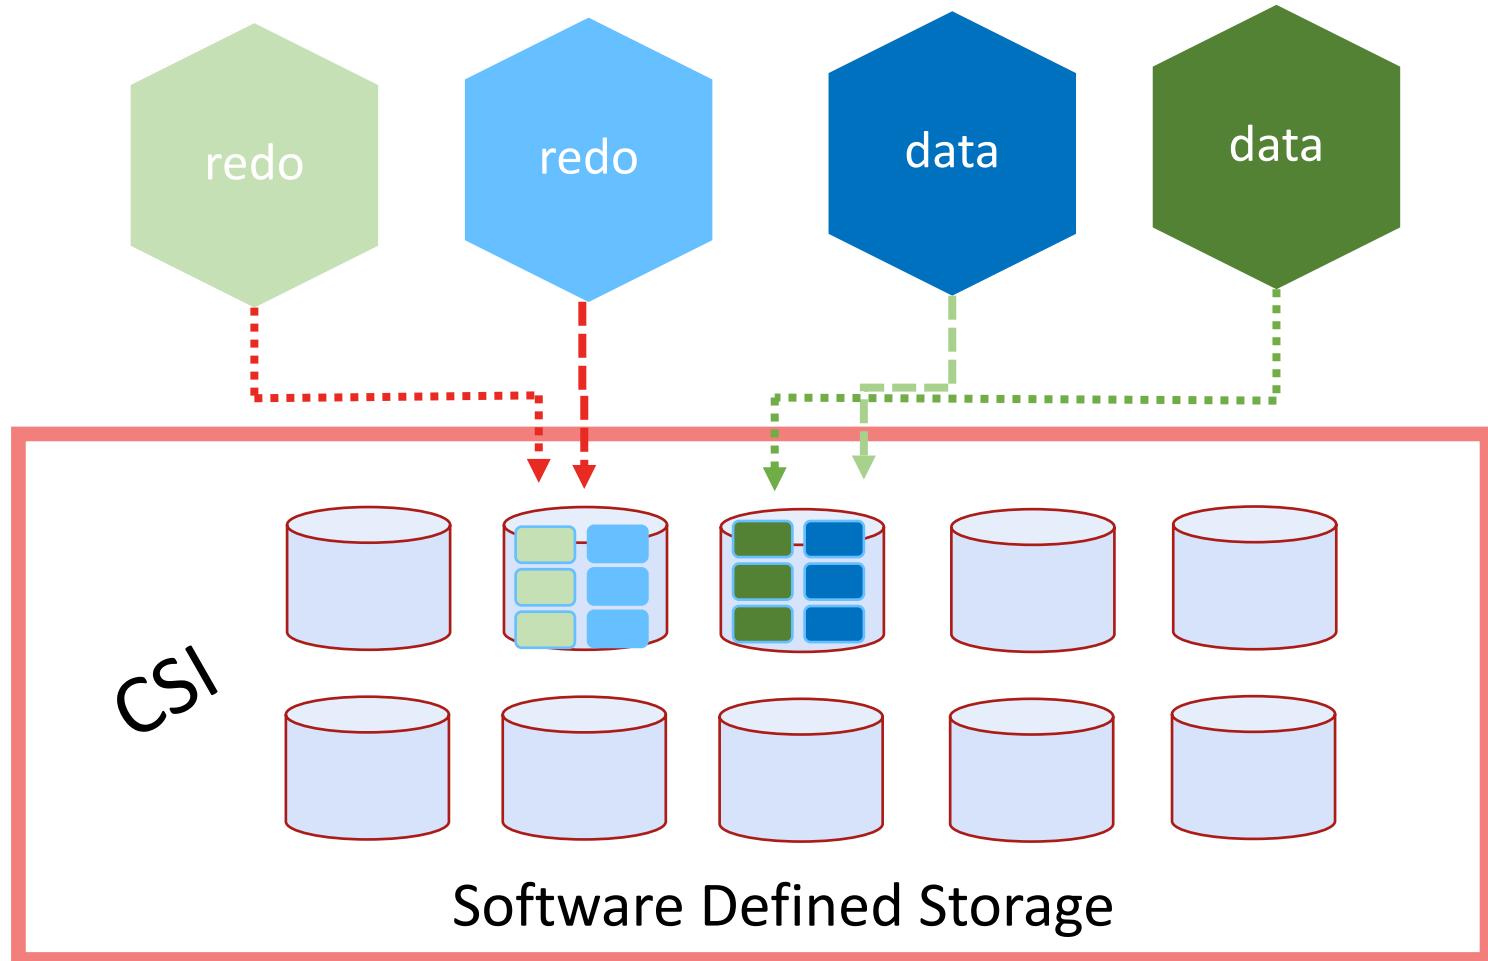
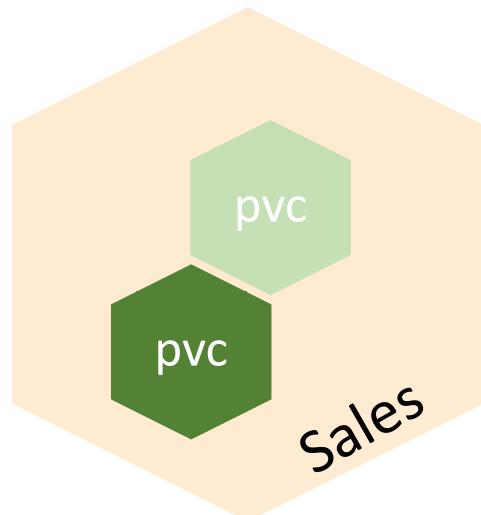
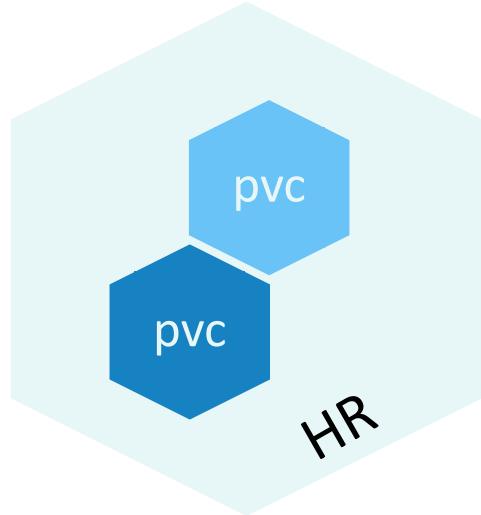
# Let us fix it ..



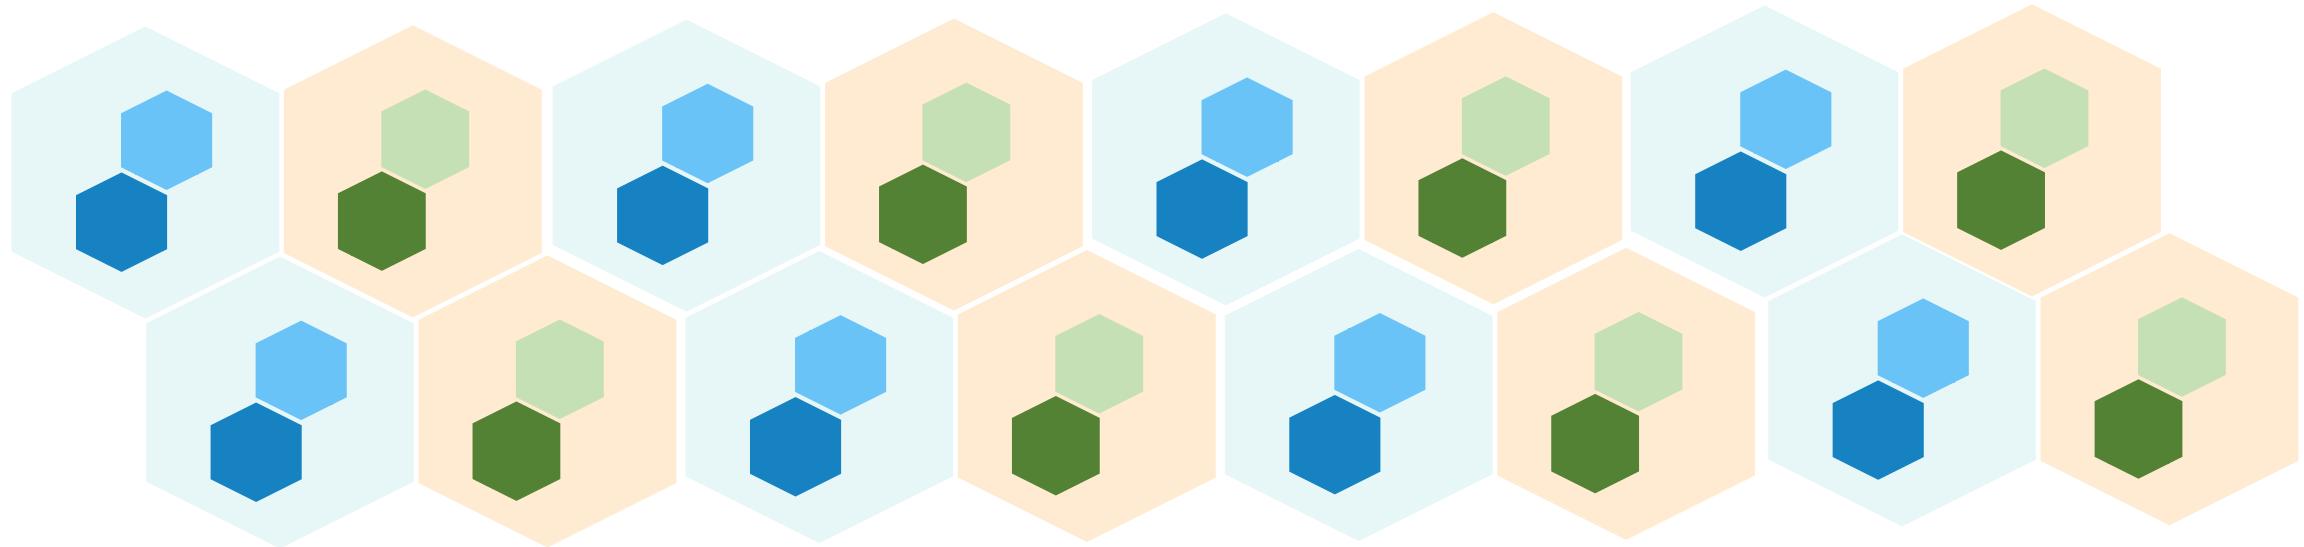
Are we good ??

Not yet.. Not completely

# Perils of consolidation ..



# Perils of consolidation ..



Imagine managing 100s of PODs and  
1000s of PVCs allocated from 100s of Drives

Are we good ??

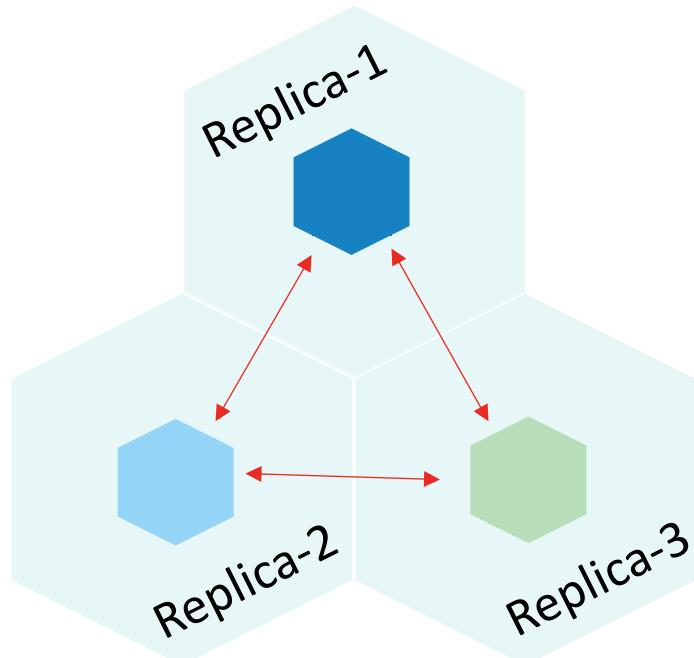


MySQL Admin is Happy

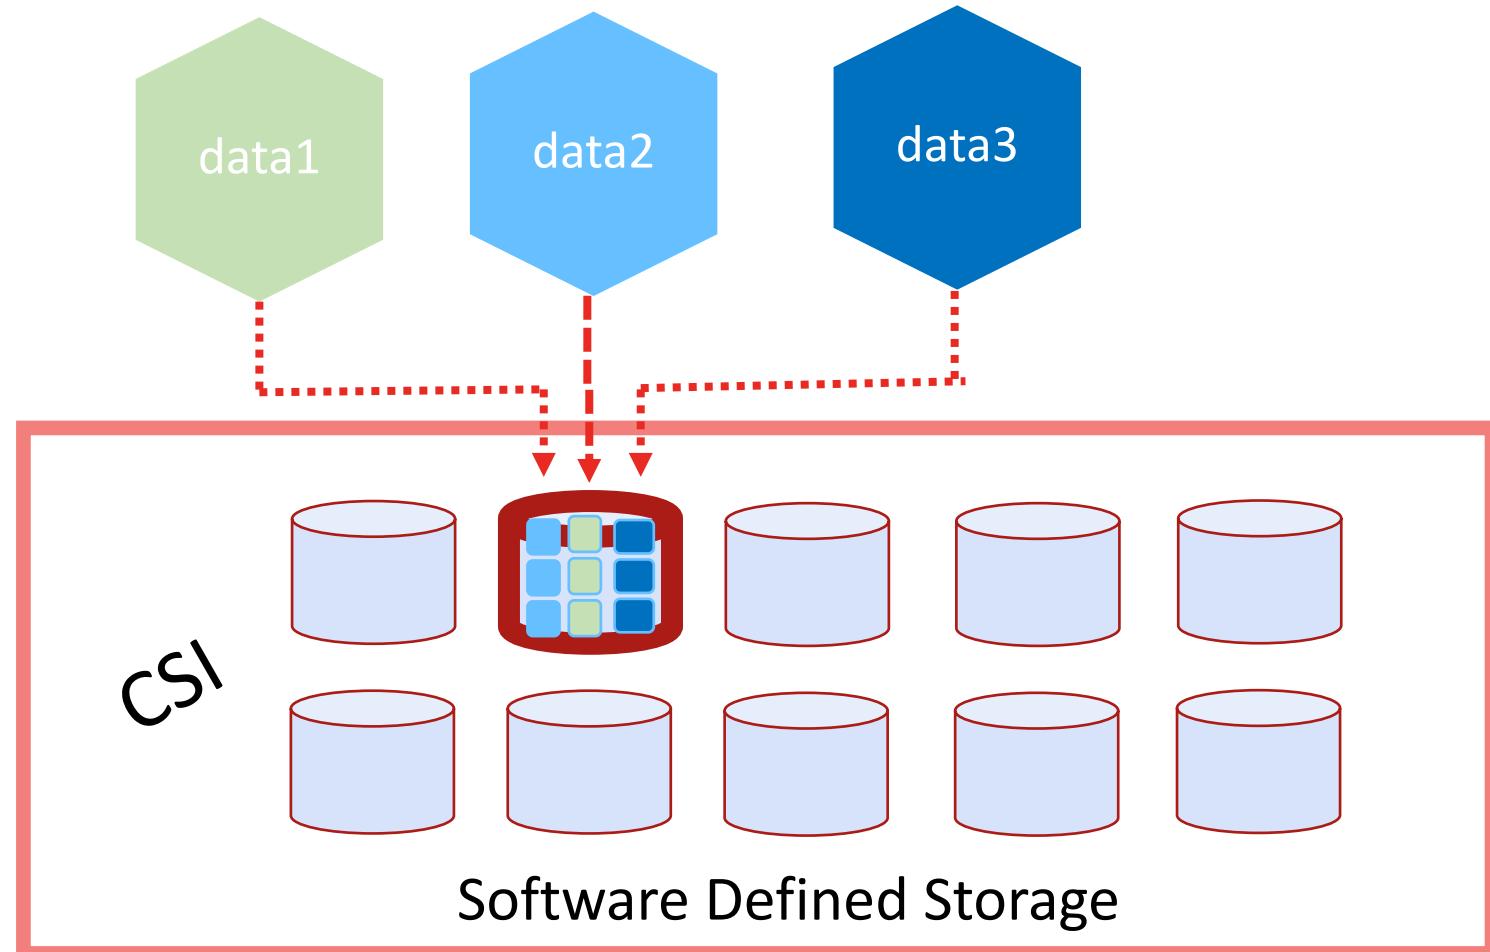
Cassandra Admin is NOT



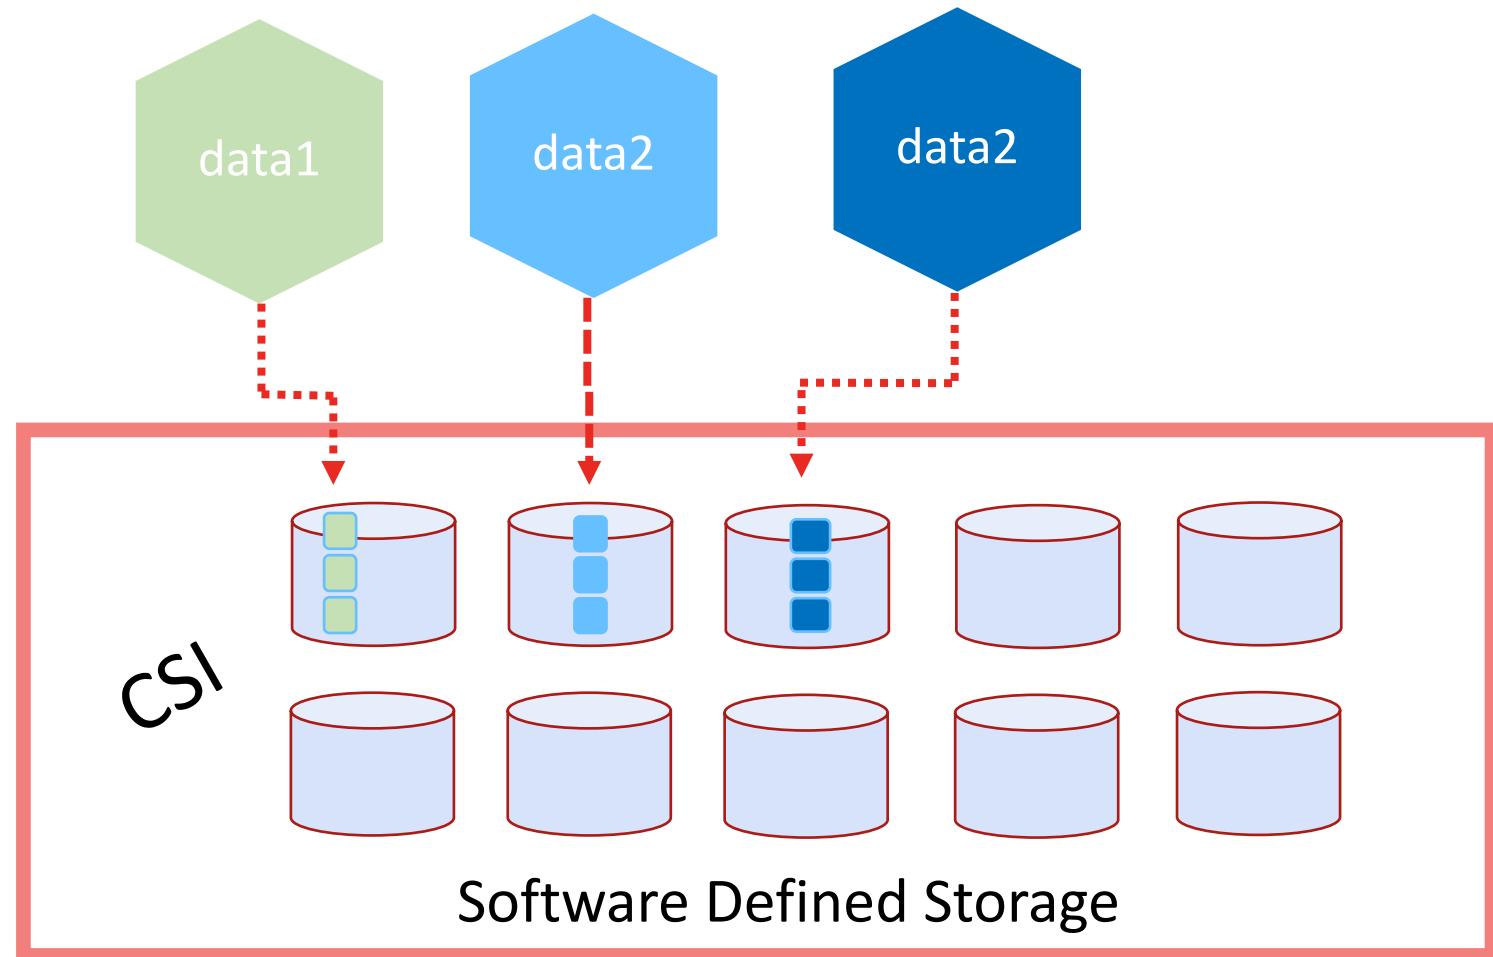
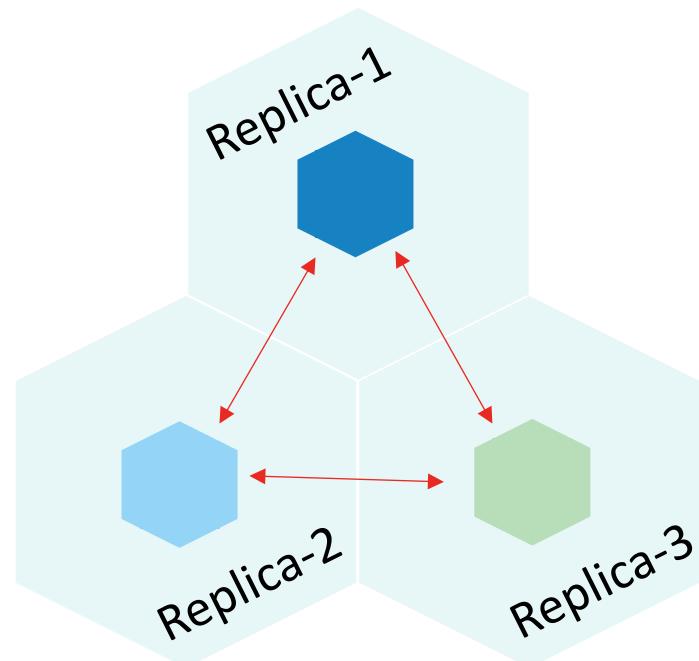
# What's with Cassandra ??



**Still resilient to disk failure ???**

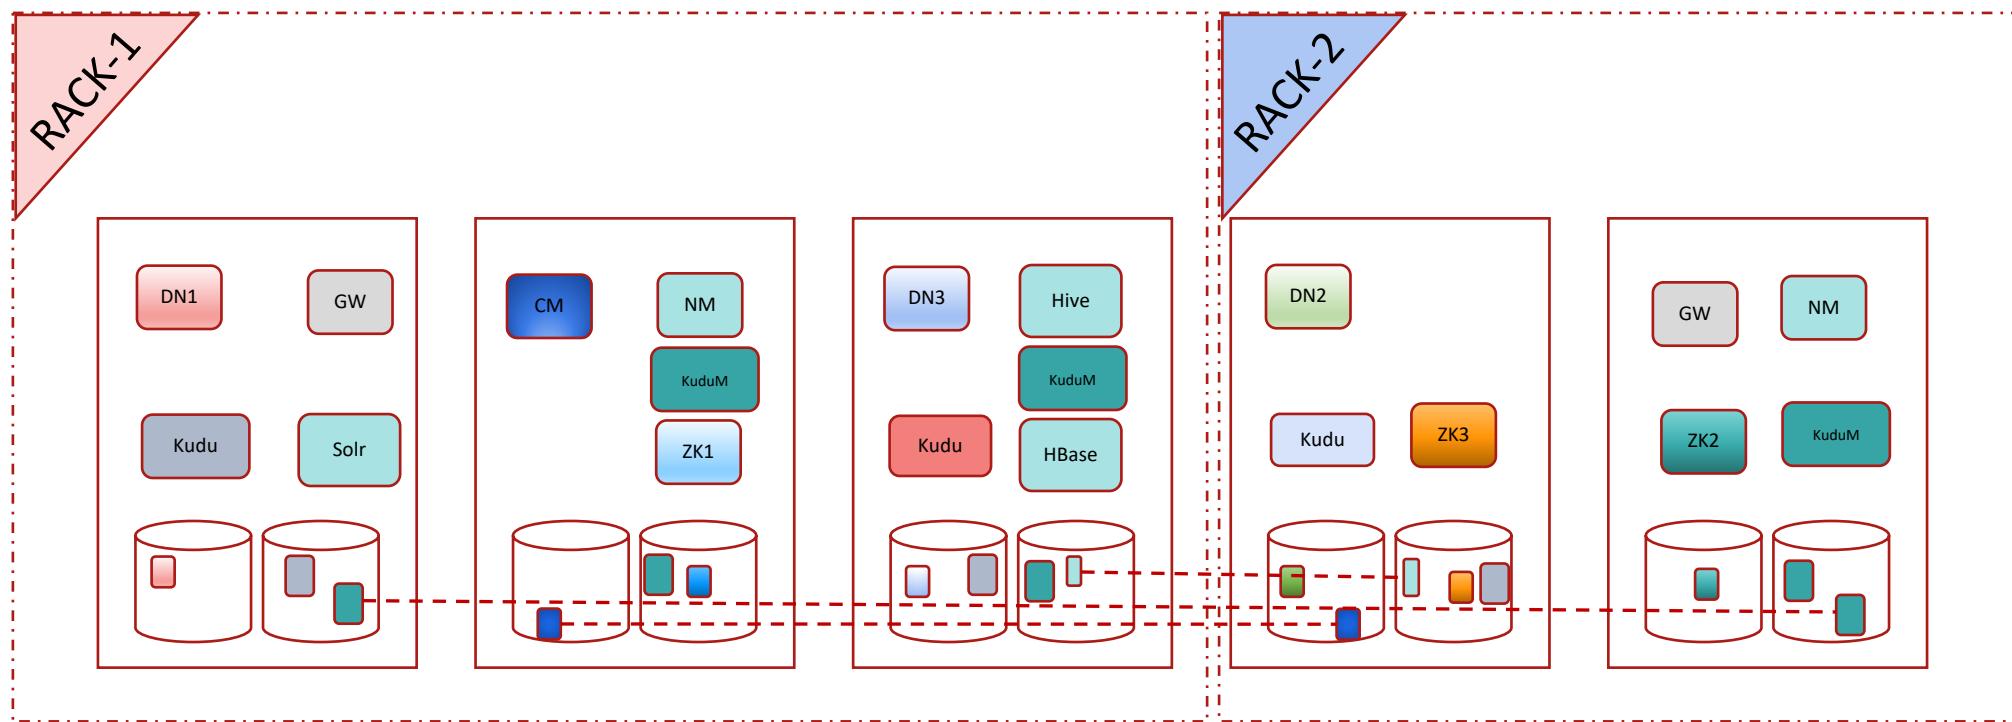


# Let us fix it ...



# What about Hadoop (HBase, Kudu) ?

# Hadoop components



Compute  
anti-affinity

Location  
Awareness  
Rack / DC

Storage &  
Compute  
Affinity

IO patterns  
QoS

High  
Availability

# Big data deployment and management challenges



**Storage workload types  
(IO patterns / QoS)**

Compute  
anti-affinity

**Location Awareness  
Rack / DC**

**Storage & Compute  
Affinity**

**High Availability**

**Scale-out compute  
and storage**

**Data Protection  
(Backup / DR)**

**Snapshot / Rollback**

# Recap

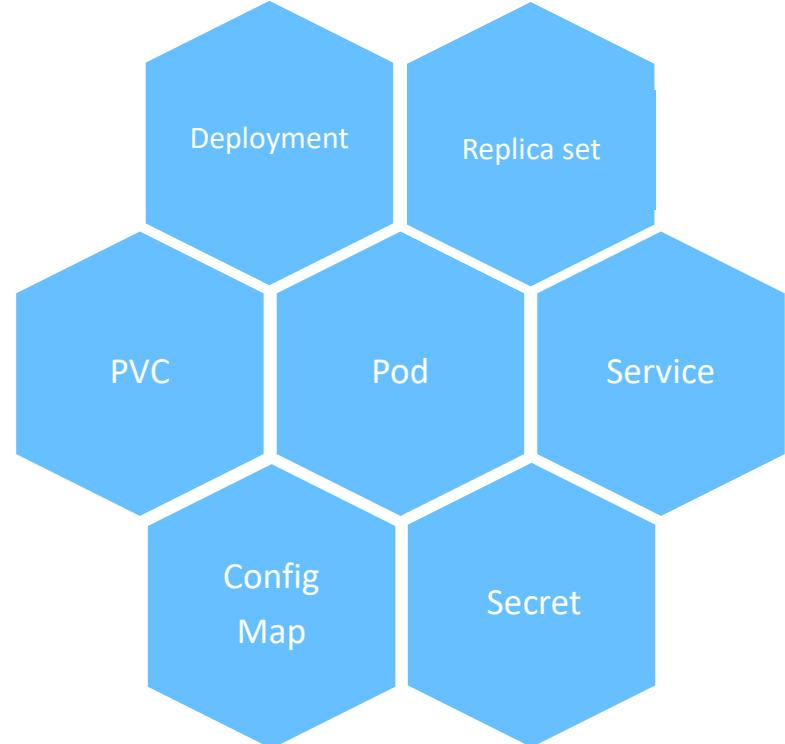
- › Data-heavy applications deal with Multiple volumes
- › Every volume will have different IO characteristics
- › Consolidation (packing) makes the problem even harder
- › Application Replication (Cassandra / Mongo) makes the allocation tricky

What are we looking for.....???

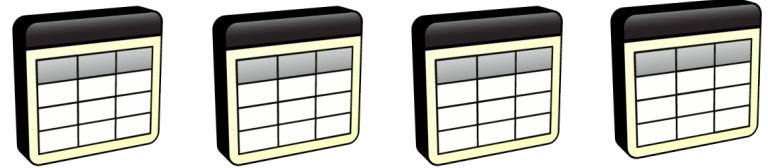
Application Aware Storage Provisioning

# Data Protection and Cloud Portability

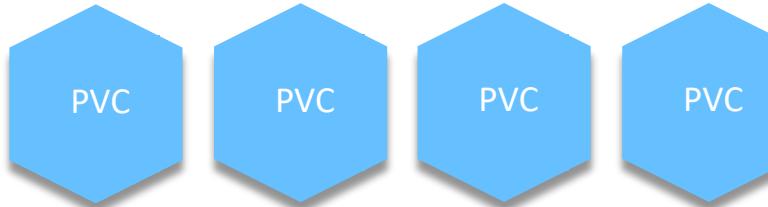
# Let us talk Data Protection



DB Checkpoints

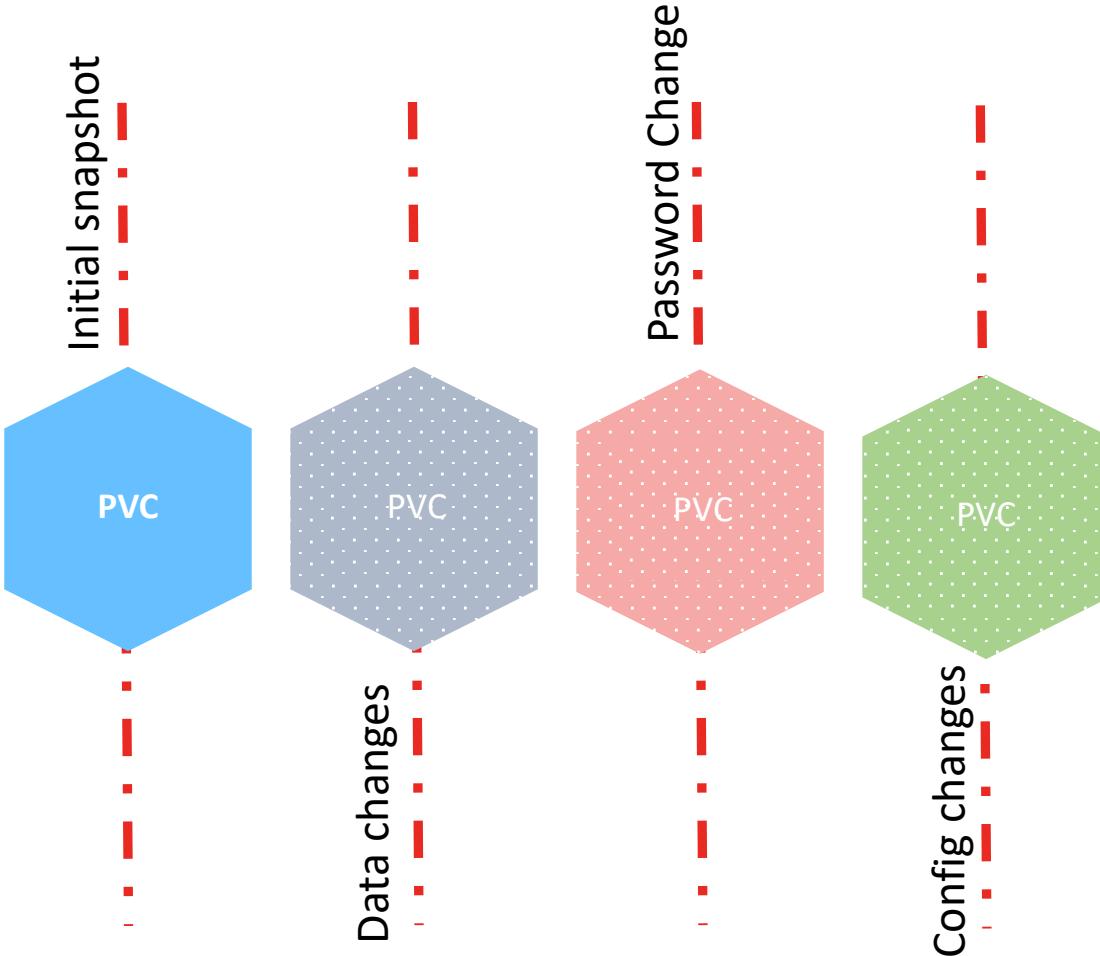
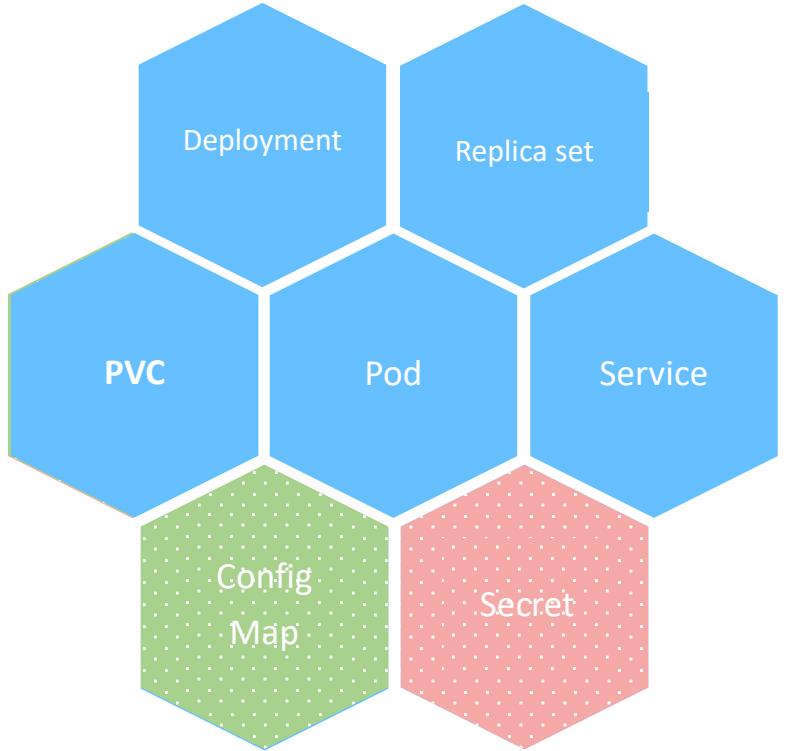


Timeline



Volume Checkpoints

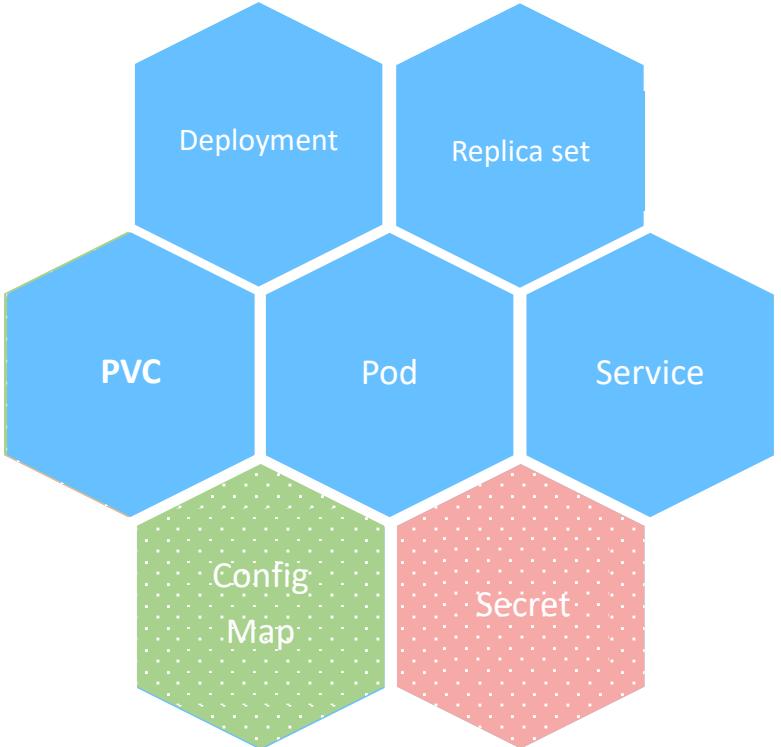
# Volume snapshots



Rollback to this snapshot

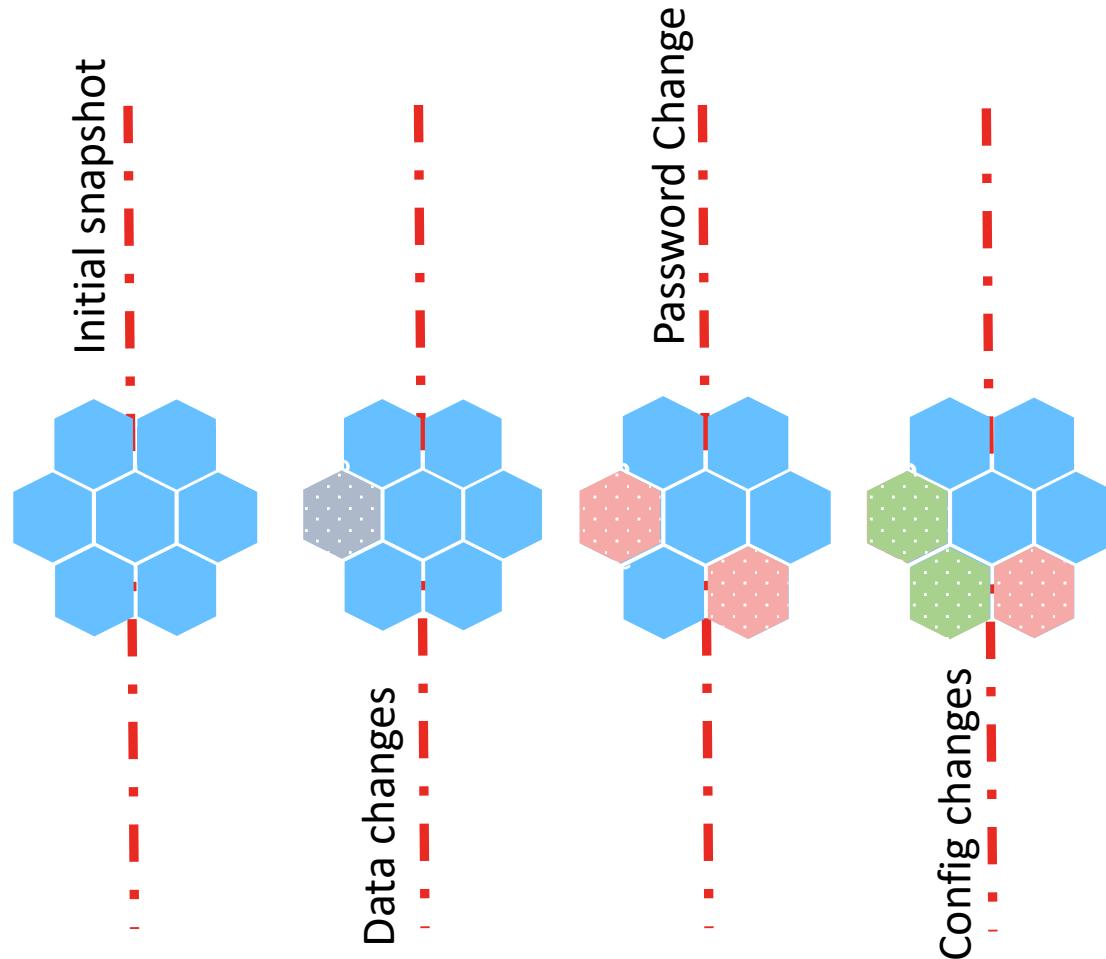
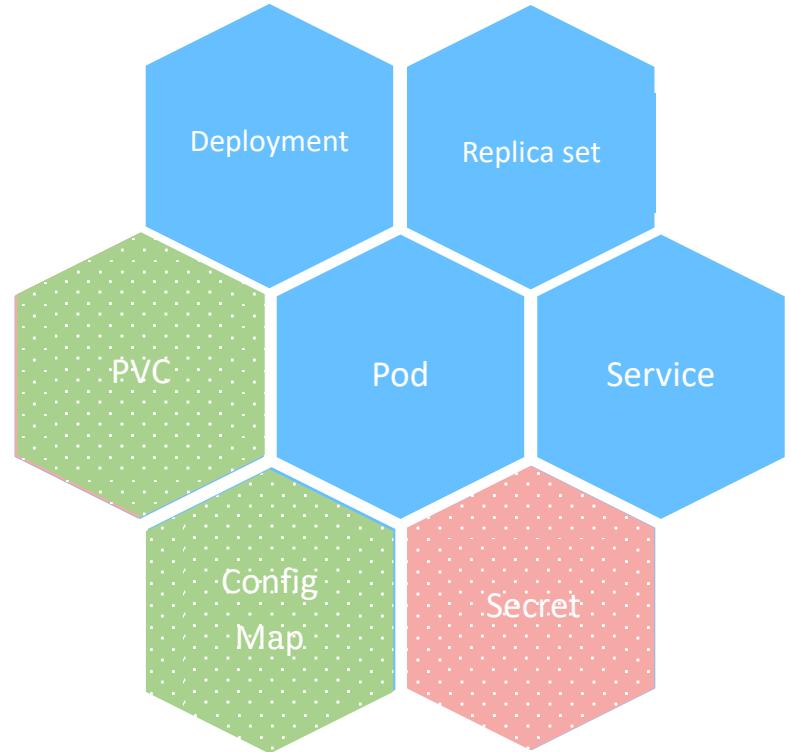
What is the problem here ?

# Volume snapshots



**Config Drift !!!**

# Let us fix it ...



# Recap (Data Protection)

- › Snapshots and backups are not just data dumps
- › Not all application have checkpoints and snapshots
- › Data snapshots are prone to config drift issues
- › Consistency group is a very critical construct
- › Application buffers / FS page cache will need to be flushed to disk

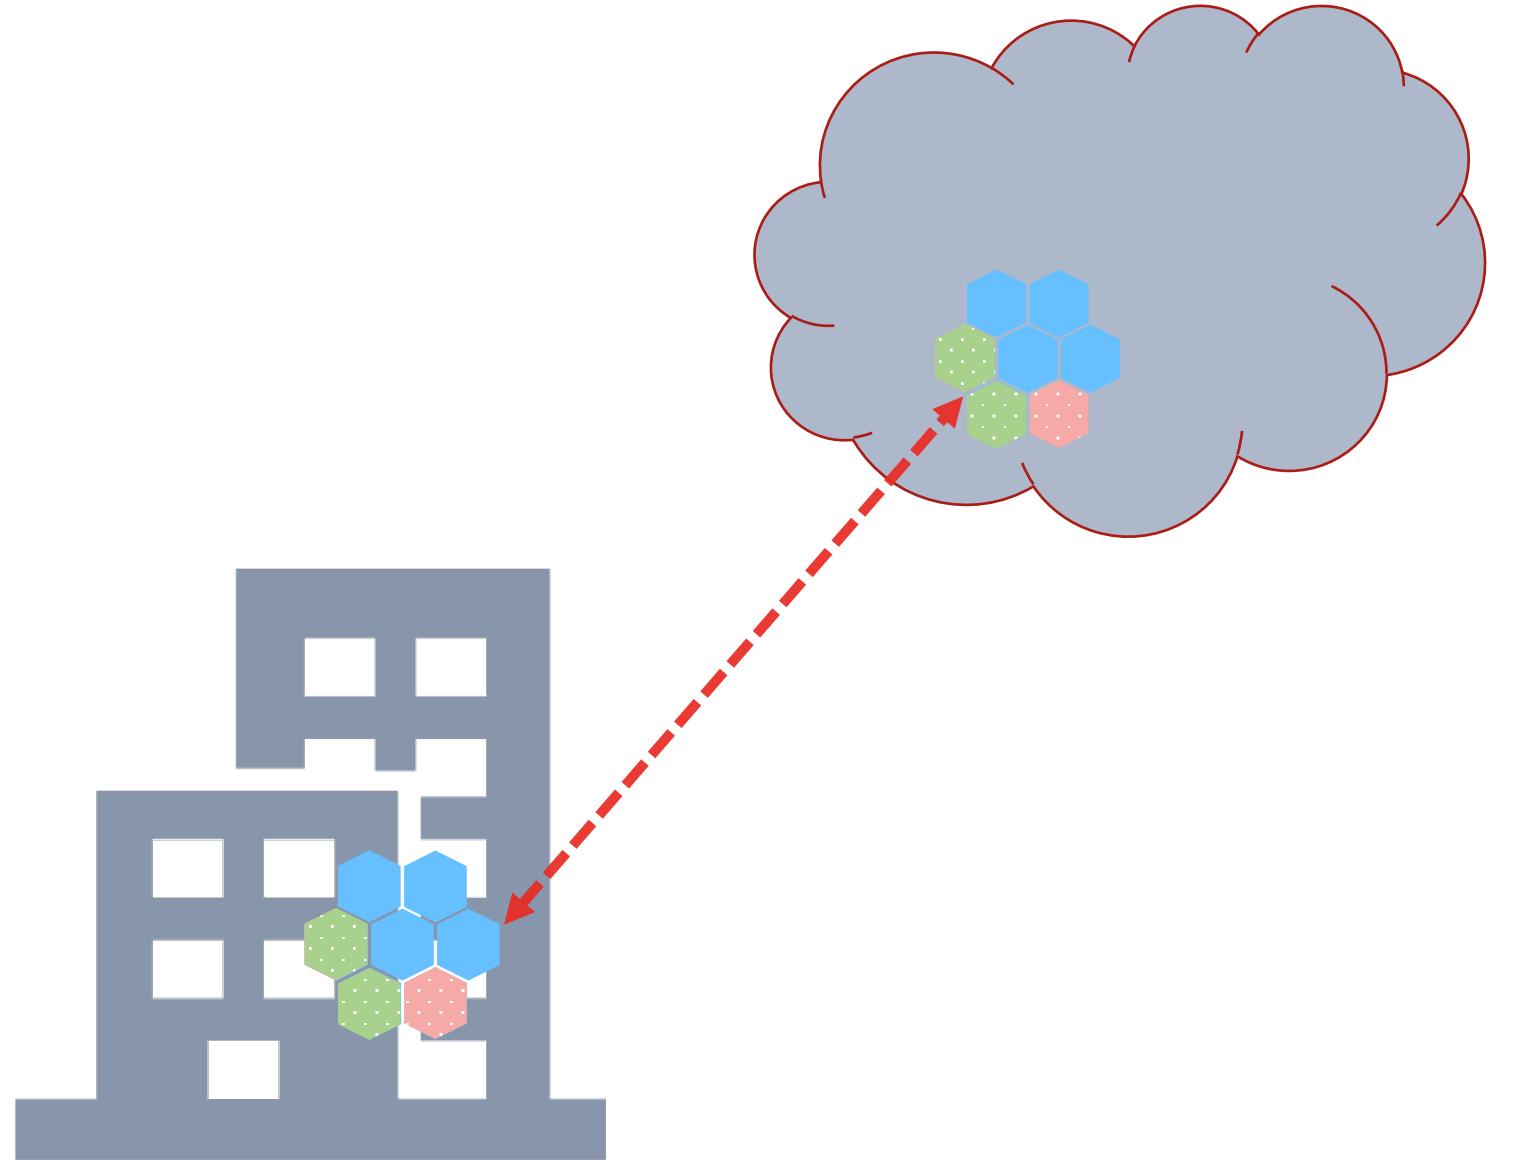
What are we looking for.....???

Application Snapshots

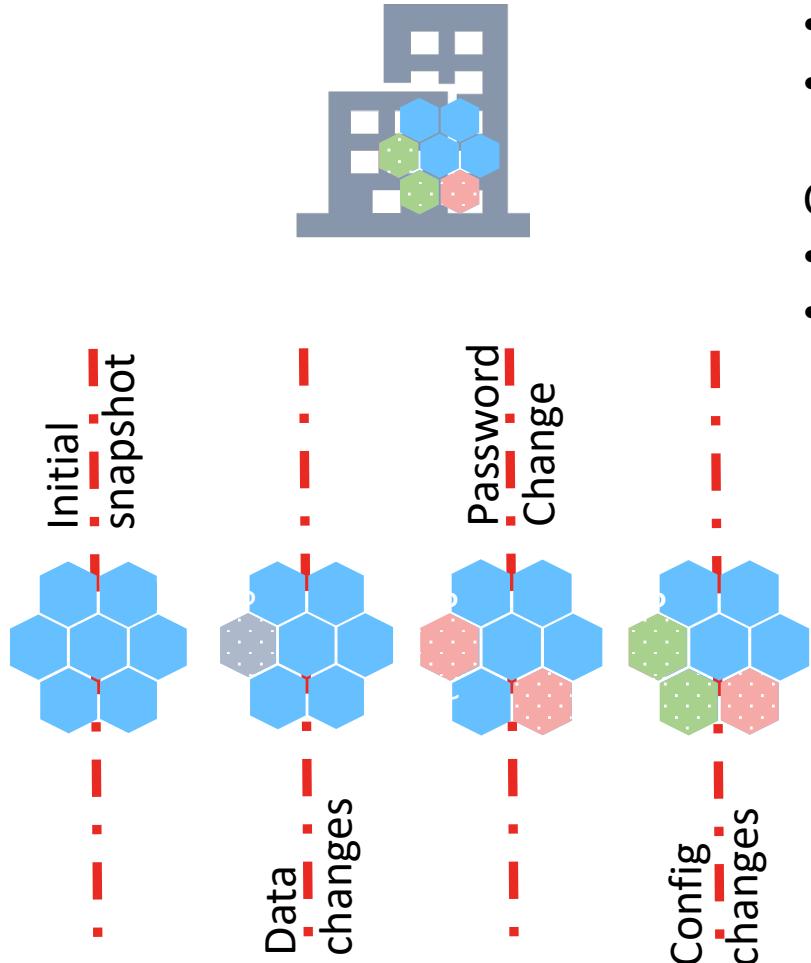
# Any other challenges

# Let us talk about portability

- › Why do we need this?
  - › Hardware refresh
  - › Datacenter migration
  - › Vendor lock-in
  - › Performance
  - › Test / Dev setups
  - › Upgrade firedrills



# Application backups

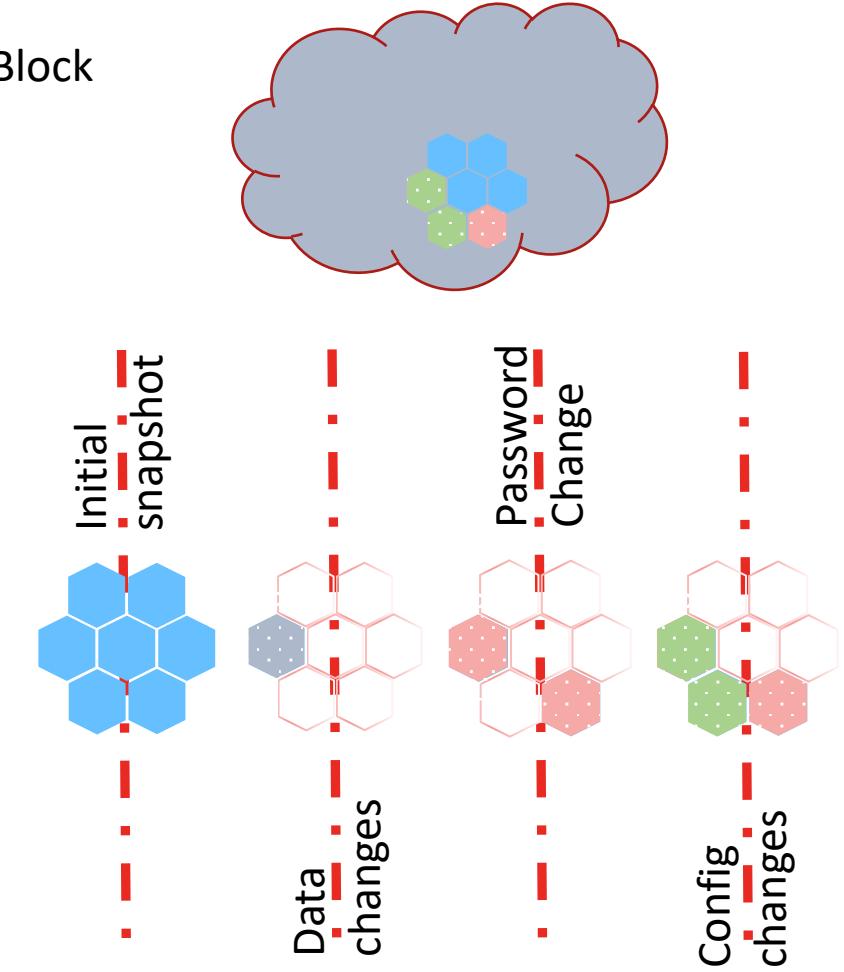


Time:

- Avoid full rehydration to Block
- Rehydrate on demand

Cost:

- Use Object store (Cheap)
- Send differentials



# Collaborate on Applications using a Git-like workflow

## Use Cases:

- Clone databases from prod to dev/test for running reports
- Validate upgrades before applying to production
- Enable git like push/pull for geo-dispersed teams to collaborate

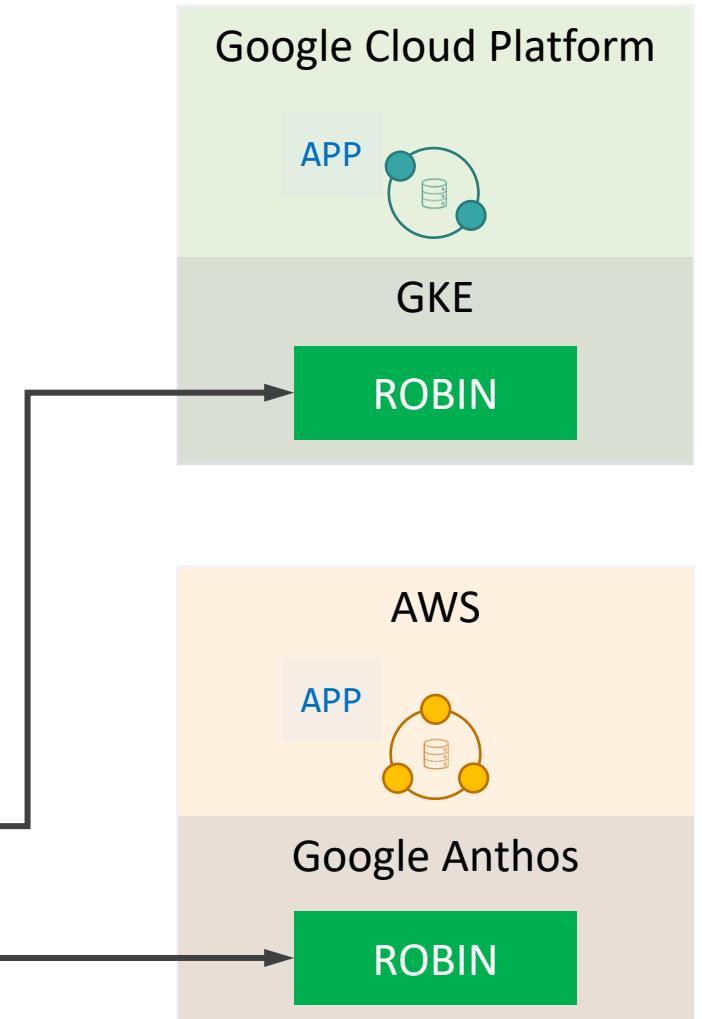


**STEP1: robin snapshot mysql mysql-snap**

**STEP2: robin clone mysql-snap testdev-mysql**

**STEP3: robin push mysql-snap gcs://bucket**

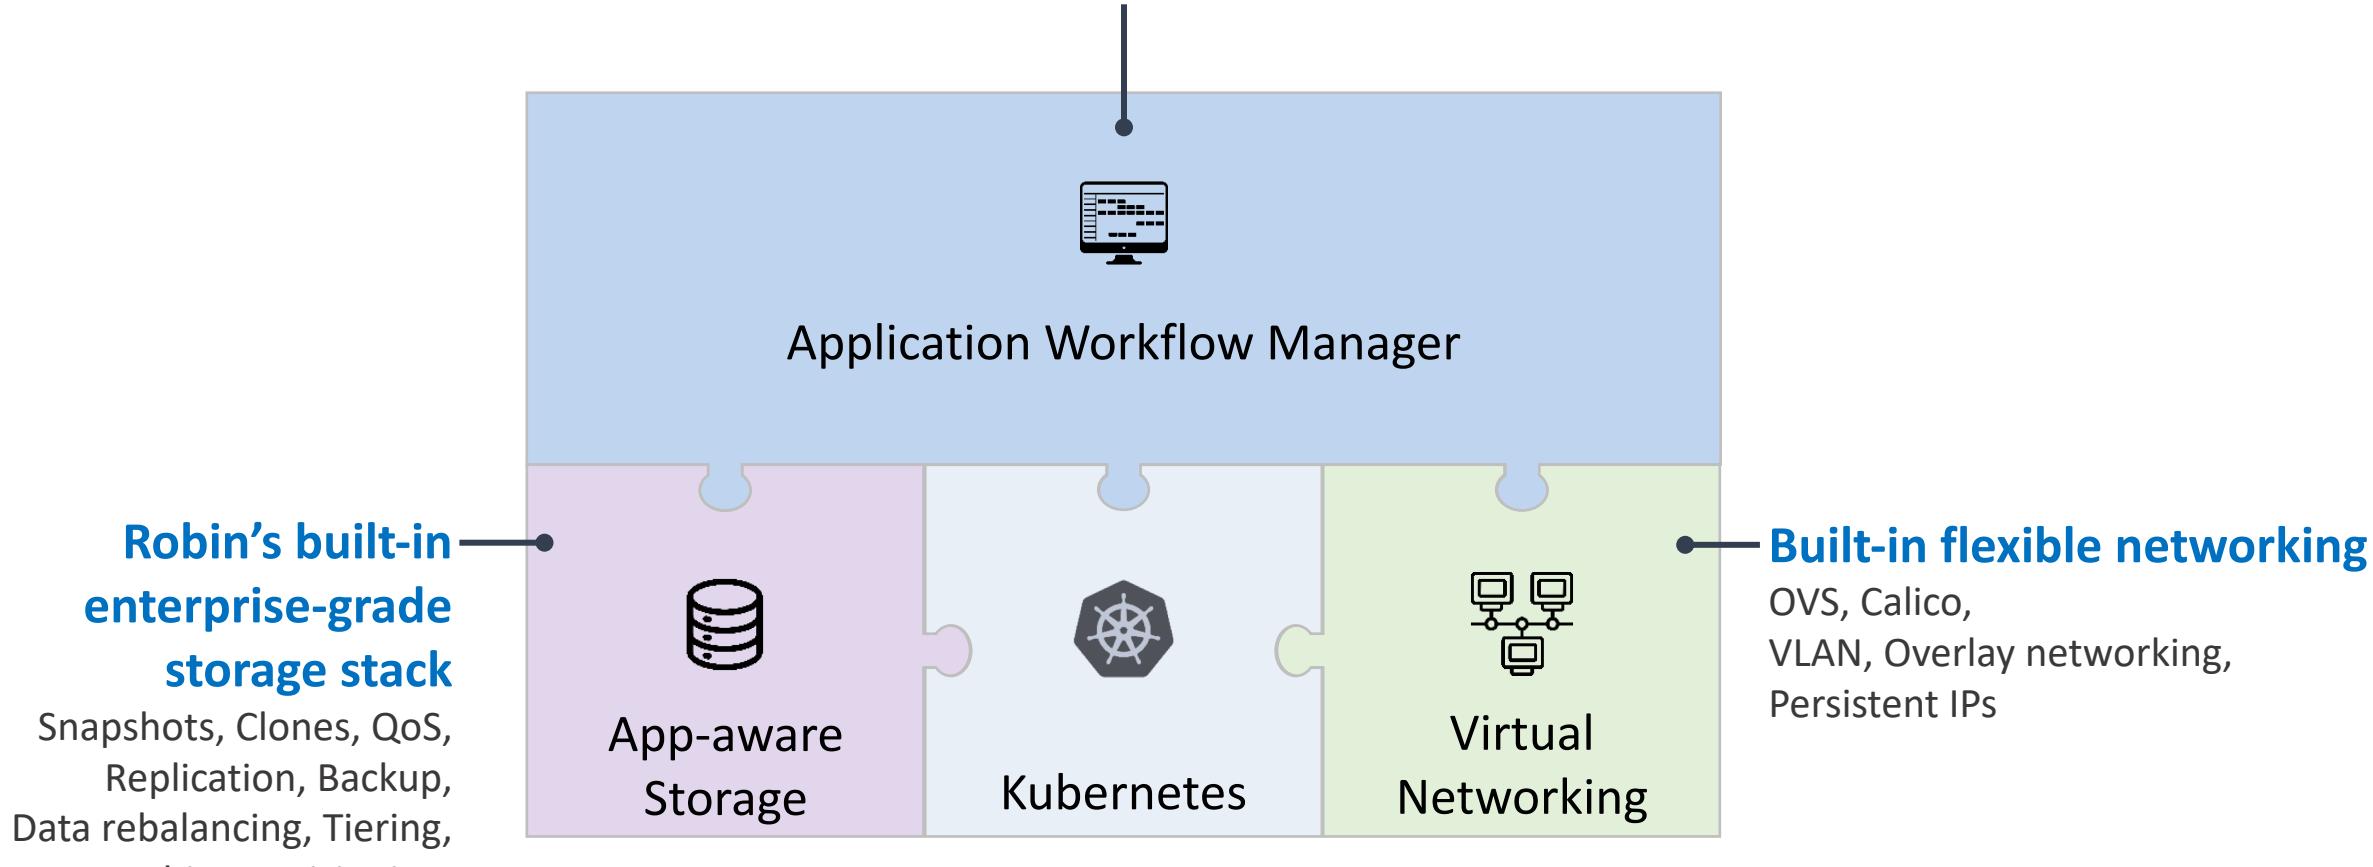
**STEP4: robin pull gcs://bucket/mysql-snap mysql**



# Robin Architecture Overview

**1-click application Deploy, Snapshot, Clone, Scale, Upgrade, Backup**

Application workflows configure Kubernetes, Storage & Networking



**Works any where**



**ROBIN software allows you run complex Big Data and Databases on Kubernetes**  
**(Storage + Networking + Application Workflow Management + Kubernetes)**

## DEPLOYMENT PROOF POINTS

**11 billion security events ingested and analyzed a day**  
**(Elasticsearch, Logstash, Kibana, Kafka)**

**6 petabytes under active management in a single ROBIN cluster**  
**(Cloudera, Impala, Kafka, Druid)**

**400 Oracle RAC databases managed by a single ROBIN cluster**  
**(Oracle, Oracle RAC)**

# Product Demos

## Demo Series 1

ROBIN's Kubernetes-native Storage and Data Management stack to deploy and manage day-2 operations of Stateful Apps

[Snapshot and Rollback](#)

[Clone](#)

[Backup](#)

[Cloud-portability](#)

## Demo Series 2

ROBIN's Super-Operator for Enterprise Apps

[1-click Deploy, Cloud-Sync and Cloud-Motion](#)

# [get.robin.io](https://get.robin.io) Portal

- › [get.robin.io](https://get.robin.io): Download Robin Storage for free
- › [docs.robin.io](https://docs.robin.io): Read product documentation
- › [slack.robin.io](https://slack.robin.io): Interact with Robin engineering

The screenshot shows the get.robin.io portal homepage. At the top, there's a navigation bar with the ROBIN logo, DOWNLOAD, ACTIVATE, MY LICENSES, DOCS, SLACK, and a user profile icon. Below the navigation, a main heading reads "Supercharge Kubernetes to Run Postgres" with a subtitle "1-click Deploy, Scale, Snapshot, Clone, Backup, Migrate". A blue button says "Download and experience for free". Below this, there's a large image of a laptop screen displaying a dashboard with a grid of application icons. The icons represent various services and databases, including cloudera, druid, hortonworks, mapr, spark, cassandra, datastax, elasticsearch, elk, kafka, mongo, centos, nginx, redhat, ubuntu, db2, mariadb, mysql, oracle, oracle-rac, saphana, sqserver, weblogic, and couchdb. At the bottom, a section titled "What is ROBIN?" explains that ROBIN extends the agility, efficiency, and portability of Kubernetes to all applications, even complex Big Data, Databases, AI/ML and Custom Apps, on any infrastructure, On-Premise, Hybrid Cloud or Multi-Cloud.

# ROBIN.IO

Supercharge Kubernetes to Deliver Big Data and Databases as-a-Service

1-click Deploy, Scale, Snapshot, Clone, Upgrade, Backup, Migrate

